

The Cost of Assuming Ionisation Balance

Transient Error in Balmer-Ratio Diagnostics of
Divertor Hydrogen Plasmas

Nikhil Jain

Department of Chemical Engineering
Indian Institute of Technology Kanpur

Supervisor: Prof. R. G. S. Pala

September 2026

Certificate

It is certified that the work contained in this thesis, entitled *The Cost of Assuming Ionisation Balance: Transient Error in Balmer-Ratio Diagnostics of Divertor Hydrogen Plasmas*, by Nikhil Jain, has been carried out under my supervision and that it has not been submitted elsewhere for a degree.

Prof. R. G. S. Pala
Department of Chemical Engineering
Indian Institute of Technology Kanpur
September 2026

Declaration

The work reported in this thesis is my own. The standard it is written to is that every numerical result is produced by a script in the accompanying repository and quoted with the script and the data file that produced it, and that every value taken from the literature is cited.

Nikhil Jain
September 2026

Abstract

A Balmer line ratio is inverted for divertor electron temperature at known density against a collisional–radiative table at fixed (T_e, n_e) . Indexing on two axes requires the neutral-to-ion ratio to be a function of those axes alone, true only after ionisation balance settles, which a divertor during an edge-localised mode has not. That escape, one channel dominating so the ratio cancels, is unavailable for H_α and H_β . **What follows is the cost of that closure where it is adopted, not a demonstration that divertor spectroscopy fails.**

A 43-state time-dependent model of atomic hydrogen, resolved in ℓ to $n = 8$ and bundled to $n = 15$, is integrated through abrupt temperature steps on a 50 by 8 grid over 1 to 10 eV and 10^{12} to $10^{15}/\text{cm}^3$.

The approximation usually blamed is not the one that fails: the quasi-steady-state closure for the excited manifold holds to 8.66×10^{-6} at the reference point. What lags is the ground state, holding almost every atom and turning over four orders of magnitude slower.

That cost factorises exactly as $\varepsilon_{\text{plateau}} = |\exp(\bar{S} G \Delta \ln T_e) - 1|$, with $\bar{S}$ the reservoir sensitivity of the line ratio and G the reservoir gain. The sensitivity obeys a ceiling, $|f_3 - f_4| < 1$, set by the structure of the linear system, not by any rate coefficient.

Over 100 μs , 45 of the 448 pairs above 2 eV exceed 10 % error, the worst 17.5 %; timescale separation licenses the reduction and predicts none of it. Inverted, the ratio has no solution for 166, reaches one only across a fold for 44, and over the remaining 238 reads a 4.81 % step back with 48 % median error.

Two measured boundaries exclude the model below 2 eV. Above it, the closed-parcel assumption every magnitude depends on fails at all 45 breakdown pairs. The structural results survive that; the magnitudes are conditional on it.

Acknowledgements

I thank Prof. R. G. S. Pala for supervising this work, and in particular for insisting that a check which has been made to pass is a check that no longer verifies anything. Several results in this thesis exist because a test was allowed to fail and then explained.

Contents

Abstract	3
Acknowledgements	4
Nomenclature	14
1 Reading the Light from a Divertor	16
1.1 Why the divertor is hard	16
1.2 Light from a hydrogen atom	19
1.3 What a spectrometer records	20
1.4 From a ratio to a temperature	21
1.5 The assumption inside the table	23
1.5.1 A worked example, because the distinction is easy to assert and hard to demonstrate	24
1.6 A plasma that never finishes changing	26
1.7 What is already known	27
1.8 The research question and the hypothesis	30
1.9 Outline of the thesis	31
2 The Atoms	32
2.1 What happens to a hydrogen atom in a plasma	32
2.1.1 The atom, from scratch	32
2.1.2 Six processes, and how fast they run	33
2.2 The state space	34
2.3 From a cross section to a rate coefficient	36
2.3.1 Detailed balance, and why de-excitation is never computed directly	37
2.3.2 Two hydrogen energies, and why both are right	38
2.4 Electron-impact excitation	38
2.5 Ionisation and recombination	40
2.5.1 Ionisation: the only true loss	40
2.5.2 Recombination: two mechanisms, two density dependences	41
2.6 Radiative decay	42
2.6.1 The photons are assumed to leave	44
2.7 Proton-impact ℓ -mixing, and why it is not optional	45
2.7.1 Why proton collisions mix ℓ	45
2.7.2 The implemented formula	46
2.7.3 Without it, $n = 2$ is not a physical manifold	47
2.8 How well do we know these rates?	48
2.8.1 Where the $n = 5$ disagreement is localised	51

2.8.2	Does any of it reach the result?	53
2.9	Two errors found by audit	56
2.9.1	The ℓ -mixing Coulomb logarithm	56
2.9.2	The eigenvalue magnitude filter	57
2.9.3	What both errors have in common	57
2.10	What this model assumes	57
2.11	Summary	58
3	Two Clocks	61
3.1	The rate equation	61
3.1.1	A balance on every level	61
3.1.2	Matrix form	63
3.1.3	Why recombination is a feed stream, not a matrix element	64
3.1.4	Units, and a source that is quadratic in density	64
3.2	Conservation, and how we know the matrix is right	64
3.2.1	A worked example	65
3.2.2	The measured residual	65
3.3	Two clocks	66
3.3.1	Why the system has natural timescales at all	66
3.3.2	The spectrum	66
3.3.3	What the two modes are	67
3.3.4	The separation across the operating range	69
3.3.5	Is the definition robust?	70
3.3.6	How much of the separation is atomic data?	71
3.3.7	What the separation does and does not license	72
3.4	The quasi-steady-state approximation	72
3.4.1	The shortcut	72
3.4.2	Partition and elimination	72
3.4.3	The reduced dynamics	73
3.4.4	When is it valid?	73
3.5	Two reference states	75
3.5.1	Why ratios	75
3.5.2	The hidden variable	75
3.5.3	Does the observable actually depend on it?	78
3.5.4	The ground-fed fraction, and where the sensitivity vanishes	79
3.5.5	The same result for what a spectrometer measures	82
3.5.6	Two reference states, and they are not the same	85
3.5.7	Two error measures	86
3.5.8	What a two-parameter table must assume	86
3.6	A caveat: the modes are not orthogonal	87
3.7	Summary	88
4	Does the Model Work?	89
4.1	What a check can and cannot show	89
4.1.1	Three things that get called validation	89
4.1.2	Severity, and how to measure it	90
4.1.3	The three grades used below	90
4.2	The conservation check is exact by construction	91

4.2.1	Four injected faults	91
4.2.2	Why, and what the check is actually worth	91
4.3	The gate suite, read one gate at a time	92
4.3.1	Gate A: detailed balance, and where its physics content actually lives	92
4.3.2	One energy ladder on both sides, which is severe	93
4.3.3	Gate B: the coronal limit	93
4.3.4	Gate C: what the Saha gate tests, which is not the Saha limit	94
4.4	Gate E: the timescale hierarchy	94
4.4.1	The gates do not stop anything	95
4.5	The checks that could have failed	95
4.5.1	The superposition residual	95
4.5.2	The reduced-versus-full ratio test	95
4.5.3	Pinned hashes on the inputs	96
4.5.4	Guards that recompute rather than trust	96
4.5.5	Two checks that fired, and one that was withdrawn	96
4.5.6	The bound that is a theorem	97
4.6	Conditioning, precision, and state-space convergence	97
4.6.1	The inverse is well behaved, and here is how well	97
4.6.2	Where the ladder was cut	98
4.7	Is the shell ratio the line ratio?	100
4.7.1	And is the plasma transparent to its own Balmer light?	100
4.8	Two tests of the reduction, and their expected outcomes	101
4.8.1	Does the open boundary manufacture the slow clock?	101
4.8.2	Is the single-slow-mode estimate a bound on the time-averaged error?	102
4.9	External comparison: Fujimoto’s population coefficients	104
4.9.1	Timescales, at a point that lies exactly on this grid	104
4.9.2	The recombining coefficients agree	104
4.9.3	The ionising coefficients at low p disagree, and the target is not verified	105
4.10	Gate D fails, and the failure is in the gate	107
4.10.1	The diagnosis, and the test that settles it	107
4.10.2	Three further defects in the gate itself	109
4.10.3	What is claimed, and what is not	110
4.11	Errors found by our own audit	110
4.11.1	The ℓ -mixing Coulomb logarithm	110
4.11.2	And ℓ -mixing is saturated	111
4.11.3	The eigenvalue magnitude filter	111
4.11.4	The central quantity was misnamed	112
4.11.5	A statement in Chapter 3 that measurement contradicted	112
4.12	Reproduction, and what a document cannot certify	112
4.12.1	Regeneration from the canonical matrix	112
4.12.2	Independent re-derivation	113
4.12.3	A prediction that was falsified, and improved the claim	113
4.12.4	One item demoted: the May/August cross-validation	114
4.12.5	Two provenance defects that are not repaired	114
4.13	What this chapter establishes	115
4.13.1	What the model may be used for	115
4.13.2	What it may not be used for	117

5	What the Model Says	118
5.1	What happens after a temperature step	119
5.1.1	Two error measures, and why the step size is part of each	119
5.1.2	The closure everyone worries about is not the dominant error	121
5.2	The error maps	122
5.2.1	The density structure is a broad crest, not a line	125
5.3	Why a percentage is not a result	126
5.4	The exact response to a finite step	127
5.4.1	The split has exactly two channels	127
5.4.2	A better variable than the departure coefficient	128
5.4.3	The decomposition	128
5.4.4	The reservoir gain, and why it is the object to report	129
5.4.5	A worked number at the benchmark	131
5.4.6	Use the exact form, not the linearisation	133
5.5	A bound that no atomic dataset can exceed	133
5.5.1	What the bound does not constrain	135
5.5.2	The inversion, measured	136
5.6	Where the error is largest, and why	139
5.6.1	Across density: the operating point, not the span	139
5.6.2	Three tests of the mechanism	141
5.6.3	Alternatives that would produce the same appearance	141
5.6.4	Across temperature: a displacement effect	142
5.6.5	Three explanations this work does not support	143
5.7	Does the transient visit the plateau?	143
5.8	How long the error lasts, and what an ELM averages	144
5.8.1	The census, inside the defensible scope	146
5.8.2	Two structural objections, tested	148
5.8.3	Stepping temperature and density together as a two-parameter perturbation stress test	149
5.9	Timescale separation does not predict the reservoir-closure error	151
5.10	Where the quantitative scope begins	155
5.10.1	Holding the electron density fixed stops being consistent	155
5.10.2	The plasma is not optically thin at the cold edge	156
5.10.3	The boundaries coincide	158
5.11	What this chapter establishes	158
6	What This Model Cannot Say	162
6.1	Optical thickness: when the plasma cannot let its own light out	163
6.1.1	The optically thin assumption and how it reaches the slow clock	163
6.1.2	Where the assumption fails, in numbers	163
6.1.3	Self-consistent Lyman escape factors applied to the matrix	164
6.1.4	Above two electron-volts, nothing moves	165
6.1.5	Below two electron-volts, the cold-edge number is inflated threefold	166
6.1.6	Below 1.15 electron-volts the ridge location is not determined	167
6.1.7	A contradiction resolved, and what is still open	167
6.2	Transport: the atoms do not stay where the model puts them	168
6.2.1	An objection the model cannot answer from inside	168
6.2.2	Batch-reactor parcel against a recycling divertor	168

6.2.3	Why this bears on the result specifically	168
6.2.4	The assumption fails hardest where the result is generated	169
6.2.5	What this reaches, and what it does not	170
6.2.6	The claim, restated as the conditional it is	170
6.2.7	One chord, one point: the line-of-sight integration	171
6.3	No molecular channels, in a gas that is mostly molecules	171
6.3.1	An atomic-only state space	171
6.3.2	Molecular activated recombination as a third feed stream	171
6.3.3	How much gas is involved	172
6.3.4	What survives it, and what does not	172
6.4	Where the ridge sits depends on what the model assumed about neutrals	173
6.4.1	Mechanism and location as separate claims	173
6.4.2	The neutral density as the controlling variable	173
6.4.3	Ridge location under a rescaled neutral density	173
6.5	The model does not identify the crest with detachment	174
6.5.1	The verbal match between the ridge and detachment	174
6.5.2	The ridge density against the detached-divertor band	174
6.5.3	A structure without a named tokamak region	175
6.6	The open ion boundary: a closed-parcel test	175
6.6.1	The fixed ion density in the rate equation	175
6.6.2	Closing the manifold with a 44th ion state	175
6.6.3	The caveat that must not be softened	176
6.7	An external benchmark this model fails, and what the failure measures .	176
6.7.1	One population coefficient agrees and one does not	176
6.7.2	The recombination-fed and ground-fed coefficients r_0 and r_1	176
6.7.3	The disagreement, and what it measures	176
6.7.4	The source's closure, imposed and tested	177
6.7.5	Two limitations of the benchmark, stated precisely	177
6.7.6	What the benchmark establishes, and what it does not	177
6.8	The top of the level ladder, and what the bundling costs	178
6.8.1	Truncation at $n_{\max} = 15$ and bundling above $n = 8$	178
6.8.2	Three consequences, in decreasing order of severity	178
6.8.3	Collisional ℓ -mixing against escape in the bundled shells	178
6.9	Remaining assumptions, and the sensitivities that are known	181
6.10	Scope: five questions this thesis leaves open	182
6.11	What is left standing	183
6.11.1	What this model cannot say	183
6.11.2	What this model can say, with the objections computed	183
6.11.3	The boundary of the claim	184
7	What It Means	185
7.1	The excited-state closure holds; the reservoir closure does not	185
7.2	A forward criterion for the equilibrium lookup	186
7.2.1	Five things that follow, and one that does not	187
7.3	What is new, at the smallest size that is defensible	189
7.4	What this work cannot claim	192
7.5	What would settle what is open	194

A	State Ordering	196
A.1	What the state space is, and where it is cut	196
A.2	The assumption the ordering embeds	196
A.3	Two hydrogen ionisation energies	198
A.4	Provenance	198
B	Atomic Data Sources	199
B.1	The stack, end to end	199
B.2	Which checks can fail, and which cannot	199
B.3	The seam at $n = 10$	199
B.4	Two constants, and why both appear	201
C	Numerical Methods and Convergence	202
C.1	How the systems are solved	202
C.2	Conditioning	203
C.3	Convergence tests with a stamped artifact	203
C.4	State-space convergence, and what is not established	204

List of Figures

1.1	Where the hidden assumption enters divertor spectroscopy	22
2.1	The 43-state space and the five processes the rate matrix contains	35
2.2	CCC benchmarked against RMPS	49
2.3	Which calculation is anomalous at $n = 5$	52
2.4	What the atomic-data disagreement costs	54
3.1	The collisional–radiative network and its slow/fast partition	62
3.2	Relaxation timescales of the 43 collisional–radiative modes	67
3.3	The observable as a function of the hidden variable	79
3.4	Ground-fed fractions and the sensitivity of the observable	80
5.1	What the plateau looks like: the excited states keep up, the reservoir does not	120
5.2	The reversal: the approximation that is doubted holds, the assumption that is not examined fails	123
5.3	Where the equilibrium-reservoir closure costs most	124
5.4	The density dependence, all eight columns	125
5.5	The two structural coefficients, and the invariant they multiply to	132
5.6	Timescale separation does not predict reservoir-closure error	152
5.7	Lyman trapping and the restricted range	157
5.8	Two independent constraints put the same floor under the quantitative scope	159

List of Tables

2.1	Process rates at the benchmark point	34
2.2	Provenance of the excitation rate coefficients	40
2.3	CCC benchmarked against Anderson RMPS	50
2.4	Rydberg-series test of the $n = 5$ disagreement	51
2.5	Which transitions carry the atomic-data sensitivity	55
2.6	Atomic data sources	59
3.1	The five slowest relaxation modes at the benchmark point	66
4.1	Fault injection into the conservation check	91
4.2	Effect of closing the ion boundary on the slow timescale	102
4.3	Accuracy of the single-slow-mode estimate	103
4.4	The validation ladder, graded by severity and provenance	116
5.1	$\varepsilon_{\text{plateau}}$ against the size of the imposed heating step, in grid intervals: the grid maximum with its location, and the same-point series at three fixed operating points. Values from <code>validation/reservoir_gain/reservoir_gain.csv</code> (<code>verify_reservoir_gain.py</code> , heating direction). The one- and two-interval maxima sit at $[0, 4]$ and the four-interval maximum at $[0, 5]$, all on the cold edge, and none is quantitative (Section 5.10).	126
5.2	The reservoir gain G at three grid points, for steps of one, two and four grid intervals. G moves by 5 to 6 % across a fourfold change in step size, the secant closing on its tangent, quantified below, while $\varepsilon_{\text{plateau}}$ itself changes by a factor 4.5 to 5.2 (Table 5.1). All nine values are the column G of <code>validation/reservoir_gain/reservoir_gain.csv</code> (<code>verify_reservoir_gain.py</code>), heating direction, and are reproduced from it to the digits printed by <code>verify_headline_provenance.py</code>	130
5.3	Density profiles of the plateau error and of the two factors of Eq. (5.4), each normalised to its own row maximum and then taken as the median over the 32 temperature rows with $T_e = 2\text{--}9\text{ eV}$, for a one-interval temperature step. From <code>verify_ridge_mechanism.py</code>	139

5.4	Why the sensitivity has an interior maximum in density, at the post-step temperature $T_e = 3.089\text{ eV}$ (temperature index 24) of the benchmark heating step. The cap $\tanh(\Delta /4)$ is nearly constant while the operating point $\ln(u^{\text{CRE}}/u^{\text{peak}})$ sweeps through zero. $\varepsilon_{\text{plateau}}$ is for a one-interval, +4.81% step. Every value is reduced from stamped artifacts by <code>verify_headline_provenance.py</code> : Δ and $\tanh(\Delta /4)$ from <code>validation/molecular_channel/</code> at the post-step index 24, and $ \bar{S} $ from <code>validation/reservoir_gain/</code> , where it is the secant of the step $23 \rightarrow 24$ and is stored against index 23. The $\ln(u^{\text{CRE}}/u^{\text{peak}})$ column takes its numerator at index 23 and its denominator at index 24; evaluating it consistently at either index shifts every entry by about 0.25 and moves neither the sign change, which stays between $j = 2$ and $j = 3$, nor the maximum of $ \bar{S} $ at $j = 3$	140
5.5	Breakdown census on the $100\text{ }\mu\text{s}$ time average of the model-defined instantaneous CR-equilibrium-reference error, counted on the single-slow-mode estimate $\hat{\varepsilon}$ of Eq. (5.16) with a one-interval step. The average is taken over a chosen duration τ_{drive} ; it is not a measurement model of an ELM. Each row states the set it counts over.	146
5.6	Duration sensitivity of the census, over the 448 pairs above 2 eV at a one-interval step. The maps throughout this chapter are computed at $100\text{ }\mu\text{s}$; $506\text{ }\mu\text{s}$ is the extrapolation to the ITER pedestal.	146
5.7	The breakdown census under a joint step, at $\tau_{\text{drive}} = 100\text{ }\mu\text{s}$. The middle row is Table 5.5's middle row, recomputed by an independent script and reproducing it exactly.	150
5.8	Partial correlation of $\log M$ with $\log \varepsilon_{\text{plateau}}$ after residualising both on a basis in $(\log T_e, \log n_e)$. Linear control is too weak to remove the shared temperature trend. Under quadratic control or richer the partial is negative in every one of the twelve combinations. Stamped in <code>validation/partial_correlation/</code> .153	
5.9	Partial correlation of $\log M$ with $\log \varepsilon_{\text{plateau}}$ over the analysed pairs at $T_e \geq 2\text{ eV}$, under the quadratic control of Table 5.8 and under saturated two-way fixed effects (one dummy per T_e row and one per n_e column). The null columns give the mean and standard deviation of the partial over 2000 permutations of $\log \varepsilon_{\text{plateau}}$ within each T_e row; no draw reached the observed value in any cell. Stamped in <code>validation/partial_fe/</code>	154
5.10	Electron-density change implied by conservation of nuclei across one step	155
6.1	Breakdown census under Lyman trapping	165
6.2	Lyman trapping at the cold-edge worst point	166
6.3	Ratio of neutral escape time to τ_{slow} , over the 45 defended breakdown pairs and over the other 635 analysed pairs. L is the distance from the emitting parcel to the plasma boundary. "Satisfies" means $\tau_{\text{esc}} > \tau_{\text{slow}}$, that is, the atom is ionised before it leaves, which is what the closed parcel assumes.	169
A.1	The 43-state ordering	197
B.1	Provenance of every rate in the matrix	200
C.1	Conditioning and non-normality over the grid	203
C.2	Reproducible numerical checks	203

Nomenclature

symbol	meaning	unit
<i>Operator and state vector</i>		
$\mathbf{n}$	population vector, one entry per level	$1/\text{cm}^3$
$\mathbf{L}$	collisional–radiative rate matrix	$1/\text{s}$
$\mathbf{S}$	recombination source, per unit ion density	$1/\text{s}$
Ω_{QSS}	Schur complement of $\mathbf{L}$ on the fast block	$1/\text{s}$
$\mathcal{S}, \mathcal{F}$	slow subspace $\{1s\}$ and the 42 fast levels	
<i>Conditions</i>		
T_e	electron temperature	eV
n_e	electron density	$1/\text{cm}^3$
n_{ion}	ion (proton) density	$1/\text{cm}^3$
n_{1s}	ground-state atom density	$1/\text{cm}^3$
T_n	neutral-atom temperature	eV
u	reservoir variable, n_{1s}/n_{ion}	
b_1	ground-state departure coefficient	
<i>Observable and its two supply channels</i>		
R	the observable ratio; n_3/n_4 where a shell ratio is meant, and the A -weighted H_α/H_β ratio where a line ratio is, the derivation covering both (Sec. 3.5.5)	
R^{CRE}	R at collisional–radiative equilibrium	
R^{QSS}	R from the excited block at a given n_{1s}	
R^{PE}	R on the plateau: new operator, old reservoir	
$\mathbf{n}^{(0)}, \mathbf{n}^{(1)}$	recombination-fed and ground-fed populations	$1/\text{cm}^3$
f_n	ground-fed fraction of shell n	
Δ	logistic displacement between shells 3 and 4	
$\bar{S}$	mean sensitivity of R to the reservoir	
$\Delta \ln u$	reservoir displacement, $\ln(u^+/u^-)$, post- over pre-step	
G	reservoir gain, secant $\Delta \ln u / \Delta \ln T_e$	
x	ground-density scale factor across a step	

Timescales

τ_{relax}	excited-manifold relaxation time, $1/ \lambda_1 $
τ_{slow}	reservoir turnover time, $1/ \lambda_0 $
τ_{drive}	duration of the driving event
τ_{esc}	neutral escape time from the emitting parcel
M	timescale separation, $\tau_{\text{slow}}/\tau_{\text{relax}}$

Error measures

ε^{QSS}	residual of the excited-state closure
ε^{CRE}	distance to collisional–radiative equilibrium
$\varepsilon_{\text{step}}$	distance between the two tabulated answers
$\varepsilon_{\text{plateau}}$	plateau error, the quantity this thesis reports
$\hat{\varepsilon}$	single-slow-mode estimate of the time-averaged error over τ_{drive} , Eq. (5.16); an estimate, not a

Abbreviations

CR	collisional–radiative
CRE	collisional–radiative equilibrium
QSS	quasi-steady state
PE	partial equilibrium, the frozen-reservoir state
ELM	edge-localised mode
ADAS	Atomic Data and Analysis Structure

Chapter 1

Reading the Light from a Divertor

A physicist sits at a console and looks at a spectrum. An optical fibre viewing the bottom of a fusion reactor has collected light along a narrow line of sight, and a spectrometer has spread it out by wavelength. Two lines stand out: a red one at 656 nm and a blue-green one at 486 nm, both emitted by hydrogen atoms. She divides the brightness of the first by the brightness of the second, looks the resulting number up in a table at the density another instrument has supplied, and reads off an electron temperature. That is how one of the least accessible regions of the machine gets measured. It is not the hottest part; the core is far hotter. It is the part where the plasma meets a solid surface, and the one whose conditions decide whether that surface survives.

This thesis is about the table. It was built by solving a model of hydrogen's excited states to steady state, and it is indexed on two numbers: temperature and density. A table can be indexed on two numbers only if every other quantity the answer depends on has been eliminated by an assumption. The assumption in this case is that ionisation and recombination in the plasma have come into balance, not that they are nearly balanced, or changing slowly, but that the plasma has *finished changing*.

The plasma in question never finishes changing. It is the most violently unsteady region of the machine. So the question is whether the table still holds, and if not, by how much it is wrong and where.

That is all this chapter does: start from the spectrum, show where the table comes from, and expose the assumption buried inside it. Sections 1.1, 1.2, 1.3 and 1.4 follow the chain from the divertor, through the hydrogen atom and the spectrometer, to the inversion of a measured ratio. Section 1.5 is the hinge: it identifies the quantity the table has eliminated and separates two approximations that are routinely confused with one another. Section 1.6 asks whether the plasma is ever in the state the table assumes, Section 1.7 says what is already known and by whom, and Section 1.8 states the question this thesis answers and the hypothesis it set out to test.

1.1 Why the divertor is hard

Scope of the measurement problem: what is the thing being measured, and why is measuring it hard enough to be worth a thesis?

Start with the machine. A tokamak is a doughnut-shaped vessel in which a gas of hydrogen isotopes is heated until its electrons separate from its nuclei. The resulting mixture of free electrons and bare nuclei is a plasma, and it is held away from the walls by a magnetic field wound around the doughnut. If the plasma gets hot enough and stays

confined long enough, deuterium and tritium nuclei fuse and release energy. ITER, under construction in southern France, is the machine built to demonstrate this at reactor scale.

Confinement is never perfect, and this is the useful part. Particles and heat leak steadily outward across the magnetic field, and once they cross the last magnetic surface that closes on itself they are no longer trapped: they stream along field lines instead. The layer outside that surface is called the scrape-off layer (SOL), and in a modern tokamak the field lines in it are not allowed to strike the wall at random. They are steered magnetically to a dedicated set of water-cooled plates at the bottom of the vessel. Those plates, together with the volume of plasma and gas in front of them, are the *divertor*.

In chemical-engineering terms, the divertor is the quench-and-separation section bolted onto the outlet of a very hot reactor. It receives the reactor's waste heat and its ash, the helium that fusion produces. It cools the stream and pumps the ash away, and it must do so without being destroyed by the duty.

Being destroyed by the duty is the entire difficulty. The exhaust power does not spread itself over the plate; it arrives concentrated in a narrow strip where the field lines land.

How concentrated, and what the target can take, are two different questions, and the ITER divertor physics basis is careful to separate them in a way worth repeating. The engineering standard is 10 MW/m^2 : it has long been asserted, in that document's words, that this "is a value which must be managed by the actively cooled divertor technology", and it is the qualification standard for laboratory testing of the components [1, Sec. 1]. **That figure is not a prediction of what arrives.** The same section states plainly that "the value of target loading in ITER is not therefore a prediction of SOLPS modelling, but rather the constraint imposed to restrict the acceptable range of core plasma performance". It is a budget, and the modelling is run so as to stay inside it.

What arrives when the divertor is not held inside it is larger. At the lowest divertor neutral pressure in the ITER SOLPS-4.3 database, about 2 Pa, the peak load on a shaped target is about 25 MW/m^2 , rising to about 40 MW/m^2 if the heat flux channel turns out to be as narrow as 1 mm [1, Sec. 4]. Against those, a criterion based on avoiding tungsten recrystallisation over the roughly 2000 hours of high-performance exposure expected of the first divertor gives $16 \pm 2.5 \text{ MW/m}^2$, and critical-heat-flux tests scaled to the final monoblock design put the ceiling on the top surface at about 20 MW/m^2 [1, Sec. 4]. The distance between what arrives and what the surface tolerates is the reason a divertor is a research problem and not a plumbing one, and closing it means removing the power before it reaches the surface at all.

The way out is not a better plate material. It is to stop the plasma before it arrives. If enough neutral hydrogen gas is allowed to build up in front of the target, the plasma ions collide with it: in a charge-exchange collision an ion takes an electron from a neutral atom and becomes a fast neutral itself, carrying its momentum out of the plasma stream, and the energy that remains is radiated away by excited atoms rather than deposited on the plate. Over the last few centimetres the plasma pressure and the heat flux both collapse. The plasma is then said to be *detached* from the target: the region where ionisation happens lifts off the plate and sits in the volume in front of it. Detachment is not a fault to be avoided. It is the intended operating regime of a reactor divertor.

Detachment is also what makes the atomic physics of hydrogen unavoidable. A detached divertor plasma is cold, of order a few electronvolts, where one electronvolt corresponds to about 11 600 K. At a few electronvolts hydrogen is only partly ionised. Neutral atoms and ions coexist in comparable numbers, recombination competes with ionisation on equal terms, and the light emitted by neutral hydrogen becomes the brightest

thing in the visible spectrum. The plasma one most needs to characterise is therefore exactly the plasma whose light carries the most information.

What conditions are those, in numbers? Edge modelling of the ITER baseline burning plasma gives target electron temperatures spanning roughly one to a few tens of electronvolts across the design database, with the large majority of cases below 15 eV [2], and falling below 1 eV across half a metre of each target in the most strongly detached solutions obtained so far [3]. The near-target electron density under partial detachment is expected to reach about $10^{15}/\text{cm}^3$ [4]. The same work makes the point this thesis depends on, that at such densities radiative decay is no longer fast compared with electron-impact processes, so the excited states must be treated with a full collisional–radiative calculation rather than a simple equilibrium formula.

Every calculation in this thesis is performed on a fixed grid spanning $T_e = 1\text{--}10\text{ eV}$ in electron temperature and $n_e = 10^{12}\text{--}10^{15}/\text{cm}^3$ in electron density. The upper density matches the expected near-target value cited above and the temperature range covers the cold, detached end of the modelled database, which is where the atomic physics of neutral hydrogen governs the emission. It does not cover the hotter attached cases, and it is a modelling choice rather than a prediction: it should be read as the range over which the model was exercised, not as a claim about what any particular machine will do. The lower end deserves a sharper statement than that, because it is below anything the sources read for this chapter describe. Across the 53 ITER SOLPS-4.3 cases spanning a wide range of fuel and impurity throughput, the upstream separatrix density is 3×10^{19} to $6 \times 10^{19}/\text{m}^3$ [1, Sec. 3.1.1], that is 3×10^{13} to $6 \times 10^{13}/\text{cm}^3$, and a divertor density is higher than an upstream one rather than lower. **The bottom of this grid, $10^{12}/\text{cm}^3$, is therefore an extension below the modelled range and is not defended as an ITER divertor condition.** It is kept because the coronal limit is approached there and the atomic physics is at its cleanest, which makes it useful for testing the model even where it is not useful for predicting a machine. Results at that edge are labelled throughout.

It would be simpler to put an instrument into the plasma, and ITER will do exactly that. A Langmuir probe is a small electrode exposed to the plasma and biased to a range of voltages; the current it draws as a function of bias gives the local electron temperature and density. ITER carries a designed divertor Langmuir probe system whose stated purpose is to measure plasma parameters at the target plates, and whose thermomechanical analysis concludes that the sensors survive the environment [5]. The claim that probes cannot survive in a reactor divertor is often made and is not supported by the design analysis.

The limitation is not survival but what a probe can see. It is mounted flush in the target surface, so it reports conditions in the thin charged layer at the plate rather than in the volume in front of it, and the volume in front of it is precisely where a detached plasma does its recombining. It perturbs what it measures. Its interpretation relies on sheath theory that is least trustworthy in exactly the cold, strongly recycling conditions of interest. And the coverage is necessarily sparse: a set of sensors embedded in a handful of divertor cassettes samples points on a surface, not a volume.

Light has neither the perturbation nor the surface-only limitation, though it localises nothing along its chord, and it has one decisive advantage: it leaves the plasma by itself. A window, a fibre, a spectrometer and a detector make a measurement that touches nothing, perturbs nothing, and samples the emitting volume rather than the surface. Spectroscopy is therefore complementary to probe measurement rather than a replacement for it. This is why hydrogen line emission is a workhorse divertor diagnostic, and why the chain of

reasoning that converts photons into T_e and n_e carries so much weight.

1.2 Light from a hydrogen atom

Origin of the two spectral lines: where do the two lines in the opening paragraph come from, and what would make their brightnesses differ?

A hydrogen atom is one electron bound to one proton. The electron cannot take any energy it likes: the bound states form a discrete ladder, labelled by an integer $n = 1, 2, 3, \dots$ called the principal quantum number, with energies

$$E_n = -\frac{13.6 \text{ eV}}{n^2}. \quad (1.1)$$

Here E_n is the energy of the level in electronvolts, measured from the free-electron state at $E = 0$, so all bound levels are negative and $n = 1$, the ground state, is the deepest at -13.6 eV . What Eq. (1.1) does physically is fix the whole spectrum of hydrogen with one constant: the levels crowd together as n grows, and 13.6 eV is the energy needed to remove the electron entirely from the ground state, which is called the ionisation energy.

Each level n is not a single state. The electron's orbital angular momentum is also quantised, by an integer $\ell = 0, 1, \dots, n-1$, and states of different ℓ have historically been labelled by letters: $\ell = 0$ is written S, $\ell = 1$ is P, $\ell = 2$ is D, $\ell = 3$ is F. So the $n = 2$ level contains a $2s$ state and a $2p$ state, the $n = 3$ level contains $3s$, $3p$ and $3d$, and so on. In an isolated hydrogen atom all the ℓ states of a given n have almost exactly the same energy, which is why they can often be lumped together, and Chapter 2 explains the one place in this problem where lumping them destroys the physics. The number of independent quantum states in a level is its statistical weight g , equal to $2n^2$ for hydrogen when electron spin is counted; a level with a larger g has more ways of being occupied and is correspondingly favoured in thermal equilibrium.

An atom in an excited level does not stay there. It can fall to a lower level and emit the energy difference as a single photon. This is spontaneous emission, and it is a first-order rate process exactly like a unimolecular decomposition: the probability per unit time that an atom in level q decays to a lower level p is a constant A_{qp} , in $1/\text{s}$, called the Einstein coefficient for spontaneous emission. The number of such decays per unit volume per unit time is therefore $A_{qp}n_q$, where n_q is the population of level q in $1/\text{cm}^3$. Not every pair of levels is connected: quantum mechanics permits a strong (electric-dipole) transition only when ℓ changes by exactly one, a restriction called a selection rule, and Chapter 2 shows that the levels which *cannot* decay this way are the ones that dominate this problem.

Because the energy released is fixed by Eq. (1.1), every transition emits at one wavelength, and the emission appears in a spectrograph as a narrow *spectral line*. Its brightness is proportional to the population of the upper level. That single sentence is why spectroscopy is a diagnostic at all: the light records how the atoms are distributed over their levels, and the distribution is set by the plasma conditions.

The transitions that end on $n = 2$ fall in the visible, which is convenient for windows, fibres and detectors, and they are collectively called the Balmer series. The two brightest are the subject of this thesis:

$$H_\alpha : \quad n = 3 \rightarrow 2, \quad \Delta E = 13.6 \text{ eV} \left(\frac{1}{4} - \frac{1}{9} \right) = 1.89 \text{ eV} \Rightarrow 656 \text{ nm}, \quad (1.2)$$

$$H_\beta : \quad n = 4 \rightarrow 2, \quad \Delta E = 13.6 \text{ eV} \left(\frac{1}{4} - \frac{1}{16} \right) = 2.55 \text{ eV} \Rightarrow 486 \text{ nm}. \quad (1.3)$$

The arithmetic is worth doing once by hand, because it shows that the red line and the blue-green line of the opening paragraph are not empirical facts to be looked up: they follow from Eq. (1.1) and the relation $\lambda = hc/\Delta E$, with $hc = 1240 \text{ eV nm}$.

What these two lines have in common is their lower level. Both end on $n = 2$, so the ratio of their brightnesses does not depend on anything about $n = 2$: it depends only on how the atoms are shared between the $n = 3$ and $n = 4$ levels. The whole of this thesis is a statement about that one ratio, and Section 1.3 explains why it, and not the brightness of either line by itself, is what a divertor spectroscopist can actually use.

1.3 What a spectrometer records

Content of a brightness measurement: given that a spectrometer measures brightness, what can and cannot be extracted from that measurement?

Picture the geometry first. The instrument looks along a straight line through the plasma and collects the light emitted anywhere along it, from a small solid angle. It cannot tell where along the line a photon came from. Whatever it reports is a sum over the whole chord.

Formally, if each decay from level q to level p releases a photon of energy $h\nu_{qp}$ in joules and the decays occur at $A_{qp}n_q$ per unit volume per unit time, the detector records the brightness

$$I_{qp} = \frac{h\nu_{qp}}{4\pi} \int_{\text{chord}} A_{qp} n_q(s) ds, \quad (1.4)$$

in $\text{W}/(\text{m}^2 \text{ sr})$, where s is distance along the line of sight in cm and the factor 4π accounts for the fact that an atom radiates into all directions equally, so only that fraction of the emission reaches the detector. Equation (1.4) says the brightness is the emission rate per unit volume, integrated over the length of plasma the instrument happens to be looking through.

It also assumes that a photon, once emitted, escapes without being reabsorbed by another atom on its way out. A plasma for which that is true is called *optically thin* at that wavelength. For the Balmer lines in a divertor this is usually a fair assumption, because the lower level $n = 2$ holds very few atoms to absorb them; for transitions that end on the ground state it is often badly false, since the ground state holds almost all of them. Chapter 6 takes that up, quantifies it for this model, and it is the single most important restriction on what follows.

Now the discouraging part. Equation (1.4) has one measured number on the left and, on the right, the population we want multiplied by a path length nobody knows, integrated over an emitting region whose extent nobody knows either. Worse, n_q is itself proportional to the amount of neutral hydrogen in the viewing volume, and that is set by how much gas recycles off the target and how it moves, neither of which the spectrometer sees. An absolutely calibrated brightness of a single line therefore determines a product of unknowns and none of them individually.

Take the ratio of two lines and most of this cancels. Define the observable used throughout this thesis:

$$R \equiv \frac{I_{H\alpha}}{I_{H\beta}} = \frac{\nu_{32} A_{32}}{\nu_{42} A_{42}} \cdot \frac{\int n_3 ds}{\int n_4 ds}, \quad (1.5)$$

where n_3 and n_4 are the populations of the $n = 3$ and $n = 4$ levels in $1/\text{cm}^3$ and R is dimensionless. The prefactor is a ratio of atomic constants, known to high accuracy. What

Eq. (1.5) does is remove the geometry: provided the two lines are emitted by the same volume, the absolute neutral density and the absolute calibration of the instrument divide out, and so does the path length as a scale. What does not divide out is any variation of n_3/n_4 along the chord: the measured ratio is the n_4 -weighted average of the local ratio, and it equals the local ratio only where the emitting region is uniform, which is the zero-dimensional case this thesis treats. That proviso is the substantive assumption here, and a good one for two lines of the same series from the same species. Only the *relative* sensitivity of the detector at two nearby visible wavelengths is needed, which is a far easier measurement to make and to maintain.

What survives is a pure ratio of two excited-state populations in a uniform emitting volume, which is why R , rather than any single brightness, is the quantity a divertor spectroscopist works with. A line ratio contains no geometry and no absolute density, so every piece of physics in it is compressed into one question: *what sets the ratio of the $n = 3$ and $n = 4$ populations?*

1.4 From a ratio to a temperature

The model between measurement and inference: what stands between the number the spectrometer produces and the (T_e, n_e) the physicist writes down?

An atom in a plasma is pulled in several directions at once. Free electrons collide with it and knock its bound electron up the ladder of Eq. (1.1), or right off it, which is *ionisation*. Ions capture free electrons back into bound levels, which is *recombination*; this happens either with the excess energy carried off by a photon (radiative recombination) or with it carried off by a second electron that happens to be present (three-body recombination, which therefore needs two electrons and becomes important at high density). Meanwhile every excited atom is decaying radiatively on its own clock, and every decay from a high level cascades down through the intermediate ones. No single one of these processes dominates in a divertor. All of them run at once, and their rate coefficients span more than ten orders of magnitude, so which of them matters depends on the level and on the conditions rather than being settled once.

The object that keeps track of all of it is a *collisional–radiative* (CR) model. It is nothing more exotic than a population balance written for every level of the atom: accumulation equals everything arriving minus everything leaving, with one term per atomic process and a measured or computed rate coefficient attached to each. The framework was set out for hydrogenic ions by Bates, Kingston and McWhirter in 1962 [6], whose central point was that neither of the two classical limits describes a laboratory plasma: not coronal equilibrium, where the density is so low that every collisional excitation is followed by a radiative decay before anything else happens, and not local thermodynamic equilibrium (LTE), where collisions are so frequent that the level populations take their thermal equilibrium values. Real plasmas sit between the two, and the populations there have to be computed.

In chemical-engineering terms this object is a reaction network in a well-mixed vessel: one species per energy level, one reaction per atomic process, a feed stream, and a set of coupled first-order rate equations with a rate-constant matrix. Chapter 3 builds it explicitly and pushes that correspondence considerably further than a metaphor would go.

Solve the CR model at a fixed T_e and n_e and it returns the population of every level, hence n_3 , hence n_4 , hence R through Eq. (1.5). Do that on a grid of (T_e, n_e) and the




Figure 1.1: **Where the hidden assumption enters divertor spectroscopy.** The inference runs left to right. A spectrometer records the Balmer emission; the ratio of two line intensities, here $I_{H\alpha}/I_{H\beta}$ at 656.1 and 486.0 nm, reduces the spectrum to a single number per line of sight; with the electron density supplied by an independent measurement, that number is looked up in a table of the ratio computed as a function of (T_e, n_e) , and the temperature that reproduces it is reported. One ratio is one constraint, so the table's second axis is fixed from outside; Chapter 5 measures this inversion at known density. The step that is rarely stated is the third. The table is built from the CR-equilibrium populations, so it assumes the ionisation balance has already settled, $n_g/n_{ion} = u^{CRE}(T_e, n_e)$. A plasma still in transit after a change in temperature does not satisfy that, and the inversion then returns a temperature that no plasma had. This thesis is about the size of that error and about which of the approximations in the chain is actually responsible for it. This is a schematic. No measured or computed quantity is plotted; the two wavelengths are the only numbers on it, and they follow from the ionisation energies in the model's own state index. The drawn line heights are arbitrary and carry no information about the real intensity ratio.

result is a table, or equivalently a family of contours in the (T_e, n_e) plane along which R is constant. Inverting a measurement means locating the measured value of R on that map. One ratio is one constraint on two unknowns, so practice supplies a second: another line ratio, the Stark broadening of a line by the plasma's own microscopic electric fields, continuum emission, or an independent estimate of one of the two parameters. The intersection of two contours fixes the point. Figure 1.1 draws the chain end to end, from the recorded spectrum to the reported (T_e, n_e) , and marks the step at which the table's own assumption enters.

This is not a hypothetical procedure. Balmer-spectrum analysis against a CR model has been used to infer divertor conditions since at least the Alcator C-Mod measurements of Pigarov, Terry and Lipschultz, who matched calculated and measured Balmer spectra, with populations of the excited states supplied by a collisional-radiative model, to determine the plasma parameters in the recombining region of a divertor at electron temperatures near one electronvolt [7]. The modern versions of the technique are more elaborate but the chain is the same one.

In practice the table is rarely built by hand. The Atomic Data and Analysis Structure

(ADAS) distributes precomputed coefficients for exactly this purpose, tabulated over (T_e, n_e) [8, 9]. What matters for the argument of this thesis is a detail of how ADAS stores them: for each line it supplies *two* emissivity coefficients, one for the emission driven by excitation out of the neutral ground state and one for the emission driven by recombination from the ion. The database keeps the two apart deliberately. The next section is about what happens when they are combined.

1.5 The assumption inside the table

Number of independent arguments of the ratio: a table has two axes; how many quantities does the answer actually depend on?

Take the physical picture before any algebra. An atom sitting in the $n = 3$ level got there by one of two routes. Either it started as a neutral atom in the ground state and an electron collision kicked it up, fed *from below*, or it started as an ion and captured a free electron into a high level, from which it cascaded down, fed *from above*. Those are two different supply lines drawing on two different reservoirs: the population of ground-state neutral atoms, n_{1s} , and the population of ions, n_{ion} . Every excited level is filled by some mixture of the two.

Because the population balances are linear in the populations, the mixture is exact rather than approximate. For every level p ,

$$n_p = \underbrace{a_p n_{1s}}_{\text{fed from below}} + \underbrace{c_p n_{\text{ion}}}_{\text{fed from above}}, \quad (1.6)$$

where n_p , n_{1s} and n_{ion} are in $1/\text{cm}^3$, the dimensionless coefficient a_p is the number of atoms in level p supported per ground-state atom, and c_p is the number supported per ion. What Eq. (1.6) does is separate the atomic physics from the plasma composition: a_p and c_p depend on T_e and n_e and on nothing else, while n_{1s} and n_{ion} describe how much of each reservoir is present. This split is as old as the CR model itself. It is the population-coefficient decomposition of Bates, Kingston and McWhirter, written by Fujimoto and McWhirter in the normalised form Chapter 3 adopts [6, 10, 11]. The two terms have names that will recur: Fujimoto calls the first the *ionising plasma component* and the second the *recombining plasma component* [11, Sec. 4.1], and a divertor shifts the balance between them as its conditions change. Those names belong to the *components*. Whether the plasma itself is ionising or recombining is a statement about the sign of its net stage-to-stage flux, which is a different quantity, and this thesis keeps the two apart by calling a level *ground-fed* or *recombination-fed* according to which component supplies it. The crossover is not identified with any named operating regime here; Chapter 6 gives the densities and explains why that identification does not hold.

Now put Eq. (1.6) into the observable. For the two Balmer levels,

$$R \propto \frac{a_3 n_{1s} + c_3 n_{\text{ion}}}{a_4 n_{1s} + c_4 n_{\text{ion}}} = \frac{a_3 (n_{1s}/n_{\text{ion}}) + c_3}{a_4 (n_{1s}/n_{\text{ion}}) + c_4}. \quad (1.7)$$

The second form is the important one: dividing through by n_{ion} shows that the observable depends on the atomic coefficients, which are functions of (T_e, n_e) , *and* on the ratio n_{1s}/n_{ion} , which is not. That ratio is a third, independent quantity. It says how much neutral hydrogen is present per ion, and it is fixed by the plasma's recent history and by transport, not by the local temperature and density.

A table with two axes must therefore either dispose of the third quantity or carry it, and there are three ways to proceed. **Route one** is to work where a single supply dominates *both* levels, so that the third argument cancels between numerator and denominator of Eq. (1.7) and the answer is right whatever the neutral fraction does. **Route two** is to adopt a closure, a rule that fixes n_{1s}/n_{ion} in terms of T_e and n_e . One such rule is *equilibrium ionisation balance*: solve the CR system all the way to steady state and take the neutral-to-ion ratio it predicts for the local conditions. That state is collisional-radiative equilibrium (CRE), and it is a much weaker condition than full thermodynamic equilibrium: it means only that the CR rate equations have stopped changing, not that the populations are thermal. **Route three** declines to dispose of the third quantity at all: it carries n_{1s}/n_{ion} as an unknown, brackets it over a plausible range, and reports the range of temperatures the ratio then admits, which is what the next subsection finds careful practice doing. **This thesis measures what route two costs where it is taken.**

1.5.1 A worked example, because the distinction is easy to assert and hard to demonstrate

It is fair to ask which route the literature actually takes, and the question deserves a specific paper rather than an impression. Lomanowski and co-workers infer divertor plasma properties in JET-ILW from hydrogen Balmer and Paschen spectroscopy, and their temperature inference is a line ratio of exactly the form of Eq. (1.7): the $D9 \rightarrow 2$ to $D5 \rightarrow 2$ intensity ratio, with the level populations built from ADAS photon emissivity coefficients [12, Sec. 3.2].

They take route one, and they do not eliminate the third quantity. The figure that defines their method plots the ratio against T_e and the neutral fraction, and its companion panel is drawn at a stated $n_0/n_e = 0.50$ rather than at a value derived from ionisation balance. The consequence is reported plainly: “it is possible for the measured line ratio value to correspond to different temperatures, if the extent of the excitation contribution is not known”. Bracketing the neutral fraction over 0.1 to 0.8 turns one measured ratio into $T_e = 1.8$ to 2.5 eV, quoted as an *upper limit* rather than a value [12, Sec. 3.2].

What lets them proceed is the cancellation of route one. They adopt the assumption that the ratio is “mainly driven by recombination processes in detached plasmas”, following earlier measurements at Alcator C-Mod and JET [13, 14], and that assumption is available to them because, in their words, “the contribution due to excitation is negligible for $n_{\text{upper}} > 5$ for $T_e < 2$ eV for any reasonable values of the neutral fraction” [12, Sec. 3.2].

Three things follow, and they set the terms of this thesis.

First, the criticism must be aimed correctly. No procedure read for this thesis eliminates n_{1s}/n_{ion} by imposing ionisation balance on a Balmer ratio and reading off a temperature. The careful practice is route one. **This thesis is therefore not a demonstration that standard divertor Balmer diagnostics fail. It is a measurement of what the equilibrium-reservoir closure costs where route one is unavailable,** and the thesis is to be held to that narrower claim throughout.

Second, route one is unavailable for exactly the lines studied here. The cancellation works because $a_p n_{1s} \ll c_p n_{\text{ion}}$ for both levels, which is the statement that the ground-fed fraction of each shell is near zero. That holds for $n \geq 5$ in a detached plasma and fails for $n = 3$ and $n = 4$, where neither supply dominates. The two lines this thesis studies are the two where the third quantity cannot be cancelled away, which is why, wherever it is

neither carried as an unknown nor supplied from outside, the closure is what remains, and why its cost is worth measuring.

Third, the non-uniqueness is not this thesis's invention. That one measured ratio can correspond to more than one temperature is stated in [12] and is found independently here from the model's own operator, in Section 5.5.1. Two routes to the same defect is better evidence than either alone.

That closure is an assumption about the plasma, not a piece of atomic physics, and the databases are not at fault. ADAS keeps the excitation-driven and recombination-driven coefficients separate precisely so that a user can weight them by the n_{1s} and n_{ion} appropriate to the plasma in hand [8]. The most careful divertor analyses do exactly that, and treat the neutral content as an unknown rather than a given: Verhaegh and co-workers, analysing Balmer emission from the divertor of the TCV tokamak, take the neutral fraction n_{1s}/n_e to be a quantity that is estimated rather than measured, assign it a broad prior spanning more than an order of magnitude, and propagate that uncertainty through the whole inference by Monte Carlo so that every inferred quantity emerges as a probability distribution rather than a number [15]. In that community, the dependence written in Eq. (1.7) is known and is handled, and this thesis does not discover it.

The same paper is worth one more sentence, because it draws a boundary this thesis has to respect. Verhaegh and co-workers restrict their analysis to the higher Balmer lines, $n \geq 5$, on the grounds that the lower ones, H_α above all, are contaminated by emission from molecular processes that their atomic treatment does not model [15]. This thesis studies H_α and H_β , which is to say precisely the two lines that careful practice avoids in detached conditions, and it does so with a model containing no molecules. Section 1.7 returns to this and Chapter 6 quantifies what it costs; it is stated here so that no reader reaches those chapters feeling it was kept back. The assumption enters at the moment a measured ratio is inverted against a two-parameter lookup, because a two-parameter lookup returns exactly one value of R per (T_e, n_e) pair and therefore requires the two supplies to have been combined in some fixed proportion.

There are now two approximations in play, and keeping them apart is the single most important thing this thesis does.

The quasi-steady-state (QSS) approximation. That the *excited* levels have already reached steady state with whatever reservoirs are present at that instant. This is the standard shortcut of CR theory, and it is what makes a_p and c_p in Eq. (1.6) functions of (T_e, n_e) alone. **The approximation being tested in this thesis is the Bodenstein steady-state approximation, applied to atomic level populations instead of radical concentrations.** It is the same approximation used in reaction kinetics for reactive intermediates, with the same justification, that the intermediates turn over far faster than anything else in the system. In principle it has the same failure mode too.

The ionisation-balance assumption. That the *reservoirs themselves*, n_{1s} and n_{ion} , stand in the proportion that CRE would give for the local T_e and n_e . This is what removes the third argument from Eq. (1.7) and allows the table to have two axes.

These are statements about different populations relaxing at different speeds, and only the first is what the literature normally means by “the validity of QSS”. A two-parameter table needs both to hold. Chapter 3 makes the distinction exact, defines a separate error measure for each, ε^{QSS} and ε^{CRE} , and shows why conflating them makes it possible to

measure the second and report it as the first. Whether either holds in a real divertor is the subject of the next section.

1.6 A plasma that never finishes changing

Time available between disturbances: how long does the plasma get between disturbances, and is that long enough?

The ionisation-balance assumption would be harmless if divertor plasmas sat still. They do not, and they are not meant to.

The largest regular disturbance is the edge-localised mode (ELM). The confined plasma maintains a steep pressure gradient at its edge; periodically that gradient becomes unstable, the edge relaxes, and a burst of particles and energy is expelled into the SOL and flushed down the field lines onto the divertor.

The numbers that matter here are the ones describing a single event, and they are short. In present-day machines an ELM heat pulse at the target lasts between about 0.75 ms and 3 ms, and the scaling to ITER gives a deposition rise time of 250 μ s with a decay roughly twice that, the pulse being well approximated by a triangular waveform; each event carries away between two and ten per cent of the energy stored in the plasma edge [16]. So the conditions in the divertor are reset over a few hundred microseconds, repeatedly.

There is a second and independent extrapolation of that duration, and it is worth having because it rests on a physical mechanism rather than on a fit. Measuring the ELM power pulse at the target by the rise time of the surface temperature, $\tau_{\text{IR}}^{\text{ELM}}$, across ASDEX Upgrade, JET and JT-60U, Loarte *et al.* found it correlates with the time an ion takes to travel from the pedestal to the target along the field, $\tau_{\parallel}^{\text{Front}} = 2\pi R q_{95} / c_{s,\text{ped}}$, and *not* with the duration of the ELM magnetohydrodynamic event itself [17, Sec. 4]. The measured scaling is $\tau_{\text{IR}}^{\text{ELM}} [\mu\text{s}] = 0.29 [\tau_{\parallel}^{\text{Front}} (\mu\text{s})]^{1.38}$, and at the expected ITER pedestal, $n_{\text{ped}} = 8 \times 10^{19} / \text{m}^3$ and $T_{\text{ped}} = 3.5 \text{ keV}$, it gives $\tau_{\text{IR}}^{\text{ELM}} = 506 \mu\text{s}$ [17, Sec. 5]. **The divertor is therefore disturbed for several hundred microseconds, and the figure has a mechanism behind it: it is set by parallel ion transport, so it lengthens as the pedestal cools.** Chapter 5 uses 100 μ s for its maps and reports what changes at 506 μ s, which is not nothing.

How often they are reset is a more delicate question, and the usual figure is weaker than it looks. The expected natural ELM frequency in ITER is of order 1 Hz at the full plasma current of 15 MA, rising to a few hertz at reduced current and scaling roughly as the inverse square of the current [18]; present-day machines, being smaller, run at tens to hundreds of hertz. The frequently quoted range of one to two hertz for ITER traces, as far as could be established for this chapter, only to a conference paper that was not peer reviewed, and it is not repeated here for that reason. Loarte *et al.* [17], which is the refereed extrapolation of ELM properties to ITER used above, extrapolates the ELM energy loss and the pulse duration and does not state a frequency at all; the word hertz does not appear in it. **Nothing in this thesis depends on the frequency.** An ELM is not a gentle modulation that the plasma tracks quasi-statically; it is idealised in this thesis as a step, because its rise of a few hundred microseconds is short against the reservoir clock over part of the operating range and long against it over the rest; Chapter 5 states at which points each holds and tests the step against a finite ramp. The relevant question is therefore not how often the steps come but how long the recovery takes relative to a

single step, and that is a question about the plasma's own internal clocks.

An ELM is not the only transient. The detachment front itself moves: as the neutral pressure in the divertor is raised, the ionisation region lifts off the plate and travels back up the field lines, and where it settles is controlled by feedback on a gas valve. Operating near detachment onset means operating near a threshold, where the plasma responds to a small change in the control parameter by moving a long way. Add the current ramps at the start and end of a discharge, periodic core relaxations, and deliberate density modulation, and the picture is of a system that is repeatedly stepped rather than one that is ever allowed to settle. Whether the intervals between steps are milliseconds or seconds does not by itself decide the question; what decides it is whether the plasma reaches its new balance before the next step arrives.

Now set that against the clock the table assumes, and note carefully which clock that is. It is not the lifetime of an excited state. Those are nanoseconds; Eq. (1.4)'s Einstein coefficients are of order 10^8 1/s, and nothing a tokamak does on a millisecond disturbs them. The relevant clock is the time the *ground-state* population takes to come back into balance with the ions after conditions change, and that is a much slower thing.

One can see why it is slower without any model at all. To ionise a hydrogen atom out of the ground state an electron must bring 13.6 eV. In a plasma at $T_e = 3$ eV, the fraction of electrons in the Maxwellian tail with at least that energy is of order $\exp(-13.6/3) \approx 1.1 \times 10^{-2}$: about one electron in a hundred. To excite the same atom to $n = 2$ costs 10.2 eV, still several times $k_B T_e$, and to move an atom that is *already* in $n = 3$ takes a fraction of an electronvolt, which essentially every electron can supply. The ground state is therefore separated from everything else by a barrier, the first excitation and the ionisation energy at 10.2 and 13.6 eV, both several times $k_B T_e$, which only the tail of the distribution can cross, while the excited levels exchange population freely among themselves. That is the physical origin of the two very different clocks, and Chapter 3 turns the argument into a calculation.

In chemical-engineering terms this is a Damköhler-number question. If the characteristic time of the disturbance is short compared with the slowest chemical relaxation in the system, then the composition during the disturbance is not the equilibrium composition, and a chart that assumes equilibrium describes a state the system has not reached. Here the disturbance rises in a few hundred microseconds; the fast levels relax in nanoseconds; and the slow reservoir relaxes on a timescale Chapter 3 computes and Chapter 5 compares against the event.

So the complication, stated in full: a table built on two assumptions is being applied to a plasma that plainly violates the second one, and possibly the first. Which of the two actually fails, and what it costs, cannot be decided from anything said so far.

1.7 What is already known

Prior work and the scope of the contribution: how much of this is new, and how much has been known for decades?

This question has been asked before, in several forms, and the contribution of this thesis is narrower than the question it asks; the narrowing is stated here.

The time-dependent hydrogen problem has been solved before, and the overlap is close enough to state plainly. Sawada and Fujimoto studied exactly this system:

atomic hydrogen, a CR model, and an abrupt change of plasma parameters. Their paper is titled, in part, for the validity range of the quasi-steady-state solution of the coupled rate equations [19]. What they established is a statement about *which level controls the approach to steady state*. Solving the coupled rate equations after a step, they found that in an ionising plasma the levels below Griem’s boundary, the level at which radiative and collisional depopulation rates are equal, relax at their own natural lifetimes while every level above it relaxes at the rate of the boundary level itself; in a recombining plasma the roles invert, the levels above the boundary relaxing on their own timescales and those below following the boundary. Either way the overall response of the excited manifold is set by the relaxation time of Griem’s boundary level.

That is a result about the excited levels, and it is the reason this thesis does not claim to have discovered that the excited manifold responds on a finite timescale. The difference is in what is being asked, and it is visible in how the calculation is set up: under ionising conditions they keep $n(1)$ constant and set the ion density to zero [19, p. 5566]. The reservoir is an input to their problem, not a variable of it. Sawada and Fujimoto asked when the excited-level solution is valid, and answered it with a timescale internal to that manifold. The question here is what a *diagnostic inversion* costs when the ground-state reservoir underneath that solution has not settled. That is a different quantity, controlled by a different and far slower clock, and one that can be wrong even when the excited-level solution is perfectly accurate. Chapter 5 returns to Griem’s boundary level, which the present model locates independently, and uses it as a check.

The separation of timescales that both papers rest on is older than either, and Sawada and Fujimoto quote it in its original form: a quasi-equilibrium population of the excited system is established almost instantaneously “without the number densities of free electrons and bare nuclei being appreciably altered”, because “the relaxation times for the excited levels are very much shorter than the relaxation time for the ground levels or for the free electrons” [19, p. 5566]. Chapter 3 measures that ratio as M and Chapter 5 shows that it does not predict the diagnostic error, which is the point at which this thesis departs from the tradition it is quoting.

The general validity criterion was stated twenty-five years ago, and it is a criterion to be checked, not a verdict. Greenland analysed the reduction of CR systems in general form, in two papers of the same year, and reached two conclusions that bear directly on this work. The first is that the conditions for a valid reduced CR model are conditions on the *eigenvectors* of the full matrix of rates, and that, in his words and in contrast to the conventional view, the eigenvalues have secondary importance [20]. The second is that the validity criteria he derives are *not related to the equilibrium timescales* of the system [21]. The two statements come from different papers and are cited separately here because they are separately attributable. Those statements are general, they are prior, and this thesis does not improve on them as statements. What it does is quantify them for one specific diagnostic: Chapter 5 shows that the standard timescale-separation criterion does not predict the diagnostic error, measures the error it fails to predict, and maps where in the operating range that error is large. Turning a general criterion into a number a diagnostician can apply is the contribution; discovering the criterion is not.

That the ratio depends on the neutral reservoir is current best practice, not a discovery. As Section 1.5 noted, Verhaegh and co-workers already treat the neutral fraction as a free assumed parameter and propagate its uncertainty by Monte Carlo [15].

The dependence in Eq. (1.7) is therefore known. What this thesis adds is a *closed form* for the sensitivity, an exact expression for how strongly the ratio responds to a change in the reservoir, derived in Chapter 3, together with a bound showing that the response can never exceed a certain size; and a *map* of where in the (T_e, n_e) plane that sensitivity is large enough to matter. A closed form does something a Monte Carlo cannot: it says which conditions are safe without requiring the calculation to be repeated.

The classical validity criteria answer a different question. Griem asked at what principal quantum number collisions become fast enough to hold a level in thermal equilibrium with the free electrons, and defined a boundary level above which the upper part of the manifold is thermalised [22, 23]. That is a statement about the internal structure of the excited levels at one instant. It says nothing about whether the ground state has settled, which is the question here. Chapter 5 nonetheless finds Griem’s boundary level turning up where the present model says it should, which is a useful consistency check rather than a new result.

The nonequilibrium Balmer thermometer is fifty years old. Park proposed the H_α/H_β intensity ratio as an indicator of electron temperature in a plasma out of Saha equilibrium, and tabulated the conversion between the excitation temperature such a ratio appears to show and the true electron temperature [24]. The apparatus is the one used here. He assumed a quasi-steady-state distribution within the excited levels, wrote the normalised population of level p in the same two-coefficient form, $\rho(p) = r_0(p) + r_1(p)\rho(1)$, and carried the normalised ground-state population $\rho(1)$, which is b_1 in the notation of Chapter 3, as the third argument. That a population away from equilibrium makes this ratio imply a temperature other than the equilibrium one is therefore established, and nothing in this thesis claims it.

What Park evaluated is the pair of limits he called *terminal* nonequilibrium: decaying, where $r_0(p)$ dominates, and rising, where $r_1(p)\rho(1)$ dominates. He named them terminal because they are the conditions in which any further departure from equilibrium does not affect the emission ratio of the hydrogen lines. In the language of Section 3.5.4 that definition is the vanishing of the very sensitivity this thesis measures: both are limits in which the two shells are fed in the same proportion, the ratio goes blind to the reservoir, and a table on two axes is well defined again. Between them the ratio is not blind, and it is there that a divertor sits. What Park leaves open is the size of the error a table incurs in that interval, and whether a plasma passing through it is described by the equilibrium closure at all.

Real divertor Balmer light is not purely atomic. One line of work must be flagged now because it bounds what any conclusion of this thesis can claim. In detached conditions a substantial part of the observed Balmer emission does not come from the atomic processes modelled here at all. Hydrogen in a cold divertor is largely molecular, and molecules open additional routes to excited atomic levels, among them molecular-assisted recombination (MAR), in which a molecular ion recombines and dissociates into an excited atom. Analyses of discharges on the TCV and MAST-U tokamaks attribute a large share of the low- n Balmer light to these channels near detachment onset [25, 26]. The model built in this thesis contains no molecules at all. Chapter 6 states the consequence plainly, and it is the principal reason the results here are presented as statements about a class of CR model rather than as predictions for a particular machine.

The system has been modelled, the general validity criterion has been stated, and the dependence on the neutral reservoir is known and handled in practice. What is missing is a quantitative, closed-form answer to the specific question of how much the temperature inferred from the Balmer ratio at known density is in error when the reservoir has not settled, and a map of where that error lives. That is the gap this thesis fills.

1.8 The research question and the hypothesis

The computation and its expected outcome: what is the thesis actually going to compute, and what does the simplest physical argument predict it will find?

The question, stated plainly:

When a hydrogen plasma at divertor conditions is disturbed on a timescale short compared with the relaxation of its ionisation balance, how large is the error in the temperature inferred from the H_α/H_β ratio at known density, what physical quantity sets that error, and where in the operating range is it largest?

Answering it needs a model that does not assume the answer. This work therefore builds a time-dependent CR model of atomic hydrogen, every level tracked as a dynamical variable with orbital angular momentum resolved wherever that resolution changes the physics, subjects it to an idealised step in electron temperature, the limiting form of the change an ELM delivers, and follows the line ratio in time. Chapter 2 assembles the atomic data it runs on and Chapter 3 assembles the equations.

Expected outcome. Section 1.5 identified two assumptions, and the hypothesis at the outset was that the first one fails. The QSS approximation is the step that eliminates several dozen dynamical variables by fiat; it is the step that every review of CR modelling flags; and it is the step whose chemical-engineering analogue, the Bodenstein approximation, every undergraduate is warned can fail when the intermediates are not fast enough. In a plasma whose conditions are reset over a few hundred microseconds by every ELM, the natural expectation is that some of the excited levels will simply fail to keep up, the population distribution will lag behind the conditions, and R will be wrong for that reason. That was the expectation this work set out to test, and it is why the model tracks every level in time rather than assuming any of them are stationary.

What is at stake if that expectation is right. If the excited levels lag, then the coefficients a_p and c_p of Eq. (1.6) are not functions of (T_e, n_e) during a transient, the whole population-coefficient framework loses its footing during exactly the events one most wants to diagnose, and the repair is a genuinely time-dependent inversion.

And if it is wrong? Then the second assumption is still standing there, untested, and it is the one the table needs just as much.

Which of the two fails, by how much, and what physical quantity controls it, is the subject of Chapters 3 and 5. The answer is not the one stated above as the expectation; that much can be said here, because the shape of the argument depends on it. The reason it is not is the reason this thesis exists.

Two things should be said in the same breath, before any of it. First, the claim this work defends is conditional: *if* the ground-state neutral population is frozen for the duration of the disturbance, which requires that neutral transport and recycling do not renew it faster than the atomic physics does, then the inversion carries the error computed here. Second, the model is zero-dimensional, optically thin, and contains no molecules. Chapter 6 examines every one of those boundaries and quantifies what it costs, rather than listing them as caveats. A result stated without its scope is the characteristic failure of this subject, and it is not repeated here.

1.9 Outline of the thesis

Route through the remaining chapters: where does the argument go from here?

Chapter 2, *The Atoms*, establishes the inputs: the processes and their rates, the levels tracked and why orbital angular momentum is resolved for the lower ones and bundled above them, how a measured cross section becomes a rate coefficient, the case for including a process the standard treatments often omit, and the provenance and accuracy of every dataset used, including where the atomic data are least secure.

Chapter 3, *Two Clocks*, assembles the CR rate equation from a population balance on each level, diagonalises the rate matrix to expose the timescale structure the Boltzmann-factor argument of Section 1.6 anticipated, presents the QSS approximation as the Bodenstein approximation applied to the fast band, derives the exact dependence of R on the ground-state reservoir, and defines the two reference states and the two error measures the rest of the thesis keeps apart.

Chapter 4, *Does the Model Work?*, applies the tests a wrong model would fail, detailed balance, the coronal and LTE limits, the timescale hierarchy and an external benchmark against independently published population coefficients, and records the errors this project’s own audit found, what each cost, and how each was caught.

Chapter 5, *What the Model Says*, reports the relaxation in time, the size of each of the two errors, the mechanism that sets the larger one, a map of where in the (T_e, n_e) plane it lives, whether it survives averaging over a transient event, and whether the standard timescale-separation criterion predicts any of it.

Chapter 6, *What This Model Cannot Say*, bounds the claim: optical thickness, molecular channels, the truncation of the level manifold, the fixed ion reservoir and the remaining assumptions, each quantified where quantification is possible and stated as open where it is not.

Chapter 7 states the criterion in a form another modeller can apply to their own rate matrix, says where in the divertor operating range a temperature inversion at known density should be distrusted, and sets out what a time-dependent inversion would require.

The argument begins with the atoms, and with how well their rates are known, in Chapter 2.

Chapter 2

The Atoms

Chapter 1 left us with a question and no instrument to answer it. A spectroscopist measures the ratio of two hydrogen lines in a divertor and reads a temperature and a density off a table computed on the assumption that the plasma has stopped changing; a divertor never stops changing. Everyone asked whether the table still holds worries first about the excited states, the fast, delicate part of the atom, which surely cannot keep up with a plasma torn about on microsecond timescales. Testing that suspicion needs a model that follows every level of the atom in time, assuming in advance that none of them has settled.

This chapter builds the raw material for that model. There is no time derivative anywhere in it. It is about the *atoms*: what can happen to a hydrogen atom in a plasma, how fast, where the numbers come from, and how far they can be trusted. Chapter 3 puts them in motion.

One sentence belongs before any physics, because it is the fastest route into everything that follows. **The approximation this thesis tests is the Bodenstein steady-state approximation, applied to atomic energy levels instead of radical concentrations.** Anyone who has written $d[R^\bullet]/dt \approx 0$ for a reactive intermediate has used it. The 43 energy levels below are the chemical species, the atomic processes are the elementary reactions, and this chapter measures the rate constants before that approximation is made.

2.1 What happens to a hydrogen atom in a plasma

2.1.1 The atom, from scratch

A hydrogen atom is one electron bound to one proton. Quantum mechanics allows the electron only certain energies, labelled by the *principal quantum number* $n = 1, 2, 3, \dots$, with binding energy

$$I_n = \frac{I_H}{n^2}, \quad I_H = 13.6058 \text{ eV}. \quad (2.1)$$

I_n (eV) is the energy needed to pull the electron free and leave a bare proton. Both sides are an energy, so the dimensional check is trivial; the *shape* of Eq. (2.1) is what matters. The levels crowd together as n grows: $n = 1$ lies 10.20 eV below $n = 2$, but $n = 14$ and $n = 15$ are separated by 0.0089 eV. High levels are easy to move between and easy to strip off; low levels are not. That asymmetry generates the two timescales of Chapter 3.

Within a shell the electron also carries *orbital angular momentum*, quantised by $\ell = 0, 1, \dots, n-1$. In hydrogen, and only in hydrogen, all ℓ states of a given n have the same

energy: they are *degenerate*. Spectroscopy labels them $\ell = 0, 1, 2, 3, \dots \equiv \text{S, P, D, F, } \dots$, so “ $2p$ ” means $n = 2, \ell = 1$. The number of quantum states sharing a label is the *statistical weight*

$$g_{n\ell} = 2(2\ell + 1), \quad g_n = \sum_{\ell=0}^{n-1} g_{n\ell} = 2n^2, \quad (2.2)$$

the factor 2 counting electron spin. g is a dimensionless counting number and appears only where a thermodynamic argument needs to know how many ways a state can be occupied: $g = 6$ for $2p$, $g = 2$ for $2s$.

An atom dropping to a lower level can emit the energy difference as one photon. The levels are discrete, so the photon energy is too, and the light appears at one sharp wavelength: a *spectral line*. Lines ending on $n = 2$ form the Balmer series and lie in the visible, which is why a fibre-coupled spectrometer sees them; this thesis uses H_α ($n = 3 \rightarrow 2$, 656 nm) and H_β ($n = 4 \rightarrow 2$, 486 nm). A line’s brightness is proportional to the population of the level it comes from, so measuring light measures populations, which is why the populations are worth computing.

2.1.2 Six processes, and how fast they run

Six things can happen to an atom in a plasma of electron density n_e ($1/\text{cm}^3$) and temperature T_e (eV), and the model is built from these and nothing else. *Electron-impact excitation* pushes the bound electron up, and needs a free electron carrying at least the energy gap. *Electron-impact de-excitation* is the reverse and has no threshold, so it is generally the faster of the pair. *Ionisation* removes the electron altogether, and is the only process that takes an atom out of the set of levels we track. *Recombination* does the reverse, by two distinct mechanisms treated in Section 2.5. *Spontaneous emission* drops the atom to a lower level with no collision partner at all, the only process whose rate is independent of density. *Proton-impact ℓ -mixing* shuffles population between the ℓ sublevels of one shell; it costs no energy, because those sublevels are degenerate, and Section 2.7 shows it is the fastest process in the model.

The chemical-engineering reading is exact. Each level is a species and each process an elementary reaction. Excitation, de-excitation, ionisation and ℓ -mixing need a collision partner, so their rates carry a density factor exactly as a second-order rate law does. Spontaneous emission is unimolecular, like a first-order decomposition. Recombination is the odd one out, because it creates atoms out of a species we have chosen not to track, and Chapter 3 shows it therefore behaves like a constant feed stream into a stirred tank rather than a reaction inside it.

Before quoting numbers we state what we expect, so the table can disagree. An allowed optical transition should run at $\sim 10^8$ 1/s; collisions at $n_e \sim 10^{14}/\text{cm}^3$ should be slower, since a coefficient of order $10^{-9} \text{ cm}^3/\text{s}$ times that density is only 10^5 1/s; and ground-state ionisation at a few electronvolts should be slower still, needing an electron above 13.6 eV from a distribution whose mean energy is far below it. What ℓ -mixing will do cannot be estimated this simply.

Table 2.1 confirms the expectations and exceeds one of them by three orders of magnitude. Radiative decay of $2p$ takes 1.60 ns and ground-state ionisation 71.0 μs , both as anticipated, but ℓ -mixing empties $2s$ in under 7 ps, two hundred times faster than the fastest radiative decay in the atom. That number is why Section 2.7 exists, and the ten-decade spread is the stiff network that motivates the steady-state approximation for intermediates in the first place.

Table 2.1: Representative rates for the six processes at the benchmark point $T_e = 2.947$ eV, $n_e = 1.389 \times 10^{14}/\text{cm}^3$ (grid indices [23, 5]). Rate coefficients are read from the stored arrays of Table 2.6 at temperature index 23; a rate in 1/s is the coefficient times n_e , except for spontaneous emission, which needs no partner. The recombination entry is a feed per unit proton density, not a rate an atom experiences. The spread is ten orders of magnitude, which is why the problem is numerically stiff.

Process	Example	Rate (1/s)	1/rate
Proton ℓ -mixing	$2s \rightarrow 2p$	1.475×10^{11}	6.78 ps
Spontaneous emission	$2p \rightarrow 1s$	6.268×10^8	1.60 ns
Ionisation	$2p \rightarrow \text{H}^+$	8.639×10^6	116 ns
Collisional de-excitation	$2p \rightarrow 1s$	8.173×10^5	1.22 μs
Collisional excitation	$1s \rightarrow 2p$	7.686×10^4	13.0 μs
Ionisation	$1s \rightarrow \text{H}^+$	1.409×10^4	71.0 μs
Recombination feed	$\text{H}^+ \rightarrow 1s$	1.095×10^1	n/a

2.2 The state space

Hydrogen has infinitely many bound levels, so a model must stop somewhere, and must also decide how finely to describe each level it keeps. Too few levels and the cascade feeding the Balmer lines is cut short; too much detail at high n and the matrix grows without the extra states carrying independent information.

This model tracks **43 states**, drawn in Figure 2.1 together with the five processes that connect them. Levels $n = 1\text{--}8$ are resolved in ℓ (indices 0–35, ordered $1S, 2S, 2P, 3S, \dots, 8J$); levels $n = 9\text{--}15$ are *bundled*, one state per shell carrying the full shell weight $g_n = 2n^2$ (indices 36–42). The ordering is fixed by `data/processed/collisions/K_exc_full/state_index.csv` and every array in the pipeline is indexed in it. The ion H^+ is deliberately *not* a 44th state: it is an external reservoir, a choice whose consequences are listed in Section 2.10.

Why ℓ must be resolved at low n . A first guess is that ℓ is a refinement, nice to have, not load-bearing. At $n = 2$ it is wrong, and a selection rule is the reason. Electric-dipole (E1) radiation, much the strongest emission mechanism, requires $\Delta\ell = \pm 1$. The transition $2p \rightarrow 1s$ satisfies this and runs at $A_{2p \rightarrow 1s} = 6.2684 \times 10^8/\text{s}$ (`Radiative/A_resolved.npy`, element [0, 2]). The transition $2s \rightarrow 1s$ does not: both states have $\ell = 0$, and the E1 rate is exactly zero. In this dataset the total radiative loss from $2s$ is $\gamma_{2s} = 0$ identically; the two-photon decay hydrogen really possesses, of order 10/s, is not included, and would be negligible beside the collisional rates of Table 2.1 in any case. So the two sublevels of $n = 2$ are not refinements of each other: one empties in 1.6 ns and the other cannot empty radiatively at all. Bundling them and weighting the decay statistically would hand the metastable $2s$ level a decay rate it does not have.

Why bundling is safe at high n . Above $n = 8$ the argument reverses. Proton ℓ -mixing scales roughly as n^4 (Section 2.7), so by $n \approx 9$ it redistributes population within a shell far faster than anything else happens to that shell. The sublevels then sit permanently in the ratio of their statistical weights and one number per shell holds all the information.

That is an argument, not a measurement, and it can be tested: do the ℓ sublevels of a shell the model *does* resolve come out statistical anyway? For $n = 4$ the statistical




Figure 2.1: **The 43-state space and the five processes the rate matrix contains.** Left: the level structure. Hydrogen is ℓ -resolved up to $n = 8$, which is 36 states, and bundled into one level per n from $n = 9$ to 15 , which is 7 more. Vertical spacing is schematic; the right-hand panel gives the true energies, and shows why, with 37 of the 43 states lying within 0.85 eV of the continuum, a to-scale ladder would be unreadable. Five processes appear in the matrix: electron-impact excitation and de-excitation between levels, radiative decay ($A_{2p \rightarrow 1s} = 6.27 \times 10^8 \text{ s}^{-1}$, and 34 of the 36 resolved states carry a non-zero decay rate), electron-impact ionisation from every state, recombination, which enters as the source vector rather than as a matrix element and feeds every state, and proton-impact ℓ -mixing, which is $\Delta n = 0$ and acts on 35 states, every resolved level with $n \geq 2$. Only the $n = 2$ and $n = 3$ mixing arrows are drawn; at $n = 4$ the sublevels are too close together for an arrow to render legibly. The dashed boxes mark the partition the analysis rests on, and 2s is why it can be drawn where it is. 2s has no E1 decay to the ground state, its radiative rate to 1s being exactly zero in the matrix, so on radiative grounds alone it would be metastable and would belong with the slow states. It is nonetheless fast, because proton ℓ -mixing to 2p carries 99.97% of its total loss rate of $1.47 \times 10^{11} \text{ s}^{-1}$ at the benchmark, and never less than 99.92% anywhere on the grid, rising to 99.99% at the cold edge. That single channel is what leaves $\mathcal{S} = \{1s\}$ as the only slow state and the remaining 42 levels in the fast manifold. The state list, including which levels are bundled and every ionisation energy, is read from the model's own state index, not written into this figure.

share of the $4f$ sublevel is $g_{4f}/g_4 = 14/32 = 0.4375$. Solving the model across the grid gives 0.4361–0.4375 (register entry A6), statistical to four figures. The largest departure, 0.3%, falls at $n_e = 10^{12}/\text{cm}^3$, the lowest density on the grid and exactly where ℓ -mixing is weakest. Two caveats belong in the same breath. The bundled block carries no ℓ -mixing of its own: it is *assumed* statistical rather than driven to it, and that assumption has not been tested directly. And the truncation at $n_{\text{max}} = 15$ is a real limitation whose size Chapter 6 reports as measured rather than assumed.

Where the ordering and the grids come from. A 43-state model has one correct row ordering, and every array must agree on it; getting it wrong produces no error, only plausible numbers about the wrong states. The pipeline therefore never lets a script declare the ordering for itself. `src/validation/cr_context.py` loads the grids, the matrix and the state ordering from the files the pipeline wrote, cross-checks their shapes, and raises if any is missing or inconsistent, so that a script which hardcoded them cannot keep running and silently mislabel every result after the pipeline changes. Every verification result in Chapters 4 and 5 comes from a script that imports it. The grid itself is declared in `src/rates/assemble_cr_matrix.py` and written to disk so downstream code reads rather than redeclares it: 50 temperatures logarithmically spaced over $T_e = 1\text{--}10\text{ eV}$ and 8 densities over $n_e = 10^{12}\text{--}10^{15}/\text{cm}^3$, giving 400 points. The benchmark point carried through this thesis is the one nearest $T_e = 3\text{ eV}$, $n_e = 10^{14}/\text{cm}^3$: index [23, 5], $T_e = 2.947\text{ eV}$, $n_e = 1.389 \times 10^{14}/\text{cm}^3$. It is an illustrative anchor, a single place in parameter space that becomes familiar through repetition. No conclusion rests on it; every general claim is carried by the full grid.

2.3 From a cross section to a rate coefficient

Atomic collision theory does not deliver rate constants; it delivers *cross sections*. The cross section $\sigma(E)$ (cm^2) is the effective target area an atom presents to an electron of kinetic energy E : a beam of flux Γ ($1/(\text{cm}^2\text{ s})$) drives the process $\Gamma\sigma$ times per second. For excitation σ vanishes below the threshold ΔE , rises above it, and falls again at high energy as the electron passes too quickly to interact.

A plasma is not a beam. What a kinetic model needs is one number per temperature, the *rate coefficient* $C(T_e) = \langle \sigma v \rangle$, which for a Maxwell–Boltzmann electron distribution reduces to

$$C(T_e) = \sqrt{\frac{8}{\pi m_e}} (k_B T_e)^{-3/2} \int_{\Delta E}^{\infty} \sigma(E) E e^{-E/k_B T_e} dE, \quad (2.3)$$

with m_e the electron mass (kg), $k_B T_e$ the temperature as an energy (J), E the incident energy (J), ΔE the threshold (J) and σ in m^2 ; the result is converted to cm^3/s . Dimensionally $\text{cm}^2 \times \text{cm}/\text{s} = \text{cm}^3/\text{s}$, and $C n_e n_p$ is then $1/(\text{cm}^3\text{ s})$, a number of events per unit volume per unit time: this is mass-action kinetics with C as a second-order rate constant. What Eq. (2.3) does physically is weight the cross section by how many electrons arrive at each energy and how fast, with the lower limit enforcing that electrons below threshold contribute nothing.

One reading of it is a common trap. The prefactor carries $T_e^{-3/2}$, but $C(T_e)$ does not fall as $T_e^{-3/2}$: the integral grows as more of the distribution clears the threshold, and for a threshold process that growth wins easily. The ground-state excitation coefficient

$C_{1s \rightarrow 2p}$ is $5.5316 \times 10^{-10} \text{ cm}^3/\text{s}$ at the benchmark temperature and nearly three orders of magnitude smaller at 1 eV.

The implementation is `compute_K_CCC.py`, function `maxwell_average`, and four choices in it could have gone wrong silently. The integral is trapezoidal on a 5000-point uniform grid; a 500-point grid was tested and gives $\sim 17\%$ error at 1 eV, so the grid density is measured in the coarse direction, not defaulted. The fine direction was measured later: on a 100000-point grid the ground-state column of $\mathbf{L}$ moves by 1.9% (median) at 2 eV and 4.3% at 1 eV, τ_{slow} by 1.0% and 3.3%, because the raw cross sections begin up to 0.062 eV above the threshold and the trapezoid places that step to within one grid spacing of 0.19 eV; no quantity Chapter 5 quotes moves by more than 1.5% at $T_e \geq 2 \text{ eV}$ (τ_{slow} 1.5%, u_{CRE} 1.3%, $f_3 - f_4$ 0.8%, $\varepsilon_{\text{plateau}}$ 1.0%), and the grid is left at 5000 points with this as the size of its contribution (`verify_maxwell_grid_convergence.py`, `validation/maxwell_grid_convergence/`). The threshold is taken from Eq. (2.1) rather than from the data file, avoiding the noisy near-threshold region. Interpolation returns zero outside the tabulated range. And the integral stops at the top of the data rather than at infinity. Every one of the 1320 excitation blocks is tabulated to between 955 and 969 eV; with the integrand $\sigma E e^{-E/k_B T_e}$ and σE allowed to grow logarithmically beyond the data, the largest missing tail at 10 eV is 3×10^{-41} of the integral (`validation/maxwell_tail/`). No block stops at 100 eV, and had one done so its tail at 10 eV would have been 0.05% for a ground-state dipole transition, because σE and not the Boltzmann factor sets it; a bound built from the Boltzmann factor alone understates the tail.

2.3.1 Detailed balance, and why de-excitation is never computed directly

Excitation and de-excitation are not independent. This is the same constraint that ties the forward and reverse rate constants of an elementary step to its equilibrium constant, and the argument has the same shape. Put the plasma in complete thermodynamic equilibrium, where every process is balanced by its reverse. For levels 1 (lower) and 2 (upper),

$$n_1 n_e C_{12} = n_2 n_e C_{21} \implies \frac{C_{12}}{C_{21}} = \frac{n_2}{n_1} \Big|_{\text{eq}} = \frac{g_2}{g_1} \exp\left(-\frac{\Delta E}{k_B T_e}\right), \quad (2.4)$$

with C_{12}, C_{21} in cm^3/s , g_1, g_2 from Eq. (2.2), and $\Delta E = I_1 - I_2 > 0$. Both sides are dimensionless. Physically: going up needs a rare fast electron and coming down does not, so de-excitation wins by the Boltzmann factor, partly offset by there being more upper states to fall out of.

The electron density cancels, because the same electron is the partner in both directions. So do the populations: C_{12} and C_{21} are Maxwellian averages of cross sections, which depend on T_e and atomic structure and nothing else. Equilibrium was only a device to fix the ratio; the ratio holds in any Maxwellian plasma, however far the populations are from equilibrium. That is the entire justification for using Eq. (2.4) in a model whose subject is departure from equilibrium. It fails only if the electrons themselves are non-Maxwellian.

Worked number. For $1s \rightarrow 2p$ at the benchmark temperature, $g_1 = 2$, $g_2 = 6$, $\Delta E = 10.20435 \text{ eV}$, so Eq. (2.4) predicts $3 \exp(-10.20435/2.947052) = 9.404800 \times 10^{-2}$. The stored arrays give $C_{1s \rightarrow 2p}/C_{2p \rightarrow 1s} = 5.53155 \times 10^{-10}/5.88163 \times 10^{-9} = 9.404800 \times 10^{-2}$,

identical to seven figures. Over all 819 transition pairs at that temperature, using g and I_n from `state_index.csv`, the largest relative deviation is 2.48×10^{-9} .

That agreement checks the implementation and nothing more, and must not be reported as validating the atomic data. It is machine precision because it is true by construction: de-excitation coefficients are never Maxwell-averaged from de-excitation cross sections but derived from the excitation coefficients through Eq. (2.4) (`compute_K_CCC.py`, function `detailed_balance`). The reason for building it that way is that it guarantees Saha–Boltzmann populations in the high-density limit, which Chapter 4 uses as a gate, and which is therefore a weaker gate than one the construction did not guarantee. The independent statement, that the underlying cross sections themselves respect microscopic reversibility, is a separate check reported in Section 2.8.

2.3.2 Two hydrogen energies, and why both are right

A reader auditing the source will find two different numbers near 13.6 eV and should not treat the difference as sloppiness. They are two distinct physical quantities.

The first is the *Rydberg energy* $R_\infty hc = 13.605\,693\,122\,994$ eV, the ionisation energy of an idealised hydrogen atom with an infinitely heavy, point-like, non-relativistic nucleus. It is a defined constant of atomic theory, and it is the unit in which collision calculations tabulate their energies. The second is the *ionisation energy of real hydrogen*, 13.598 434 599 702 eV, the measured energy needed to remove the electron from a real H atom in its ground state. The second is smaller by 0.0534 %.

The pipeline uses each where it belongs, and the rule is simple. Cross sections and thresholds supplied in Rydbergs must be converted back with the constant they were tabulated in, so $R_\infty hc$ is carried by every collisional quantity: the excitation thresholds of Section 2.4, the ionisation thresholds of Section 2.5, and I_H inside the ℓ -mixing formula Eq. (2.10). Photon energies and spectroscopic level energies describe real atoms and carry the real ionisation energy. Both constants were checked against their published values and both are correct for what they are.

Two warnings go with that. First, the difference is *not* simply the reduced-mass correction: scaling $R_\infty hc$ by μ/m_e reproduces the ionisation energy for about five digits and then stops, and the residual 1.47×10^{-4} eV is the Dirac fine-structure and Lamb-shift contribution. Writing one as the other times m_e/m_H is wrong at the 10^{-5} level. Second, the Rydberg energy is stored in the source under a name, `CHI_H`, conventionally reserved for the ionisation potential. Anything that reads that symbol expecting an ionisation potential, a Saha or Boltzmann exponent in particular, is 0.0534 % too large. Nothing in the assembled matrix does so, but the name is a trap and is flagged here rather than renamed.

2.4 Electron-impact excitation

Excitation lifts atoms out of the ground state and populates everything the spectrometer sees, so its cross sections matter more than any other input. They come from the *convergent close-coupling* (CCC) method of Bray and co-workers [27, 28], in which the collision problem is solved by expanding the target in a large square-integrable basis and enlarging that basis until the answer stops moving. The “convergent” is the claim that the limit is reached rather than estimated. The dataset is the 3115 raw files supplied by Prof. Igor Bray (personal communication, 22 February 2026) in

data/raw/ccc/e-H_XSEC_LS/, parsed by `src/parsers/parse_ccc.py`. The acquisition report describes the database requested as 3117 files; 3115 arrived, and they are the 3115 the provider’s own `file_list.txt` names, so which two are absent cannot be settled without the provider’s manifest and is left open. Everything downstream of the files is counted and asserted in `validation/ccc_accounting/`: 2025 files carry two resolved-state labels, of which 870 are excitation, 870 de-excitation and 285 the excluded $\Delta n = 0$, and the remaining 1090 carry a bundled-shell label.

Three instructions came with the data and all three are load-bearing.

The filenames read backwards. The convention is `FINAL.INITIAL`, so `1S.2P` holds $2p \rightarrow 1s$ de-excitation and the excitation $1s \rightarrow 2p$ lives in `2P.1S`. Reading it the wrong way round would transpose the dataset with no immediate symptom. The model’s rate coefficients are Maxwell-averaged from the *excitation*-direction files, selected in `compute_K_CCC.py` by the test $n_f > n_i$, and de-excitation follows by Eq. (2.4). The de-excitation-direction files are parsed and stored but never enter the rate tables.

Transitions with $\Delta n = 0$ were excluded, on the data provider’s instruction. The note reads: “Transitions with $n_{\text{final}} = n_{\text{initial}}$ (e.g. $2S \leftrightarrow 2P$, $3S \leftrightarrow 3P$) have unphysically large cross sections due to non-relativistic degeneracy in the CCC calculation. All $\Delta n = 0$ transitions are excluded from the database.” The reason is that in a strictly non-relativistic treatment the ℓ sublevels of a shell are exactly degenerate, the transition energy is zero, and the long-range dipole interaction gives a divergent cross section. After the exclusion 2190 blocks remain, all with $\Delta n \neq 0$: 1320 in the excitation direction and 870 in the de-excitation direction. Collapsed over ℓ , the 1320 become the 618 CCC pairs of Table 2.2, which with the 201 Vriens–Smeets pairs make the 819, the number of $\Delta n \neq 0$ pairs among 36 resolved and 7 bundled states once the 84 same-shell pairs are removed; every integer in that chain is asserted by the script. The acquisition record’s count of 1740 (a historical project value with no producing script, not used as evidence here and superseded by the accounting in `validation/ccc_accounting/`) is the 870 of each direction among the resolved-state files alone: the 450 excitation blocks into $n = 10$ had not yet been parsed, and the 324 into $n = 9$ were already among the 870. The accuracy the provider states for what remains is $\sim 5\%$, which the benchmark of Section 2.8 tests and does not confirm: it measures 13% for $n_{\text{upper}} \leq 4$.

That exclusion leaves no electron-driven route between $2s$ and $2p$ at all, a consequence easy to miss and one that Section 2.7 turns into the central argument of this chapter.

Covering the shells CCC does not reach. CCC data extend to $n = 8$ resolved in ℓ , with resolved data reaching $n = 9$ and $n = 10$. Transitions involving higher bundled shells are filled from the semi-empirical formulae of Vriens and Smeets [29], coded in `compute_K_VS.py`; the implemented expressions are their Eqs. (17)–(19) directly, not the equivalent rearrangement of Hartgers *et al.* [30] that the file’s header docstring describes. Their stated accuracy is $\sim 20\%$ against the $\sim 5\%$ CCC states for itself, so the merge uses CCC wherever it exists; neither figure is one this thesis has measured, and the one it can test comes out at 13% (Section 2.8). Those formulae need one further input, the shell-to-shell absorption oscillator strength, built from the same Einstein coefficients used in Section 2.6, with a Kramers fallback where a value is missing that is recorded per transition in the metadata. Table 2.2 records which source supplies which block,

Table 2.2: Which source supplies which excitation transitions, from the coverage record in `src/rates/assemble_K_exc.py` and the per-transition attribution in `K_exc_meta.csv`. “res” is an ℓ -resolved state ($n \leq 8$), “bund” a bundled shell ($n = 9\text{--}15$). The collapse rows are ℓ -resolved CCC data summed over final ℓ to give a shell total, which keeps CCC accuracy for the first two bundled shells instead of falling back to Vriens–Smeets. De-excitation coefficients are not listed because they are not independent data: each follows from its excitation partner by Eq. (2.4).

Block	Source	Pairs	Stated accuracy (provider’s)
res $\rightarrow$ res, $n \leq 8$	CCC	546	$\sim 5\%$
res $\rightarrow n = 9$ bundled	CCC, collapsed over ℓ	36	$\sim 5\%$
res $\rightarrow n = 10$ bundled	CCC, collapsed over ℓ	36	$\sim 5\%$
res $\rightarrow n = 11\text{--}15$	Vriens–Smeets	180	$\sim 20\%$
bund $\leftrightarrow$ bund	Vriens–Smeets	21	$\sim 20\%$
total		819	

because “the excitation data are from CCC” would misdescribe a table that is a quarter Vriens–Smeets.

2.5 Ionisation and recombination

2.5.1 Ionisation: the only true loss

Ionisation is the only process that removes an atom from the bound manifold altogether; everything else moves an atom from one tracked level to another. In reactor terms it is the irreversible step consuming reactant to product, and the rest is internal shuffling. The ionisation rate coefficient Q_p (cm^3/s) is obtained by Maxwell-averaging the CCC *total ionisation cross section* (TICS) with Eq. (2.3), the lower limit being the binding energy of Eq. (2.1) (`compute_K_TICS.py`). The threshold is ℓ -independent in hydrogen but the cross sections are not, so the coefficients are ℓ -resolved for $n \leq 8$.

The stored coefficients are not uniform in origin and the metadata records this state by state. The 36 resolved levels carry `CCC_TICS`; the $n = 9$ bundle carries `CCC_TICS_direct`, taken from the bundled `TICS.9` file rather than assembled from sublevels; the six shells $n = 10\text{--}15$ carry `Lotz1968`, the semi-empirical formula of Lotz [31], with a note recorded against every one of them: “overestimates CCC by factor $\sim 4\text{--}8$; negligible impact” (`collisions/tics/K_ion_final_meta.csv`). That note is a flag, not a resolved question, and Section 2.8 returns to it, because the overestimate does not stay confined to ionisation.

Worked number. At the benchmark point $Q_{1s} = 1.0140 \times 10^{-10} \text{ cm}^3/\text{s}$ and $Q_{2p} = 6.2171 \times 10^{-8} \text{ cm}^3/\text{s}$ (`K_ion_final.npy`, temperature index 23): about 600 times easier to ionise an atom already in $n = 2$, which has to supply only 3.40 eV instead of 13.61 eV. This step-ladder character is why the excited levels matter to a quantity as apparently unrelated as the total ionisation rate.

2.5.2 Recombination: two mechanisms, two density dependences

Recombination runs the other way, and telling its two mechanisms apart matters because they scale differently with density. In *radiative recombination* the excess energy leaves as a photon: one electron and one proton are needed, so events go as $n_e n_{\text{ion}}$ and the coefficient α_{RR} has units cm^3/s . In *three-body recombination* a second free electron carries the energy away: three particles must meet, so events go as $n_e^2 n_{\text{ion}}$ and $\alpha_{3\text{BR}}$ has units cm^6/s . Nothing here is special to plasmas. It is the difference between a bimolecular and a termolecular elementary step, with the same consequence: the second mechanism takes over as density rises.

Radiative recombination is computed shell by shell from Johnson’s three-term Gaunt-factor expansion [32], itself a fit to the treatment of Seaton [33],

$$\alpha_{\text{RR}}(n, T_e) = D y^{3/2} e^y [g_0(n)E_1(y) + g_1(n)E_2(y) + g_2(n)E_3(y)], \quad y = \frac{I_n}{k_B T_e}, \quad (2.5)$$

with $D = 5.197 \times 10^{-14} \text{cm}^3/\text{s}$, E_i the exponential integrals and g_0, g_1, g_2 the Gaunt polynomials of Johnson’s Table 1 (`recombination_rates.py`). y and the bracket are dimensionless and D carries the units, so both sides are cm^3/s . Physically, capture into a shell is easiest so capture is favoured at low electron energy. $\alpha_{\text{RR}}(n, T_e)$ decreases monotonically with T_e for every level and has no maximum at any temperature. The two limits follow from the exponential integrals in Eq. (2.5). As $T_e \rightarrow 0$, $y \rightarrow \infty$ and $e^y E_k(y) \rightarrow 1/y$ for every k , so $\alpha \rightarrow D (g_0 + g_1 + g_2) y^{1/2} \propto T_e^{-1/2}$ with a positive coefficient, $g_0 + g_1 + g_2 = 0.876, 0.908$ and 0.925 at $n = 2, 3$ and 4 ; as $T_e \rightarrow \infty$, $y \rightarrow 0$ and $e^y E_1(y) \rightarrow -\gamma - \ln y$, so $\alpha \propto T_e^{-3/2} \ln(T_e/I_n) \rightarrow 0$. Evaluating the module’s own function at fixed y gives a logarithmic slope of -0.50 at $y = 340$ and -1.36 at $y = 3.4 \times 10^{-4}$ for every shell from $n = 2$ to 8 (`validation/rr_limits/`). The shell total is split over ℓ statistically, $\alpha_{\text{RR}}(n, \ell) = \alpha_{\text{RR}}(n)(2\ell + 1)/n^2$ [34]. Radiative capture itself favours low ℓ , so the split is not a statement about capture; it is justified by what happens next. Proton ℓ -mixing redistributes a captured population within its shell faster than any competing loss: over the whole grid the mixing rate out of a sublevel exceeds its total radiative decay by at least a factor 9 at $n = 3$ and 86 at $n = 4$, the two shells the observable is built from (`validation/molecular_channel/`, item 1), so the ℓ -distribution is set by mixing to within about one part in that ratio whatever the capture distribution was. The condition is the timescale ratio, not a density threshold.

$g_1(n)$ in Eq. (2.5) is negative and grows with n , so the three-term sum can in principle turn negative for high shells at low temperature, which would argue for the single-term form instead. The code keeps the three-term expansion behind a clamp at zero, which would silently replace a negative value rather than report it, but neither `alpha_RR_resolved.npy` nor `alpha_RR_bundled.npy` contains an exact zero anywhere on the 50-point temperature grid: the clamp never fires over the range this thesis uses, though it would need revisiting for a colder grid or a longer ladder.

Three-body recombination is not computed independently at all. It comes from ionisation by detailed balance [23], exactly as de-excitation came from excitation:

$$\alpha_{3\text{BR}}(n, \ell, T_e) = Q_{n\ell}(T_e) \frac{g_{n\ell}}{2} \left(\frac{h^2}{2\pi m_e k_B T_e} \right)^{3/2} \exp\left(\frac{I_n}{k_B T_e}\right), \quad (2.6)$$

where h is Planck’s constant and the bracketed factor is the cube of the electron thermal de Broglie wavelength, a volume, so the units read $\text{cm}^3/\text{s} \times \text{cm}^3 = \text{cm}^6/\text{s}$, as required.

Equation (2.6) imposes Saha equilibrium between a level and the continuum: in complete equilibrium the ionisation flux $Q n_{n\ell} n_e$ and the recombination flux $\alpha_{3\text{BR}} n_e^2 n_{\text{ion}}$ must balance, and rearranging Saha for $n_e n_{\text{ion}} / n_{n\ell}$ gives it directly. Like detailed balance for de-excitation, this guarantees the high-density limit rather than testing it.

Assembled, the two mechanisms give the source used throughout the thesis, $\mathbf{S} = n_e(\alpha_{\text{RR}} + n_e \alpha_{3\text{BR}})$, whose units are $1/\text{cm}^3 \times \text{cm}^3/\text{s} = 1/\text{s}$, so that $\mathbf{S} n_{\text{ion}}$ is atoms created per unit volume per unit time. Chapter 3 explains why this is a feed stream rather than an element of the rate matrix.

Worked numbers, and a surprise. The first guess is that three-body recombination is a high-density correction, negligible at $n_e = 10^{12}/\text{cm}^3$. That is true only for the ground state. Forming the three-body fraction $n_e \alpha_{3\text{BR}} / (\alpha_{\text{RR}} + n_e \alpha_{3\text{BR}})$ from the stored coefficients at temperature index 23 gives, for $1s$, below 10^{-4} at $n_e = 10^{12}/\text{cm}^3$ and 0.0085 at $10^{15}/\text{cm}^3$; but for $n = 15$ it gives 0.9996 at the *lower* density and above 0.9999 at the higher. Capture into the ground state is radiative essentially everywhere on the grid; capture into the top shell is three-body essentially everywhere on the grid. Equation (2.6) shows why: $\alpha_{3\text{BR}}$ inherits an ionisation coefficient that is enormous for a shell bound by 0.060 eV, and the exponential that suppresses it for the ground state is close to unity at the top of the ladder. A loosely bound electron is easy to knock off and, by the same token, easy to capture in a three-body encounter.

The three-body coefficients for $n = 10$ – 15 therefore inherit the Lotz overestimate through Eq. (2.6), which Section 2.8 takes up.

2.6 Radiative decay

An excited atom need not be disturbed to decay. The *Einstein coefficient for spontaneous emission* A_{qp} (1/s) is the probability per unit time that an atom in upper level q drops to lower level p , and the total radiative loss from q is $\gamma_q = \sum_{p < q} A_{qp}$, also in 1/s, with $1/\gamma_q$ the radiative lifetime. This is a first-order rate constant in the ordinary chemical sense, the one process with no concentration dependence, so the only one whose relative importance rises as the plasma thins.

Not every pair of levels is connected. E1 radiation requires $\Delta\ell = \pm 1$; transitions violating it are *forbidden*, meaning slower by many orders of magnitude rather than impossible. The dataset here contains E1 transitions only, so a forbidden transition has $A = 0$ exactly. The related *oscillator strength*, a dimensionless measure of how strongly a transition couples to the radiation field, feeds the Vriens–Smeets formulae of Section 2.4.

Two consequences of the selection rule matter to this thesis. First, $2s \rightarrow 1s$ is forbidden, so $\gamma_{2s} = 0$ (Section 2.2). Second, $4f \rightarrow 2p$ is forbidden because ℓ would change by two, so the $4f$ sublevel holds 43.75% of the $n = 4$ shell by weight and contributes nothing to H_β , decaying only to $3d$. The $n = 4$ *shell* population and the H_β *line* intensity are therefore not proportional, and Chapter 4 measures how much that matters to the observable. The three channels that do make H_α are $3s \rightarrow 2p$ ($6.3172 \times 10^6/\text{s}$), $3p \rightarrow 2s$ ($2.2461 \times 10^7/\text{s}$) and $3d \rightarrow 2p$ ($6.4688 \times 10^7/\text{s}$); the three that make H_β are $4s \rightarrow 2p$ ($2.5796 \times 10^6/\text{s}$), $4p \rightarrow 2s$ ($9.6735 \times 10^6/\text{s}$) and $4d \rightarrow 2p$ ($2.0637 \times 10^7/\text{s}$).

The coefficients come from the exact hydrogenic code of Hoang-Binh [35], distributed through the Computer Physics Communications Program Library as ADUU v1.0, which evaluates the radial integrals analytically rather than fitting them [36]. Its 1015 transitions

for $n \leq 15$ are stored in `Radiative/H_A_E1_LS_n1_15_physical.csv`; every entry has $|\Delta\ell| = 1$ and $n_{\text{upper}} > n_{\text{lower}}$, and `radiative_rates.py` sweeps every allowed combination on load and raises if one is missing rather than assuming completeness. For the bundled shells the A values are averaged over ℓ with statistical weights, the same assumption that justified bundling in Section 2.2, applied consistently. The pattern to notice is that γ falls steeply with n , from $6.27 \times 10^8/\text{s}$ for $2p$ to $8.16 \times 10^4/\text{s}$ for the $n = 15$ bundle: radiative decay dominates at low n and collisions at high n , and the level where they cross is a measurable feature of the model that Chapter 4 locates.

These are not the recommended values, and it matters to say so. An external check of the four Lyman-series coefficients finds them 0.056 to 0.060% *above* the critically evaluated values recommended by Wiese and Fuhr [37]: for Lyman- α , $6.2684 \times 10^8/\text{s}$ stored against $6.2649 \times 10^8/\text{s}$ recommended. The offset is systematic, not scattered, and it has a cause. What this code computes is the *exact non-relativistic, infinite-nuclear-mass hydrogenic* transition probability, reproduced to 1–4 parts in 10^5 ; the recommended values include the finite proton mass and the small relativistic and quantum-electrodynamic corrections that the non-relativistic Schrödinger problem does not contain. Both sets are correct for what they are, and using the exact hydrogenic set throughout is internally self-consistent, which is what a rate matrix needs. **But this dataset must not be described as the recommended or evaluated set, and no evaluated compilation should be cited as its source.** It is the analytic hydrogenic solution, and that is what is written here.

The largest silent error available in a rate matrix. One number in this section can be wrong by a third without anything complaining, and the trap is worth spelling out because it is the exact mistake this chapter’s ℓ -resolution is designed to avoid. The value $A_{2p \rightarrow 1s} = 6.2684 \times 10^8/\text{s}$ is an *ℓ -resolved level rate*: it is the decay probability of an atom that is in $2p$. The corresponding *n -resolved shell rate* for $n = 2 \rightarrow n = 1$ is smaller by the fraction of the shell that can radiate at all,

$$A_{n=2 \rightarrow n=1} = \frac{g_{2p}}{g_2} A_{2p \rightarrow 1s} = \frac{6}{8} \times 6.2684 \times 10^8/\text{s} = 4.7013 \times 10^8/\text{s}, \quad (2.7)$$

the factor $3/4$ being the statistical weight of $2p$ within $n = 2$, and the missing quarter being the $2s$ sublevel that has no E1 route to the ground state at all. Both sides are in $1/\text{s}$. Placing the ℓ -resolved number on an n -resolved matrix element therefore overestimates Lyman- α emission by 33.3%, and nothing in a conservation check or a sign check would notice, because the error is in a magnitude and magnitudes cancel in column sums (Section 2.9).

This model is ℓ -resolved for $n \leq 8$, so the ℓ -resolved value is the correct one to use here, but that was audited rather than assumed. Reading the stored block directly: the element $A[1s, 2s]$ is exactly zero, the element $A[1s, 2p]$ carries $6.2684 \times 10^8/\text{s}$, $\gamma_{2s} = 0$ and $\gamma_{2p} = 6.2684 \times 10^8/\text{s}$, and across all 36×36 resolved elements there is not one non-zero entry with $|\Delta\ell| \neq 1$. The shell-averaged value of Eq. (2.7) appears nowhere in the resolved block. The bundled block $n = 9\text{--}15$ is n -resolved by construction and carries statistically weighted shell rates throughout, which is the consistent treatment for states that are always in statistical ℓ -equilibrium.

2.6.1 The photons are assumed to leave

Writing γ_q as a bare sum of Einstein coefficients contains an assumption: that every emitted photon escapes and never returns. A plasma for which this holds is *optically thin*. If it is optically thick, a photon emitted by one atom is absorbed by another, which returns to the upper level, and the effective decay rate falls. The chemical-engineering analogue is a reaction slowed because the product cannot leave and keeps re-adsorbing: the kinetics are unchanged, the effective rate constant is not.

The size of the effect is set by the *optical depth* τ_c , the number of absorption lengths a photon must cross to escape, and carried by the *population escape factor* Θ_P , defined so that the effective coefficient is $\Theta_P A$. For the resonance line the model is most exposed to, Lyman- α ($2p \rightarrow 1s$), `src/analysis/escape_factor.py` evaluates the escape factor of Holstein [38] in the formulation of Behringer [39] adopted as routine ADAS214 of the Atomic Data and Analysis Structure (ADAS) [8], by direct quadrature and with no fitting formula:

$$\Theta_P(\tau_c) = \frac{1}{\sqrt{\pi}} \int_{-\infty}^{\infty} e^{-x^2} \exp(-\tau_c e^{-x^2}) dx, \quad \tau_c = n_{1s} \sigma_0 \frac{D}{2}, \quad (2.8)$$

with x the frequency offset in Doppler widths (dimensionless), n_{1s} the ground-state density ($1/\text{cm}^3$), σ_0 the line-centre absorption cross section (cm^2) from the Ladenburg relation with oscillator strength $f_{12} = 0.4162$ and $\lambda_0 = 121.567 \text{ nm}$, and D twice the escape path length (cm). Both sides of the first expression are dimensionless. It says: weight the escape probability by the fraction of photons emitted at each frequency, and note that photons in the line wings get out when photons at the line centre cannot.

The geometry is mislabelled, and the label matters. The module documents Eq. (2.8) as applying to “a homogeneous slab of half-thickness b ”. Read directly from the ADAS214 manual, equation 3.14.14 is the *isotropic* case (the escape probability at the centre of a sphere, ADAS geometry g1) and not the slab case, g2. The distinction is not cosmetic in the thick regime: a true slab requires the further angular average $\Theta_{\text{slab}}(\tau) = \int_0^1 \Theta_{\text{iso}}(\tau/\mu) d\mu$, which is smaller by a factor approaching two, measured as 0.529 at $\tau_c = 1$ and 0.477–0.484 at $\tau_c = 100$ –1000. Everything else in the module was checked against the primary source and is correct: the convention $\tau_c = \kappa_0 b$ with b the edge-to-centre distance is exactly ADAS’s own, the Ladenburg constant $\pi e^2/m_e c = 2.654\,008\,85 \times 10^{-2} \text{ cm}^2 \text{ Hz}$ is right, the Doppler profile normalisation and thermal speed reproduce ADAS’s own tabulated widths, and the leading Holstein constant is 1 rather than the 1.6 sometimes quoted. **This is reported, not repaired.** Where the plasma is thin, $\Theta_P \rightarrow 1$, the two geometries agree and nothing depends on the choice, and Section 2.10 and Chapter 3 show that this is the whole region carrying this thesis’s quantitative results. Where the plasma is thick, at the cold edge, the escape factor computed here is roughly twice too large if a slab is the right picture, and Chapter 6 carries that caveat rather than resolving it.

Two constants in this module are wrong, and both are immaterial. An external audit of the same file found the elementary charge entered as $e = 4.803\,26 \times 10^{-10} \text{ esu}$ against the CODATA value $4.803\,204\,713 \times 10^{-10}$ [40], high by 0.001 15 % and entering σ_0 as e^2 , so 0.0023 %; and the Doppler width computed with $m = 1.672\,62 \times 10^{-24} \text{ g}$, which is the *proton* mass where the physically correct object is the neutral hydrogen *atom*, $1.673\,532\,840\,7 \times 10^{-24} \text{ g}$. A neutral atom Doppler-broadens with its own mass, so the line is 0.027 % too wide. Neither shifts any conclusion in this thesis by a detectable amount,

and neither has been edited, because a constant corrected quietly is a constant nobody can find again.

Every rate matrix in this thesis is built with bare A coefficients, $\Theta_P \equiv 1$. Equation (2.8) is a diagnostic applied afterwards, to find where that choice is defensible. Chapter 3 carries out that evaluation across the grid and Chapter 6 states the consequences where the plasma is not thin. The question can be deferred because τ_c depends on n_{1s} , an output of the model rather than an input, and on a path length D the model does not contain at all.

2.7 Proton-impact ℓ -mixing, and why it is not optional

2.7.1 Why proton collisions mix ℓ

Collisional–radiative treatments of hydrogen often omit proton collisions on the grounds that protons are slow and heavy and therefore ineffective. For transitions that change n that is right. For transitions that only change ℓ it is exactly backwards, and the chain is followed here link by link, because it explains a rate that would otherwise be merely a large number in a table.

Degeneracy makes the mixing free, and slowness helps. Within a shell, ns , np , nd have identical energy, so moving population among them costs nothing: there is no threshold to clear, unlike excitation $n \rightarrow n + 1$. And a proton, 1836 times heavier than an electron, moves 43 times more slowly at the same temperature, so the encounter lasts far longer. For a process needing a sustained push rather than a sharp kick, a long gentle encounter is worth more.

The mechanism is the linear Stark effect. A passing proton is a moving charge; while nearby, the atom sits in a nearly static electric field. Because s , p and d are degenerate, hydrogen responds linearly, and the field’s own eigenstates are mixtures of s , p and d . An atom in a pure np state is therefore no longer in an eigenstate once the field arrives; it evolves as a superposition, and when the proton leaves the population has been redistributed across ℓ . No photon and no dipole selection rule is involved.

Distant collisions dominate, logarithmically, and Debye screening cuts them off. The field at impact parameter b falls as $1/b^2$, so the transition probability does too; but the number of protons passing near b grows as $b db$. The product is db/b , so every decade in b contributes equally and the integral produces a logarithm, which is why the process is strong, and why it needs a cut-off. That cut-off is physical: a proton is dressed by a screening cloud, its potential falls as $e^{-R/\lambda_D}/R$, and beyond about one Debye length λ_D the atom feels nothing. So $b_{\max} \approx \lambda_D$ and the Coulomb logarithm is $\ln(\lambda_D/b_{\min})$.

Assembling those links by an impact-parameter calculation gives the scaling

$$\nu_{\ell\text{-mix}} \propto n_{\text{ion}} n^4 T_i^{-1/2} \ln \Lambda, \quad (2.9)$$

with ν in $1/\text{s}$, n_{ion} the proton density ($1/\text{cm}^3$), T_i the *proton* temperature (eV) and $\ln \Lambda$ the Coulomb logarithm. The n^4 comes from a dipole matrix element scaling as n^2 and entering squared; the $T_i^{-1/2}$ from slower protons interacting for longer. Equation (2.9) tells us in advance that ℓ -mixing will be fast, will get faster with density, and will grow explosively up the ladder, which is the quantitative content of the bundling argument in Section 2.2.

2.7.2 The implemented formula

The rates actually used come from the Debye-cutoff form of the Pengelly–Seaton treatment [41] given by Badnell *et al.* [42], called PSM20 in the code and in this thesis, coded in `src/rates/compute_lmix.py`. For the downward transition within shell n , writing $\ell_{>} = \max(\ell, \ell')$,

$$q_{ji} = \frac{a_0^3}{\tau_0} \left(\frac{\pi \mu I_H}{k_B T_e} \right)^{1/2} \frac{D_{ji}}{\omega_\ell} F(U_m), \quad D_{ji} = 6n^2 \ell_{>} (n^2 - \ell_{>}^2), \quad (2.10)$$

with a_0 the Bohr radius, $\tau_0 = a_0/\alpha c$ the atomic unit of time, $\mu = m_p/2m_e = 918.08$, $\omega_\ell = 2\ell_{>} - 1$ the weight of the lower sublevel, and D_{ji} the dimensionless dipole coupling. The prefactor a_0^3/τ_0 carries the units and everything else is dimensionless, so q_{ji} is a rate coefficient. Note $D_{ji} \sim 6n^4$ for $\ell_{>} \ll n$, recovering the n^4 of Eq. (2.9), and that temperature enters as $T_e^{-1/2}$, recovering the $T_i^{-1/2}$ under $T_i = T_e$. The upward rate follows by reciprocity, $q_{\text{up}}/q_{\text{down}} = (2\ell_{\text{hi}} + 1)/(2\ell_{\text{lo}} + 1)$, which is Eq. (2.4) for degenerate states, where the Boltzmann factor is unity and only the weights survive.

$F(U_m)$ is the Coulomb logarithm of the physical picture, in the form Badnell’s Eq. 9 gives it:

$$F(U_m) = \frac{\sqrt{\pi}}{2} U_m^{-3/2} \text{erf}\left(\sqrt{U_m}\right) - \frac{e^{-U_m}}{U_m} + E_1(U_m), \quad U_m = \frac{E_{\min}}{k_B T_e}, \quad (2.11)$$

with $E_{\min}$ the minimum energy transfer set by the cut-off radius $R_c = \lambda_D$. F is dimensionless and *diverges* logarithmically as $U_m \rightarrow 0$: a smaller minimum energy transfer means a larger cut-off radius means more decades of impact parameter contributing. That limit is the subject of Section 2.9, where the audit found it implemented backwards.

Three approximations enter and are carried through the thesis: $T_i = T_e$, $n_{\text{ion}} = n_e$ by quasi-neutrality, and $\Delta\ell = \pm 1$ only. A fourth is easy to miss on reading the code. The Debye length inside U_m is evaluated at a fixed reference density $10^{14}/\text{cm}^3$ rather than at each grid density, on the stated grounds that F depends on density only logarithmically. That stated ground is wrong, and quoting one number per shell understates by how much. Evaluating the module’s own expression across the eight density columns at 1 eV, F varies by a factor 3.7 for $n = 2$, which has only the $\ell = 1$ channel, and by between 50 and 687 across the seven ℓ channels of $n = 8$, the largest belonging to $\ell = 1$ and the smallest to $\ell = 7$. Over the whole (T_e, n_e) grid the $n = 8$, $\ell = 1$ spread reaches 1564 (`compute_lmix.py`, functions `_psm20_Um` and `_psm20_F`). The practice survives for a different reason, measured rather than assumed. Rebuilding $\mathbf{L}$ with a density-consistent cutoff at every grid point moves $f_3 - f_4$ by at most 0.42%, and scaling every ℓ -mixing rate over a range of $\times 0.1$ to $\times 10$ moves it by less than 0.3% at the benchmark and the density maximum. ℓ -mixing is saturated, so the frozen table costs nothing that can be measured. The stored table has no density axis at all, so the linear $\propto n_e$ scaling applied at assembly is exact while the Coulomb-logarithm factor inside it is frozen at mid-grid. Badnell’s formulation offers two competing cut-offs to the Debye radius, one set by the level lifetime and one by any residual splitting of the ℓ degeneracy, and the correct R_c is the smallest of the three. Only the Debye cut-off is implemented, and the cost of that choice is now measured rather than expected (`verify_lmix_cutoffs.py`, `validation/lmix_cutoffs/`). The cut-off enters only through the Coulomb-logarithm factor $F(U_m)$ with $U_m \propto R_c^{-2}$, so a competing radius smaller than λ_D by a factor r changes F by about $\ln r$. The lifetime radius $v\tau$ exceeds λ_D at every channel and grid

point, by a factor 6.7 at its closest, and never binds. The splitting radius $\hbar v/\Delta E$, taken conservatively at the largest fine-structure splitting between the two coupled sublevels, does bind at low density: for the $2s \leftrightarrow 2p$ channel it is smaller than λ_D at 300 of the 400 grid points, everywhere below $n_e = 3.7 \times 10^{14}/\text{cm}^3$, and at $10^{12}/\text{cm}^3$ it would reduce F from about 9 to 14 down to 0.41 of that value; if the smaller Lamb-shift splitting were the relevant one instead, the reduction is to 0.90. The same cut-off binds at $n = 3$ and $n = 4$ at the lowest densities, reducing F to 0.56 and 0.72 of the Debye value at worst, and never above $n = 5$ or above $3.7 \times 10^{14}/\text{cm}^3$. The smallest factor anywhere on the grid is 0.41, inside the $\times 0.1$ to $\times 10$ scan above that moves $f_3 - f_4$ by less than 0.3%, so which cut-off is correct is an open question about the mixing rate and a closed one for the result.

2.7.3 Without it, $n = 2$ is not a physical manifold

We can now close the loop that Sections 2.2, 2.4 and 2.6 each opened, and the result is the strongest single statement in this chapter.

Ask what an atom in $2s$ can do. Its only radiative exit is two-photon emission at 8.229/s: $2s \rightarrow 1s$ is E1-forbidden, and $\gamma_{2s} = 0$ in this dataset because the two-photon channel is not carried in it. That rate is ten orders of magnitude below the collisional rates that follow, which is why its omission is harmless here and why it nonetheless matters: it is the one depopulation route that does not vanish as $n_e \rightarrow 0$. It cannot be moved to $2p$ by an electron: $\Delta n = 0$ transitions were excluded from the CCC dataset on the provider's instruction. What is left is electron-impact excitation to $n \geq 3$, electron-impact ionisation, and proton ℓ -mixing. Decomposing the total loss from $2s$ into those channels, from the stored coefficient arrays at the benchmark point: proton ℓ -mixing $1.4745 \times 10^{11}/\text{s}$; electron-impact collisions to $n \geq 3$ $3.4030 \times 10^7/\text{s}$; electron-impact ionisation $6.9216 \times 10^6/\text{s}$; radiative decay zero. ℓ -mixing carries 99.972% of everything that happens to a $2s$ atom.

Because γ_{2s} is exactly zero, every remaining channel is collisional and scales with n_e , so that share is independent of density: it depends only on temperature, running from 99.994% at $T_e = 1 \text{ eV}$ to 99.921% at 10 eV. **Over the whole 400-point grid, proton ℓ -mixing carries at least 99.92% of the $2s$ loss rate**, the minimum falling at the hot edge. There the contrast is starkest: without ℓ -mixing $2s$ would empty at $6.17 \times 10^5/\text{s}$; with it, at $7.79 \times 10^8/\text{s}$, a factor of 1262.

The conclusion is not that omitting ℓ -mixing perturbs the $n = 2$ manifold. It is that omitting it changes which process controls the manifold, by three orders of magnitude, and leaves $2s$ an artificial metastable trap. Any model of hydrogen that resolves ℓ and excludes $\Delta n = 0$ electron impact *must* include proton ℓ -mixing or it is not describing the $n = 2$ shell at all.

Worked numbers. At the benchmark point $q_{2s \rightarrow 2p} = 1.0612 \times 10^{-3} \text{ cm}^3/\text{s}$ (`K_1mix.npy`, element [2, 1, 23]), giving $1.4745 \times 10^{11}/\text{s}$ and a $2s$ lifetime of 6.78 ps. Within $n = 5$ the ℓ -mixing loss rates run from $3.25 \times 10^{12}/\text{s}$ for $5s$ (0.31 ps) to $3.47 \times 10^{11}/\text{s}$ for $5g$ (2.88 ps). An independent hand estimate from the impact-parameter picture, with the Debye length as cut-off and $b_{\min} = n^2 a_0$, gives 0.16 ps for $n = 5$: high by roughly six, in the expected direction, because the crude estimate uses the full orbital size for the dipole matrix element and drops the angular factors that D_{ji} supplies, so the large number in the table is physics and not a coding error. Finally, the largest single entry anywhere in the assembled rate matrix at the benchmark point is the $8s$ diagonal, $-9.5235 \times 10^{12}/\text{s}$ (`L_grid.npy`, element [23, 5, 28, 28]): again ℓ -mixing, at the highest resolved shell, exactly

where the n^4 of Eq. (2.9) puts it.

2.8 How well do we know these rates?

Internal consistency checks cannot establish that a rate is the right size. Detailed balance holds to 2.5×10^{-9} and column-sum conservation, which Chapter 3 uses, to 3.6×10^{-11} , both by construction, so both would hold for a dataset uniformly wrong by a factor of three. The only test with real severity is comparison against an independent calculation.

That comparison is against Anderson *et al.*, who computed effective collision strengths Υ for all $n \leq 5$, $\Delta n \neq 0$ transitions in hydrogen using the R-matrix with pseudostates (RMPS) method, a different formulation, basis and code from CCC. Two papers carry that calculation and they are not interchangeable. The original tabulation [43] was superseded by a corrigendum [44] after Dere and Mason pointed out that its dipole effective collision strengths did not approach the correct high-energy asymptotic form. The underlying collision strengths were sound; the fault was in the Maxwell averaging. The high-energy tail had been fitted by least squares using the infinite-energy limit point, but RMPS collision strengths oscillate shallowly at all energies because of high-lying pseudo-thresholds, so the fit degraded progressively beyond the last computed energy. The corrigendum recomputes Υ from the same collision strengths using linear interpolation of the reduced collision strength between the highest finite energy and the infinite-energy limit. **This work uses the corrigendum values throughout.**

How much that choice matters depends entirely on temperature. Comparing the two tabulations transition by transition, the mean change over $T_e = 1$ to 10 eV is 1.4 % and no transition moves by more than 13.1 %; at 15 eV the largest change is 33 % and at 25 eV it reaches 86 %, on $1s \rightarrow 3p$ in both cases. The revision is therefore almost irrelevant to this thesis, which never leaves $T_e \leq 10$ eV, and decisive for any extension of this model to attached-divertor or scrape-off-layer temperatures. The distinction is recorded here because the superseded table remains in circulation.

The comparison itself is done in `src/validation/anderson2002_benchmark.py`, which reads Υ directly from the published corrigendum table rather than from a transcription, and converts it to a rate coefficient through the authors' Eq. (3). That conversion has one trap worth naming: the statistical weight in the denominator is that of the *lower* state and is $(2S+1)(2L+1) = 2(2\ell+1)$ for hydrogen doublets, so using $2\ell+1$ introduces a systematic factor of two. The table's row convention, upper state first, is verified at run time rather than assumed, by checking that the radiative rate listed against the $2p-1s$ row is 6.27×10^8 /s. Comparison is made at all 50 temperatures of the working grid, with Υ interpolated logarithmically in T_e ; since the grid spans 1 to 10 eV and the RMPS table spans 0.5 to 25 eV, this interpolates and never extrapolates. Table 2.3 gives the outcome over the resulting 4250 comparisons.

Taken at face value the last row is a failure: fewer than half the comparisons agree within 20 %, most of the shortfall in transitions reaching $n = 5$. But the degradation with upper shell is graded rather than abrupt, and the disagreement does not lie *entirely* at $n = 5$: $n_{\text{upper}} = 4$ already falls to 74.0 % within 20 % against 96.9 % for $n_{\text{upper}} = 3$. Within the $n_{\text{upper}} = 4$ block the spread is carried by $\Delta n = 1$ and $\Delta n = 3$; the $\Delta n = 2$ transitions agree within 20 % at every point tested. Nor is the agreement uniform in temperature. For $n_{\text{upper}} \leq 4$ the mean absolute deviation runs from 7.2 % at 1 eV to 17.9 % at 10 eV, so the figure of 13 % in Table 2.3 is a grid average and not a bound. Figure 2.2 shows both




Figure 2.2: **The CCC excitation rate coefficients benchmarked against the independent R-matrix-with-pseudostates calculation.** (a) The ratio of the two, rather than one against the other: over twelve decades a $\pm 20\%$ band is thinner than the plotted line, and any pair of datasets looks perfect on a 1:1 log-log plot. The horizontal axis is $\Delta E/T_e$, how far below threshold the Maxwell average samples the cross section; each transition spans one decade, its position is its threshold and its shape the temperature dependence of the disagreement. Colour is the upper shell. Transitions into $n' \leq 3$ sit inside the band almost everywhere; those into $n' = 5$ scatter from a factor 0.2 to a factor 3 in two modes, the near-degenerate $4 \rightarrow 5$ set high at small $\Delta E/T_e$ and the ground-state set low at large. The disagreement grows toward threshold for every transition into $n = 5$ from $n = 2$ and 3 and shrinks toward it for most into $n = 4$; opposite signs in neighbouring shells rule out a defect in the averaging, which would act on all alike, and place the difference in the cross sections. (b) Mean absolute deviation against temperature, over the full range the RMPS table covers. Agreement degrades smoothly with the upper shell rather than holding and then failing: $n' = 4$ crosses 20% inside the range this work uses, which is why the 13% quoted for $n' \leq 4$ is a grid average and not a bound. From `ccc_vs_anderson2002_thesis_Te_grid.csv` with thresholds from `ccc_vs_anderson2002_benchmark.csv`, and `ccc_vs_anderson2002_full_Te_range.csv`.

Table 2.3: The CCC excitation rate coefficients used in this model, compared against the R-matrix with pseudostates calculation of Anderson *et al.* [43] as revised by the corrigendum [44], over 85 matched transitions at all 50 temperatures of the working grid (4250 comparisons). The error is $(K_{\text{CCC}}/K_{\text{RMPS}} - 1)$; a negative median means CCC is the smaller of the two. Source: `data/processed/collisions/ccc_vs_anderson2002_thesis_Te_grid.csv` and the report `ccc_anderson.md`. The agreement degrades smoothly with the upper shell rather than failing abruptly at $n = 5$.

Subset	comparisons	within 20 %	mean error	median error
$n_{\text{upper}} = 2$	100	100.0 %	10.9 %	−10.5 %
$n_{\text{upper}} = 3$	450	96.9 %	8.5 %	−5.6 %
$n_{\text{upper}} = 4$	1200	74.0 %	14.9 %	−9.9 %
$n_{\text{upper}} = 5$	2500	12.3 %	41.1 %	−34.2 %
$n_{\text{upper}} \leq 4$	1750	81.4 %	13.0 %	−8.8 %
all transitions	4250	40.7 %	29.5 %	−21.5 %

features: the $n' = 5$ points scatter over more than a factor of ten while the $n' \leq 3$ points stay inside the 20 % band. The horizontal axis is $\Delta E/T_e$, how far below threshold the Maxwell average is sampling the cross section, because that is the one variable on which a defect in the averaging would have to show: an interpolation or integration error near threshold would grow with $\Delta E/T_e$ for every transition alike. It does not. Over the 25 transitions into $n = 5$ from $n = 2$ and $n = 3$ the disagreement grows toward threshold in every one (rank correlation of |error| with $\Delta E/T_e$ of +0.65 and +0.73 for the two lower shells); over the 24 transitions into $n = 4$ it shrinks toward threshold in 17 and grows in 3 (pooled −0.3 to −0.5); and the 20 transitions between the nearly degenerate $n = 4$ and $n = 5$ shells split ten to seven. Opposite signs in neighbouring shells are not what a shared averaging defect produces. They are what a difference in the cross sections themselves produces, and the class in which the disagreement grows toward threshold is the class whose upper shell is the top of the RMPS basis, which is where Section 2.8.1 finds that basis misbehaving. Plotting the same ratio against the size of the rate coefficient instead does not rule out a threshold artefact: the rate mixes strength, threshold and temperature, and the correlation it gives, 0.02, is flat because two opposing trends cancel.

At $n = 5$ the disagreement has two modes with opposite signs. Where the transition is dipole-allowed between the nearly degenerate $n = 4$ and $n = 5$ shells, as for $4s \rightarrow 5p$ with its gap of 0.31 eV, CCC gives the larger rate, by up to a factor of 2.97; every one of the 105 comparisons in which CCC exceeds RMPS by more than a factor of two belongs to the four transitions $4s \rightarrow 5p$, $4p \rightarrow 5s$, $4p \rightarrow 5d$ and $4d \rightarrow 5f$, and the excess grows with temperature, from 1.4 at 1 eV to 3.0 at 5 eV for $4s \rightarrow 5p$. Where the transition is a high-multipole one with no dipole matrix element, as for $1s \rightarrow 5g$ in which ℓ changes by four, CCC gives rates several times smaller: a factor of 5 for $1s \rightarrow 5g$ at every temperature on the grid, and 275 comparisons across eleven transitions fall below half the RMPS value. This second mode dominates the aggregate, which is why the median error at $n_{\text{upper}} = 5$ is −34.2 %. The dipole transitions out of the ground state sit between the two, with CCC low by a factor that shrinks with temperature, 0.54 at 1 eV to 0.75 at 10 eV for $1s \rightarrow 5p$. A count over four temperatures of one direction only gives ten comparisons beyond a factor of two, all with CCC the larger; the figure plots all fifty.

Table 2.4: Successive percentage changes in the reduced rate coefficient $R(n_{\text{up}})$ of Eq. (2.12), threshold factor removed, along three dipole series at 1 eV and 10 eV. At 1 eV the RMPS step into $n = 5$, the top shell of their target basis, is a *rise* in all three series, the only positive entries in the table; at 10 eV neither calculation kinks and the RMPS steps are uniformly the flatter. Dashes mark shells the RMPS tabulation does not contain. Source: `data/processed/collisions/anderson_validity_dipole_series.csv`, columns `R_CCC_bz` and `R_And_bz`.

Series	T_e	Source	2 $\rightarrow$ 3	3 $\rightarrow$ 4	4 $\rightarrow$ 5	5 $\rightarrow$ 6	6 $\rightarrow$ 7
$1s \rightarrow np$	1 eV	CCC	-19.5	-9.2	-20.6	-1.2	-2.5
		RMPS	-19.6	-9.9	+28.4	–	–
	10 eV	CCC	-41.4	-17.9	-9.9	-5.2	-3.8
		RMPS	-34.0	-10.9	-6.9	–	–
$2p \rightarrow nd$	1 eV	CCC	–	-42.6	-27.3	-14.9	-8.3
		RMPS	–	-42.6	+12.4	–	–
	10 eV	CCC	–	-64.6	-33.2	-18.5	-11.8
		RMPS	–	-59.1	-29.3	–	–
$2s \rightarrow np$	1 eV	CCC	–	-28.5	-24.2	-10.2	-10.1
		RMPS	–	-29.8	+12.0	–	–
	10 eV	CCC	–	-54.4	-26.6	-14.4	-9.6
		RMPS	–	-47.1	-21.5	–	–

2.8.1 Where the $n = 5$ disagreement is localised

Both calculations are approximate and neither is a measurement, so the disagreement cannot be settled by preferring one author. It can, however, be *localised*, using an asymmetry between the two datasets: the RMPS tabulation stops at $n = 5$, while CCC continues to $n = 10$. Only one of them can be asked what happens next.

The test uses a property of hydrogen that holds regardless of which code computes it. Along a Rydberg series of dipole transitions $n_{\text{lo}} \rightarrow n_{\text{up}}$ the absorption oscillator strength falls as $f \sim C/n_{\text{up}}^3$ as $n_{\text{up}} \rightarrow \infty$, and at a temperature well above threshold the rate coefficient is dominated by the Bethe term, the leading high-energy contribution to a dipole-allowed cross section and itself proportional to f , so it inherits that scaling. The rate coefficient also carries the threshold factor $e^{-\Delta E/T_e}$, with ΔE the excitation energy of the transition, which both calculations share and which at 1 eV dominates everything else: along $1s \rightarrow np$ it alone falls by 26, 15 and 10 % at the steps into $n = 5$, 6 and 7. Dividing it out leaves the reduced quantity

$$R(n_{\text{up}}) = K(n_{\text{lo}} \rightarrow n_{\text{up}}) n_{\text{up}}^3 e^{\Delta E/T_e}, \quad (2.12)$$

proportional to Υn_{up}^3 over the weight of the lower state, which must flatten as the series is climbed. Equation (2.12) does not say what the constant is, and it is not used to compute anything; its only role is that a dataset tracking hydrogenic behaviour produces a smooth $R(n_{\text{up}})$, so a dataset that kinks at one particular n_{up} is departing from that behaviour there. Leaving the threshold factor in gives steps of -41, -16 and -12 % for CCC, which are that factor rather than the scaling under test. Table 2.4 gives the successive changes in R for three series at both ends of the temperature range.

Figure 2.3 plots the same quantity against shell at both ends of the temperature range.




Figure 2.3: **Which of the two calculations is anomalous at $n = 5$.** The reduced rate coefficient $K n_{\text{upper}}^3 e^{\Delta E/T_e}$ of Eq. (2.12) along three dipole series, with the threshold factor the two calculations share divided out. For hydrogen the oscillator strength falls as n_{upper}^{-3} asymptotically, so this quantity must flatten as the series is climbed; the constant it approaches is not used, only the smoothness. Dark curves are $T_e = 1$ eV, faded curves 10 eV. CCC continues to $n = 10$. At 1 eV the RMPS values track CCC to $n = 4$ and then *rise* into $n = 5$, the top shell of their target basis, in all three series; at 10 eV neither calculation kinks and the RMPS values are uniformly the flatter. The kink belongs to the temperatures at which the Maxwell average samples the cross section close to the $n = 5$ threshold. It localises the departure from hydrogenic behaviour at one shell of one calculation without proving CCC correct, which is not smooth there either. From `anderson_validity_dipole_series.csv`, columns `R_CCC_bz` and `R_And_bz`.

The two datasets agree on the shape of each series until the last step the RMPS calculation can take, and what happens at that step depends on the temperature. At 1 eV the RMPS reduced rate *rises* into $n = 5$, by 28, 12 and 12 % in the three series, while every other step of either calculation in Table 2.4 is a fall; the ratio of the RMPS rate to the CCC rate along $1s \rightarrow np$ is 1.15, 1.15, 1.14 and then 1.85. CCC is not smooth there either, dropping by 20.6 % into $n = 5$ on $1s \rightarrow np$ between neighbouring steps of -9 and -1 %, so both calculations carry structure at that shell and the localisation rests on the sign: a rise in Υn_{up}^3 along a Rydberg series is consistent with a truncated basis, and a smooth continuation is not. At 10 eV there is no kink. The RMPS steps are three to seven points flatter than the CCC ones at $3 \rightarrow 4$ and $4 \rightarrow 5$ alike, the step into $n = 5$ is where the two calculations differ *least*, and the same ratio climbs steadily, 1.06, 1.19, 1.29, 1.34. The kink is a 1 eV feature.

The sign at 1 eV matches the bulk sign of the disagreement in Table 2.3: the RMPS $n = 5$ rates sit high, so CCC reads low against them. The temperature dependence constrains the mechanism. A top shell that absorbed flux belonging to $n \geq 6$ and the continuum would do so at every collision energy, and the excess would appear at every temperature; it appears only where the Maxwell average samples the cross section within about an electron-volt of the $n = 5$ threshold, which is where the $n \geq 6$ thresholds, the ionisation limit and the pseudo-state thresholds of a truncated expansion crowd together. If the kink is a basis artefact, the evidence places it near threshold rather than at high energy; that it is one has not been shown.

This localises the anomaly; it does not prove CCC correct. The scaling argument is asymptotic, and 1 eV is not a temperature at which the Bethe term cleanly dominates, so the test is suggestive rather than decisive. What it establishes is narrower and still useful: the departure from hydrogenic behaviour occurs at one specific shell of one specific calculation, at the temperatures where the two disagree most, rather than being distributed through the CCC dataset. **This work therefore treats the $n = 5$ disagreement as a limitation of the benchmark rather than as a measured error in the data being benchmarked**, and claims no more than that.

2.8.2 Does any of it reach the result?

The preceding argument is about rate coefficients. A discrepancy in a rate coefficient matters only insofar as it moves something the thesis reports, and the way to find out is to make the substitution and look. The CR matrix is linear in electron density, verified rather than assumed, to a relative 3×10^{-16} , so it separates exactly into a radiative part and a collisional part,

$$\mathbf{L}(T_e, n_e) = \mathbf{L}_{\text{rad}}(T_e) + n_e \mathbf{L}_{\text{coll}}(T_e), \quad (2.13)$$

so the collisional rate coefficients can be replaced without disturbing anything else. For all 85 transitions the RMPS calculation covers, the CCC rate coefficient in $\mathbf{L}_{\text{coll}}$ was replaced by the RMPS one and the eigenvalue problem of Chapter 3 re-solved at every one of the 400 grid points. Excitation and de-excitation were scaled by the same factor, so detailed balance survives the substitution, and the diagonal was compensated so that column sums, and with them the loss to the continuum, are unchanged. Before any substitution the procedure reproduces the unperturbed spectrum, including all three ranges of Eq. (3.18), which is the check that makes the perturbed numbers worth reading.

Replacing every rate coefficient the benchmark covers changes τ_{slow} by 11.4 % on average across the grid and by at most 18.7 %; τ_{relax} moves by 2.4 %; and M , which




Figure 2.4: **What the atomic-data disagreement costs, and which coefficient each dataset moves.** (a) Shift in the timescale separation M when the RMPS rate coefficients replace the CCC ones at every transition the RMPS calculation covers, over all 400 grid operators. The shift grows with density and saturates above $n_e \approx 10^{13}/\text{cm}^3$, where the collisional part of $\mathbf{L}$ has taken over from the radiative part, and it is largest at the cold end even though Section 2.8 finds the rates best constrained there: at 1 eV the excitation rates sit deep in the Boltzmann exponential, so a given fractional change in a rate produces a larger one in the eigenvalue. (b) Mean absolute change in each structural coefficient under the two substitutions, over the 448 defended pairs ($T_e \geq 2$ eV, window-resolvable, one grid interval). The two datasets act on different objects. Replacing the excitation data moves Δ and with it the cap $\tanh(|\Delta|/4)$, and $\bar{S}$ follows the cap. Replacing the ionisation data barely moves the cap – Δ is the ratio of ratios $(a_3/a_4)/(c_3/c_4)$, in which a change similar at $n = 3$ and $n = 4$ largely cancels – and instead moves the operating point beneath it, pair by pair and in both directions, which is what carries $\bar{S}$ and $\epsilon_{\text{plateau}}$. The reservoir gain G is robust to both, at under 2%, which is a second reason to tabulate it rather than a percentage error. From `ccc_anderson_grid_impact.csv` and `recombination_substitution.csv`.

inherits both, moves by 11.1% on average. At the benchmark point M falls from 9982 to 8855. This is the test that Section 2.8.1 could not perform.

Figure 2.4(a) maps that shift over the grid. It grows with density and saturates above $n_e \approx 10^{13}/\text{cm}^3$, and it is largest at the cold end, the opposite of where the rate coefficients are best known, because at 1 eV the excitation rates sit deep in the Boltzmann exponential and a given fractional change in a rate produces a larger one in the eigenvalue it feeds.

Its outcome is not the expected one, and the useful part is the disagreement with what was expected. Decomposing the substitution by which transitions are replaced:

The $n = 5$ transitions barely move M and the ground-state ones move it most (Table 2.5), for a structural rather than a numerical reason: τ_{slow} is the timescale of the slowest mode, and Section 3.3 identifies that mode as the depletion of the ground state, so the rates out of 1s set it almost by themselves. Rates into $n = 5$ enter only through cascades that have already been diluted by everything they pass through. The 0.9% is also a net: the five $1s \rightarrow 5\ell$ transitions alone move M by -3.3% and the forty-five excited-fed transitions into $n = 5$ by 2.4% , so part of the reason $n = 5$ does not reach the result is that its two halves pull in opposite directions (`validation/anderson_subset_impact/`).

Table 2.5: Change in the timescale separation M when the CCC rate coefficients are replaced by the RMPS values, for subsets of the 85 benchmarked transitions. The 50 transitions reaching $n = 5$, which carry almost all of the benchmark disagreement, move M by under 1 %; the 14 ground-state transitions, which carry much less of it, move M by more than the full substitution does. Source: `validation/anderson_subset_impact/`, which repeats the substitution of `ccc_anderson_grid_impact.py` for each subset.

Transitions replaced	count	mean $ \Delta\tau_{\text{slow}} $	ΔM at benchmark
all benchmarked ($n_{\text{upper}} \leq 5$)	85	11.4 %	−11.3 %
$n_{\text{upper}} = 5$ only	50	4.7 %	−0.9 %
$n_{\text{upper}} \leq 4$ only	35	7.2 %	−10.8 %
ground-state $1s \rightarrow n\ell$ only	14	10.4 %	−12.4 %

The same conclusion does not follow from the general ground that $n = 5$ channels are weak; they are weak *in the specific eigenvalue the thesis reports*, which had to be computed rather than asserted. The same calculation shows that the transitions which do matter are the ones the benchmark constrains least well among $n_{\text{upper}} \leq 4$: ground-state excitation agrees within 20 % in only 61.8 % of comparisons at $n_{\text{upper}} \leq 4$, against 81.4 % for that block as a whole. Chapter 3 carries the consequence, which is that M is reported with the 11 % by which the R-matrix substitution moves it: the sensitivity to one dataset substitution, not an uncertainty.

Three further soft spots are worth naming rather than leaving to be discovered. *Detailed balance in the raw cross sections*: separately from the enforced consistency above, the CCC cross sections were checked for microscopic reversibility as supplied (`src/parsers/qc_ccc.py`) and satisfy it at the 0.05 % level. That is a genuine physics check on the provider’s calculation, and it passes comfortably. *Lotz ionisation for the top six shells, and its echo in recombination*: the metadata records that Lotz “overestimates CCC by factor $\sim 4\text{--}8$ ” for $n = 10\text{--}15$, and because Eq. (2.6) builds $\alpha_{3\text{BR}}$ from Q , that overestimate propagates straight into the three-body coefficients of the same shells, the shells whose feed is essentially all three-body. Chapter 6 reports, from an independent comparison against published population coefficients, that the terminal shell is overpopulated by 4.9–6.2 times while $n = 10$ agrees to 4 to 20 %. The two observations agree in size and location, but the test that would establish the link has not been done, so the connection is a hypothesis, not a finding. What is measured is what the block does to the reservoir, which is what a reader of Chapter 5 needs. Scaling the six Lotz coefficients and the three-body recombination built from them together by 1/2, 1/4 and 1/8, the far end of the metadata’s own factor, moves u^{CRE} by at most 0.63 % (at $10^{12}/\text{cm}^3$, and by 0.03 % above $10^{14}/\text{cm}^3$), τ_{slow} by at most 1.7 %, the reservoir displacement $\Delta \ln u$ by 0.03 %, and $\varepsilon_{\text{plateau}}$ by 0.7 % at the median and 3.9 % at worst over the 448 defended pairs, with the census moving from 45 to 43 (`verify_high_n_ionisation_scaling.py`, `validation/high_n_ionisation_scaling/`). Scaling both directions of the exchange preserves the Saha ratio of those shells with the continuum and only slows it, and the block is weakly coupled where three-body recombination is small, which is why the reservoir does not see it; the over-population of the terminal shell is a separate question this scan does not answer. No second dataset exists above $n = 9$, so this is a scan and not a substitution. *Vriens–Smeets for the bundled block*: 201 of the 819 excitation pairs come from a semi-empirical fit of quoted accuracy $\sim 20\%$ rather than a close-coupling

calculation, which is the price of the ladder reaching $n = 15$. And one gap that the benchmark simply cannot close: the RMPS tabulation stops at $n = 5$, so every CCC transition reaching $n \geq 6$, the whole upper half of the ladder, has no independent check in this thesis at all.

2.9 Two errors found by audit

Two implementation errors were found in this pipeline by auditing it against its own published sources. Neither was found by a test failing; both were found by reading the code against the paper it claimed to implement, and by recomputing a quantity unconditionally and comparing with the stored value. They belong in the body of the thesis for two reasons: a reader assessing the numbers is entitled to know what the audit turned up, and the *size* of each error, once bounded, is a physical result in its own right.

2.9.1 The ℓ -mixing Coulomb logarithm

The first implementation of the factor $F(U_m)$ of Eq. (2.11) in `src/rates/compute_lmix.py` set it to 1, justified in a comment as “the low-density limit $U_m \rightarrow 0$, $F \rightarrow 1$ ”.

The justification inverts the limit. Reading Badnell’s Eq. 9 [42] directly from the source, $F(U_m)$ is the Coulomb logarithm of the impact-parameter picture, and $F(U_m \rightarrow 0)$ diverges logarithmically rather than tending to one, as it must, since a vanishing minimum energy transfer means an unbounded cut-off radius. The values of U_m at divertor conditions are indeed small, 10^{-3} to 10^{-1} , which is precisely where F is *large*. At the benchmark temperature the correct factor for $2p \rightarrow 2s$ is $F = 6.95$ at $U_m = 1.05 \times 10^{-3}$; across transitions it runs from about 7 at low n to about 3 at $n = 8$. That setting under-counted every ℓ -mixing rate by 3 to 7.

That the correction is present in the data this thesis uses can be checked rather than taken on trust. The canonical matrix is `cr_matrix/L_grid.npy`, written 21 July 2026, SHA-256 beginning `2d92b58e1693`; the ℓ -mixing coefficient it was built from is $q_{2p \rightarrow 2s} = 3.53722 \times 10^{-4} \text{ cm}^3/\text{s}$ at temperature index 23, while evaluating Eq. (2.10) independently with F set to 1 gives $5.08977 \times 10^{-5} \text{ cm}^3/\text{s}$. The ratio is 6.9497 and the independently computed $F(U_m)$ at the same point is 6.94961, so the stored coefficient equals $q F$ to 8 parts in 10^6 .

What the correction changed is the interesting part. Scaling every intra-shell off-diagonal by the correct F , restoring the exact column sums and recomputing the spectrum moves the excited-state relaxation time by -0.12% and the ionisation-balance time by 0.07% at the benchmark point; across the grid the effect on the relaxation time is bounded by 0.85% , and scaling the ℓ -mixing block by 3 and by 7 moves it by at most 0.151% over the 400 points (`verify_residual_claims.py`, `validation/residual_claims/`). A factor of 3 to 7 in the fastest rates in the matrix moves the slow eigenvalues by less than one percent.

That insensitivity is not luck; it is the timescale separation of Chapter 3. ℓ -mixing runs at picoseconds; the relaxation being measured runs at nanoseconds, and on that timescale the ℓ sublevels are fully equilibrated whether the mixing is three times faster or seven, so the slow eigenvalue cannot see the difference. It licenses nothing elsewhere: absolute ℓ -resolved populations and any ℓ -resolved line ratio depend on the correction directly, which is why it was fixed rather than merely bounded.

2.9.2 The eigenvalue magnitude filter

Two analysis scripts extracted timescales by sorting the eigenvalues of the rate matrix and then discarding every eigenvalue of magnitude below $1/s$, with the literal `eigs = eigs[eigs < -1.0]` and a comment calling the discarded values “near-zero numerical noise”.

The slowest genuine eigenvalue of this system is not noise. It is the ionisation-balance rate, and in the cold, thin corner of the grid its magnitude falls below $1/s$: the corresponding timescale reaches tens of seconds. Wherever that happened the filter deleted the physical slow mode and promoted the whole eigenvalue ladder by one rung, so the reported ionisation-balance time was really the relaxation time of the level above it.

19 of the 400 grid points were affected, all at $T_e \leq 1.389 \text{ eV}$ and $n_e \in [10^{12}, 1.93 \times 10^{13}] \text{ 1/cm}^3$. At every one the corrupted relaxation time equalled the corrupted ionisation-balance time of the unfiltered calculation exactly, the signature of a one-rung shift, and how the mechanism was confirmed rather than guessed. The reported range of the timescale-separation ratio changed from $1.34\text{--}1.01 \times 10^8$ to $86.77\text{--}1.72928 \times 10^9$. The benchmark point was bit-identical before and after, and so were the step-error maps at all 400 points: the corruption was real, confined, and did not touch the central result. Every check the filtered code ran continued to pass, so it was found not by a check failing but by recomputing the spectrum unconditionally, storing it separately, and comparing bit-for-bit. The filtered outputs were kept rather than deleted, under names ending `_FILTERED_20260721`, so the comparison remains reproducible.

The last statement needs to be precise rather than tidy. The filter was first removed from the analysis scripts that write the timescale grids, which select on `eigs < 0.0`. The literal `eigs[eigs < -1.0]` survived for some weeks afterwards in two other files, `src/rates/solve_cr.py` line 269 and `src/rates/check_mz.py` line 10, neither of which produces any number quoted in this thesis: the first is a superseded solver and the second a diagnostic print. Both were corrected on 16 September 2026, once it had been established that nothing imported either, and the literal no longer appears anywhere in `src/`. The gap between the two repairs is recorded rather than smoothed over, because during it a reader who searched the repository for the literal would have found it and been entitled to ask which numbers passed through it.

2.9.3 What both errors have in common

Both were invisible to every consistency test the pipeline ran. The ℓ -mixing error preserved column sums exactly, because ℓ -mixing is a conservative redistribution and its magnitude cancels in the sum. The eigenvalue filter produced eigenvalues that were all negative, all finite and correctly ordered. A structurally consistent, physically wrong result is the characteristic failure mode of this kind of model, and it is why Chapter 4 separates checks a wrong result would also have passed from checks it would not. The structural response is the loading module of Section 2.2, which raises on any inconsistency between grids, state ordering and matrices.

2.10 What this model assumes

Every item here is a deliberate choice, gathered so that they can be found in one place and quantified where the thesis is able to quantify them.

Maxwellian electrons. Every rate coefficient is an average over a Maxwell–Boltzmann distribution, and Eq. (2.4), on which every de-excitation and every three-body coefficient rests, holds exactly for a Maxwellian and not otherwise. A suprathermal tail, which a real transient can produce, would break it.

$T_i = T_e$ **and** $n_{\text{ion}} = n_e$. The proton temperature in the ℓ -mixing rate is set to the electron temperature; since $\nu_{\ell\text{-mix}} \propto T_i^{-1/2}$, a factor of two error in T_i is 40% in the rate, inside the insensitivity of Section 2.9, but stated rather than assumed away. Quasi-neutrality gives $n_{\text{ion}} = n_e$; no impurities are modelled.

A fixed reference density in the Debye cut-off. U_m in Eq. (2.11) is evaluated at $10^{14}/\text{cm}^3$ rather than at each grid density, freezing a factor that varies across the density columns at 1 eV by 3.7 for $n = 2$ and by 50 to 687 across the ℓ channels of $n = 8$ (Section 2.7.2); the cost to $f_3 - f_4$ is at most 0.42%.

Optically thin. All Einstein coefficients enter bare, $\Theta_P \equiv 1$. Chapter 3 evaluates Eq. (2.8) across the grid and finds the assumption excellent over the hot half and failing along the $T_e = 1$ eV edge; Chapter 6 quantifies the consequence.

$n_{\text{max}} = 15$, **with** $n \geq 9$ **bundled.** Seven shells are one state each, with no ℓ -mixing inside the bundled block: they are *assumed* in statistical ℓ -equilibrium rather than driven to it, and that has not been tested directly. Chapter 6 reports the measured excess in the terminal shell.

Uniform plasma, no transport, no molecules. A single well-mixed volume of atomic hydrogen only, in reactor terms a perfectly stirred tank rather than a plug-flow reactor. At the cold end of the grid the gas is largely neutral and in reality substantially molecular, and molecular-assisted recombination would feed $n = 3$ directly. That is a physics limit rather than a data limit, and it is why the grid stops at 1 eV.

An open system. The ion H^+ is a fixed external reservoir: there is no 44th state, the model does not conserve nuclei, and ionisation is a one-way leak out of the tracked manifold. This is why the ionisation-balance timescale is boundary-dependent, and why the effective coefficients this model produces cannot be compared directly with those tabulated by ADAS. No ADAS dataset enters the rate matrix at any point; ADAS appears in this thesis only as an external comparison.

2.11 Summary

This chapter has built the input side of a collisional–radiative model of atomic hydrogen: 43 levels, six elementary processes, and a rate coefficient for every pair of levels at every point of a 50×8 grid in temperature and density. Table 2.6 is the provenance record for all of it.

One provenance item on that table was resolved by test rather than inspection. The merged excitation arrays appeared to predate a CCC file they are assembled from; re-running `assemble_K_exc.py` into a scratch directory returns all five arrays identical by SHA-256. The file whose date raised the question, `K_CCC_metadata.csv`, carries no rate data, and the rate tables themselves predate the merge. The stored arrays are current.

Table 2.6: Every atomic dataset entering the model, with the file it is stored in and the script that produced it. Paths are relative to `data/processed/`; the arrays are those named in the `PATHS` dictionary of `src/rates/assemble_cr_matrix.py`. All rate-coefficient arrays are tabulated on the 50-point temperature grid $T_e = 1\text{--}10\text{ eV}$. Two hydrogen energies appear in the pipeline and they are two *different physical quantities*, not two estimates of one; Section 2.3.2 says which is used where.

Quantity	Source	Stored file	Producing script
Excitation and de-excitation, $n \leq 10$	CCC (Bray, priv. comm. 2026); $\Delta n = 0$ excluded	<code>collisions/K_exc_full/</code> <code>compute_K_full.py,</code> <code>K_deexc_full.npy</code>	<code>assemble_K_exc.py</code>
Excitation, bundled shells	Vriens–Smeets Eqs. (17)–(19)	merged into the same arrays	<code>compute_K_VS.py</code>
Ionisation, $n \leq 9$	CCC total ionisation cross sections	<code>collisions/tics/K_ionization.npy</code>	<code>compute_K_TICS.py</code>
Ionisation, $n = 10\text{--}15$	Lotz semi-empirical formula	same array	<code>ionization_rates.py</code>
Radiative recombination	Johnson three-term Gaunt expansion; ℓ -split statistically	<code>recombination/alpha_Recombined.npy</code>	<code>recombination_rates.py</code>
Three-body recombination	detailed balance from ionisation, via Saha	<code>recombination/alpha_Recombined.npy</code>	<code>recombination_rates.py</code>
Einstein A coefficients	Hoang-Binh, ADUU v1.0	<code>Radiative/A_resolved.npy,</code> <code>A_bund_res.npy,</code> <code>A_bund_bund.npy,</code> <code>gamma_*.npy</code>	<code>radiative_rates.py</code>
Proton-impact ℓ -mixing	PSM20 Debye-cutoff formula, Badnell <i>et al.</i>	<code>lmix/K_lmix.npy</code>	<code>compute_lmix.py</code>
State ordering and grids	the pipeline’s own output	<code>collisions/K_exc_full/</code> <code>state_ordering_matrix.py,</code> <code>cr_matrix/Te_grid_L.npy</code> <code>ne_grid_L.npy</code>	<code>state_ordering_matrix.py,</code> via <code>cr_context.py</code>

Four things carry forward. The **state space is not arbitrary**: ℓ is resolved below $n = 9$ because $2s$ and $2p$ behave completely differently, and bundled above it because ℓ -mixing makes the sublevels of a high shell indistinguishable, which was checked, not asserted. **Proton-impact ℓ -mixing is structural**: it carries at least 99.92% of the loss from $2s$ everywhere on the grid. **The data are good where it matters**: the $n_{\text{upper}} \leq 4$ transitions that drive the kinetics agree with an independent RMPS calculation at a mean absolute error of 13.0%, and the factor 2–5 disagreement at $n = 5$ is localised to the top shell of the RMPS target basis (Section 2.8.1). Substituting the RMPS rates wholesale moves the timescale separation of Chapter 3 by 11%, carried by ground-state excitation and not by $n = 5$. And **two implementation errors were found by audit and both were bounded**.

What none of this has done is put the atoms in motion. Chapter 3 assembles the 43×43 rate coefficients and the feed term into a single matrix equation and asks whether its fast clock is fast enough for the standard shortcut to be safe.

Chapter 3

Two Clocks

Chapter 2 established what can happen to a single hydrogen atom in a plasma and how fast each process runs. This chapter asks the next question: if we know what happens to one atom, how do we follow 10^{12} or more of them per cubic centimetre, distributed over 43 energy levels, and what makes the answer hard?

The answer to the second half is that the system has two clocks running at once, separated by four orders of magnitude. Everything in this thesis follows from that separation, including a result that the separation itself does not predict, as Chapter 5 shows.

What this chapter takes as given. Three things are established elsewhere and used here without re-derivation. From Chapter 1: that divertor conditions are inferred spectroscopically, that the Balmer H_α/H_β ratio is the working diagnostic, and the operating range $T_e = 1\text{--}10\text{ eV}$, $n_e = 10^{12}\text{--}10^{15}/\text{cm}^3$ over which everything is computed. From Chapter 2: the 43-level state space and the reason ℓ is resolved for $n \leq 8$ and bundled above it; the origin and accuracy of every rate coefficient; and the Maxwellian averaging that turns a cross section into a rate. This chapter assumes those rates exist and are correct, and asks only what the resulting equations do.

3.1 The rate equation

3.1.1 A balance on every level

Consider one energy level of the hydrogen atom, say $2p$, and count atoms. Atoms arrive in $2p$ when an electron collision excites a ground-state atom into it, when an atom in a higher level decays into it radiatively, when a collision knocks an atom down from above, and when a proton captures an electron directly into it. Atoms leave $2p$ when they are excited upward, knocked downward, ionised, or radiate to $1s$.

The rate of change of the $2p$ population is the sum of everything arriving minus the sum of everything leaving. That is all a rate equation is.

A rate equation of this kind is a species balance on a well-mixed vessel, the form written for any reaction network, with one species per energy level and one reaction per atomic process. Figure 3.1 shows the network and the partition that Section 3.4 will exploit. The only unusual feature is the number of species, 43, and the fact that some of the “reactions” are radiative and therefore proceed without a collision partner.




Figure 3.1: The collisional–radiative network of atomic hydrogen and the partition used in Section 3.4. Levels with $n \leq 8$ are resolved in orbital angular momentum; $n = 9–15$ are bundled with a statistical ℓ -distribution. Proton-impact ℓ -mixing (orange) couples sublevels within every resolved shell and is the dominant loss channel for the metastable $2s$ level, which has no electric-dipole route to the ground state; arrows are drawn for $n = 2$ and $n = 3$ only, the $n = 4$ sublevels being too closely spaced on this schematic for a legible arrow. The broad arrows are not elementary transitions: $\mathbf{a} = -\mathbf{L}_{FF}^{-1}\mathbf{L}_{FS}$ and $\mathbf{c} = -\mathbf{L}_{FF}^{-1}\mathbf{S}_F$ are *network responses*, each summing every cascade route from its reservoir into a given level (Section 3.5.3). Vertical spacing is schematic and not drawn to energy scale.

Write it out for a general level p :

$$\frac{dn_p}{dt} = \underbrace{\sum_{q \neq p} C_{qp} n_e n_q}_{\text{collisions in}} - \underbrace{\sum_{q \neq p} C_{pq} n_e n_p}_{\text{collisions out}} + \underbrace{\sum_{q > p} A_{qp} n_q}_{\text{radiation in}} - \underbrace{\gamma_p n_p}_{\text{radiation out}} - \underbrace{Q_p n_e n_p}_{\text{ionisation}} + \underbrace{S_p n_{\text{ion}}}_{\text{recombination}} \quad (3.1)$$

where n_p is the population of level p in $1/\text{cm}^3$; C_{qp} is the rate coefficient for the collisional transition $q \rightarrow p$ in cm^3/s ; A_{qp} is the Einstein coefficient for spontaneous emission $q \rightarrow p$ in $1/\text{s}$; $\gamma_p = \sum_{q < p} A_{pq}$ is the total radiative loss rate from p ; Q_p is the ionisation rate coefficient out of p in cm^3/s ; and S_p is the recombination source per unit ion density in $1/\text{s}$, given in Eq. (3.5) below.

Collisional terms carry a density factor because a collision requires a partner; radiative terms do not, because an excited atom decays whether or not anything else is nearby. Which partner matters: electron-impact terms scale with n_e , proton-impact terms with n_{ion} . For the pure hydrogen plasma considered here quasi-neutrality gives $n_{\text{ion}} = n_e$, so both scale together and the distinction is one of physics rather than of parametric dependence. This is the first hint of the two clocks: radiative rates are fixed by atomic structure, collisional rates scale with density, and the competition between them is what makes the answer depend on position in the plasma.

3.1.2 Matrix form

Equation (3.1) holds for each of the 43 levels. Stacking them and collecting the terms that are linear in the tracked populations gives

$$\boxed{\frac{d\mathbf{n}}{dt} = \mathbf{L} \mathbf{n} + \mathbf{S} n_{\text{ion}}} \quad (3.2)$$

with $\mathbf{n}$ the 43-component population vector ($1/\text{cm}^3$), $\mathbf{L}$ the collisional–radiative rate matrix ($1/\text{s}$), and $\mathbf{S} n_{\text{ion}}$ the recombination feed ($1/(\text{cm}^3 \text{s})$). The matrix elements are

$$L_{pq} = C_{qp}^{(e)} n_e + C_{qp}^{(i)} n_{\text{ion}} + A_{qp} \quad (p \neq q), \quad (3.3)$$

$$L_{pp} = - \left(\sum_{q \neq p} [C_{pq}^{(e)} n_e + C_{pq}^{(i)} n_{\text{ion}}] + \gamma_p + Q_p n_e \right). \quad (3.4)$$

The superscripts distinguish the two collision partners. $C^{(e)}$ denotes electron-impact transitions and is the dominant term for changes of principal quantum number. $C^{(i)}$ denotes proton-impact transitions, which in this model are non-zero only for ℓ -changing collisions within a shell ($\Delta n = 0$, $\Delta \ell = \pm 1$, $n \leq 8$). That term is easy to overlook and it is not optional. Chapter 2 shows it carries 99.92% of the total loss rate from the metastable $2s$ level. The reason is that $2s$ has no other fast exit: the transition $2s \rightarrow 1s$ is forbidden to electric-dipole radiation, and no two-photon rate is included in the present radiative dataset, so the only route out other than proton ℓ -mixing is electron-impact de-excitation, some three orders of magnitude slower. Removing proton ℓ -mixing therefore does not merely perturb the $n = 2$ manifold; it changes which process controls it.

Note the index order: L_{pq} with $p \neq q$ is the rate of the transition $q \rightarrow p$, a *gain* into p . Off-diagonal elements are therefore non-negative and diagonal elements are negative. This convention is the opposite of the one some texts use, and getting it backwards transposes the physics; it is stated here once and used throughout.

3.1.3 Why recombination is a feed stream, not a matrix element

There is one term in Eq. (3.1) that refuses to become a matrix element, and understanding why is worth a paragraph.

$\mathbf{L}$ multiplies $\mathbf{n}$. Every entry of $\mathbf{L}\mathbf{n}$ must therefore be proportional to the population of some *tracked* level. Excitation, de-excitation, radiative decay and ionisation all satisfy this: they take an atom that is in the manifold and move it, or remove it. Recombination does not. It is proportional to $n_e n_{\text{ion}}$, the density of a species outside the neutral manifold entirely, and so it enters as a constant feed, independent of $\mathbf{n}$.

The chemical-engineering analogy is exact and worth holding onto for the rest of the thesis. $\mathbf{L}$ is the reaction-rate matrix of a network of 43 species in a well-mixed vessel. $\mathbf{S}n_{\text{ion}}$ is the inlet stream. And the ion population is the reservoir feeding that inlet, a quantity this model does not evolve, which becomes important in Section 3.5 and again in Chapter 6.

3.1.4 Units, and a source that is quadratic in density

The source deserves care because its units are easy to state wrongly. Writing it out,

$$\mathbf{S} = n_e (\boldsymbol{\alpha}_{\text{RR}} + n_e \boldsymbol{\alpha}_{\text{3BR}}) \quad (3.5)$$

where $\boldsymbol{\alpha}_{\text{RR}}$ is the radiative recombination coefficient (cm^3/s) and $\boldsymbol{\alpha}_{\text{3BR}}$ the three-body coefficient (cm^6/s). Three-body recombination requires a second electron to carry away the excess energy, so it carries an extra factor of n_e ; the source as a whole is therefore *quadratic* in electron density, and $\mathbf{S}$ has units of $1/\text{s}$, not cm^3/s .

This matters physically, not just dimensionally. At $n_e = 10^{12}/\text{cm}^3$ the three-body channel already supplies more than 10% of the feed into 30 of the 36 resolved levels, and is negligible only for $n \leq 3$; at $n_e = 10^{15}/\text{cm}^3$ it supplies more than 99% into the high- n shells. The same equation therefore describes two physically different regimes at the two ends of the density grid. It would be convenient if that allowed the radiative and three-body datasets to be tested separately, one at each end. Section 4.11.5 shows the measurement does not support this: the three-body channel is already dominant at the bottom of the density grid, so there is no radiative-only end to test against.

Dimensional check on Eq. (3.2): $[\mathbf{L}\mathbf{n}] = 1/\text{s} \cdot 1/\text{cm}^3 = 1/(\text{cm}^3 \text{s})$, and $[\mathbf{S}n_{\text{ion}}] = 1/\text{s} \cdot 1/\text{cm}^3$, matching $[d\mathbf{n}/dt]$. Every term is a number of atoms per unit volume per unit time, as a balance equation must be.

3.2 Conservation, and how we know the matrix is right

A 43×43 matrix assembled from six separate atomic datasets offers a great many opportunities to make a mistake that produces plausible numbers. Before using $\mathbf{L}$ for anything, we need a test that a wrong matrix would fail.

The test comes from counting atoms. Sum Eq. (3.2) over all levels with the feed switched off. Every collisional and radiative transition moves an atom *within* the manifold: it is a loss from one level and a gain to another, so the two contributions cancel in the sum. Only ionisation genuinely removes an atom from the bound states. Therefore

$$\boxed{\sum_p L_{pq} = -Q_q n_e \quad \text{for every column } q.} \quad (3.6)$$

Each column of $\mathbf{L}$ must sum to the ionisation leak from the level that column represents, no more, no less.

The check is worth having, and what it can detect is narrower than it looks: each block of the assembly writes its diagonal as minus the column sum of the off-diagonal entries it has just written, so the identity holds by construction whatever those entries contain, and of the four faults Section 4.2 injects, only an inconsistency between the two separately stored radiative arrays moves the residual. Chapter 4 grades it on that basis.

3.2.1 A worked example

The rule is easiest to see in a two-level system. Take $1s$ and $2p$ only, at $T_e = 3\text{ eV}$ and $n_e = 10^{14}/\text{cm}^3$, with representative coefficients $C_{12} = 1 \times 10^{-11} \text{ cm}^3/\text{s}$, $C_{21} = 1 \times 10^{-10} \text{ cm}^3/\text{s}$, $A_{21} = 6.27 \times 10^8/\text{s}$, $Q_1 = 1 \times 10^{-13} \text{ cm}^3/\text{s}$ and $Q_2 = 1 \times 10^{-10} \text{ cm}^3/\text{s}$. Then

$$\mathbf{L} = \begin{pmatrix} -1010 & 6.2701 \times 10^8 \\ 1000 & -6.2702 \times 10^8 \end{pmatrix} 1/\text{s}, \quad \sum_p L_{pq} = \begin{pmatrix} -10 \\ -10^4 \end{pmatrix} 1/\text{s} \quad (3.7)$$

and $Q_1 n_e = 10/\text{s}$, $Q_2 n_e = 10^4/\text{s}$. The rule holds.

Two features of this small matrix survive into the full one and are worth noticing now.

First, the second column is a difference of two numbers that agree to four significant figures: 6.2701×10^8 against 6.2702×10^8 , leaving 10^4 . Almost all of the $2p$ loss rate is radiative decay to the ground state, and the physically interesting part, the ionisation leak, is four orders of magnitude smaller. Conservation is therefore always quoted as a *relative* residual; an absolute one would be meaningless.

Second, that same near-cancellation is the seed of the two clocks. The $2p$ level empties in nanoseconds by radiating, but the atom that arrives back in $1s$ is still in the manifold. Only the small ionisation term actually removes it. Section 3.3 makes this quantitative.

3.2.2 The measured residual

For the full matrix, evaluated at all 400 grid points, the largest relative column-sum residual is

$$\max_{q, T_e, n_e} \left| \frac{\sum_p L_{pq} + Q_q n_e}{Q_q n_e} \right| = 3.61 \times 10^{-11}, \quad (3.8)$$

which is round-off for a matrix whose non-zero entries span about twelve decades. The assembly is arithmetically sound.

This establishes that every transition entered as a loss also entered as a gain somewhere. It says nothing whatever about whether the *magnitude* of any rate is correct: ℓ -mixing, for instance, cancels exactly in the column sum and so is invisible to this test. Checking magnitudes requires comparison against independent sources, which is the subject of Chapter 4. A test that a wrong result would also have passed is not a test.

3.3 Two clocks

3.3.1 Why the system has natural timescales at all

Equation (3.2) is linear, so with $\mathbf{L}$ and the feed held fixed its solution is a steady state plus a sum of decaying exponentials:

$$\mathbf{n}(t) = \mathbf{n}^{\text{ss}} + \sum_k c_k e^{\lambda_k t} \mathbf{v}_k, \quad \mathbf{L}\mathbf{v}_k = \lambda_k \mathbf{v}_k, \quad \mathbf{n}^{\text{ss}} = -\mathbf{L}^{-1} \mathbf{S} n_{\text{ion}}. \quad (3.9)$$

Each eigenvalue λ_k has a negative real part; the system is stable, as a system of decaying populations must be. Each eigenvalue contributes a mode decaying with time constant $\tau_k = 1/|\lambda_k|$. The set of τ_k is the set of clocks the system runs on.

The whole chapter turns on what Eq. (3.9) has achieved. In the populations n_p the 43 balance equations are hopelessly coupled: nothing happens to $3d$ without also happening to $2p$, $4f$, and everything else. Diagonalising $\mathbf{L}$ is a change of coordinates. Instead of asking how many atoms are in each *level*, ask how much of each *eigenvector* the state contains. In those coordinates the coupling disappears entirely. The k th coordinate obeys $\dot{c}_k = \lambda_k c_k$, which is the first-order decay $c_k(t) = c_k(0)e^{-t/\tau_k}$ familiar from single-step kinetics. Forty-three coupled equations become forty-three independent ones, each with its own clock.

The eigenvectors are therefore not abstract: each is a *pattern of populations* that decays without changing shape, and any state of the system is a superposition of those patterns fading at their own rates. Asking which physical process a mode corresponds to is then a matter of reading which levels carry its weight, which is what Section 3.3 does next. One caveat, developed in Section 3.6: this decomposition is exact, but the patterns are not orthogonal, so a small state can be the sum of two large and nearly cancelling ones.

For hydrogen at divertor conditions there are 43 of them. The striking fact, and the one this thesis is built on, is that they are not spread out evenly.

3.3.2 The spectrum

Figure 3.2 shows all 43 relaxation times, and Table 3.1 lists the five slowest at the benchmark point $T_e = 2.947$ eV, $n_e = 1.389 \times 10^{14}/\text{cm}^3$.

Table 3.1: The five slowest relaxation modes at the benchmark point. The gap between the first and second is 9982; the largest of the remaining 41 gaps is the 2.80 in this table, at $k = 2$, so the slow mode is isolated from the fast band by a factor of 3567.

k	λ_k (1/s)	τ_k (s)	τ_k/τ_{k+1}
0	-4.3999×10^4	2.2728×10^{-5}	9982
1	-4.3919×10^8	2.2769×10^{-9}	1.23
2	-5.3994×10^8	1.8521×10^{-9}	2.80
3	-1.5110×10^9	6.6181×10^{-10}	2.71
4	-4.0922×10^9	2.4436×10^{-10}	—

There is one gap and then a continuum.

The slowest mode is separated from the next by a factor of ten thousand; after that the modes crowd together, each within a factor of about two or three of its neighbour. It




Figure 3.2: Relaxation timescales at the benchmark point. (a) All 43 modes. The slowest, τ_{slow} , is the relaxation of the ground-state population; the next, τ_{relax} , is the slowest-decaying mode of the excited manifold, the bottleneck of the fast set, and their ratio is the separation M . (b) The ratio of consecutive timescales, which shows that the first gap is exceptional: the largest of the other 41 is 2.80, so the first exceeds every competitor by a factor of 3567 and the 42 fast modes form a dense band rather than a hierarchy.

is not a hierarchy of three or four well-separated groups, as one might expect from the n -shell structure. It is one very slow mode sitting far above a dense quasi-continuum of fast ones.

3.3.3 What the two modes are

The eigenvectors identify them.

The slow mode $\mathbf{v}_0$ is almost entirely ground state: its $1s$ component is 1.0000 and every excited component is below 10^{-3} . Its participation ratio,

$$\text{PR}(\mathbf{v}) = \frac{(\sum_i |v_i|^2)^2}{\sum_i |v_i|^4}, \quad (3.10)$$

which counts the effective number of states a mode involves, unity for a mode confined to a single level and 43 for one spread evenly over all of them, is 1.00. This mode is the relaxation of the ground-state population against the prescribed ion reservoir. It is often called the ionisation balance, and that is the right physical picture, but the term should be used with care here: this model holds n_{ion} fixed (Section 3.1), so what relaxes is n_{1s} alone and not a self-consistent neutral-ion partition. Chapter 6 assesses what that costs.

Its eigenvalue is closely related to, but not equal to, the direct ionisation rate of the ground state. Writing

$$|\lambda_0| = \eta Q_{1s} n_e, \quad (3.11)$$

the factor η ranges from 1.16 to 29.5 across the grid, and equals 3.12 at the benchmark point. It is nowhere below unity on the grid, and it is the *stepwise-ionisation enhancement*: ground-state atoms are not only ionised directly but also excited into the manifold and

ionised from there. Equation (3.21) shows exactly this: η is the ratio of the effective ionisation rate to the direct one,

$$\eta = \frac{|\mathbf{\Omega}_{\text{QSS}}|}{Q_{1s} n_e}, \quad (3.12)$$

the numerator being the full Schur complement and the denominator the direct ground-state ionisation rate alone. The denominator is emphatically *not* $|\mathbf{L}_{SS}|$: by Eq. (3.1) the ground state has no radiative loss and no ℓ -mixing partner, so $|\mathbf{L}_{SS}| = (\sum_q C_{1s \rightarrow q} + Q_{1s}) n_e$, which exceeds $Q_{1s} n_e$ by the whole of the excitation flux. Since the Schur complement returns part of that flux to the ground state, $|\mathbf{\Omega}_{\text{QSS}}| < |\mathbf{L}_{SS}|$ always, and the ratio $\mathbf{\Omega}_{\text{QSS}}/\mathbf{L}_{SS}$ is bounded above by unity, whereas η is bounded below by it. Note also that η is defined through $|\lambda_0|$ while Eq. (3.12) is written with $|\mathbf{\Omega}_{\text{QSS}}|$; the two agree only to $O(M^{-1})$, which at the benchmark is one part in 10^4 and at the weakest grid point one part in 87. Consistently, η grows with density, where collisional climbing of the ladder is easier, and falls with temperature, where direct ionisation is already fast: at $T_e = 10$ eV, $n_e = 10^{12}/\text{cm}^3$ it is 1.16, and at $T_e = 1$ eV, $n_e = 10^{15}/\text{cm}^3$ it is 29.5; almost all ionisation there proceeds stepwise. We write

$$\tau_{\text{slow}} \equiv \frac{1}{|\lambda_0|} = 22.73 \mu\text{s}. \quad (3.13)$$

A note on the name. Calling Eq. (3.13) τ_{QSS} would attach the label of the *fast*-manifold elimination of Section 3.4 to the *slow* reservoir at the opposite end of the spectrum; the symbol is therefore τ_{slow} throughout. The alternative τ_{res} was rejected because Sawada and Fujimoto use it for the *response* time of the excited levels [19], this thesis's τ_{relax} , and Chapter 7 compares the two works directly.

The next mode $\mathbf{v}_1$ is the excited states rearranging among themselves. Its participation ratio is 2.64, so it lives on two or three levels rather than across the manifold. Its ground-state weight depends on the measure and both are worth stating: in the raw eigenvector the 1s component is the largest entry, at 1.0000, while scaled by the equilibrium populations it is 1.9×10^{-4} . The second measure is the one that matters for the populations, and it is the reason this mode is not a ground-state mode:

$$\tau_{\text{relax}} \equiv \frac{1}{|\lambda_1|} = 2.277 \text{ ns}. \quad (3.14)$$

Their ratio is the *separation*

$$M \equiv \frac{\tau_{\text{slow}}}{\tau_{\text{relax}}} = 9982 \quad \text{at the benchmark point.} \quad (3.15)$$

The excited manifold settles ten thousand times faster than the ionisation balance moves. Put in units a reader can hold: if the ionisation balance took a year to equilibrate, the excited states would be done in under an hour.

Which operator M belongs to

One point of notation has to be fixed here, because it is the source of an apparent contradiction between this chapter and Chapter 5.

M is a property of an *operator*, not of an experiment. Nothing in $|\lambda_1|/|\lambda_0|$ refers to a disturbance, an initial condition, or a history. Written with its argument exposed,

$$M[\mathbf{L}] \equiv \frac{\tau_{\text{slow}}[\mathbf{L}]}{\tau_{\text{relax}}[\mathbf{L}]} = \frac{|\lambda_1[\mathbf{L}]|}{|\lambda_0[\mathbf{L}]|}, \quad M(T_e, n_e) \equiv M[\mathbf{L}(T_e, n_e)], \quad (3.16)$$

and Eq. (3.15) is $M(2.947 \text{ eV}, 1.389 \times 10^{14} / \text{cm}^3)$ while Eq. (3.18) is the range of $M(T_e, n_e)$ over the 400 grid operators. Both are statements about matrices, and no notion of a step enters either.

A temperature step brings a second operator into scope. Writing $\mathbf{L}^- = \mathbf{L}(T_e^-, n_e)$ and $\mathbf{L}^+ = \mathbf{L}(T_e^- + \Delta T_e, n_e)$, the initial populations are those of the CR equilibrium of $\mathbf{L}^-$, but every relaxation for $t > 0$ is governed by $\mathbf{L}^+$. The two separations

$$M^- \equiv M[\mathbf{L}^-], \quad M^+ \equiv M[\mathbf{L}^+] \quad (3.17)$$

are therefore both meaningful, and it is M^+ that bounds the post-step relaxation. At the benchmark point a step of one grid interval in T_e gives $M^- = 9982$ and $M^+ = 8243$. The difference is not a disagreement between chapters: M falls steeply with temperature, with a measured local slope $d \ln M / d \ln T_e \approx -4$, so a 4.8 % step moves it by 17 %.

The alternative convention, defining M to mean the post-step operator always, was considered and rejected, for a reason worth recording. The T_e grid is logarithmic with a ratio of $10^{1/49} = 1.0481$, so a step of one grid interval from the benchmark lands exactly on the neighbouring grid point. The post-step operator is $\mathbf{L}(T_e[24], n_e[5])$, one of the 400 operators whose separation Eq. (3.18) already reports. This is not special to that step. Because $\mathbf{L}$ is tabulated only on the grid, every step used in this thesis is an integer number of grid intervals: the larger step reported elsewhere as “+0.6 eV” is four intervals, $T_e[23] \rightarrow T_e[27]$, an increment of 0.6094 eV and a factor $1.0481^4 = 1.2068$, for which $M^+ = 4856$. *Every* post-step operator in this work is therefore also a static grid operator. Under that convention the same matrix would carry one name in this chapter and another in Chapter 5. Being pre-step or post-step is a role a matrix plays in an experiment, not a property it has; the notation should not pretend otherwise.

3.3.4 The separation across the operating range

Table 3.1 is one grid point. Across all 400,

$$0.87 \text{ ns} \leq \tau_{\text{relax}} \leq 38.9 \text{ ns}, \quad 75.4 \text{ ns} \leq \tau_{\text{slow}} \leq 67.2 \text{ s}, \quad 86.8 \leq M \leq 1.73 \times 10^9. \quad (3.18)$$

τ_{relax} varies by a factor of 45; τ_{slow} by nine orders of magnitude, from 75.4 ns to 67.2 s. The separation is never small: in the hottest, densest corner of the grid, where it is weakest, the excited states are almost ninety times faster than the ionisation balance.

The scope of that lower figure needs stating, because it is easy to quote in a way it does not support. $M = 86.8$ is the minimum over all 400 grid points, and the point carrying it is excluded from the transient analysis of Chapter 5 by the plateau-window criterion, which requires $M > 900$ for a window to exist at all. The floor over the *analysed* set is therefore $M \geq 902$, and that floor is imposed by the criterion rather than measured. The same distinction applies to τ_{slow} : Eq. (3.18) reports the minimum over all 400 points, 75.4 ns at the same [49, 7] corner, so that its three ranges share one scope, whereas the floor over the 346 points that satisfy $M > 900$ is $\tau_{\text{slow}} \geq 1.18 \text{ } \mu\text{s}$. Neither number licenses a statement of the form “*the separation is at least X wherever the model is used*”: the

first is a minimum over a set that includes points the analysis discards, and the second is a restatement of the cut. Both are reported here as descriptions of the spectrum, and the argument of Section 3.4 does not rest on either.

The one-gap structure of Table 3.1 also survives across the grid, though with far less margin than the benchmark suggests. Define the *isolation* of the slow mode as the first gap divided by the largest of the other 41. At the benchmark it is 3567. Its minimum over all 400 points is 24.3, at $T_e = 10$ eV, $n_e = 10^{15}/\text{cm}^3$, where the first gap is 86.8 and its nearest competitor 3.57. A factor of 24 is still an unambiguous separation, one slow mode and then a band, but it is 147 times less margin than Figure 3.2 displays, and the benchmark value must not be quoted as though it described the grid.

That corner is worth naming, because it is where every part of this argument is weakest at once: the minimum τ_{relax} , the minimum τ_{slow} , the minimum M and the minimum isolation all occur at the same grid point, [49, 7]. It is the single stress point of the two-clock picture. It is also, by the plateau-window criterion, outside the set analysed in Chapter 5, which is a limitation to state rather than a convenience, since it means the weakest separation on the grid is never tested against a transient.

τ_{relax} falls with density with a fitted logarithmic slope $d \ln \tau_{\text{relax}} / d \ln n_e$ between -0.46 and -0.53 across the temperature range, rather than the -1 that a purely collisional relaxation would give. The difference is radiative decay, which is density-independent and increasingly dominates the excited-state loss budget as the density falls.

3.3.5 Is the definition robust?

Two definitions of τ_{relax} are available and can be compared. The first, used above, takes the second eigenvalue of the full 43×43 matrix. The second computes the slowest eigenvalue of the excited block $\mathbf{L}_{FF}$ alone: the relaxation of the excited manifold with the ground state held fixed. These are different calculations and there is no guarantee they coincide.

They agree to 0.0098 % at the benchmark point and to better than 0.346 % everywhere on the grid. The excited manifold relaxes at essentially the same rate whether or not the ground state is allowed to participate.

It would be a mistake to read that agreement as independent evidence that the subsystems separate, and the arithmetic says why. Block decoupling is a perturbation in the coupling between the two blocks, and its leading correction is of relative size M^{-1} . At the benchmark $M^{-1} = 0.0100$ % against the measured 0.0098 %. The two agree to 2 %. The disagreement between the two definitions of τ_{relax} is the first-order coupling correction; it is a measurement of the spectral gap, not a second, independent demonstration of it. Quoting it as corroboration would be the error Section 3.2 warns against: a test that a wrong result would also have passed.

Read correctly it is still worth having, and is a sharper statement than the one it replaces. Across the grid the product $M \times (\text{relative disagreement})$ runs from 0.074 to 7.94 with median 0.494, order unity as first-order perturbation theory requires, with a coefficient that varies by two orders across the operating range. The two definitions differ by exactly the amount the block structure predicts, which confirms the perturbative picture rather than assuming it. The formal argument for separability is made in Section 3.4 and rests on the Schur complement, not on this comparison.

3.3.6 How much of the separation is atomic data?

The previous subsection asked whether M survives a change of definition. It does. A separate question is whether it survives a change of *input*, and the two are not related: both definitions of τ_{relax} read the same matrix, so both would move together if the matrix were wrong. Every number in Eq. (3.18) descends from rate coefficients that Chapter 2 shows are known to about 13 % for the transitions that dominate, and a separation quoted to four figures from data known to two is quoting the arithmetic rather than the physics.

Section 2.8.2 measures the sensitivity directly, by substituting the independent R-matrix rate coefficients for the CCC ones wherever the independent calculation exists and re-solving the spectrum at all 400 grid points. The result is that τ_{slow} shifts by 11.4 % on average and M by 11.1 %, while τ_{relax} shifts by only 2.4 %. At the benchmark point M falls from 9982 to 8855. That displacement must not be confused with the $M^+ = 8243$ of Eq. (3.17), which is the same atomic data evaluated at a different temperature. Here the temperature is held fixed and the data are changed instead, so the two numbers measure unrelated things and neither constrains the other.

The asymmetry between the two clocks is the informative part. τ_{relax} is set by excited-state relaxation, and at these densities that budget is dominated by radiative decay, which the substitution does not touch: swapping collisional rate coefficients moves a term that is not in control. τ_{slow} is set by ground-state depletion, which is collisional almost entirely, so it absorbs the full change. The fast clock is anchored by spontaneous emission and the slow clock is not, and that is why an atomic-data uncertainty of 13 % in the rates becomes 11 % in M rather than cancelling between numerator and denominator.

The sensitivity is not uniform over the grid, and it runs opposite to where the rate coefficients are best known. It grows with density, from 6.3 % at $10^{12}/\text{cm}^3$ to 13.7 % at $10^{15}/\text{cm}^3$, saturating above $n_e \approx 10^{13}/\text{cm}^3$ where the collisional part of $\mathbf{L}$ has taken over from the radiative part and the rate uncertainty passes through undiluted. It falls with temperature, from 17.0 % at 1 eV to 5.6 % at 10 eV, even though Chapter 2 finds the *rates* best constrained at the cold end. At 1 eV the excitation rates sit deep in the Boltzmann exponential, where a given fractional change in a rate coefficient produces a larger fractional change in the eigenvalue it feeds. Good data in a sensitive regime is not the same as a robust answer.

What this does and does not disturb is worth separating. M spans more than seven orders of magnitude across the grid, so a shift of 11 % changes no statement about its order of magnitude anywhere: the weakest separation on the grid moves from 86.8 to 80.8, and the two-clock picture is not in question at any grid point. **M is therefore reported throughout this thesis with the 11 % by which one independent excitation dataset moves it**, a sensitivity to a substitution and not an uncertainty, and values such as the benchmark 9982 should be read as $\sim 10^4$ rather than as four significant figures.

One consequence is sharper and less comfortable. The plateau-window criterion of Chapter 5 admits a grid point to the transient analysis only if $M > 900$, and six of the 400 points change side under the substitution. The analysed set is therefore itself weakly dependent on which atomic database was used, by about 2 % of its membership. That does not affect any conclusion drawn from the 346 points, all of which sit far from the cut, but it does mean the cut is not a property of the plasma alone, and it should not be described as one.

3.3.7 What the separation does and does not license

A separation of 10^4 invites the quasi-steady-state approximation, which Section 3.4 derives. Two warnings, both of which matter later.

First, M compares two *rates*. Chapter 5 measures a *distance*, how far a diagnostic reading sits from the truth. Knowing how quickly a system recovers from a displacement says nothing about how far it was displaced, and these are independent quantities.

Second, eigenvalues describe the long-time behaviour of each mode in isolation. The matrix $\mathbf{L}$ is strongly asymmetric, downward radiative rates exceeding upward collisional ones by some five orders of magnitude at the benchmark density, which makes non-normality plausible without proving it: a non-symmetric matrix may still be normal and have orthogonal eigenvectors. Section 3.6 settles it numerically, finding the numerical abscissa positive at every grid point while every eigenvalue is negative. For a non-normal matrix the modal amplitudes can be large and opposing, and a perturbation can grow before it decays even though every eigenvalue is negative. The spectrum bounds the asymptotic decay; it does not bound the finite-time error.

Both warnings point the same way: the two clocks tell us the system *can* be simplified, and they say nothing about what the simplification costs. Working that out is the business of the rest of this chapter.

3.4 The quasi-steady-state approximation

3.4.1 The shortcut

A separation of 10^4 invites a shortcut. If the excited states settle in nanoseconds while the ionisation balance takes microseconds, then on any timescale relevant to the second, the first has already finished. So do not solve for the excited states at all: assume they have reached whatever configuration the current conditions demand, and solve only for the slow subsystem.

This is the quasi-steady-state approximation, and every practical collisional–radiative model uses some form of it. Without it, a plasma simulation that resolves ns timesteps in order to follow the excited states cannot afford to run for the ms an experiment lasts.

The chemical-engineering version of this argument is the Bodenstein steady-state approximation. For a reaction $A \rightarrow I \rightarrow P$ with a short-lived intermediate, one sets $d[I]/dt = 0$ and solves algebraically for $[I] = (\text{production})/(\text{loss})$. The intermediate stops being a differential unknown and becomes slaved to whatever $[A]$ currently is. What follows is the same move applied to 42 intermediates at once, coupled to each other.

3.4.2 Partition and elimination

Split the 43 levels into a *slow* subspace $\mathcal{S} = \{1s\}$ and a *fast* subspace $\mathcal{F}$ containing the 42 excited levels, and partition Eq. (3.2) conformably:

$$\frac{d}{dt} \begin{pmatrix} \mathbf{n}_S \\ \mathbf{n}_F \end{pmatrix} = \begin{pmatrix} \mathbf{L}_{SS} & \mathbf{L}_{SF} \\ \mathbf{L}_{FS} & \mathbf{L}_{FF} \end{pmatrix} \begin{pmatrix} \mathbf{n}_S \\ \mathbf{n}_F \end{pmatrix} + \begin{pmatrix} \mathbf{S}_S \\ \mathbf{S}_F \end{pmatrix} n_{\text{ion}}. \quad (3.19)$$

Section 3.3 justifies this particular partition rather than some other: the slow eigenvector is 1.0000 ground state with every excited component below 10^{-3} , so the slow mode and the ground state are the same object to the precision that matters.

Setting $d\mathbf{n}_F/dt = 0$ in the lower block gives an algebraic system for the excited populations,

$$\boxed{\mathbf{n}_F^{\text{QSS}} = -\mathbf{L}_{FF}^{-1} (\mathbf{L}_{FS} \mathbf{n}_S + \mathbf{S}_F n_{\text{ion}}).} \quad (3.20)$$

Read the right-hand side as a sentence. The bracket is everything arriving in the excited manifold from outside it: $\mathbf{L}_{FS} \mathbf{n}_S$ is excitation up from the ground state, $\mathbf{S}_F n_{\text{ion}}$ is recombination down from the ions. The factor $-\mathbf{L}_{FF}^{-1}$ divides by the loss network. It is a matrix rather than a number only because a level's losses feed other levels, which have their own losses, and so on; the division is over the whole network at once. For a single Bodenstein intermediate it would collapse to $1/k_{\text{loss}}$.

Two things make Eq. (3.20) usable. $\mathbf{L}_{FF}$ is invertible because it is a strictly column-diagonally-dominant matrix with negative diagonal: every excited level loses population to the ground state or to the continuum, so each column sum is strictly negative rather than zero, and the Levy–Desplanques theorem applies. Physically, there are no perfectly trapped states within the excited manifold. Its condition number lies between 1.48×10^3 and 1.74×10^5 across the grid, so the inversion is numerically well behaved despite the twelve decades spanned by the entries. Both the range and the condition number are produced by `verify_operator_conditioning.py` and stamped in `validation/operator_conditioning/`.

3.4.3 The reduced dynamics

Substituting Eq. (3.20) back into the slow block eliminates $\mathbf{n}_F$ entirely:

$$\frac{d\mathbf{n}_S}{dt} = \underbrace{(\mathbf{L}_{SS} - \mathbf{L}_{SF} \mathbf{L}_{FF}^{-1} \mathbf{L}_{FS})}_{\Omega_{\text{QSS}}} \mathbf{n}_S + (\mathbf{S}_S - \mathbf{L}_{SF} \mathbf{L}_{FF}^{-1} \mathbf{S}_F) n_{\text{ion}}. \quad (3.21)$$

A 43-dimensional differential system has become a one-dimensional one. The matrix Ω_{QSS} is the *Schur complement* of $\mathbf{L}_{FF}$ in $\mathbf{L}$, and it has a direct physical reading: $\mathbf{L}_{SS}$ is the direct loss from the ground state, and $-\mathbf{L}_{SF} \mathbf{L}_{FF}^{-1} \mathbf{L}_{FS}$ is the correction for atoms that leave the ground state, wander through the excited manifold, and either come back or are ionised on the way. It is stepwise ionisation, summed over every possible path.

This is the object the rest of the thesis is about. Ω_{QSS} is precisely the effective ionisation rate that a plasma modeller using QSS would construct, and it is the algebraic form of the effective coefficients tabulated in frameworks such as ADAS [8]. When Chapter 5 asks what the approximation costs, it is asking what this quantity costs.

3.4.4 When is it valid?

The elimination assumed the excited states are always at their instantaneous steady value. The error in that assumption can be estimated directly. Write the deviation from the instantaneous target as $\boldsymbol{\delta}(t) = \mathbf{n}_F(t) - \mathbf{n}_F^{\text{QSS}}(t)$ and subtract Eq. (3.20) from the lower block of Eq. (3.19):

$$\frac{d\boldsymbol{\delta}}{dt} = \mathbf{L}_{FF} \boldsymbol{\delta} - \frac{d\mathbf{n}_F^{\text{QSS}}}{dt}. \quad (3.22)$$

The deviation relaxes at the rate set by $\mathbf{L}_{FF}$, whose slowest mode is τ_{relax} , and is forced by the motion of the target itself. If the driving conditions change on a timescale τ_{drive} ,

then $|\mathbf{d}\mathbf{n}_F^{\text{QSS}}/\mathbf{d}t| \sim \|\mathbf{n}_F^{\text{QSS}}\|/\tau_{\text{drive}}$, and balancing the two terms in Eq. (3.22) at quasi-static equilibrium gives $\|\boldsymbol{\delta}\| \sim \tau_{\text{relax}}\|\mathbf{n}_F^{\text{QSS}}\|/\tau_{\text{drive}}$, that is

$$\frac{\|\mathbf{n}_F - \mathbf{n}_F^{\text{QSS}}\|}{\|\mathbf{n}_F^{\text{QSS}}\|} \sim \frac{\tau_{\text{relax}}}{\tau_{\text{drive}}}, \quad (3.23)$$

so the approximation is accurate whenever $\tau_{\text{drive}} \gg \tau_{\text{relax}}$. For the self-consistent problem, where the only thing driving the excited states is the slow subsystem itself, $\tau_{\text{drive}} \rightarrow \tau_{\text{slow}}$ and the condition becomes

$$M = \frac{\tau_{\text{slow}}}{\tau_{\text{relax}}} \gg 1. \quad (3.24)$$

Section 3.3 measured $M = 9982$ at the benchmark and $M \geq 86.8$ at every one of the 400 grid points, the floor carrying the qualification stated there. **The condition is satisfied everywhere on the grid, comfortably.**

Two qualifications must travel with that statement.

Equation (3.23) is a spectral estimate, and a bound is available. The step from Eq. (3.22) to Eq. (3.23) identified the relaxation rate of $\boldsymbol{\delta}$ with the slowest eigenvalue of $\mathbf{L}_{FF}$. That identification is exact only for a normal operator. The quasi-static balance itself gives more: $\boldsymbol{\delta} \simeq \mathbf{L}_{FF}^{-1}\dot{\mathbf{n}}_F^{\text{QSS}}$, hence

$$\|\boldsymbol{\delta}\| \leq \|\mathbf{L}_{FF}^{-1}\|_2 \|\dot{\mathbf{n}}_F^{\text{QSS}}\|, \quad \|\mathbf{L}_{FF}^{-1}\|_2 = \frac{1}{\sigma_{\min}(\mathbf{L}_{FF})}, \quad (3.25)$$

and since $\sigma_{\min} \leq \min_k |\lambda_k|$ for every matrix, the operative timescale $\|\mathbf{L}_{FF}^{-1}\|_2$ is never shorter than the spectral one. Measured across the grid, the ratio $\|\mathbf{L}_{FF}^{-1}\|_2 |\lambda_{\min}(\mathbf{L}_{FF})|$ lies between 1.118 and 1.526, with median 1.151: the spectral estimate understates the rigorous timescale by at most 53%. Defining the corresponding bounded separation $M^{\text{eff}} = \tau_{\text{slow}}/\|\mathbf{L}_{FF}^{-1}\|_2$,

$$M^{\text{eff}} = 8659 \text{ at the benchmark,} \quad M^{\text{eff}} \geq 77.3 \text{ at all 400 grid points,} \quad (3.26)$$

so Eq. (3.24) survives non-normality with a margin reduced by 13% and not by orders of magnitude. The penalty is this small even though $\kappa(\mathbf{L}_{FF})$ reaches 1.74×10^5 because the ill-conditioning lives in the fast subspace, among modes separated by many decades in rate, and barely touches the slowest mode of $\mathbf{L}_{FF}$, which is the one that sets QSS validity. The operator is strongly non-normal in a way that is almost irrelevant to this particular question.

That said, Eq. (3.25) is still a statement about the settled deviation, and the warning of Section 3.3 applies to the approach as opposed to the endpoint: the balance in Eq. (3.22) is a *quasi-static* argument, valid once the fast transients have decayed, and it says nothing about what $\boldsymbol{\delta}$ does in the first few τ_{relax} after a disturbance, when the non-orthogonality of the eigenvectors can allow transient growth. Whether that matters for the quantity actually measured is not decidable from the spectrum, and Chapter 5 settles it by integrating the full system rather than by argument.

The condition is not sharp. At the benchmark point Eq. (3.23) gives $(M^+)^{-1} = 1.21 \times 10^{-4}$, the post-step operator being the one that governs the relaxation being compared, whereas direct integration returns a residual of order 10^{-8} : the approximation is four orders of magnitude better than the scaling argument requires. The unwritten prefactor, set by how strongly the observable projects onto the moving part of the target,

is therefore small. Equation (3.23) predicts the trend correctly and the magnitude not at all, and it should be quoted only for the former.

That would appear to end the matter. It does not, and the reason is the subject of the next section.

3.5 Two reference states

3.5.1 Why ratios

A spectrometer viewing a divertor records a line-integrated brightness. Turning that into an absolute population requires the neutral density along the sightline and the path length, neither of which is known to better than a factor of a few. Practitioners therefore work with *ratios*, in which the instrument calibration and the common geometric scale divide out. What a ratio does not remove is the chord integration itself: as Section 1.3 sets out, the measured ratio is the n_4 -weighted average of the local ratio, equal to the local ratio only where that ratio is constant over the emitting region. Everything below is written for the local ratio.

The ratio used throughout this thesis is

$$R \equiv \frac{n_3}{n_4} = \frac{n_{3s} + n_{3p} + n_{3d}}{n_{4s} + n_{4p} + n_{4d} + n_{4f}}, \quad (3.27)$$

the population ratio of the $n = 3$ and $n = 4$ shells. Up to Einstein coefficients it is the Balmer H_α/H_β ratio, since H_α is the $3 \rightarrow 2$ transition and H_β the $4 \rightarrow 2$. The shell ratio is used here because it is the cleaner object to derive with, not because the results are restricted to it: Section 3.5.5 shows every result of this chapter holds for the A -weighted line ratio itself, and for any non-negative linear functional of the state vector. The substitution is therefore not free but it is also not load-bearing, and Section 4.7 measures what it costs: the transient error of Chapter 5, computed from the A -weighted line ratio and from the shell ratio, agrees to 0.065% at the benchmark point (0.063571 against 0.063612), and the line-ratio value lies below the shell-ratio value at every grid point and step direction, with a worst case of 5.2% at the low-density edge (`validation/weighted_census/`). The difference is not the darkness of the $4f$ level alone. The ℓ -distribution within a shell differs between the ground-fed and the recombination-fed channels, so the A -weighted sum and the shell sum of Eq. (3.27) respond differently to the transient, and the H_α numerator contributes more to the difference than the H_β denominator: substituting one side at a time at that corner gives 0.967 for the numerator against 0.981 for the denominator, and 0.948 with both substituted (`validation/weighted_census/`); the scaling of the 5.2% with the proton ℓ -mixing rate is in Section 4.7. It is convenient to normalise to the ground state as well, writing $r_p = n_p/n_{1s}$, since that is the form in which the QSS solution is usually tabulated.

3.5.2 The hidden variable

Dividing Eq. (3.20) through by the ground-state population gives the QSS ratios

$$r_p^{\text{QSS}} = \frac{\left[-\mathbf{L}_{FF}^{-1} (\mathbf{L}_{FS} \mathbf{n}_S + \mathbf{S}_F n_{\text{ion}}) \right]_p}{n_{1s}}. \quad (3.28)$$

It is tempting, and in the literature common, to treat Eq. (3.28) as a function of T_e and n_e alone, tabulate it, and invert a measured ratio against the table. That step deserves more scrutiny than it usually receives.

Look at what the numerator contains. The term $\mathbf{L}_{FS}\mathbf{n}_S$ is proportional to the ground-state density; the term $\mathbf{S}_F n_{\text{ion}}$ is proportional to the ion density. What fixes r_p^{QSS} is not either supply on its own but the *ratio* of the two.

It is convenient to measure that ratio against its equilibrium value. Define the ground-state *departure coefficient*

$$b_1 \equiv \frac{n_{1s}}{Z_1 n_e n_{\text{ion}}} \quad (3.29)$$

where Z_1 is the Saha–Boltzmann coefficient of the ground state,

$$Z_p [\text{cm}^3] = 10^6 \frac{g_p}{2g_i} \left(\frac{h^2}{2\pi m_e e T_e} \right)^{3/2} \exp\left(\frac{\chi_p}{T_e}\right), \quad \chi_p = \frac{13.606}{p^2}, \quad (3.30)$$

where T_e and χ_p are both in eV, e is the elementary charge in C (converting eV to J), the bracket evaluates in m^3 , and the factor 10^6 converts to cm^3 . Here $g_p = 2p^2$ is the statistical weight of level p , $g_i = 1$ that of the bare proton, and the factor 2 in the denominator is the spin degeneracy of the free electron. Temperatures are quoted in electron-volts throughout this thesis; no Boltzmann constant appears, and T_e should everywhere be read as the electron thermal energy.

Z_p is the factor for which $n_p = Z_p n_e n_{\text{ion}}$ holds in complete thermodynamic equilibrium, so b_1 is dimensionless and equals unity there. Departures above or below unity measure over- and under-population relative to that thermodynamic reference; no more should be read into the sign than that.

The name “departure coefficient” suggests a small correction, and the numbers show it is nothing of the kind. Evaluated on the collisional–radiative steady state, b_1 runs from 76.6 at $[0, 7]$, the coldest and densest corner of the grid, to 2.683×10^5 at $[49, 0]$, the hottest and thinnest, and equals 956 at the benchmark point. The ground state is over-populated relative to Saha–Boltzmann by two to five orders of magnitude throughout, which is the statement that a plasma at these densities is far from thermodynamic equilibrium: radiative decay returns atoms to the ground state faster than collisions can lift them out, and where the resonance lines escape the photons are not replaced. A quantity that varies over three decades across the operating range is not a small correction, and treating it as fixed is the assumption this thesis examines.

The location of those extremes is not where a first guess puts it, and the correction matters later. One expects the largest departure where the gas is coldest and most neutral. It is the opposite: b_1 is largest in the *hottest, thinnest* corner and smallest in the coldest, densest one. The reason is that b_1 measures distance from Saha–Boltzmann, and Saha–Boltzmann at 10 eV predicts a nearly fully ionised gas, so any surviving neutral population is an enormous departure from it. At 1 eV Saha already predicts a mostly neutral gas, so a predominantly neutral plasma is only a modest departure. The neutral *fraction* therefore runs the other way: solving $\mathbf{L}\mathbf{n} + \mathbf{S} = 0$ directly and restoring $n_{\text{ion}} = n_e$, the ratio n_{1s}/n_{ion} falls from 40.4 at $[0, 0]$ and 20.5 at $[0, 7]$, gases that are 95 to 98% neutral, to 5.3×10^{-6} at $[49, 7]$. **A large departure coefficient does not mean many neutrals**, and the two quantities peak at opposite corners of the grid.

Where the optically thin assumption holds

That last observation forces a declaration that the preceding sections left implicit. The Einstein coefficients in Eqs. (3.1), (3.3) and (3.4) enter bare: every emitted photon is assumed to leave the plasma. Where the ground-state density is high that is false: a Lyman- α photon is reabsorbed, the atom returns to $2p$, and the effective $A_{2p \rightarrow 1s}$ is reduced by the population escape factor Θ_P of Holstein [38].

Whether this matters is a question about n_{1s} , and the numbers above make it answerable. Taking a homogeneous slab of thickness D and a Doppler-broadened Lyman- α line at $T_{\text{at}} = T_e$:

point	T_e (eV)	n_e (1/cm ³)	n_{1s} (1/cm ³)	τ_c	Θ_P
[49, 0]	10	10^{12}	1.1×10^7	9×10^{-7}	1.000
[49, 7]	10	10^{15}	5.3×10^9	5×10^{-4}	1.000
[23, 5]	2.95	1.4×10^{14}	1.2×10^{11}	0.020	0.986
[0, 0]	1	10^{12}	4.0×10^{13}	11.1	0.032
[0, 7]	1	10^{15}	2.0×10^{16}	5.6×10^3	3×10^{-5}

with τ_c and Θ_P quoted at $D = 10$ cm; Θ_P at $[0, 0]$ is 0.48, 0.032 and 0.0023 for $D = 1, 10$ and 100 cm respectively, so the qualitative conclusion does not depend on the choice, but no single figure should be quoted without it. The model itself contains no geometry, and D is an input.

The reading is clean. Over the hot half of the grid the plasma is optically thin to Lyman- α to better than one part in 10^3 , and at the benchmark point $\Theta_P = 0.986$: the bare- A treatment is justified where the quantitative results of this thesis are established. Along the $T_e = 1$ eV edge it is not. There Θ_P falls to 3×10^{-5} , the effective $A_{2p \rightarrow 1s}$ collapses, and with it b_1^{CRE} , the switching points c_m/a_m , and the ground-fed fractions of Section 3.5.4.

The results of this chapter are therefore stated for $T_e \gtrsim 2$ eV, and the cold edge of the grid is retained for the asymptotic behaviour it exposes rather than as a prediction. It is also the corner where the atomic-only model is weakest for an independent reason: at 95 to 98 % neutral and 1 eV the gas is molecular, and neither H_2 nor molecular-assisted recombination appears in **L**. Chapter 6 takes both limitations up quantitatively.

The structural results of this chapter are statements about whatever **L** is supplied, and the condition they need must be stated exactly, because it is easy to claim too much here. The quasi-steady-state closure, the three-argument form of R^{QSS} and the identity $d \ln R / d \ln b_1 = f_3 - f_4$ require that $\mathbf{L}_{FF}$, $\mathbf{L}_{FS}$ and $\mathbf{S}_F$ contain no n_{1s} . That is what makes Eq. (3.33) affine in the reservoir, and the affinity is what makes f_p a logistic of unit width.

For a matrix with the Lyman transitions trapped at a *fixed* escape factor the condition holds, and the identities carry over with a different b_1^{CRE} . For the self-consistent trapped problem it does not: Θ_P depends on n_{1s} , so the effective transition probabilities do too, **a** is no longer independent of the reservoir, and the unit width is lost. The ceiling then becomes $\tanh(m|\Delta|/4)$ with $m \neq 1$ and not known in advance. Section 6.1 of Chapter 6 treats the self-consistent case and quantifies where it matters.

A molecular reservoir breaks the condition more seriously, because H_2 set by recycling is a *third* independent supply. Then $n_p = a_p n_{1s} + c_p n_{\text{ion}} + m_p n_{\text{H}_2}$, f_p is not a one-variable logistic, and the ceiling does not follow at all. The framework of this chapter is two-channel

by construction, and that is a statement about its form rather than about the accuracy of its coefficients.

b_1 is a **dynamical variable** with its own equation of motion. Differentiating Eq. (3.29) at fixed n_{ion} ,

$$\dot{b}_1 = \frac{\dot{n}_{1s}}{Z_1 n_e n_{\text{ion}}} - b_1 \left(\frac{\dot{Z}_1}{Z_1} + \frac{\dot{n}_e}{n_e} \right). \quad (3.31)$$

Only the first term comes from Eq. (3.21), which governs n_{1s} ; the second comes from the normalisation moving. This distinction matters for the temperature-step experiment of Chapter 5. At the instant of a step, n_{1s} is continuous but T_e is not, so $Z_1(T_e)$ jumps and b_1 changes discontinuously purely because its reference has changed. Once T_e and n_e are again held fixed, $b_1 \propto n_{1s}$ and its subsequent relaxation is that of the ground-state population, on τ_{slow} .

Either way, $b_1(t)$ is not determined by the instantaneous (T_e, n_e) ; it carries the plasma's history. Nothing in the QSS derivation fixes it.

3.5.3 Does the observable actually depend on it?

The previous subsection shows that each *individual* normalised population depends on b_1 . That is not yet an argument about R . A sceptical reader is entitled to object: perhaps every level depends on b_1 and the dependence cancels in the ratio, leaving H_α/H_β a clean function of (T_e, n_e) after all. Ratios are chosen precisely because things cancel in them. The objection must be answered, not assumed away.

It can be answered exactly, because Eq. (3.20) is linear in the two supplies. Define

$$\mathbf{a} \equiv -\mathbf{L}_{FF}^{-1} \mathbf{L}_{FS}, \quad \mathbf{c} \equiv -\mathbf{L}_{FF}^{-1} \mathbf{S}_F, \quad (3.32)$$

so that Eq. (3.20) reads

$$\mathbf{n}_F^{\text{QSS}} = \mathbf{a} n_{1s} + \mathbf{c} n_{\text{ion}}. \quad (3.33)$$

Both vectors are dimensionless and both have a plain reading. a_p is the number of atoms that end up sitting in level p per ground-state atom present: not only those excited into p directly, but those excited anywhere and arriving at p afterwards by cascade or by a further collision. The matrix inverse is what performs that sum over every possible route. Likewise c_p is the population of p per ion, from electrons captured into any level and settling into p on the way down. Neither depends on how many ground-state atoms or ions are actually present; both are properties of $\mathbf{L}_{FF}$, $\mathbf{L}_{FS}$ and $\mathbf{S}_F$ alone, and therefore functions of (T_e, n_e) only.

Every excited level is thus fed from two directions, and what distinguishes one level from another is the *proportion* in which it draws on the two.

Writing a_m and c_m for the sums over the ℓ -resolved states of shell m , and using $n_{1s} = b_1 Z_1 n_e n_{\text{ion}}$ from Eq. (3.29),

$$n_m^{\text{QSS}} = n_{\text{ion}} [a_m b_1 Z_1 n_e + c_m] = a_m n_{1s} + c_m n_{\text{ion}}. \quad (3.34)$$

The prefactor n_{ion} is common to both shells and cancels in the ratio, as does the instrument calibration. Equation (3.34) is a relation at a point, and reading it as a statement about a measured chord requires in addition that n_3/n_4 be constant over the emitting region (Section 1.3). The neutral density does *not* cancel, and it is important not to read the previous sentence that way: what an intensity ratio removes is the need to *know* the




Figure 3.3: The shell ratio R^{QSS} as a function of the ground-state departure coefficient, at the benchmark point, from Eq. (3.35). The curve is a one-parameter family: it asks what R would be if b_1 took other values at fixed (T_e, n_e) . Between the two supply limits the observable changes by a factor of several; a table indexed on (T_e, n_e) alone must select one point on this curve. Both axes are logarithmic, so the visible slope is $d \ln R / d \ln b_1$ directly, the quantity plotted in Figure 3.4. The marker shows the value the CR-equilibrium closure selects, and the shaded strip the range of b_1^{CRE} over the whole grid.

absolute neutral density and the path length in order to form it, not the dependence of the ratio on how much neutral hydrogen is present relative to the ions. That dependence is the subject of this chapter and it survives, inside the bracket, as b_1 :

$$R^{\text{QSS}}(b_1) = \frac{a_3 b_1 Z_1 n_e + c_3}{a_4 b_1 Z_1 n_e + c_4} \quad (3.35)$$

Differentiating,

$$\frac{\partial R^{\text{QSS}}}{\partial b_1} = \frac{Z_1 n_e (a_3 c_4 - a_4 c_3)}{(a_4 b_1 Z_1 n_e + c_4)^2} \quad (3.36)$$

The observable is independent of b_1 if and only if

$$a_3 c_4 = a_4 c_3 \quad \Longleftrightarrow \quad \frac{a_3}{c_3} = \frac{a_4}{c_4}, \quad (3.37)$$

that is, if and only if the $n = 3$ and $n = 4$ shells draw on the two supplies in *identical proportion*. There is no reason for them to. The $n = 4$ shell lies closer to the continuum and is more readily fed by recombination; the $n = 3$ shell is more readily reached by excitation from the ground state. Chapter 5 measures the difference and finds it substantial across most of the operating range.

3.5.4 The ground-fed fraction, and where the sensitivity vanishes

Equation (3.36) settles whether the dependence exists. It is not, however, the right measure of how *strong* it is, because R itself varies over orders of magnitude across the




Figure 3.4: The ground-fed fractions of the two shells and their difference, at the benchmark point; the dotted vertical line marks b_1^{CRE} there. Both fractions rise from zero to unity as the ground-state reservoir grows, but they do so at different rates, and their difference, which by Eq. (3.39) is exactly $d \ln R / d \ln b_1$, vanishes in both limits and peaks where the two supplies compete. Note that $f_3 - f_4$ is a difference of fractions, not itself a fraction.

operating range and an absolute derivative is not comparable between grid points. The dimensionless measure is the logarithmic sensitivity, and it takes a form with a direct physical reading.

Define the *ground-fed fraction* of shell m ,

$$f_m \equiv \frac{a_m b_1 Z_1 n_e}{a_m b_1 Z_1 n_e + c_m} \quad (3.38)$$

the share of shell m 's population that arrived from below, by excitation out of the ground state, as opposed to from above by recombination. That f_m lies in $[0, 1]$ is not automatic: it requires $a_m \geq 0$ and $c_m \geq 0$, and both are components of $-\mathbf{L}_{FF}^{-1}$ acting on non-negative vectors. The claim therefore needs $-\mathbf{L}_{FF}^{-1}$ to be non-negative *elementwise*, which holds for the same reason $\mathbf{L}_{FF}$ is invertible (Section 3.4): $-\mathbf{L}_{FF}$ has non-positive off-diagonal entries and strictly positive column sums, making it a non-singular M-matrix, and the inverse of a non-singular M-matrix is elementwise non-negative. Physically, adding atoms to one level cannot reduce the steady population of another, so the statement is unsurprising; it is recorded because the two limits below depend on it. Direct evaluation confirms it: at none of the 400 grid points does $-\mathbf{L}_{FF}^{-1}$ contain a negative entry.

Differentiating Eq. (3.35) logarithmically then gives

$$\frac{d \ln R}{d \ln b_1} = f_3 - f_4 \quad (3.39)$$

The Balmer ratio's sensitivity to the ground-state reservoir is exactly the difference in ground-fed fraction between the two shells. This is the mathematical core of the thesis, and Eq. (3.39) is where the physics becomes visible.

It vanishes in two limits, and only there. When $b_1 Z_1 n_e \gg c_m / a_m$ for both shells, both are fed almost entirely from the ground state, $f_3 = f_4 = 1$, and the ratio is blind to

the reservoir. When $b_1 Z_1 n_e \ll c_m/a_m$, both are fed almost entirely by recombination, $f_3 = f_4 = 0$, and the ratio is blind again. *In between, where the two supplies compete, the shells are fed in different proportion and the sensitivity is largest.*

How large, and where

Equation (3.39) says the sensitivity peaks somewhere between the two limits. It does not yet say how high, or at what reservoir. Both follow in closed form, and the answer is more restrictive than one might expect.

Work in the logarithmic reservoir variable $x = \ln(b_1 Z_1 n_e)$. That combination is the neutral-to-ion ratio itself: $b_1 Z_1 n_e = n_{1s}/n_{\text{ion}}$, so $x = \ln u$ in the notation Chapter 5 uses throughout, and the Saha coefficient cancels out of it entirely. The two chapters work in the same variable under two names. In that variable Eq. (3.38) is exactly a logistic,

$$f_m(x) = \frac{1}{1 + e^{-(x-x_m)}}, \quad x_m \equiv \ln\left(\frac{c_m}{a_m}\right), \quad (3.40)$$

so the two curves in Figure 3.4 are not merely similar in shape: they are the *same* curve, of the same width, displaced along the axis. Every property of the 42-state network enters only through where each sits. Their separation is

$$\Delta \equiv x_4 - x_3 = \ln\left(\frac{a_3/a_4}{c_3/c_4}\right) = \ln\left(\frac{R^{\text{grd}}}{R^{\text{rec}}}\right), \quad (3.41)$$

the logarithm of the span between the two asymptotes of Figure 3.3. Equation (3.40) also makes explicit a distinction that is easy to lose. **The switching point of shell m is $u_m = c_m/a_m$, a ratio of two network coefficients and nothing else.** In the b_1 coordinate the same point reads $b_1 = c_m/(a_m Z_1 n_e)$, which carries $Z_1 n_e$ and is therefore a Saha-normalised label for a quantity that does not depend on the Saha balance. A plasma may lie far above the Saha–Boltzmann ground-state population, $b_1 \gg 1$, and still have both shells fed predominantly from above. The benchmark point is exactly such a case: $b_1^{\text{CRE}} = 956$, while the two shells do not become half ground-fed until $b_1 = 2.73 \times 10^3$ and 1.91×10^4 respectively, so $f_3 = 0.260$ and $f_4 = 0.048$: three-quarters of $n = 3$ and 95% of $n = 4$ arrive by recombination. A large departure coefficient and a predominantly ground-fed manifold are different statements, and only the second controls Eq. (3.39). Neither should be conflated with calling the plasma *ionising*, which is properly a statement about the direction of the net population flux and is not fixed by the magnitude of b_1 alone. Writing $s = x - \frac{1}{2}(x_3 + x_4)$ and using $f(-y) = 1 - f(y)$, the difference $f_3 - f_4$ is *even* in s . Its maximum therefore sits at the midpoint, and takes the value

$$\boxed{\max |f_3 - f_4| = \tanh\left(\frac{|\Delta|}{4}\right) = \tanh\left[\frac{1}{4} \left| \ln\left(\frac{R^{\text{grd}}}{R^{\text{rec}}}\right) \right| \right]} \quad (3.42)$$

where the modulus matters only for the generalisation below: for the n_3/n_4 ratio considered here $R^{\text{grd}} > R^{\text{rec}}$ and $\Delta > 0$, but for a pair of shells drawn in the opposite order Δ changes sign while the left-hand side cannot. The maximum is attained at $b_1 Z_1 n_e = \sqrt{(c_3/a_3)(c_4/a_4)}$, the geometric mean of the two switching points. The area is fixed by the same quantity: integrating Eq. (3.39) over the *whole* reservoir axis gives $\int_{-\infty}^{+\infty} (f_3 - f_4) d \ln b_1 = \ln R^{\text{grd}} - \ln R^{\text{rec}} = \Delta$. The limits are not decoration: the identity is a sum rule over the full excursion from purely recombination-fed to purely ground-fed, and an

integral over any finite interval is smaller *in magnitude*. The magnitude matters because the integrand keeps a fixed sign only while $\Delta > 0$, which holds for the (3, 4) pair; for a pair taken in the opposite order the integrand is negative throughout and every finite-interval integral is larger, not smaller. Chapter 5 uses a finite-interval version of the same integral for a different purpose, to define a mean sensitivity over the excursion a real step produces, and the two must not be confused. That second relation is only the fundamental theorem of calculus, and it is quoted here for what it is, a consistency check tying the two figures together, not an independent result.

Three things follow, and the third is the useful one.

First, the peak sensitivity is set *entirely* by how far apart the two supply limits lie. Two plasmas with the same span between R^{rec} and R^{grd} have the same worst-case sensitivity, whatever the details of their cascade networks.

Second, $\tanh$ is bounded by unity, so $|\text{d} \ln R / \text{d} \ln b_1| < 1$ everywhere and for every pair of shells. **A shell ratio can never respond to the ground-state reservoir faster, in relative terms, than the reservoir itself moves.** Integrating a logarithmic sensitivity whose magnitude is everywhere below unity gives the finite-excursion statement $|\Delta \ln R| \leq |\Delta \ln b_1|$, so a stale reservoir cannot be amplified in logarithmic space, only attenuated.

3.5.5 The same result for what a spectrometer measures

Everything above is derived for the shell ratio $R = n_3/n_4$, and a spectrometer does not measure shell populations. It measures line intensities, which are sums over the dipole-allowed sublevel transitions weighted by their Einstein coefficients. The gap is worth closing rather than measuring, because the derivation never used the fact that the weights are unity. It used only that they are non-negative.

Let $\mathbf{w}_\alpha$ and $\mathbf{w}_\beta$ be the weight vectors of the two lines, carrying $h\nu A$ on the levels each line drains and zero elsewhere. The local emissivity is $j_\alpha = \mathbf{w}_\alpha^\top \mathbf{n}_F$, and substituting the two-channel form $\mathbf{n}_F = \mathbf{a}n_{1s} + \mathbf{c}n_{\text{ion}}$ of Eq. (3.33) gives

$$j_\alpha = A_\alpha n_{1s} + C_\alpha n_{\text{ion}}, \quad A_\alpha \equiv \mathbf{w}_\alpha^\top \mathbf{a}, \quad C_\alpha \equiv \mathbf{w}_\alpha^\top \mathbf{c}, \quad (3.43)$$

where $\mathbf{w}_\alpha$ carries $h\nu A$ on the levels the line drains, in energy per unit time, and $A_\alpha, C_\alpha \geq 0$ because $\mathbf{a}$ and $\mathbf{c}$ are non-negative elementwise and $\mathbf{w}_\alpha$ is non-negative. The logistic form needs the *strict* inequality, since $A_\alpha = 0$ would put the switching point at infinity, and strictness follows from $-\mathbf{L}_{FF}^{-1}$ being strictly positive here rather than merely non-negative; it is confirmed numerically below. **That is the same affine form with the same positivity**, so the line ratio $j_\alpha/j_\beta = (A_\alpha u + C_\alpha)/(A_\beta u + C_\beta)$ is affine-over-affine in the reservoir exactly as the shell ratio is, and every result of this section carries over unchanged: the ground-fed fraction of a *line* is the same unit-width logistic, its switching point is C_α/A_α , the logarithmic sensitivity is the difference of two such fractions, and the peak is $\tanh(|\Delta_j|/4)$ with $\Delta_j = \ln[(A_\alpha/A_\beta)/(C_\alpha/C_\beta)]$. Nothing is re-derived; the substitution $a_m \rightarrow \mathbf{w}^\top \mathbf{a}$ is the whole of it. The result holds for any non-negative linear functional of the fast-state vector, of which a shell sum and an A -weighted line are two instances.

Checked, not asserted. Evaluating Eq. (3.43) at all 400 grid points with the pipeline's own resolved A -values, H_α from $3s \rightarrow 2p$, $3p \rightarrow 2s$ and $3d \rightarrow 2p$ and H_β from $4s \rightarrow 2p$, $4p \rightarrow 2s$ and $4d \rightarrow 2p$ (`verify_emissivity_generalisation.py`): all four coefficients

are strictly positive everywhere, and maximising $|F_\alpha - F_\beta|$ over u numerically reproduces $\tanh(|\Delta_j|/4)$ to 9×10^{-16} , so the bound is found rather than imposed.

One check is worth more than the others. A constant factor on either weight vector cancels out of Δ_j , so a photon-counting detector and an energy-weighted one must return the *same* Δ_j even though their emissivities differ by $h\nu$. Measured, the two agree to 7×10^{-16} . **The displacement that sets the bound is a property of the atomic network and not of the detector.** The cancellation is exact because every channel of a given line carries the same $h\nu$ in this model, the levels being degenerate in ℓ ; with fine structure resolved the two weightings would differ within a line and the agreement would be approximate rather than exact.

How far the shell ratio was from the observable. Over the grid Δ_j runs from 0.673 to 2.196 against 0.711 to 2.199 for the shell Δ , a median disagreement of 0.134 % and a worst case of 5.4 %; the caps differ by 0.116 % at the median and 5.3 % at worst. At the benchmark, $\Delta_j = 1.9444$ against 1.9456, and the cap is 0.4511 against 0.4514. The worst case sits at $T_e = 10$ eV, $n_e = 10^{12}/\text{cm}^3$, the low-density corner, for the reason Section 4.7 gives: $4f$ contributes exactly zero to H_β while carrying a large share of the $n = 4$ shell, $14/32 = 43.75$ % by statistical weight and 43.66 % as actually populated at that corner. The two differ precisely because ℓ -mixing is weakest there, which is also why the shell and the line disagree most about what $n = 4$ contains.

The shell ratio was therefore a good proxy, and this subsection means the thesis no longer needs it to be one. **Section 4.7's measurement of the substitution cost becomes a consistency check on two routes to the same answer rather than a caveat carried by the analytic results.**

What this bounds is the error in the *observable*. A diagnostician does not suffer an error in R directly: they measure R and invert it for (T_e, n_e) . Take the simplest case first, in which n_e has been fixed by an independent measurement so that only T_e is being inferred. Then the error in the inferred temperature is

$$\delta \ln T_e = \frac{(\partial \ln R / \partial \ln b_1)_{T_e, n_e}}{(d \ln R^{\text{CRE}} / d \ln T_e)_{n_e}} \delta \ln b_1, \quad n_e \text{ known.} \quad (3.44)$$

Both derivatives are taken at fixed n_e , and the relation is a local linearisation: it is exact only in the limit of small $\delta \ln b_1$, and Chapter 5 uses the exact finite-step form in its place wherever a number is quoted (Section 5.5.2). It also requires the denominator to be non-zero, which is not a formality: Section 5.5.1 finds $d \ln R^{\text{CRE}} / d \ln T_e$ passing through zero inside the operating range on six of the eight density columns, and there the quotient diverges and the inversion has no temperature resolution at all.

If n_e is not known independently, one ratio is not enough. Two unknowns need two observables, and the local statement is then a Jacobian rather than a quotient: with two measured ratios $R^{(1)}$ and $R^{(2)}$,

$$\begin{pmatrix} \delta \ln R^{(1)} \\ \delta \ln R^{(2)} \end{pmatrix} = \mathbf{J} \begin{pmatrix} \delta \ln T_e \\ \delta \ln n_e \end{pmatrix}, \quad \begin{pmatrix} \delta \ln T_e \\ \delta \ln n_e \end{pmatrix} = \mathbf{J}^{-1} \begin{pmatrix} \delta \ln R^{(1)} \\ \delta \ln R^{(2)} \end{pmatrix}, \quad (3.45)$$

where $\mathbf{J}$ is the Jacobian of $(\ln R^{(1)}, \ln R^{(2)})$ with respect to $(\ln T_e, \ln n_e)$, and the inversion exists only where $\det \mathbf{J} \neq 0$. That condition is the classical failure mode of two-ratio diagnostics rather than a technicality, because the two families of contours in the (T_e, n_e) plane

can run nearly parallel over wide regions. To parallel Eq. (3.44), a reservoir displacement $\delta \ln b_1$ enters as $(\delta \ln T_e, \delta \ln n_e)^\top = \mathbf{J}^{-1}(\partial \ln R^{(1)}/\partial \ln b_1, \partial \ln R^{(2)}/\partial \ln b_1)^\top \delta \ln b_1$, and the amplification is governed by the conditioning of $\mathbf{J}$ rather than by a single quotient. Chapter 1 already noted that one ratio is one constraint on two unknowns and that practice supplies a second; Eq. (3.44) is what that second constraint buys. This thesis measures the numerator and does not measure $\mathbf{J}$, so no claim about a two-parameter inversion is made anywhere in it.

Where the ratio is itself weakly sensitive to temperature, which is where Balmer diagnostics are known to be least useful, a bounded error in R can still become a large error in the inferred T_e ; the inversion Jacobian in the denominator is supplied in Chapter 5, and without it no statement about diagnostic error follows from this section.

These closed forms also locate the benchmark point on the curve rather than leaving it as an anchor of unstated significance. Its sensitivity is $f_3 - f_4 = 0.212$, some 47% of the maximum 0.4514, which is attained nearly one decade higher at $b_1 = 7.2 \times 10^3$. The benchmark is therefore on the rising flank and not at the worst case, which is the reading to attach to every error reported against it in Chapter 5. Two identities follow from Eq. (3.38) and fix the marker in Figure 3.3 in terms of the fractions of Figure 3.4,

$$R = R^{\text{grd}} \frac{f_4}{f_3} = R^{\text{rec}} \frac{1 - f_4}{1 - f_3}, \quad (3.46)$$

both giving $R = 0.7778$ at the benchmark. They are algebraic rearrangements rather than new physics, but they are the cheapest available cross-check between the two figures and the figure script enforces them.

Third, the attenuation is severe unless the span is large. For $\Delta \ll 1$ the peak is $\Delta/4$: a diagnostic whose two supply limits differ by only 10% has a worst-case log-sensitivity of 0.024. Sensitivity to the hidden variable is bought with span, and $R = n_3/n_4$ has span because $n = 3$ and $n = 4$ draw on the two supplies in genuinely different proportion. That is why this ratio is diagnostically useful and also why it is exposed.

Note that this is the logarithmic sensitivity, not the absolute derivative $\partial R/\partial b_1$ of Eq. (3.36). The latter does *not* vanish as $b_1 \rightarrow 0$; it tends to $Z_1 n_e (a_3 c_4 - a_4 c_3)/c_4^2$, which is finite. The two limits are statements about the *fractional* response of the observable, which is what a diagnostic error is.

The two supply channels are Fujimoto's, who names the first term of Eq. (3.33) the *recombining plasma component* and the second the *ionising plasma component* [11, Sec. 4.1]. Naming the components that way is his and is standard. **Identifying which component dominates one shell with the state of the plasma is a different claim, and this thesis does not make it.** Call a shell *ground-fed* when excitation out of the ground state supplies most of its population and *recombination-fed* when capture from the continuum does. That is a statement about one level's supply budget. Whether the plasma is *ionising* is a statement about the sign of the net stage-to-stage flux, as Section 3.5.4 already insisted, and a level can be predominantly ground-fed in a plasma that is recombining on balance. The crossover Eq. (3.39) singles out is the ground-fed one: between the two limits both supply channels contribute, the two shells draw on them in different proportion, and the sensitivity is largest. That crossover is a statement about the switching points c_m/a_m of the network, and where it falls in (T_e, n_e) is measured rather than assumed; identifying it with any named divertor operating regime is a separate claim, requiring densities this section does not supply, and is not made here. Chapter 5 maps $f_3 - f_4$ over the operating range and shows the diagnostic error tracking it.

Equation (3.35) is not a fit, an expansion, or an approximation: it is an exact rewriting of Eq. (3.20), valid wherever $\mathbf{L}_{FF}$ is invertible. Nor is it new. It is the standard population-coefficient decomposition of collisional–radiative theory, introduced by Bates, Kingston and McWhirter [6] and written by Fujimoto and McWhirter [10] in the Saha-normalised form $b(p) = r_0(p) + r_1(p) b(1)$, with r_0 the recombining and r_1 the ionising population coefficient. Substituting $b(p) = n_p/[Z_p n_e n_{\text{ion}}]$ into that expression and clearing the normalisation gives the correspondence

$$a_p = r_1(p) \frac{Z_p}{Z_1}, \quad c_p = r_0(p) Z_p n_e, \quad (3.47)$$

both sides dimensionless. The two forms are the same linear solution in different coordinates; what is done here is to ask what the ratio of two such populations depends on, which Eq. (3.39) answers.

Section 4.9 tests the decomposition against Fujimoto’s published population coefficients. A second test, against the ADAS effective coefficients [8], cannot be made level by level because r_0 and r_1 are published as curves rather than tables: $\mathbf{a}$ and $\mathbf{c}$ are instead summed against the ionisation rates to form the effective stage-to-stage coefficients, which *are* tabulated, and compared with SCD96 and ACD96. That comparison, Gate D of Section 4.10, fails at every grid point and is left unrepaired; Section 4.10.1 traces the failure to the gate summing over both supply channels where ADAS keeps them apart, and rebuilt on the two channels separately both halves agree with ADAS within a factor two at all 400 points. The size of the disagreement, the gate’s unimplemented recombination half and the second reading that was open are recorded there.

The contraction is severe in any case, since forty-two coefficients are collapsed into one, but it is not merely a rate check either: a_p and c_p enter the two coefficients with different weights, so an error in the *split* between the two supplies displaces the ionisation and recombination coefficients in opposite directions, which no single-channel comparison detects. Chapter 5 verifies the superposition $\mathbf{n}_F = \mathbf{a}n_{1s} + \mathbf{c}n_{\text{ion}}$ to a relative error of 3×10^{-14} .

3.5.6 Two reference states, and they are not the same

Equation (3.35) shows that the QSS ratio is a function of three arguments, not two. This gives two distinct things a measurement can be compared against.

The QSS reference. $R^{\text{QSS}}(T_e, n_e; b_1(t))$, evaluated at the *instantaneous* departure coefficient. This is what adiabatic elimination claims: the excited manifold is in steady state with whatever reservoir is present at that moment. Its accuracy is governed by Section 3.4’s validity condition.

The CR-equilibrium reference.

$$R^{\text{CRE}}(T_e, n_e) \equiv R^{\text{QSS}}(T_e, n_e; b_1^{\text{CRE}}(T_e, n_e)), \quad (3.48)$$

where b_1^{CRE} is obtained by additionally setting $d\mathbf{n}_S/dt = 0$ in Eq. (3.21) and solving for the steady ground-state population. This *is* a function of (T_e, n_e) alone and can therefore be tabulated.

The distinction is easy to miss because both are called “the QSS solution” in common usage, and in a steady plasma they are the same number. They separate whenever the

ground-state population has not finished settling, which, in a divertor disturbed by an edge-localised mode every few milliseconds, is most of the time.

Two remarks on terminology, both of which matter for what follows. First, $b_1 = 1$ denotes complete thermodynamic equilibrium, whereas b_1^{CRE} denotes the steady state of the collisional–radiative system, which is a different and much weaker condition; where excitation out of the ground state dominates and the plasma is optically thin, b_1^{CRE} is very much larger than unity, and, as Section 3.5.4 showed at the benchmark, that does not by itself make the shells ground-fed. Second, this model holds the ion density fixed (Section 3.1), so the quantity relaxing on τ_{slow} is the ground-state population against a prescribed reservoir, not a self-consistent ionisation balance; what that restriction costs is assessed in Chapter 6.

3.5.7 Two error measures

Each reference defines an error:

$$\varepsilon^{\text{QSS}}(t) = \left| \frac{R(t)}{R^{\text{QSS}}(T_e, n_e; b_1(t))} - 1 \right|, \quad (3.49)$$

$$\varepsilon^{\text{CRE}}(t) = \left| \frac{R(t)}{R^{\text{CRE}}(T_e, n_e)} - 1 \right|. \quad (3.50)$$

ε^{QSS} asks whether the excited states have kept up with the reservoir they actually see. ε^{CRE} asks whether the tabulated answer is right. The first is the question the literature worries about; the second is the one a diagnostician suffers.

They are not the same question, and Chapter 5 shows they have very different answers. Naming them apart is the single most important definitional step in this thesis: conflating them makes it possible to measure the second and report it as the first.

3.5.8 What a two-parameter table must assume

It remains to say precisely what is being assumed when a measured ratio is inverted against a table indexed on (T_e, n_e) .

Equation (3.35) establishes that $R = R(T_e, n_e, b_1)$. A two-parameter table must therefore do one of three things: adopt a closure $b_1 = b_1(T_e, n_e)$; operate in a regime where one supply dominates both shells, so that Eq. (3.36) is negligible and the third argument does not matter; or hold some other quantity fixed that determines b_1 implicitly. Equilibrium ionisation balance is a closure, not a mathematical necessity, and it is important not to overstate the point.

It is important to be precise about which practice is being criticised, because the modern effective-coefficient databases are not themselves at fault. ADAS [8], for instance, tabulates photon emissivity coefficients separately for the excitation-driven and recombination-driven components. Used as intended, those two components are weighted by n_{1s} and n_{ion} respectively, which is exactly the information b_1 carries; the database does not impose a closure at all. Indeed the excitation/recombination split of Eq. (3.33) is conceptually the same decomposition.

The assumption enters at the *inversion* step, when a measured ratio is mapped to (T_e, n_e) through a two-parameter lookup. Such a lookup returns a single ratio per

(T_e, n_e) pair and therefore requires the two components to have been combined in a fixed proportion; that is, it requires a closure for b_1 . The usual choice is $b_1 = b_1^{\text{CRE}}$, obtained by solving the collisional–radiative system to complete steady state. A ratio inverted this way carries the assumption that the plasma’s ground-state population has settled at its steady value for the local conditions: not an assumption about how quickly excited states equilibrate, but an assumption that the plasma has *finished changing*.

The second possibility is not merely hypothetical, and it is the reason Eq. (3.36) is worth having. Where one supply dominates both shells, the table is safe regardless of what the ground state is doing. Chapter 5 maps where in the divertor operating range that is the case, and where it is not.

The remainder of this thesis quantifies what the assumption costs.

3.6 A caveat: the modes are not orthogonal

Section 3.3 presented the eigen-decomposition as a change of coordinates in which the 43 coupled equations become 43 independent decays. That is exact. But a change of coordinates is only as convenient as the coordinates are well behaved, and these are not.

The reason is a physical asymmetry running through the whole matrix. Radiative decay carries atoms downward at up to $6.3 \times 10^8/\text{s}$; electron-impact excitation carries them upward at $Cn_e \sim 10^4/\text{s}$ at the benchmark density. Downward and upward coupling therefore differ by nearly five orders of magnitude, and $\mathbf{L}$ is very far from symmetric. That asymmetry does not by itself establish non-normality, for the reason given in Section 3.3: a non-symmetric matrix may still be normal. The non-normality is established separately, by measurement, and the numerical abscissa is positive at all 400 grid points.

The consequence is worth stating carefully, because it is easy to overstate. Every eigenvalue has negative real part, so every mode decays and any perturbation eventually vanishes; that much the spectrum does guarantee. What it does not guarantee is that the perturbation decays *monotonically*. When eigenvectors are nearly parallel, a state of modest size can be the sum of two much larger, nearly cancelling modal contributions. If those two modes decay at different rates the cancellation is progressively spoiled, and the state can grow, sometimes by orders of magnitude, before the decay of the individual modes finally wins. The standard diagnostic for whether this is possible at all is the numerical abscissa

$$\mu(\mathbf{L}) = \lambda_{\max}\left(\frac{1}{2} [\mathbf{L} + \mathbf{L}^\top]\right), \quad (3.51)$$

the largest eigenvalue of the symmetric part of $\mathbf{L}$, which gives the initial slope of $\|\mathbf{n}(t)\|$ for the worst-case perturbation. For this matrix $\mu(\mathbf{L}) > 0$ at every grid point, so transient growth is not merely possible but guaranteed for *some* disturbance.

Whether it occurs for the disturbance and the observable that matter here is a different question, and one the spectrum cannot answer. $\mu(\mathbf{L}) > 0$ is a statement about the worst case over all perturbation directions in the Euclidean norm, and other norms carry their own logarithmic norms; the Balmer ratio is one particular scalar function of the state, and a temperature step is one particular disturbance. A quantity may be entirely well behaved even when the operator that governs it admits growth in principle. This is why Section 3.4 was careful to label Eq. (3.23) an estimate rather than a bound: the eigenvalue argument describes the system after the transients have settled, and the transients themselves are outside its reach.

Chapter 5 therefore settles the question the only way it can be settled, by integrating the full 43-state system through a temperature step and measuring what the Balmer ratio does, rather than by inferring it from the spectrum.

3.7 Summary

Starting from a population balance on each of 43 levels, this chapter assembled the collisional–radiative rate equation $\dot{\mathbf{n}} = \mathbf{L}\mathbf{n} + \mathbf{S}n_{\text{ion}}$ and checked the assembly against particle conservation, which holds to 3.6×10^{-11} across the operating range.

Diagonalising $\mathbf{L}$ revealed a spectrum with one gap and then a continuum: a single slow mode, the relaxation of the ground-state population, sitting four orders of magnitude above a dense band of fast modes whose slowest member is the rearrangement of the excited manifold. The separation $M = \tau_{\text{slow}}/\tau_{\text{relax}}$ spans two to nine orders over the grid and, as Section 3.3 records, its floor sits at a point the transient analysis excludes, so the four-orders figure is the benchmark value rather than a property of the grid.

That separation licenses the quasi-steady-state approximation: eliminate the excited populations algebraically and solve only for the slow subsystem. The elimination is the Bodenstein steady-state approximation applied to 42 coupled intermediates at once, and it produces the Schur complement $\mathbf{\Omega}_{\text{QSS}}$, the effective ionisation rate that a plasma modeller using QSS constructs. Its validity condition, $M \gg 1$, is satisfied everywhere on the grid and comfortably, subject to the two qualifications of Section 3.4: M compares rates where Chapter 5 measures a distance, and the spectrum of a non-normal operator bounds asymptotic decay rather than finite-time error.

The chapter’s central result is that this is not the end of the argument. The QSS ratio

$$R^{\text{QSS}}(b_1) = \frac{a_3 b_1 Z_1 n_e + c_3}{a_4 b_1 Z_1 n_e + c_4}$$

is a function of *three* arguments, not two. The third is the ground-state departure coefficient b_1 , which is dynamical: it carries the plasma’s history and takes a time τ_{slow} to settle. Its logarithmic sensitivity is exactly

$$\frac{d \ln R}{d \ln b_1} = f_3 - f_4,$$

the difference in ground-fed fraction between the two shells, which vanishes when either supply dominates both shells and is largest where the two compete. A table indexed on (T_e, n_e) alone must therefore adopt a closure for b_1 , and the closure adopted in practice is complete collisional–radiative equilibrium.

Two error measures follow and must not be conflated. ε^{QSS} asks whether the excited states have kept up with the reservoir they actually see; ε^{CRE} asks whether the tabulated answer is right. The first is the question the literature worries about. The second is the one a diagnostician suffers.

Everything in this chapter has been an argument about equations. Whether the equations describe hydrogen is a separate question, and Chapter 4 addresses it: what would a wrong model fail, and does this one fail it?

Chapter 4

Does the Model Work?

Chapter 3 built a 43×43 rate matrix, eliminated the fast variables, and extracted from the result a closed-form statement about how a line ratio responds to its reservoir. Chapter 5 will use that machinery to make quantitative claims about a diagnostic in use on real machines. Between the two sits a question that no amount of further algebra answers: why should anyone believe a number this model produces?

The answer is not a single verdict. It is a graded one, because the checks that have been run on this model are not of equal worth, and several of the ones that look most impressive are worth almost nothing. This chapter sorts them.

What this chapter takes as given. The state space, the atomic data and their sources, and the benchmark of the excitation rates against an independent R-matrix-with-pseudostates (RMPS) calculation are established in Chapter 2, Sections 2.8 and 2.9, and are not repeated. The construction of $\mathbf{L}$, the conservation identity, the quasi-steady-state (QSS) elimination, and the two-channel decomposition of the excited populations come from Chapter 3. This chapter asks only what the checks run on that construction are entitled to establish.

What this chapter does not do. It reports no result about divertor plasmas. Every number below is about the instrument, not about the physics the instrument is pointed at. Keeping those separate matters here more than anywhere else in the thesis: a model that predicts something interesting is not thereby a correct model, and a check that uses the interesting prediction as its evidence has verified nothing.

4.1 What a check can and cannot show

4.1.1 Three things that get called validation

Three activities are routinely reported under one heading, and they license different conclusions.

Verification asks whether the equations that were intended got solved correctly. Conservation residuals, detailed balance, matrix identities, convergence with tolerance, agreement between two independent code paths: all of these compare the model against itself or against arithmetic. They can detect a transposed index. They cannot detect a rate coefficient that is wrong by a factor of three, because a uniformly wrong dataset satisfies every one of them.

Validation asks whether the model reproduces something established independently of it. Published coefficients from another group's code, measured cross sections, a tabulated limit. Only this class of check can reach the atomic physics.

Analysis asks what the model then says. That is Chapter 5, and none of it appears here.

4.1.2 Severity, and how to measure it

The useful property of a check is not whether it passes. It is the size of the set of wrong models it would reject. Call that its severity.

The failure mode is the one a mass balance on a continuous stirred-tank reactor shows. Compute the outlet stream as inlet minus consumption, then verify that the balance closes. It closes to machine precision, always, for any rate constant whatever, because the outlet was constructed from the balance being tested. The check is real arithmetic and it is entirely uninformative about the kinetics. A model can be wired correctly and be wrong.

Severity is measurable, and two operations measure it.

1. **Provenance inspection.** Ask whether the quantity under test was derived from the quantity doing the testing. If the de-excitation coefficients were computed from the excitation coefficients through detailed balance, then a detailed-balance check on them is arithmetic, not physics. Chapter 2 already states this for the rate dataset (Section 2.3.2); this chapter carries it through to the matrix.
2. **Fault injection.** Break the model deliberately, in a way that a real mistake could plausibly break it, and see whether the check notices. A check that survives a tenfold error in a rate coefficient has told you that a tenfold error in that coefficient is invisible to it.

Fault injection is the sharper of the two, because it produces a number rather than an argument. It is used in Section 4.2 on the conservation check, and its result is the least comfortable finding in this chapter.

4.1.3 The three grades used below

Every check in this chapter is placed in one of three grades, and the grade is stated wherever the check is reported.

Tautological. Cannot fail except on an arithmetic bug, because the quantity tested was constructed from the relation being tested. It confirms that the code implements the construction. It is not evidence about the physics, and reporting it as though it were makes the validation weaker, not stronger, by displacing the checks that carry information.

Weak. Could fail in principle, but the tolerance is so much looser than the measured value that only a gross error would trip it. A gate that passes with a factor of 150 in hand has told you the answer is not absurd.

Severe. A plausible wrong version of this model fails it. Either the check has actually fired at some point in the project's history, or a specific corruption is known to trip it.

Table 4.1: Four deliberate faults injected into the assembly of $\mathbf{L}$, with the resulting change in the matrix and the conservation residual of Eq. (3.8) recomputed after each. Three faults large enough to move individual matrix elements by up to 2.7×10^{11} 1/s leave the residual unchanged at the baseline value; only an inconsistency between the Einstein coefficients A and the total radiative loss rates γ is detected. Source: `outputs/CHANGE_REPORT.md` §4.2.

injected fault	max $ \Delta\mathbf{L} $ (1/s)	residual
none (baseline)	0	2.238×10^{-12}
$K_{\text{exc}}(1s \rightarrow 2p) \times \text{ten}$	6.9×10^5	2.238×10^{-12}
all de-excitation deleted	1.9×10^{11}	2.171×10^{-12}
excitation array transposed	2.7×10^{11}	1.837×10^{-12}
A_{pq} halved, γ_p left alone	3.1×10^8	3.63×10^1 caught

Four quantities currently printed by this project’s own scripts as validation are tautological by this definition, and Table 4.4 names them: Gate A, the column-sum conservation residual, the tanh bound, and Gate E. Sections 4.2 to 4.5 give the measurements behind each.

4.2 The conservation check is exact by construction

Section 3.2 established that each column of $\mathbf{L}$ must sum to the ionisation leak from the level that column represents, and reported the measured residual of Eq. (3.8) as 3.61×10^{-11} , a relative residual (Section 3.2). The 3.61×10^{-11} is the maximum over every column at every one of the 400 grid points; the fault-injection run below has its own baseline, 2.238×10^{-12} , at the single point it was run at. The claim attached to it there is that “an error in a single off-diagonal element, a transposed index, or a missing back-reaction all break this immediately.”

That claim is testable, and it is false.

4.2.1 Four injected faults

Four faults were introduced into the assembly, one at a time, and the residual recomputed after each. The first three are the exact failure modes the claim names. Table 4.1 gives the outcome.

A tenfold error in the single most important excitation coefficient in the model, the one that supplies the entire ground-fed channel of Section 3.5.3, changes the residual in no digit. Deleting every de-excitation process in the matrix, which is not a subtle mistake, changes it in the third digit and leaves it three orders of magnitude below any tolerance. Transposing the whole excitation input array, the classic index-order mistake, does the same.

4.2.2 Why, and what the check is actually worth

The reason is in the assembly, not in the physics. The diagonal of $\mathbf{L}$ is built as minus the column sum of the same off-diagonal array that the check then sums. Whatever numbers are in that array, correct or corrupted, the column sum returns the ionisation leak. The

identity of Eq. (3.6) is imposed by the construction rather than tested by it. This is the reactor mass balance of Section 4.1 exactly.

The one fault that was caught is the one that breaks the construction. In the notation of Eq. (3.1), A_{pq} is the Einstein coefficient for spontaneous emission from level p to level q , in 1/s, and $\gamma_p = \sum_{q<p} A_{pq}$ is the total radiative loss rate from p , also in 1/s. The individual A_{pq} enter the off-diagonal elements while their sum γ_p enters the diagonal, so halving one without the other makes the two inconsistent, and the residual rises from 2.238×10^{-12} to 36.3, thirteen orders of magnitude. The check raises rather than warns, it costs nothing to run, and on this evidence that is the width of what it covers.

The rate magnitudes are reached only by Section 2.8's comparison against RMPS and by the external comparisons of Section 4.9, and by nothing internal.

The grid point of the table. Three of the four $\max |\Delta \mathbf{L}|$ values in Table 4.1 scale with n_e , so the table fixes its own grid point. Sweeping all 400 grid points and matching on those three values identifies [23, 5], the benchmark, with a worst disagreement of 1.63%, the rounding of the printed table to two significant figures; every residual in Table 4.1 then reproduces to the digits shown (`verify_fault_injection.py`, stamped in `validation/fault_injection/`). The script rebuilds $\mathbf{L}$ through the pipeline's own assembly rather than patching the stored matrix, and the rebuild reproduces `L_grid` exactly before anything is injected, which makes the faults faults in the model rather than in a re-implementation of it.

4.3 The gate suite, read one gate at a time

Five automated gates run over the grid in `src/validation/validate_gates.py` and write `validation/gate_summary.txt`. Their headline outcome is four passes and one failure. Read as a headline it says little, because the five differ in severity by a wide margin and two of them cannot fail. This section takes them one at a time; Gate D, the failure, gets Section 4.10 to itself.

4.3.1 Gate A: detailed balance, and where its physics content actually lives

Gate A takes every excitation rate coefficient $K_{\text{exc}}(i \rightarrow j)$ in the dataset, in cm^3/s , and predicts the reverse coefficient from the Boltzmann relation between forward and reverse rates in a Maxwellian electron distribution at temperature T_e ,

$$K_{\text{deexc}}(j \rightarrow i) = K_{\text{exc}}(i \rightarrow j) \frac{g_i}{g_j} \exp\left(\frac{E_j - E_i}{k_B T_e}\right), \quad (4.1)$$

then compares the prediction against the stored reverse coefficient at three temperatures spanning the grid. Here g_i is the dimensionless statistical weight of level i , $E_j - E_i$ is the excitation energy in eV, and k_B is Boltzmann's constant. The relation holds only for a Maxwellian distribution, which Section 2.3 established and which is assumed throughout. The gate passes if the maximum relative error is below 0.01 % (`validate_gates.py:118`). Measured: below 0.01 % over all 819 pairs at all three temperatures.

Grade: tautological. The stored de-excitation coefficients were produced from the excitation coefficients through Eq. (4.1) in the first place (`compute_K_CCC.py`, function

detailed_balance), as Section 2.3.2 states. Gate A confirms that the same formula was applied twice with the same energies and weights. That is a wiring check, and Chapter 2 already flagged it as one.

The physics content of detailed balance in this project sits somewhere else and is currently invisible in the outputs. `src/parsers/qc_ccc.py:87–168` tests whether Bray’s *raw* convergent-close-coupling (CCC) cross sections, which this project did not produce, satisfy microscopic reversibility on their own. The ratio of the two directions has mean 0.9995 with standard deviation 0.0032, and 98.7 % of comparisons fall within 1 % (`ccc_qc_report.png`). A 0.05 % mean bias with 0.32 % scatter on external data is a genuine external check, and it is the one to lead with.

4.3.2 One energy ladder on both sides, which is severe

There is a second detailed-balance result, and unlike Gate A it could have failed.

A model of this kind carries energies twice: once in the Boltzmann factors that relate forward and reverse collisional rates, and once in the Saha factors that relate ionisation to three-body recombination. If two different energy tables were used, or one table were indexed differently from the other, the model would still conserve particles, still satisfy Gate A, and still produce plausible populations. Nothing internal would notice.

Testing ionisation against three-body recombination through the Saha relation at every level, with CODATA constants on the Saha side, gives a ratio of 0.999997 for every one of the 43 states at every temperature, identical across all 2150 entries to 1.4×10^{-15} (`verify_saha_balance.py, validation/saha_balance/`); the residual is the rate module’s rounded values of h , m_e and the electron-volt against CODATA, and nothing else. The point is not the closeness to unity, which detailed balance imposes. It is the identity of the residual across levels whose Saha exponents span the full binding-energy ladder of hydrogen, from $e^{13.6 \text{ eV}/k_B T_e}$ for the ground state down to $e^{0.06 \text{ eV}/k_B T_e}$ for the top shell. At $T_e = 1 \text{ eV}$ those two factors differ by $e^{13.54}$, a factor of 7.6×10^5 . Two energy ladders that differed would produce a residual drifting across that range; the measured residual does not drift, and its common value is the eight-figure rounding of the stored state table. **Grade: severe.** The check could have failed and did not.

The same audit gives the collisional side over the 819 excitation pairs at each of three temperatures, 2457 comparisons in total: maximum deviation 7.31×10^{-9} , median 7.5×10^{-10} , again consistent with the state table’s rounding.

4.3.3 Gate B: the coronal limit

At low enough density an excited level is populated by electron impact from the ground state and empties radiatively before a second collision reaches it, so its population per unit ground-state density and per unit electron density becomes independent of n_e . Gate B tests that plateau: it computes $n_{2p}/(n_{1s}n_e)$ and $n_{3d}/(n_{1s}n_e)$ at the two lowest densities on the grid and asks whether they have converged.

The tolerance is 50 % per temperature (`validate_gates.py:175`) and the gate passes if 70 % of temperatures pass (line 181). Measured: all 50 temperatures pass, worst convergence 0.0034.

Grade: weak. Passing with a factor of 150 between the measured convergence and the tolerance means the gate rejects only a gross error. It does establish that the low-density end of the grid is in the coronal regime, which Chapter 5 relies on when it

reads the low-density column as ionising, and that is worth having. It is not evidence about rate magnitudes.

4.3.4 Gate C: what the Saha gate tests, which is not the Saha limit

Gate C is named for the high-density limit in which collisions dominate radiation and populations approach their Saha–Boltzmann values. It computes three quantities per temperature: whether n_{2p}/n_{1s} increases monotonically with density, whether it stays below the Saha value (with 5 % slack), and whether it approaches that value, defined as reaching 10^{-4} of it.

The pass flag uses the first two and **excludes the third** (`validate_gates.py`:250; the **approaching** column is computed at line 239, stored, printed at line 261, and never consulted). The exclusion is not what makes the gate weak, because the excluded column would pass: the measured approach fraction runs 0.0123 to 0.0155, comfortably above the 10^{-4} the criterion asks for. What makes it weak is that the criterion is 150 times slacker than the measurement. Reaching 10^{-4} of Saha is not approaching Saha, and neither is reaching 1.5 %.

The measurement itself is the useful part. At $n_e = 10^{15} \text{ 1/cm}^3$, the top of the grid, the $2p$ population sits at about 1.5 % of its Saha–Boltzmann value, so that level is nowhere near local thermodynamic equilibrium (LTE) anywhere on this grid. Only $2p$ was tested, so this is a statement about $2p$ and not about the whole distribution. It is the expected behaviour for an optically thin plasma whose ion population is an externally imposed reservoir rather than a Saha partner, and Chapter 5 depends on it: it is why the two supply channels of Section 3.5.3 never merge, and why $f_3 - f_4$ does not vanish at high density.

Grade: weak, and mis-named. What Gate C tests is monotonicity in density and a Saha ceiling. Both would fail on a sign error in the recombination source, which is worth something. Neither tests the limit the gate is named for.

4.4 Gate E: the timescale hierarchy

Gate E diagonalises $\mathbf{L}$ at every grid point, orders the negative eigenvalues, and forms $M = \tau_{\text{slow}}/\tau_{\text{relax}} = |\lambda_1|/|\lambda_0|$. It passes if $M > 1$ at 95 % of points (`validate_gates.py`:409–412).

Measured over all 400 points: M ranges from 86.77 at [49, 7] ($T_e = 10 \text{ eV}$, $n_e = 10^{15} \text{ /cm}^3$) to 1.72928×10^9 at the cold corner; τ_{slow} spans 75.4 ns to 67.2 s and τ_{relax} spans 0.87 ns to 38.9 ns (project working note; Table 4.4). The gate asks for $M > 1$. The smallest value anywhere is 86.8.

Grade: tautological as posed, and the real content is elsewhere. The threshold sits between two and seven orders of magnitude below the measured values. No matrix this construction has produced comes within a factor of 86 of failing, and none plausibly could, because the fast radiative rates and the slow ionisation rate that set the two eigenvalues are separated by construction of the physics rather than by a numerical accident. The reportable quantity is the range, not the pass, together with the fact that three independent implementations produce it bit-for-bit: `qss_analysis.py`, Gate E itself, and an unconditional recomputation stored in `timescales_unfiltered_CHECK.npz`.

That agreement is what caught the eigenvalue-filter error of Section 4.11, and a check that has actually fired is worth more than a threshold nothing approaches.

Two properties of these numbers must travel with them, and both come from Chapter 3. τ_{slow} is boundary-dependent, which Section 4.8 quantifies. And the slow eigenvalue is not an equilibration rate: measured against the ground-state ionisation rate coefficient $K_{\text{ion}}(1s)$, in cm^3/s , it satisfies $|\lambda_0| = 2.286 K_{\text{ion}}(1s) n_e$ across the grid (project working note; Table 4.4), so τ_{slow} is the collisional–radiative ionisation lifetime of the ground-state reservoir and not the time for anything to come to equilibrium.

4.4.1 The gates do not stop anything

No gate raises on failure. Each prints **PASS** or **FAIL** and returns a flag; `main` sets `all_pass = False`, prints a warning, and exits normally. **The script returns exit code zero whether or not Gate D fails.** The residual gates elsewhere in the project do raise; this one does not, and that asymmetry is a defect.

4.5 The checks that could have failed

Everything so far has been a demotion. This section is the other half, and it is short by design: a small number of severe checks is worth more than a long list of consistency tests, and these are the ones the thesis leans on.

4.5.1 The superposition residual

The entire mechanism of Chapter 5 rests on one structural claim: that the excited populations $\mathbf{n}_F$ of Section 3.4 split exactly into a ground-fed and a recombination-fed part, $\mathbf{n}_F = \mathbf{a} n_{1s} + \mathbf{c} n_{\text{ion}}$, where the network responses $\mathbf{a}$ and $\mathbf{c}$ of Section 3.5.3 are dimensionless and independent of the two reservoir densities. If that split is not exact, the ground-fed fractions f_p of Section 3.5.4 are not well defined and the closed form for the diagnostic error does not exist.

The check solves the full system three times at each grid point and each step direction, once with each reservoir alone and once with both, and measures the relative residual of the sum against the joint solve. Across the 784 grid-point and step-direction pairs it covers, out of the 800 the grid admits, the largest residual is 3.075×10^{-14} (`plateau_gridmap.txt`, line 19).

Grade: severe, for a specific reason. Feeding the check a corrupted input in which the $n = 3$ and $n = 4$ entries are swapped, which is the mistake that would most easily go unnoticed in a two-shell observable, trips it (project working note; Table 4.4). It is one of only two checks in the project reported to catch that swap.

4.5.2 The reduced-versus-full ratio test

The second is its companion. The observable $R = n_3/n_4$ can be computed two ways: from the full 43-state steady-state solve, and from the reduced two-channel form of Eq. (3.35) with the ground-fed fractions substituted in. The two routes share no code path beyond the matrix itself. At the benchmark point $T_e = 2.947 \text{ eV}$, $n_e = 1.389 \times 10^{14}/\text{cm}^3$, an independent re-implementation obtained $R = 0.777768$ by three routes against Chapter 3’s recorded 0.7778 (project working note; Table 4.4). A $3 \leftrightarrow 4$ swap breaks the agreement.

4.5.3 Pinned hashes on the inputs

The most common way for a computation like this to become wrong is not an algebra error. It is a stale array: a matrix regenerated after a correction, combined with an analysis output produced before it. This project has three retractions in its history and at least one of them had that shape.

`src/validation/preflight.py` pins the SHA-256 hashes, fixed-length fingerprints that change if any byte of a file changes, of the two canonical inputs and compares them byte for byte before any verification script runs. Recomputed for this chapter:

file	SHA-256 (first 16 hex digits)
<code>data/processed/cr_matrix/L_grid.npy</code>	2d92b58e1693107d
<code>data/processed/cr_matrix/S_grid.npy</code>	7822f536590c76ba

Both match the pinned baselines at `preflight.py:37–38` exactly, so the matrix on disk at the time of writing is the one the pins were taken from; the statement is only as good as the discipline of running the preflight before each analysis.

Grade: severe on the data, incomplete on the scripts. Of the six scripts the same preflight claims to check, `verify_ch3_groupB.py` and `verify_grid_coverage.py` do not exist on disk under those names; a missing script is downgraded to a non-blocking advisory (`preflight.py:258–262`) while the summary still prints “All checks passed”. The hash check is sound; the sentence after it overstates what ran.

4.5.4 Guards that recompute rather than trust

`src/analysis/make_ch5_figures.py` does not plot the stored results. It rebuilds every plotted quantity from $\mathbf{L}$ and the recombination source $\mathbf{S}$ of Eq. (3.5) and compares against the stored comma-separated tables at a relative tolerance of 10^{-9} , raising with the offending grid point, column, and matrix hash if any row disagrees (lines 393–432). It reads the temperature step size out of the file header rather than redeclaring it, so a figure cannot be drawn against a step convention different from the one that produced the data. It also verifies that the untrapped run of the radiation-trapping sweep reproduces the canonical matrix exactly, without which no trapped number in Chapter 6 would be interpretable as a difference.

The same script carries a negative control: it computes the correlation between M and the diagnostic error both raw and after controlling for temperature and density, and refuses to draw the figure if the controlled correlation fails to change sign (Section 5.9 draws the conclusion).

4.5.5 Two checks that fired, and one that was withdrawn

Three episodes say more than the pass/fail table, because in each the check did something.

The Lyman- α line-centre cross-section gate, an independent implementation in `escape_factor.py` against the project’s own Einstein coefficients, fired at 1156 % on its first run, located a 4π factor, and now agrees to 0.059 % against a raising 1 % tolerance (Section 6.1).

`verify_fujimoto_table41.py` carries a deliberately broken variant with proton-impact ℓ -mixing removed. It returns a recombination-fed population coefficient $r_0(2) = 104.5$, and the script flags it itself as unphysical on the ground that a level cannot exceed

Saha equilibrium. Production gives 0.7392. A check that fires on a known-bad input is the definition of a severe one.

Third, a semigroup propagator identity, $\exp(\mathbf{L}t_1)\exp(\mathbf{L}t_2) = \exp(\mathbf{L}(t_1 + t_2))$ for a matrix exponential computed by the same routine, returned 0.000×10^0 by construction and was *withdrawn by the author* (Section 5.7), the judgement this chapter applies to Gate A, applied earlier and unprompted.

4.5.6 The bound that is a theorem

Section 5.5 derives an upper bound on the diagnostic error of the form $\tanh(\Delta/4)$, where Δ is the separation between the two shells' switching points. A script reports that the bound is honoured at all 248 points tested, and that line reads as validation.

It is not. Here Δ is the dimensionless separation between the two shells' switching points, in units of the logarithm of the reservoir strength. Given $\mathbf{a}, \mathbf{c} \geq 0$, which Section 3.5.3 proves from the M-matrix structure of $-\mathbf{L}_{FF}$, the bound is a theorem: it cannot fail except on an arithmetic bug. Fed corrupted input it still passes. Swapping shells $3 \leftrightarrow 4$ sends $\Delta \rightarrow -\Delta$ and leaves both $|\Delta|$ and the logarithmic change in R that it bounds unchanged, so the gate cannot see it. Swapping $c_3 \leftrightarrow c_4$ also passes, but for a different reason: it moves $|\Delta|$ from 1.945614 to 0.939493, a factor 2.07, and the gate still passes because it compares $|f_3 - f_4|$ against $\tanh(|\Delta|/4)$ evaluated from the same corrupted $\mathbf{c}$. The bound is satisfied by whatever $\mathbf{a}, \mathbf{c} \geq 0$ it is handed. **Grade: tautological.** The script's own docstring calls it a wiring check; the line it prints does not, and the printed line is what a reader sees.

The bound is a result, and a useful one, because it holds for any atomic dataset whatever. It is not evidence that this atomic dataset is right.

4.6 Conditioning, precision, and state-space convergence

Two numerical questions sit underneath every result in this thesis. The elimination of Section 3.4 inverts a 42×42 matrix whose entries span many orders of magnitude, which invites the question of whether the inverse means anything. And the level ladder was cut at $n_{\max} = 15$, which invites the question of what the discarded levels would have changed.

4.6.1 The inverse is well behaved, and here is how well

The relevant quantity is the condition number of $\mathbf{L}_{FF}$, the block of $\mathbf{L}$ coupling the 42 excited levels to each other (Section 3.4), and it is the factor by which a relative perturbation of the matrix can be amplified in the solution. Computed from the canonical $\mathbf{L}$ at all 400 grid points by `src/validation/verify_operator_conditioning.py` and stamped to `validation/operator_conditioning/operator_conditioning.csv`: minimum 1.4821×10^3 at $[0, 0]$ ($T_e = 1.000$ eV), maximum 1.7436×10^5 at $[33, 7]$ ($T_e = 4.715$ eV), median 2.1044×10^4 . In double precision, which carries about sixteen significant digits, a condition number of 10^5 costs five of them. Eleven remain, and every conclusion in this thesis turns on the first three.

The range was first asserted in Chapter 3 on an unscripted working note, which is not a check; the stamped script reproduces it to 0.14 % and 0.21 %. The number was right

the whole time and was not evidence until it could be re-run.

Three further measurements bound the arithmetic. The eigenvalue condition number is a different quantity from the matrix condition number above: it is $1/|\mathbf{y}^H \mathbf{x}|$ for a left and right eigenvector pair, and it measures how far a single eigenvalue moves under a perturbation of the matrix. For the two eigenvalues the thesis uses it is of order unity (1.71 and 1.98 at the benchmark point, 3.26 and 2.59 at the cold corner), despite a largest absolute column sum of 2.0×10^{14} 1/s; that large norm is a statement about the fast radiative modes and does not propagate into λ_0 or λ_1 . Recomputing the two supply channels in single precision moves them by 1.3×10^{-5} , so double precision is nowhere near a cliff. And tightening the integrator changes nothing that matters: a relative tolerance of 10^{-9} against 10^{-12} shifts the result by 1.20×10^{-8} , and computing the propagator by a dense matrix exponential against a Krylov action differs by 1.13×10^{-10} .

The eigenvalue condition numbers and the single-precision test cover three of 400 points; the tolerance comparisons cover two. The condition numbers and the displacement are recomputed under the definition $1/|\mathbf{y}^H \mathbf{x}|$ with unit-norm left and right eigenvectors by `verify_residual_claims.py` (`validation/residual_claims/`). They are point measurements, not grid scans, and are quoted as such.

4.6.2 Where the ladder was cut

The quantity that has to converge is not a population but the difference in ground-fed fraction between the two Balmer shells, $f_3 - f_4$, since Chapter 5 shows the diagnostic error to be proportional to it.

Rebuilding $\mathbf{L}$ with successively lower $n_{\max}$, from 15 down to 9, with the loss channels of every removed state returned to the diagonal of the survivors, and recomputing $f_3 - f_4$ at each truncation's own u^{CRE} (`verify_nmax_downward_scan.py`, `validation/nmax_downward_scan/`) gives increments of one sign at the three named points, with the last measured ratio $d(15)/d(14)$ equal to 0.88 at the benchmark, 0.72 at the cold corner $[0, 0]$ and 0.33 at the ridge $[15, 3]$; at the 245 of the 280 defended points where the increments are one-signed and geometric the ratio has median 0.67. Summing the geometric tail with that ratio gives an extrapolated truncation error of -0.46% at the benchmark, 1.6% at the cold corner and 0.005% at the ridge; holding u at its $n_{\max} = 15$ value instead gives -0.38% , 1.4% and 0.005% , so the sign does not depend on that choice. A historical project value with no producing script, not used as evidence here and superseded by `verify_nmax_downward_scan.py`, records 0.9% , -1.4% and -0.55% with a common ratio of about 0.75; the stamped run reverses both signs and the conclusion that rested on them. Truncation makes $f_3 - f_4$ slightly too large at the benchmark and too small at the cold corner, so it is not a bias that can be corrected by a single factor. Over the defended points where the tail can be summed the extrapolated error has median 0.002% and runs from -0.82% to 1.03% , smaller than the spread between the two step conventions of Section 5.3 at every such point. At the 35 defended points of the density column $n_e = 5.2 \times 10^{13}/\text{cm}^3$ the increments change sign among $n_{\max} = 12$ to 15 and no tail can be summed; the last measured increment there is below 0.041% of $f_3 - f_4$ at every point.

A trap in this test. The first attempt dropped the top shells from the state vector without removing their loss channels from the diagonal of the surviving states, which then lost population to states that no longer existed, and $f_3 - f_4$ jumped. The stamped scan

repeats the mistake deliberately: dropping $n = 15$ alone without the correction moves $f_3 - f_4$ by 11.5 % at the benchmark, 12.6 % at the cold corner and 5.0 % at the ridge, where the correct truncation to $n_{\max} = 14$ moves it by 0.06 %, 0.64 % and 0.01 %; dropping six shells without it gives 17.9 %, 25.9 % and 13.3 % (a working note recorded 22 to 37 %). The jump is an artifact of the truncation procedure, not a dependence on $n_{\max}$. Removing state k requires adding $\sum_k \text{dropped } L_{ki}$ back to L_{ii} for every surviving i , as the production matrix does. Anyone repeating this test will reproduce the spurious number first.

The terminal shell is over-populated, and by a measured amount. Truncating the ladder at $n_{\max}$ removes, from the balance of the terminal shell, the collisional excitation out of it to $n > n_{\max}$ and the de-excitation and cascade into it from those levels (cascade is downward, so what is lost is a way out, not something “to cascade out to”). The loss removed is not small. Continuing the $12 \rightarrow 13$, $13 \rightarrow 14$ and $14 \rightarrow 15$ excitation coefficients geometrically, the missing $15 \rightarrow 16$ channel would carry 1.30 times the whole retained loss out of $n = 15$ at the benchmark and 1.20 times at $[0, 4]$, between 1.20 and 1.35 times at every grid point, while radiative decay is below 10^{-6} of that loss and ionisation is 7 to 10 % of it (`validation/terminal_shell_budget/`). The top shell sits within 6×10^{-5} of Saha–Boltzmann, so for the total population the removed exchange is nearly balanced and the sign of the bias is decided channel by channel: the ground-fed channel, whose net flux through the top of the ladder is upward, loses only an exit and can only be biased high, and it is the ground-fed coefficient $r_1(15)$ that the external comparison finds high while $r_0(15)$ agrees (Section 4.9). Comparing each shell’s ground-fed coefficient a_p as terminal shell (operator rebuilt with $n_{\max} = p$, the loss channels into the removed states kept on the diagonal) with its value in the full $n_{\max} = 15$ operator gives at the Fujimoto point $[49, 0]$ (10 eV, $10^{12}/\text{cm}^3$) excesses of 9.4 at $p = 8$, 10.7 at $p = 9$ and 6.3 at $p = 10$, then 4.4, 3.0, 2.1 and 1.4 at $p = 11$ to 14 (`verify_terminal_vs_interior.py`, `validation/terminal_vs_interior/`). At the benchmark $p = 8$ to 10 give 14.2, 12.3 and 6.5; over the 280 defended points the $p = 8$ excess spans 8.2 to 17.7 and the $p = 10$ excess 5.4 to 7.5. The excess is confined to the ground-fed channel: the shell’s total population and recombination-fed coefficient move under 1 % at the benchmark and at most 32 % and 30 % at the low-density named points. A historical project value with no producing script, not used as evidence here and superseded by `verify_terminal_vs_interior.py`, records 9.4, 10.7 and 6.3 with neither quantity nor grid point named; the stamped run reproduces them to 0.4 % as a_p at $[49, 0]$. The absence of a monotone trend over $p = 8$ to 10 holds there and at 47 of the 280 points; from $p = 9$ upward the excess falls at every grid point, at the benchmark from $p = 8$. The external excess at the terminal shell, 4.9 to 6.2 against published values (Section 4.9), lies outside the 6.3 to 10.7 band, not inside it as the working note states, but within a factor 2.2 of every value in it, and both rise with density (4.9 to 6.2 externally, 9.4 to 17.7 at $p = 8$ internally at 10 eV). That agreement in magnitude and in the sign of the density trend makes truncation a sufficient explanation for the $n = 15$ discrepancy without invoking the atomic data. The fall with p cannot be extrapolated to $p = 15$: the reference for $p = 14$ is itself a ladder with one shell above it, so the internal measurement neither predicts nor excludes 4.9 to 6.2 there. No result in this thesis rests on the top shell: the observable is built from $n = 3$ and $n = 4$, both fully resolved in orbital angular momentum ℓ and eleven shells below the cut at $n_{\max} = 15$.

What has now been checked. Above $n = 8$ the shells are bundled and their internal ℓ -distribution is assumed statistical rather than computed. The script that tests that

assumption, `verify_bundling_psm20.py`, had never been executed; two of the three defects found on its first run, the synthesised temperature grid and the identically zero ℓ -mixing array for the bundled indices, were predicted here. Section 6.8.3 reports all three defects and the result.

4.7 Is the shell ratio the line ratio?

Every analytic result in Chapter 3 is written for the shell ratio $R = n_3/n_4$, because the two-channel decomposition applies to shell populations. A spectrometer does not measure shell populations. It measures the H_α and H_β line intensities, each of which is a sum over the electric-dipole-allowed sublevel transitions, weighted by their Einstein coefficients. Section 2.6 established that these are not proportional: $4f$ holds 43.75 % of the $n = 4$ shell by statistical weight and contributes nothing to H_β , since $4f \rightarrow 2p$ would change ℓ by two and is forbidden.

The substitution therefore has to be measured, not assumed.

The line intensities are built as A -weighted sums over resolved channels, H_α from $3s \rightarrow 2p$, $3p \rightarrow 2s$ and $3d \rightarrow 2p$, and H_β from $4s \rightarrow 2p$, $4p \rightarrow 2s$ and $4d \rightarrow 2p$, with populations taken from the solve. Nothing statistical is assumed anywhere in that construction: searching the two producing scripts for statistical-weight expressions returns no matches, and the A -value lookup raises unless exactly one row matches each transition.

Two measurements follow. The $4f$ fraction of the $n = 4$ shell runs from 0.4361 to 0.4375 across the grid, against the statistical value $14/32 = 0.4375$; proton-impact ℓ -mixing drives the shell to a statistical distribution to better than 0.3 % everywhere, with the departure concentrated at $n_e = 10^{12}/\text{cm}^3$ where ℓ -mixing weakens and radiative decay competes. That fraction is a property of the full equilibrium population; the quantity that sets the line-versus-shell difference below is the ℓ -distribution within each supply channel separately, which is not statistical at the low-density edge. And the ratio of the transient error computed on the line ratio to the same error computed on the shell ratio runs from 0.9482 to 0.9999 over the 784 point-and-direction pairs, below 1 at every one: **worst case 5.2 %, at the heating step from 9.54 to 10.0 eV at the lowest density** (`verify_weighted_census.py`, stamped to `validation/weighted_census/`). At $[0, 4]$, the point carrying the largest error anywhere, the two agree to 0.06 % (0.386683 against 0.386903). Because the ratio is below 1 everywhere, the shell census of Chapter 5 is conservative for the line observable: counted on the line ratio, 44 instead of 45 of the 448 pairs above 2 eV exceed 10 % at 100 μs , with the worst case 0.1746 instead of 0.1748 at the same point. The 5.2 % scales inversely with the proton ℓ -mixing rate: with the within-shell ℓ -changing rates scaled by a common factor and the column sums of the operator preserved, the corner departure is 10.0 % at half the rate, 2.6 % at double and 1.1 % at five times, while the benchmark line-to-shell ratio stays between 0.9987 and 0.9999 (`validation/weighted_census/`, ℓ -mixing sensitivity block and `weighted_census_lmix.csv`); and it sits on the grid boundary.

That 0.06 % is a single-point figure and must not be quoted as a grid-wide one. The grid-wide statement is the 5.2 %.

4.7.1 And is the plasma transparent to its own Balmer light?

The shell-to-line substitution assumes the emitted photons escape. For the Lyman series that assumption fails at the cold edge and Chapter 6 treats it at length. For the Balmer

lines themselves it closes here, because it bears on the observable rather than on the matrix.

Section 6.1 defines the optical depth τ and the escape factor and treats the Lyman series, where transparency fails. For the Balmer lines the absorbers are the $n = 2$ atoms, $n_{2s} + n_{2p}$ in collisional-radiative equilibrium with $n_{\text{ion}} = n_e$, the Doppler width is taken at $T_n = T_e$, and the multiplet oscillator strengths are $f(2 \rightarrow 3) = 0.641$ and $f(2 \rightarrow 4) = 0.119$ [37]. The line-centre optical depth of H_α across a 10 cm slab is then 4.9×10^{-7} at the cold corner, 8.2×10^{-4} at $[0, 4]$ and 0.214 at $[0, 7]$, the single worst cell on the grid; resolving the absorbers into $2s$ and $2p$ with the model's own A coefficients changes these by at most 4%. At $[0, 7]$ the population escape factor for a photon born at the slab centre ($\kappa_0 D/2 = 0.107$) is 0.928 for H_α and 0.990 for H_β , so the observed ratio is lowered by 6.3% in that one cell; over the full path ($\tau = 0.214$) the figures are 0.861 and 12%. Above 2 eV the largest H_α optical depth is 0.010 at $[15, 7]$ and the ratio moves by less than 0.7% anywhere (`verify_balmer_optical_depth.py`, `validation/balmer_optical_depth/`). The Balmer lines are optically thin over the region where this thesis makes quantitative claims. A historical project value with no producing script, not used as evidence here and superseded by `verify_balmer_optical_depth.py`, records 6.3×10^{-7} , 1.0×10^{-3} and 0.263 for the same three cells; those values are 1.23 to 1.28 times the stamped ones and no convention tested recovers them, and the escape factor above 0.85 and ratio shift of at most about 10% recorded with them are half-slab figures attached to a full-slab optical depth. The conclusion holds with either set of values.

4.8 Two tests of the reduction, and their expected outcomes

The checks so far ask whether the model computes what it intends. This section asks something harder: whether two specific structural choices manufacture the result. In each case the size of effect that would have destroyed the claim was written down before the test was run, which is what makes the outcome informative rather than reassuring.

4.8.1 Does the open boundary manufacture the slow clock?

$\mathbf{L}$ describes an open system: atoms that ionise leave the accounting, the recombination feed $\mathbf{S}n_{\text{ion}}$ enters as a constant source rather than a matrix element, and there is no forty-fourth state for ions to accumulate in. Since the persistence claim of Section 5.8 rests entirely on τ_{slow} exceeding the duration of the disturbance, the first objection is that τ_{slow} is an artifact of the open boundary.

The falsifier, named first. Write τ_{drive} for the duration of the disturbance, fixed at 100 μs throughout this thesis as a representative edge-localised-mode (ELM) time (Section 5.8). At $[0, 4]$, the worst point on the grid, $\tau_{\text{slow}}/\tau_{\text{drive}} \approx 2300$. If closing the system shortened τ_{slow} by that factor, the slow clock would fall to τ_{drive} itself and the persistence would vanish. A factor of 2300 refutes the claim.

Two grid labels appear in the table for the same worst case: the error is evaluated at $[0, 4]$, but the operator governing the decay after the step is the post-step one, $\mathbf{L}[1, 4]$, and every timescale in this thesis names the operator it belongs to (Chapter 5).

Table 4.2: The slow timescale τ_{slow} computed with the ion treated as an external reservoir (open) and with a forty-fourth ion state appended (closed), at four grid points. Over the whole grid the ratio has minimum 1.00, median 1.00 and maximum 41.4: the boundary condition matters only where τ_{slow} already exceeds the disturbance duration by three orders of magnitude. Every entry is reduced from `validation/ion_closure/ion_closure_summary.csv` (`verify_ion_closure.py`) by `verify_headline_provenance.py`; the open column is the float64 eigensolver, which at the cold corner differs from the artifact’s 40-digit path by one part in 10^4 .

point	τ_{slow} open	τ_{slow} closed	factor
cold corner [0, 0]	67.23 s	1.6226 s	41.4
post-step operator $\mathbf{L}[1, 4]$	0.233 24 s	0.015 173 s	15.4
ridge [15, 3]	2.027 ms	1.999 ms	1.01
benchmark [23, 5]	22.728 μs	22.708 μs	1.00

The manifold was closed by appending an ion state,

$$\mathbf{L}_c = \begin{pmatrix} \mathbf{L} & \mathbf{S} \\ -\text{colsum}(\mathbf{L}) & -\sum \mathbf{S} \end{pmatrix}, \quad (4.2)$$

so that every ionising atom lands in the new state instead of being deleted. Each column of $\mathbf{L}_c$ sums to zero to machine precision, and the construction reproduces the closed-system value at the cold corner recorded in this project’s working notes, 1.6226 s, to five digits, which confirms it is the same closure rather than a new one. The measured factors are given in Table 4.2.

The shape of the distribution carries the argument: the factor is large only where τ_{slow} already exceeds τ_{drive} by about a thousand, so the error has already saturated, and it is unity everywhere τ_{slow} approaches τ_{drive} . Recomputing the entire persistence census with closed-system τ_{slow} moves the breakdown count from 202 to 200 of 680 and the worst case from 0.38682 to 0.38563, about one part in 325; restricted to $T_e \geq 2$ eV, where Chapter 5 makes quantitative claims, the count is 45 before and 45 after.

The falsifier did not appear. A factor of 2300 was needed; the measured factor at that operator is 15.

The caveat, which must not be softened. The closure adds an ion reservoir; it does not close the electron balance, since n_e is held fixed while the ion population moves, and it adds no transport. It answers the boundary condition on the atomic manifold and nothing else: Section 6.2 states the objection it cannot answer, and Section 5.10.1 the one it makes sharper rather than softer.

4.8.2 Is the single-slow-mode estimate a bound on the time-averaged error?

Chapter 5 evaluates the error at the partial-equilibrium plateau and builds from it the single-slow-mode estimate of Eq. (5.16), which multiplies the plateau error $\varepsilon_{\text{plateau}}$ of Section 5.1.1 by the fraction of τ_{drive} over which an exponential decay on τ_{slow} keeps it alive. It was first written as a lower bound, on the assumption that the observable cannot decay faster than the slowest mode. That assumption is testable, and an estimate quoted as though it were the answer is a claim about accuracy that has to be measured.

Table 4.3: The plateau error, the true error averaged over 0 to 100 μs from direct eigen-propagation of the full 43-state system, and the single-slow-mode estimate of Eq. (5.16), at five grid points. The averages are integrated at 400 time samples per decade; at 1600 per decade the five change by at most 4.7×10^{-6} relative, and the $[0, 4]$ average by 2.7×10^{-10} . Integration on a uniform 25 ns mesh, which under-resolves the nanosecond rise, reads low by 6×10^{-5} to 5×10^{-4} relative. The estimate reproduces the true average to better than 0.8 % at the five points. The value below unity at $[0, 4]$ is real: a deficit of 2.2×10^{-5} that survives the denser grid. Source: `verify_trajectory_census.py`, `validation/trajectory_census/`.

point	$\varepsilon_{\text{plateau}}$	true 100 μs average	estimate, Eq. (5.16)	true/estimate
$[0, 4]$ worst case	0.386903	0.386811	0.386820	0.99998
$[15, 3]$ ridge	0.180728	0.175283	0.174812	1.00270
$[15, 5]$	0.103550	0.072185	0.071722	1.00646
$[23, 5]$ benchmark	0.063612	0.011796	0.011712	1.00722
$[0, 0]$	0.083004	0.083084	0.083004	1.00096

Propagating the full 43-state system through the step by direct eigen-decomposition and integrating the error over 0 to 100 μs gives the comparison in Table 4.3.

At the five points tested the estimate sits within 0.8 % of the directly propagated average, and at $[0, 4]$ it sits above it, so it is an estimate and not a bound. The accuracy at these five points does not generalise: over the 448 pairs above 2 eV the ratio of true average to estimate departs from unity by up to 2.2 % at 100 μs and 3.8 % at 506 μs (Section 5.8). What matters here is that the assumption underneath it was tested rather than asserted.

The reason is structural, and it is a property the shell ratio has and another observable in this same project does not. The spectrum of $\mathbf{L}$ has one slow eigenvalue and then a band, with nothing between: at $\mathbf{L}[1, 4]$ the three slowest eigenvalues are -4.29 , -1.60×10^8 and -3.44×10^8 1/s, a gap of 3.7×10^7 . An observable that depends on the excited populations only through their ratio cannot decay faster than τ_{slow} : once the manifold is quasi-steady the excited vector is a function of u alone, so $\ln R - \ln R^{\text{CRE}} = (S/u^*) \delta u + O(\delta u^2)$ with $\delta u \propto e^{-t/\tau_{\text{slow}}}$, and the leading term vanishes only where $S = f_3 - f_4 = 0$. It does not: S is positive at all 400 grid points, with minimum 0.046 at $[44, 7]$, and the slow eigenvector’s projection onto $\ln R$ equals S/u^* to 1.2×10^{-4} at the benchmark and 1.6×10^{-8} at $[0, 4]$, the deviation being the first-order quasi-steady correction $\lambda_{\text{slow}}/\lambda_2$ itself. It exceeds 1 % only at the two hot, dense points $[48, 7]$ and $[49, 7]$, where the separation is 87-fold and which the window criterion excludes from the analysed set (`validation/slow_mode_projection/`). That decay on an intermediate mode was the failure mode being hunted, and it does occur for the ground-normalised error measure, which does not cancel the reservoir and whose 1/e time is 2.58 s against $\tau_{\text{slow}} = 67.2$ s at the cold corner $[0, 0]$. The estimate is accurate for the ratio and not for that measure; why it is nevertheless not a bound, the denominator relaxing on the same τ_{slow} , is set out in Section 5.8.

4.9 External comparison: Fujimoto’s population coefficients

Everything to this point compares the model against itself or against arithmetic. Only comparison against an independent calculation reaches the atomic physics, and this project has one such comparison in depth: against Fujimoto’s collisional–radiative treatment of atomic hydrogen [11], computed decades earlier, with different atomic data, by a different method.

4.9.1 Timescales, at a point that lies exactly on this grid

Fujimoto works a numerical example at $T_e = 10\text{ eV}$ and $n_e = 10^{18}/\text{m}^3$, which is $10^{12}/\text{cm}^3$ and lands exactly on grid point $[49, 0]$. No interpolation is involved.

quantity	Fujimoto App. 4B	this model, $[49, 0]$	ratio
excited-state response	$\sim 1 \times 10^{-7}\text{ s}$	$3.4057 \times 10^{-8}\text{ s}$	0.34
ground-state depletion	$\sim 1 \times 10^{-4}\text{ s}$	$1.7657 \times 10^{-4}\text{ s}$	1.77

The two-clock structure appears in both, separated by three orders of magnitude in both, with the fast time short and the slow time long in the same sense.

What this comparison is not. Fujimoto’s response time is $\max_p t_{\text{tr}}(p)$, the 63% rise time of the slowest-responding level, which is not $1/|\lambda_1|$. The two are different functionals of the same spectrum, so agreement to a factor of a few is the meaningful comparison and exact agreement would be suspicious. Quoting these ratios as though they were a precision test would overstate them.

One qualification on the independence of this check, found on reading the printed appendix rather than a summary of it. Fujimoto’s Figs. 4B.1 and 4B.2 are captioned as quoted from Sawada and Fujimoto (1994) [19], so the appendix and the paper Chapter 7 discusses are the same calculation, not two. **The comparison is therefore independent of this thesis but not two independent checks**, and it is counted once. The appendix also states the same construction the paper does, that $n(1)$ is held constant, which matters for Chapter 7 and is taken up there.

A second, independent route corroborates the same physics. Fujimoto’s boundary level, the principal quantum number above which collisions dominate radiative decay, sits at $p_G \approx 4\text{--}5$ at this density. Tracing the boundary descent in this model puts it at $5g$ at the low-density end of the grid, which is the crossing Chapter 2 deferred to this chapter (Section 2.6). Two independently computed models place the same structural feature within one shell of each other, at one density.

4.9.2 The recombining coefficients agree

Fujimoto’s Table 4.1 tabulates the two population coefficients of the decomposition that Section 3.5.3 derives: $r_0(p)$, how strongly level p is fed from the ion continuum, and $r_1(p)$, how strongly it is fed from the ground state. Both are dimensionless, and both are normalised to the Saha–Boltzmann population of level p , which is why r_0 is of order unity while r_1 at low density is of order 10^{-5} , and why a value of r_0 above one is unphysical.

The recombining coefficients agree to 0.1 to 1% for $n \geq 4$. The discriminating case is $r_0(2)$, where cascade and ℓ -mixing structure matter most and a hydrogenic approximation

would go wrong: this model gives 0.7392 against the tabulated 0.730, a difference of 1.3 %, a test the model could have failed and did not (the broken variant of Section 4.5 drives the same quantity to 104.5).

4.9.3 The ionising coefficients at low p disagree, and the target is not verified

At $n_e = 10^{12}/\text{cm}^3$ this model's r_1 sits a factor 8.3 low at $p = 3$ and 4.4 low at $p = 4$ against Fujimoto's Table 4.1(b). Those are precisely the two shells whose difference is the mechanism of Chapter 5, so the discrepancy cannot be dismissed as peripheral.

Reading it as a failure of the model is not supportable, for six reasons, and the status is an open comparison whose *target* is unverified.

1. **The model's r_1 is reproduced by hand.** In the coronal limit $r_1(2)$ follows from the single excitation rate coefficient $K_{\text{exc}}(1s \rightarrow n=2) = 9.2796 \times 10^{-9} \text{ cm}^3/\text{s}$ and gives 1.37×10^{-5} , against the full 43-state solve's 1.66×10^{-5} . The two agree to 21 %, which is the accuracy a coronal estimate that ignores cascades should have. Whatever produces a factor of 8 is not inside the solve.
2. **The atomic stack passes three independent external checks.** The radiative recombination coefficient into the ground state is within 2 % of the scaled standard value; the sum over excited states, $2.08 \times 10^{-13} \text{ cm}^3/\text{s}$, sits 10 % below the standard case-B coefficient $\alpha_B \approx 2.31 \times 10^{-13} \text{ cm}^3/\text{s}$, the accepted total for capture into every excited level, and that 10 % shortfall is exactly the $n > 15$ truncation of Section 4.6.2; and the ground-state ionisation coefficient at 10 eV, $4.864 \times 10^{-9} \text{ cm}^3/\text{s}$, sits 10.4 % below ADAS SCD96 read in its low-density limit, $5.43 \times 10^{-9} \text{ cm}^3/\text{s}$ at $n_e = 5 \times 10^7/\text{cm}^3$ [8, 9], where the effective coefficient reduces to ionisation out of the ground state because every excited level decays before it can be ionised. The comparison is at the low-density limit deliberately: SCD96 varies by 0.9 % over a factor four in density there, so the limit is reached. **This 10 % is not the truncation of the item above.** Direct ionisation from $1s$ is a single cross section and does not know where the level ladder stops; the difference is between this model's ionisation cross-section source and the one behind SCD96, and it is reported rather than reconciled.
3. **The tabulated asymptote demands an excitation coefficient fifteen times the accepted value.** Reproducing Fujimoto's quoted coronal asymptote requires $K_{\text{exc}}(1s \rightarrow n=2) = 1.44 \times 10^{-7} \text{ cm}^3/\text{s}$ at 11.03 eV, fifteen times the accepted value. No error in this model's implementation produces that number; only a different input does.
4. **Three candidate explanations are eliminated quantitatively.** Truncation moves $r_1(3)$ by 1.2 % between $n_{\text{max}} = 10$ and $n_{\text{max}} = 15$, which is a factor 8 short. Lyman trapping at an escape factor of 0.1 brings $r_1(2)$ to 0.92 of the table but leaves $r_1(3)$ at 0.26 of it and drives $r_0(2)$ to 8.2, a population above Saha equilibrium and therefore unphysical. The third candidate, the ℓ closure, is eliminated in the item below.
5. **The ℓ closure is a real difference, and it does not account for the gap.** Table 4.1(b) is indexed by $p = 2, 3, 4, 5, 7, 10, 15$, principal quantum numbers, and

carries no ℓ label: the published coefficients are shell quantities computed with the ℓ substructure bundled, while this model resolves ℓ below $n = 9$. That difference is real, and it is not the explanation. Imposing the source's own closure directly, by bundling this model's operator under statistical ℓ weights as $\mathbf{L}_{\text{bundle}} = \mathbf{P} \mathbf{L} \mathbf{R}$ with $\mathbf{P}_{ki} = 1$ for i in shell k and $\mathbf{R}_{ik} = g_i/g_k$, moves $r_1(3)$ by 0.42 %: the factor 8.29 becomes 8.26 (`verify_fujimoto_bundle.py`, stamped in `validation/fujimoto_bundle/`). At $p = 2$ the agreement becomes *worse*, from a factor 10.8 low to a factor 11.8 low, and $r_0(2)$ moves from 0.7392, which agrees with the tabulated 0.730, to 0.6459, which does not. The reason is structural: the bundling projector annihilates the ℓ -mixing operator identically, $\max |\mathbf{PBR}| = 8.5 \times 10^{-7}$ against a matrix scale of 1.0×10^{11} , so bundling cannot express what ℓ -mixing does; and the production model is already at the statistical- ℓ limit for $p \geq 3$, its sublevel populations sitting within 5 % of $g_{n\ell}/g_n$. The ℓ -resolved calculation was therefore already the like-for-like comparison. Removing proton ℓ -mixing altogether moves $r_1(2)$ by a factor 118, but that is the opposite limit from a statistical closure rather than a proxy for it, and the two do not bracket the tabulated value.

6. **The transcription and the row labels are correct.** Both were checked against the printed table. The $\log_{10} n_e = 18$ row reads 0.730, 0.835, 0.947 and 0.983 for r_0 at $p = 2, 3, 4, 5$, and 1.79×10^{-4} , 5.72×10^{-5} , 1.66×10^{-5} and 5.08×10^{-6} for r_1 , matching the values used here exactly. The row labels are $\log_{10}(n_e/1/\text{m}^3)$, which the table establishes internally: $r_0(2)$ reaches 0.981 at $\log_{10} n_e = 21$, and Griem's criterion places local thermodynamic equilibrium for $n = 2$ at $1.74 \times 10^{16}/\text{cm}^3$, which is $10^{21} 1/\text{m}^3$ to within the criterion's own precision. Read as $1/\text{cm}^3$ the onset would sit five orders above Griem, which is impossible; and the rows run from $\log_{10} n_e = 12$ to 24, a span of 10^6 to $10^{18} 1/\text{cm}^3$ read as $1/\text{m}^3$ and no plasma at all at the top read as $1/\text{cm}^3$. A one-decade offset in the density rows is not a tenable reading either.

Status. The r_1 deficit at low p is an open external disagreement, and this thesis does not explain it; the three candidate causes above each fail to account for a factor of 8. The ℓ -closure test carries a severity check: bundling an operator whose ℓ distribution is deliberately non-statistical moves $r_1(2)$ by a factor 129, so the null result at $p = 3$ is a measurement and not a blind instrument. What the section does *not* establish is that the disagreement leaves the quantity Chapter 5 uses untouched. Equation (3.47) makes $a_p = r_1(p) Z_p/Z_1$: r_1 is the ground-fed coefficient from which f_p is built, so a deficit in $r_1(3)$ and $r_1(4)$ is a deficit in a_3 and a_4 themselves. Bundling is not the protection, since neither shell passes through a bundled coefficient; the only protection is that Δ and the cap depend on the ratio a_3/a_4 , in which a deficit common to both shells cancels. At $10^{12}/\text{cm}^3$ the deficit is not common, 8.3 at $p = 3$ against 4.4 at $p = 4$; at $10^{15}/\text{cm}^3$ it nearly is, 1.8 against 1.9. Dividing the model's a_3 and a_4 by those ratios, so that they become what Fujimoto's coefficients would give, and recomputing at every grid point (`verify_fujimoto_r1_consequence.py`, `validation/fujimoto_r1_consequence/`) moves Δ by +0.64 under the first pair and by -0.08 under the second. Over the 280 points above 2 eV the cap then moves by 29 % at the median under the first and -4 % under the second; the crest u^{peak} falls by a factor 6.0 and 1.9; and the sensitivity at the operating point, $f_3 - f_4$ evaluated at u^{CRE} , moves by 30 % and 11 % at the median, with ranges across the points of -58 to 524 % and -38 to 68 %. Both exceed the 8 to 9 % that the R-matrix excitation substitution of Chapter 2 produces. The disagreement is therefore direct evidence about the coefficients Chapter 5 uses; what is open is whether the deficit lies in this model or in the tabulation, and the

scenario applies a deficit measured at 10 eV at every temperature, which is its stated assumption. An unexplained disagreement in an absolute population coefficient is reported here as unexplained, with its stake stated.

What the benchmark does establish. The recombination-fed channel agrees, and because r_0 and r_1 carry the same p -dependent factor $Z(p)$, a normalisation error large enough to explain the r_1 deficit is excluded by r_0 agreeing. That is a pass, and it should be read as one.

What it does not establish, and two limits on it. It does not fix the absolute normalisation of the ground-fed channel against an independent calculation, because no ℓ -resolved tabulation was available to compare against. Table 4.1(b) is computed at $T_e = 1.28 \times 10^5$ K, which is 11.03 eV, while this grid reaches 10 eV, so the comparison sits at the grid edge, with the tabulated temperature 10.3 % above it. The companion Table 4.1(a) is at 10^3 K = 0.0862 eV, two orders below this grid, so there is no second usable case. The external check is one temperature at one edge.

4.10 Gate D fails, and the failure is in the gate

The Atomic Data and Analysis Structure (ADAS) distributes effective stage-to-stage rate coefficients for hydrogen: SCD, the effective ionisation coefficient, and ACD, the effective recombination coefficient [8]. These are the coefficients a transport code uses when it does not resolve excited states, and they are the most widely used external reference this model could be compared against.

Gate D forms the model’s counterpart by taking the steady-state population n_p^{ss} of each level at each grid point, weighting it by that level’s ionisation rate coefficient Q_p of Eq. (3.1), and dividing by the total neutral population:

$$\text{SCD}_{\text{model}} = \frac{\sum_p Q_p n_p^{\text{ss}}}{\sum_p n_p^{\text{ss}}}, \quad (4.3)$$

then compares against ADAS interpolated to the same temperature within a factor-of-two density band. The gate passes a point if the ratio $\eta = \text{SCD}_{\text{model}}/\text{SCD}_{\text{ADAS}}$ lies between 0.5 and 2.0.

Gate D: FAIL (0% of points)

validation/gate_summary.txt

Measured, η ranges from 5.73 to 8346 over the full grid, and no point agrees within a factor of two. The extremes are not typical: over most of the grid the model sits high by a factor between 9 and 65, and that systematic offset peaks near $T_e \approx 2$ eV (these are the gate’s own numbers, under its linear read of the table; the read is discussed below). The disagreement has a shape, which is what makes it diagnosable rather than random.

4.10.1 The diagnosis, and the test that settles it

Two readings were open. One of them has now been tested and it is the explanation.

The first is that Eq. (4.3) is not the quantity ADAS tabulates. ADAS defines SCD as the *ionising* coefficient, the ionisation flux attributable to the ground-state population. The sum in Eq. (4.3) runs over the full steady state, which includes the recombination-fed Rydberg levels, the $\mathbf{n}^{(0)} = \mathbf{c} n_{\text{ion}}$ channel of Section 3.5.3. Ionisation out of those levels is a recombination-driven flux and belongs under ACD, not SCD. Since the high levels have

small ionisation potentials and large ionisation coefficients, moving them to the wrong side of the ledger inflates the model's SCD by a large factor, and inflation is what is measured. The test that would settle it is to recompute Eq. (4.3) over the ground-fed channel $\mathbf{n}^{(1)}$ alone, and it has now been run (`diagnose_gate_d.py`).

The two channels separate exactly, with no approximation: $\mathbf{n}_E = \mathbf{n}^{(0)} + \mathbf{n}^{(1)}$ to a relative residual below 10^{-8} at every grid point, which the script checks before using either. Rebuilding the coefficient on the ground-fed channel alone,

$$\text{SCD}_{\text{model}} = Q_{1s} + \sum_{p \neq 1s} Q_p \frac{n_p^{(1)}}{n_{1s}}, \quad (4.4)$$

and comparing against the same ADAS table with the same factor-of-two tolerance moves the agreement from 0 of 400 points to 400 of 400. The ratio η falls from a range of 5.73 to 8346 to a range of 0.702 to 1.004, with median 0.828; above 2 eV it runs from 0.721 to 0.937, median 0.831, and at the benchmark it is 0.768 (`validation/reservoir_vs_adas/`). The read matters. ADAS96 tabulates seven temperatures between 1 and 10 eV, and `diagnose_gate_d.py`, which first made this comparison, interpolates linearly in value between them (`adas_at`, lines 114 to 124); across 1 to 1.5 eV, where SCD rises by nearly two orders of magnitude, that overstates the table by up to a factor 6.65. That read gives 323 of 400, a range of 0.130 to 0.931 and a median of 0.696, with a deficit deepening toward low temperature; read logarithmically in both axes, as the table's spacing requires, the tail disappears, the row median being 0.72 at 1 eV and 0.90 by 1.2 eV. The diagnosis script is left in place; every ADAS number in this section comes from the logarithmic read.

The mechanism is visible directly. At collisional–radiative equilibrium the recombination-fed channel carries between 89.5 % and 99.99 % of $\sum_p Q_p n_p$, and its share rises monotonically with density, from 97.8 % at $10^{12}/\text{cm}^3$ to 99.99 % at $10^{15}/\text{cm}^3$ at 1 eV. Equation (4.3) is therefore reporting the recombination channel with a small ground-fed correction, which is why η tracks n_e : three-body recombination carries an extra factor of n_e and the mislabelled flux carries it into the coefficient.

The recombination half, built and compared for the first time. The same decomposition supplies the coefficient the gate never computed. The net flux into the neutral ground state carried by the recombination-fed channel, per unit $n_e n_{\text{ion}}$, is

$$\text{ACD}_{\text{model}} = \frac{S_{1s} n_{\text{ion}} + \sum_{p \neq 1s} L_{1s,p} n_p^{(0)}}{n_e n_{\text{ion}}}, \quad (4.5)$$

direct radiative recombination into 1s plus everything that cascades or is collisionally transferred down into it. Against ACD96 this agrees at *all* 400 points, with η between 0.894 and 1.009 and median 0.982; at the benchmark the model gives $2.788 \times 10^{-13} \text{ cm}^3/\text{s}$ against ADAS's 2.883×10^{-13} , a ratio of 0.967.

That is the strongest external check in this thesis. It is an absolute coefficient, compared against an independent database built from independent atomic data, agreeing to better than 10 % at the benchmark and within a factor of two everywhere. It was available all along and was not taken, because the half of the gate that would have taken it was never written.

The second reading was that the two quantities are not defined on the same system at all. ADAS coefficients describe a closed atom-plus-ion system; this model holds the

ion as a fixed external reservoir with no forty-fourth state (Section 3.5). A stage-to-stage coefficient extracted from an open system need not equal one extracted from a closed one, and Section 4.8 has already shown that the boundary condition moves τ_{slow} by up to a factor of 41. **The ACD result argues against this being the dominant effect:** ACD is extracted from the same open system through the same fixed-ion reservoir and agrees with ADAS to within 3.3 % at the benchmark, 1.8 % at the grid median and 11 % at every point; an open-system artifact large enough to produce η of 8346 in SCD would not leave ACD intact.

The residual SCD deficit, stated rather than dismissed. Rebuilding SCD on the ground-fed channel does not make the agreement perfect, and the remaining disagreement is systematic in a way worth reporting: the model’s ionisation sits below ADAS at all but one of the 400 points, by up to 30 % and by 17 % at the median, with the maximum $\eta = 1.004$ at [4, 7]. Under the linear read of the table the gap closes steeply with temperature, from a median of 0.302 over 1 to 1.5 eV to 0.817 over 5 to 10 eV, a shape that reads as the signature of truncation; the steep part is the linear read, and under the logarithmic read the deficit is between 0.72 and 0.94 at every temperature above 2 eV. A deficit of that size is still what truncation would produce, since ADAS96 carries levels above $n_{\text{max}} = 15$ and the levels nearest the continuum ionise most easily, but its temperature shape no longer singles that cause out (the $n \geq 12$ flux share the diagnosis script printed against the linear read is not carried forward). The test that would settle the cause is the n_{max} convergence study of Chapter 6 rather than this one.

The reservoir itself, against the same table. ADAS’s two coefficients fix the equilibrium reservoir: at ionisation balance $n_{1s} \text{SCD} = n_{\text{ion}} \text{ACD}$, so $u^{\text{CRE}} = \text{ACD}/\text{SCD}$, and the model obeys the same identity with its own ground-fed SCD and ACD to 4×10^{-13} at every grid point. That makes the quantity Chapter 5 is built on directly comparable, and the comparison has two parts (`verify_reservoir_vs_adas.py`, `validation/reservoir_vs_adas/`). The *level* is high: $u_{\text{model}}^{\text{CRE}}/u_{\text{ADAS}}^{\text{CRE}}$ has median 1.18 over the 400 points, runs from 0.89 at [4, 7] to 1.40 at [0, 1], and is 1.26 at the benchmark, structured about equally in T_e and in n_e . The model’s ground reservoir per ion sits a fifth above ADAS’s, which moves the operating point beneath the cap of Section 5.5 without touching the cap. The *displacement*, which is what the plateau error is built from, agrees. Over the 392 one-interval temperature steps the model’s $\Delta \ln u$ has median magnitude 0.26 and differs from ADAS’s by a median 0.014; above 2 eV the median difference is 0.011 and the worst 0.050, at [15, 0]; at the benchmark step the two are -0.270 and -0.258 . The reservoir gain G this thesis tabulates is therefore within about 5 % of the one ADAS implies, at the median over the defended range, and the refuter set before the run, a difference above 0.1 at any interior step above 2 eV, appeared at none of 264. Every number in this paragraph comes from the logarithmic read; on the linear read of `diagnose_gate_d.py`, which overstates SCD by up to a factor 6.65 over 1 to 1.5 eV, the refuter would have appeared at 9 of the 264 steps.

4.10.2 Three further defects in the gate itself

The gate is also incompletely implemented, and this must be said alongside the failure rather than instead of it.

The recombination half does not exist *in the gate*. The ACD table is loaded once, at line 299 of `validate_gates.py`, and the variable holding it is never read again anywhere in the file; only the ionisation comparison runs, despite a docstring that describes both. That half is computed in `diagnose_gate_d.py` (Section 4.10.1).

The interpolation is wrapped in a catch-all exception handler that sets the comparison value to not-a-number and continues, so a failure to find the right density band is indistinguishable from a point that was legitimately skipped.

And the gate’s status is recorded three incompatible ways: `README.md` line 119 calls it documented, the docstring at line 47 predicts a pass above $T_e = 5\text{ eV}$, and `gate_summary.txt` reports failure at 0 % of points. The correct entry is the measured one: Gate D fails.

4.10.3 What is claimed, and what is not

Gate D itself still fails, and it has not been changed. `validate_gates.py` is untouched: no tolerance was widened, no subset was selected, and `gate_summary.txt` still records 0 %. What has changed is that the failure now has a demonstrated cause rather than two undemonstrated candidates. The gate compares a mixed quantity against a table that keeps the two channels apart by design, which is a defect in the comparison and not in the matrix.

Deciding what Gate D should test belongs with whoever owns the gate. What is claimed here is narrower: rebuilt on the two channels separately, the coefficients agree with ADAS within a factor two at all 400 points for SCD, the model 0 to 30 % low with a median ratio of 0.83, consistent with the n_{max} truncation but not uniquely so, and at all 400 for ACD with a median ratio of 0.98.

The failure does not cost the thesis its only comparison of absolute effective coefficients against an external database; that comparison exists, hidden by the definition. The failure does not touch the structural results of Chapter 3, which are statements about whatever $\mathbf{L}$ is supplied and would hold for the ADAS-consistent matrix equally. It does not touch the level-resolved comparison of Section 4.9, which compares the same decomposition against tabulated coefficients rather than against a contraction of 42 coefficients into one. And the contraction is the reason the two comparisons carry different weight: many incorrect α would contract to the correct SCD, so a passing Gate D would have been weak evidence even had it passed.

4.11 Errors found by our own audit

Section 2.9 reports two implementation errors that this project found in itself, with their sizes, and observed that both were invisible to every consistency test the pipeline ran. This section states what each changes here, adds two more, and keeps the four together rather than in an appendix: a model whose authors found and bounded four of its own errors is better characterised than one reporting none, and in three of the four cases the bound is itself a result.

4.11.1 The ℓ -mixing Coulomb logarithm

Proton-impact ℓ -mixing redistributes population between sublevels of the same n and, as Section 2.7 showed, is the dominant loss channel for the metastable $2s$ level. The rate

coefficient follows Badnell’s formulation [42], whose Eq. 9 carries a factor

$$F(U_m) = \frac{\sqrt{\pi}}{2} U_m^{-3/2} \operatorname{erf}(\sqrt{U_m}) - \frac{e^{-U_m}}{U_m} + E_1(U_m), \quad (4.6)$$

in which erf is the error function, E_1 the exponential integral, and U_m the dimensionless ratio of the smallest energy transfer the Debye cutoff permits to $k_B T_e$; $F(U_m)$ is the Coulomb logarithm of the interaction, which diverges as $U_m \rightarrow 0$.

The implementation set $F = 1$, justified in its docstring as the low-density limit $U_m \rightarrow 0$; the limit was inverted, and Section 2.9 gives the finding, the factor of three to seven by which the rates were low, and the measured effect on the spectrum. Evaluated with the local Debye length at the benchmark point [23, 5], $F = 6.62$ for $2p\text{--}2s$, 5.71 for $3d\text{--}3p$, 4.26 for $5g\text{--}5f$ and 2.89 for $8k\text{--}8i$ (`validation/lmix_cutoffs/`); with the production table’s frozen density of $10^{14}/\text{cm}^3$ the $2p\text{--}2s$ value is 6.95 (Section 2.9). A derivation note records 6.65, 5.73, 4.29 and 2.92 for the same four pairs.

The insensitivity of the slow eigenvalues to that error, which Section 2.9 measures, is specific to a process already far faster than τ_{relax} ; it is not a general licence to be careless about fast rates, and the ℓ -mixing scan below is what establishes the range over which it holds.

4.11.2 And ℓ -mixing is saturated

The same insensitivity was then tested directly rather than argued. Scaling every ℓ -mixing rate in the matrix over a range of $\times 0.1$ to $\times 10$ moves $f_3 - f_4$ by less than 0.3 % at the benchmark point and at the ridge, and by less than 5 % at the cold corner. Only switching ℓ -mixing off entirely changes anything. Rebuilding the matrix with a density-consistent Debye cutoff in place of a frozen one moves $f_3 - f_4$ by at most 0.42 %.

The claim that the frozen Debye density varies the cutoff factor by about 30 % across the grid is wrong: the module’s own functions give $\times 3.7$ for $n = 2$ and at least $\times 50$ for $n = 8$ (Section 2.7); the conclusion survives because the observable does not depend on the input over that range.

4.11.3 The eigenvalue magnitude filter

The spectrum of $\mathbf{L}$ was extracted by a routine containing the line `eigs = eigs[eigs < -1.0]`, which discards every eigenvalue smaller in magnitude than 1 1/s. Where the slow eigenvalue is smaller than that, at the cold end of the grid, the filter deleted it and promoted the whole ladder by one step, so τ_{slow} , τ_{relax} and M were wrong together while all three remained finite, negative and correctly ordered. It was found by recomputing the spectrum unconditionally and comparing bit-for-bit against the stored arrays, the three-implementation agreement of Section 4.4. Section 2.9 gives the extent of the damage, the ladder shift that confirmed the mechanism, the preserved filtered outputs and the removal of the literal from two dormant files. What belongs here is the shape of the wrong answer: at the cold corner the filter turned $M = 1.729 \times 10^9$ into $M = 1.467$, and the grid minimum from 86.77 into 1.341, values that read as a physics finding, since near-breakdown of timescale separation in the coldest, thinnest plasma is exactly what a reader would expect to see. The later correction of the two dormant files changes no number in this thesis: at the benchmark the two filters return the same τ_{slow} to every digit, and they differ only where the slow mode is slower than 1 1/s.

4.11.4 The central quantity was misnamed

The largest of the four errors changed no number and inverted a sentence.

The quantity that had been reported as the quasi-steady-state error was defined as the difference between the tabulated ratio and the true one, with the tabulated ratio evaluated from a two-argument function of T_e and n_e . Section 3.5 shows that the excited-state ratio is a function of three arguments, not two, the third being the state of the ground-state reservoir. Fixing the first two therefore silently imposes equilibrium ionisation balance, so what was measured was the distance to the equilibrium answer, not the error of the quasi-steady-state closure.

Integrating the full system and evaluating both errors along the same trajectory separates them:

point	error against equilibrium	error of the QSS closure
benchmark [23, 5]	6.34×10^{-2}	8.66×10^{-6}
worst case [0, 4]	3.869×10^{-1}	6.73×10^{-9}

They differ by four to eight orders of magnitude, most at the point carrying the largest error. The physics was intact and every number reproduced; what was wrong was the name and the sentence built on it, and Chapter 5 states the consequence, that the approximation everyone worries about is not the one that fails. The error was found by an adversarial reading of the thesis text itself, not of the code, and it is why Section 5.1.1 defines four error measures separately.

4.11.5 A statement in Chapter 3 that measurement contradicted

The statement, in Section 3.1.4, was that at $n_e = 10^{12}/\text{cm}^3$ three-body recombination supplies under 10% of the feed into any level, and it was used to propose testing the radiative and three-body datasets separately at the two ends of the density grid. Measured against the matrix, 30 of 36 levels exceed 10% three-body feed at both 1.00 eV and 2.95 eV, and 26 of 36 at 10 eV. Into the bundled shells it is essentially the whole feed.

The density grid therefore does not separate the two recombination datasets, and **no such separate test was run or can be**; Section 3.1.4 now states the measured fractions and withdraws the promise. What remains is the gap: the radiative and three-body datasets enter this model only in combination, nothing in this thesis tests them apart, and a test that could would need a density below the bottom of this grid, outside the conditions the thesis is about.

4.12 Reproduction, and what a document cannot certify

4.12.1 Regeneration from the canonical matrix

Every quantitative result in Chapters 3 and 5 was regenerated from $\mathbf{L}$ and compared against the stored output. The timescale and step-error grids reproduce bit-identically, including the values that come from numerical integration and not from linear algebra. The plateau and persistence tables reproduce byte for byte, 680 of 680 rows, with 784

superposition residuals at exactly zero difference. The nineteen points corrected by the eigenvalue filter reproduce exactly, including the ladder-shift identity.

This is reproduction in the narrow sense: the same code, run again, on the same inputs, gives the same answers. It excludes a stale array and it excludes a non-deterministic solver. It says nothing about whether the algebra is right.

4.12.2 Independent re-derivation

The stronger check is re-derivation. All six central results of Chapter 3 were derived from the linear collisional–radiative system alone by a party working from the physical statement of the problem, implemented from scratch, with the predicted values written down before any project file was opened.

In the table below, b_1 is the ground-state departure coefficient of Eq. (3.29), the ratio of the actual ground-state population to its Saha–Boltzmann value, and b_1^{CRE} is the value it takes in collisional–radiative equilibrium (CRE), the steady state the lookup tables assume.

quantity	chapter3.tex	independent re-derivation
b_1^{CRE} at the benchmark	956	956.509
switching points c_m/a_m	$2.73 \times 10^3, 1.91 \times 10^4$	2727.3, 19085.5
peak b_1	7.2×10^3	7214.7
f_3, f_4 , difference	0.260, 0.048, 0.212	0.259652, 0.047725, 0.211927
maximum of $f_3 - f_4$	0.4514	0.451357
R at the benchmark	0.7778	0.777768

The timescales reproduce to four digits, $\tau_{\text{slow}} = 22.73 \mu\text{s}$, $\tau_{\text{relax}} = 2.277 \text{ ns}$, $M = 9981.9$.

Two further invariances were measured in the same pass and belong with the result they protect. The shell separation Δ that sets the bound of Section 5.5 is invariant under rescaling the ground-to-ion normalisation, verified exactly under factors of 10^{10} and 10^{-7} , while the individual channel fractions are not; and Δ varies by 0.07 % across five different ℓ -weightings of the shell sum, so the bound is not an artifact of how sublevels were weighted.

4.12.3 A prediction that was falsified, and improved the claim

The same re-derivation made one prediction that turned out to be wrong, and it is more informative than the six that were right.

The prediction was that $f_3 - f_4$ must vanish at high density, on the reasoning that in local thermodynamic equilibrium both supply channels produce the same Boltzmann distribution, so the two shells would draw identically and the difference would close. It does not happen. At $T_e = 2.947 \text{ eV}$, Δ doubles over the lowest part of the grid and then stops moving: 0.979 at $10^{12}/\text{cm}^3$, 1.946 at $1.4 \times 10^{14}/\text{cm}^3$, 1.939 at $10^{15}/\text{cm}^3$ (validation/fujimoto_r1_consequence/). Above about $7 \times 10^{12}/\text{cm}^3$ it is flat, and the bound $\tanh(\Delta/4)$ built on it spans only 0.394 to 0.450, a 14 % range, across the whole of the rest of the grid.

The reason is the open boundary again, this time in the model’s favour: the recombination source is imposed externally rather than tied to the ion density by a Saha relation, so the two channels cannot merge by construction, and Gate C’s 1.5 % of Saha says the

grid never approaches the limit in which they would. The consequence for the density maximum of the diagnostic error, a position effect rather than a span effect, is Table 5.4 in Section 5.6.1.

4.12.4 One item demoted: the May/August cross-validation

A cross-validation between two runs four months apart, on different matrices, with different observables and different numerical methods, was recorded in this project's register as verified and described as the strongest single check it had. The agreement is close: τ_{slow} to 0.064 %, τ_{relax} to 0.200 %, M to 0.264 %, $\varepsilon_{\text{step}}$ to 0.014 %. Predicting how long the error stays above 10 % from the plateau-then-decay picture gives 12.478 μs against the earlier run's full integration, 12.4794 μs .

Two corrections apply, and both weaken it.

The apparent agreement to eight parts in 10^5 on the duration is cancellation. The inputs agree only to 0.06 to 0.6 %, and carrying the earlier run's own internal numbers through gives 12.406 μs , off by 0.59 %. **The figure is agreement at the 0.6 % level**, which is still strong.

More seriously, **the earlier run has never been reproduced against the canonical matrix**. This project's own derivation record says of these same numbers that they are a reported result, not independently re-run, and should be treated as strong documented evidence rather than as self-verified. A cross-validation whose earlier half cannot be re-executed is a citation, not a check. It is reported here as documented evidence and is not counted on the ladder of Section 4.13. Under the ground-truth ordering this thesis works to, physics over mathematics over code over documents, a document is the weakest of the four, and it is where this particular claim currently rests.

4.12.5 Two provenance defects that are not repaired

Two arrays of timescales, M , τ_{slow} and τ_{relax} , are written to the same three file paths by two different scripts. Whichever ran last wins, and six downstream readers cannot tell whose numbers they are holding. This is exactly the condition under which the eigenvalue filter propagated, since the filtered and unfiltered versions of the same array have the same name.

Two things bound what that costs here, and both were checked rather than assumed. The two writers are not two calculations: each sorts the eigenvalues of the same canonical $\mathbf{L}$ and keeps the negative ones, neither applies the filter that caused the earlier incident, and the arrays on disk agree with the independently produced `verify_timescales.py` artifact to 2×10^{-16} . And none of the six readers feeds this document: every figure and table they write is absent from it. The benchmark $M = 9982$ quoted in Chapter 3 and in the abstract is not read from these arrays at all; it is $\tau_{\text{slow}}/\tau_{\text{relax}}$ computed from $\mathbf{L}$ by `verify_timescales.py` and recorded in `validation/spectrum_testpoint.csv`. The defect is therefore a latent hazard in the repository rather than an ambiguity in any number printed here, and it is left in place and named rather than quietly repaired.

An audit script exists to detect single-writer violations, and it reports an unfixed failure of this kind. It is also incomplete. Its test for a write reads the file mode from the second argument or a keyword, so the idiom `path.open("w")` passes through it unrecorded. At the tip of `main` on 23 August that idiom had four users; in the tree this thesis is built from it has 45 call sites in 29 scripts, 42 of them writing comma-

separated output, among them the producer of the plateau map that fifteen places in Chapter 5 read from (`verify_writer_census.py`, `validation/writer_census/`). The audit should not be cited as evidence that the single-writer property holds. The census confirms that the property fails for exactly the three timescale arrays, each written by both `qss_analysis.py` and `validate_gates.py`, and holds for the plateau map, which has one writer.

4.13 What this chapter establishes

Table 4.4 collects the checks in the order a reader should weight them, with the grade of Section 4.1 attached to each.

4.13.1 What the model may be used for

Three statements are supported by the evidence above, and they are the ones Chapter 5 relies on.

The equations are solved correctly. The superposition residual, the independent re-derivation of all six central results, the bit-for-bit regeneration, and the grid-wide conditioning measurement together establish that the linear algebra does what Chapter 3 says it does, to a precision many orders of magnitude finer than any quantity the thesis reports.

The input rates are traceable and benchmarked where the observable lives; one derived coefficient there is not. The rates feeding $n \leq 4$, which is where H_α and H_β originate, agree with an independent R-matrix calculation at a mean absolute error of 13.0 %, and the two transitions this thesis leans on hardest, $1s \rightarrow 2p$ and $2p \rightarrow 3d$, agree to between 2.4 % and 13.3 % across the temperature range. Above $n = 4$ the atomic data are less certain and the truncation is measurable. Section 4.6.2 tests one component of the mechanism, the sensitivity $f_3 - f_4$, and finds it converged to within 1.03 % at the defended points where the geometric tail can be summed, in a stamped run (`validation/nmax_downward_scan/`); under the same truncation u^{CRE} is converged to within 0.13 % at the 257 defended points where its increments are geometric, Δ to within 0.57 % and τ_{slow} to within 1.6 % at all 280, and $\Delta \ln u$ across the defended heating steps moves by at most 2.3×10^{-5} between $n_{\text{max}} = 14$ and 15 against a median step of 0.20 (Appendix C). What is measured for the reservoir is a scan rather than a truncation: weakening the ionisation and three-body exchange of the top six shells eightfold moves $\varepsilon_{\text{plateau}}$ by 0.7 % at the median over the defended pairs (Section 2.5). All of that concerns the rates going in. The coefficients that come out at the two shells the observable uses are a separate matter, and one of them is in open disagreement: the ground-fed $r_1(3)$ and $r_1(4)$ of this model sit factors 8.3 and 4.4 below Fujimoto’s tabulation at the lowest density, which is a deficit in a_3 and a_4 themselves and so in the sensitivity built from them (Section 4.9). Benchmarked inputs are therefore not the same thing as a benchmarked susceptibility, and this thesis can claim the first but not the second.

The structural results are robust to the choices that could have manufactured them. Closing the ion boundary changes the persistence census by one part in 325. Scaling ℓ -mixing over two decades moves the observable by under 0.3 % where the thesis makes claims. The single-slow-mode estimate of Eq. (5.16) reproduces the true time-averaged error to within 2.2 % at 100 μs over the 448 pairs above 2 eV.

Table 4.4: Every check this chapter grades, with what it tests, its severity, whether its result can be regenerated, and the measured outcome. *Tautological* means the check cannot fail except on an arithmetic bug, because the quantity tested was constructed from the relation being tested; *weak* means it can fail but only for an error far larger than any plausible one; *severe* means a plausible wrong version of this model fails it. The provenance column is the more uncomfortable of the two: *stamped* means a script in the repository regenerates the number into an artifact under `validation/`; *notes* means the number is recorded in this project’s working files with no producing script, so it cannot be re-executed and cannot fail anything today. None of the checks graded severe are in that condition, and the one *notes* entry left is the tanh bound, which is graded tautological. The counts stated in this caption are checked against the table’s own rows by `verify_ladder_counts.py`, which exits non-zero if they drift apart. Of the eighteen internal checks, twelve are severe, four are tautological and two are weak; of the five external comparisons, two pass, one passes only at the higher densities, one fails, and one is open.

check	grade	record	measured
Atomic-data provenance (Ch. 2)	severe	stamped	every coefficient traces to a named dated source; tallies exact
Excitation rates against RMPS (Ch. 2)	severe	stamped	81.4 % within 20 % for $n_{\text{upper}} \leq 4$, mean 13.0 %; $n = 5$ disagrees by factors of 2–5, localised to the top shell of the RMPS basis
Raw cross sections, microscopic reversibility	severe	figure	ratio 0.9995 ± 0.0032 , 98.7 % within 1 %
One energy ladder on both sides	severe	stamped	Saha ratio 0.999997 for all 43 states, spread 1.4×10^{-15}
Superposition of the two channels	severe	stamped	3.075×10^{-14} over 784 pairs; catches a $3 \leftrightarrow 4$ swap
Reduced against full R	severe	stamped	0.777768 by three routes; catches the same swap
Deliberate fault injection	severe	stamped	three of four injected faults invisible to the conservation residual; the fourth moves it by thirteen orders
Line-weighted against shell ratio	severe	stamped	ratio 0.948 to 1.000 over 784 pairs, below unity at every one; census 44 against 45
Escape-factor cross-check	severe	stamped	σ_0 against an independent module, 0.059 % apart against a 1 % raising tolerance
Pinned input hashes	severe	stamped	both SHA-256 match byte for byte
Recomputation guards on the figures	severe	code	every plotted row rebuilt from $\mathbf{L}$ at 10^{-9} , raises
Truncation convergence of $f_3 - f_4$	severe	stamped	increments one-signed and geometric at 245 of 280 defended points, ratio median 0.67; extrapolated -0.46 % at the benchmark, 1.6 % at the cold corner, within -0.82 to 1.03 % on the defended set. Tests one quantity, not the state space
Gate A, detailed balance	tautological	stamped	< 0.01 %, but K_{deexc} was derived from K_{exc}
Column-sum conservation	tautological	stamped	3.61×10^{-11} ; see the fault-injection row for what it cannot see
The tanh bound	tautological	notes	holds at all 248 points, and holds on corrupted input
Gate E, timescale hierarchy	tautological	stamped	$M \geq 86.8$ against a threshold of 1
Gate B, coronal limit	weak	stamped	worst convergence 0.0034 against a tolerance of 0.5
Gate C, Saha ceiling	weak	stamped	monotone and below Saha; approach

4.13.2 What it may not be used for

No effective coefficient from this model should be compared against a database without first separating its two supply channels. Gate D as written fails at every point on the grid because it compares a coefficient that mixes both channels against a table that separates them (Section 4.10); rebuilt channel by channel it agrees within a factor two at all 400 points for SCD and ACD, and the model’s total effective ionisation coefficient remains a quantity ADAS does not tabulate.

No claim rests on the terminal shell, whose population is high by a factor consistent with truncation, or on the ℓ -distribution inside the bundled shells, which is assumed rather than measured. Section 6.8.3 tests that assumption: ℓ -mixing licenses it over most of the grid, and at the two highest densities below 3 eV, where the locally evaluated mixing loses to collisional loss, detailed balance among shells within 10^{-4} of Saha–Boltzmann equilibrium licenses it instead; the adjacent-shell rates themselves are not ℓ -blind, and no grid point escapes both criteria.

Nothing here validates the physics that is absent. Every check in this chapter tests a 43-state atomic hydrogen model against itself, against arithmetic, and against other atomic hydrogen models. None of them tests transport, molecular channels, or radiation trapping, because none of those appears in $\mathbf{L}$ at all. A model can pass every check in Table 4.4 and still be the wrong model for a divertor. Chapter 6 is where that question is asked, and it is asked with numbers rather than with hedges.

The instrument is now characterised: what it computes, how precisely, on which region of the grid, and where the evidence for it runs out. What it says is the subject of Chapter 5.

Chapter 5

What the Model Says

Chapter 3 left the argument at a definition and a worry. The definition: a Balmer-ratio lookup table indexed on (T_e, n_e) is a table in *three* arguments with the third one hidden, and the hidden argument is the ground-state departure coefficient b_1 , which takes a time τ_{slow} to settle after the plasma changes. The worry, which is the one the literature carries: the excited states are what a spectrometer sees, and they may not keep up with a plasma that changes on microseconds. Chapter 4 established that the rate matrix these questions are asked of reproduces independent atomic data where such data exist, and reported the two errors that an audit of this work found and corrected.

This chapter evaluates both errors over the operating range, and the two results are worth stating before the evidence rather than after it.

The first is a reversal. The quasi-steady-state (QSS) closure, the approximation that the excited populations sit in instantaneous steady state with whatever ground-state reservoir is present, is not the problem. Measured along full time-dependent trajectories, its residual never exceeds 4.4×10^{-5} over the analysed pairs and is 6.73×10^{-9} at the grid point carrying the largest diagnostic error (Section 5.7). What fails is the assumption that the reservoir itself has finished moving, which is a different assumption, made at a different step of the inversion, and never previously separated from the first.

The second is that what that failure costs is not a number. It is a function, and the function has a closed form,

$$\varepsilon_{\text{plateau}} = \left| \exp(\bar{S} G \Delta \ln T_e) - 1 \right|,$$

in which $\Delta \ln T_e$ is the size of the transient under consideration, G measures how far the ground-state reservoir is dragged by it, and $\bar{S}$ measures how strongly the line ratio answers to the reservoir. The last of these is bounded above by a quantity that no choice of atomic data can push past unity. So the chapter reports the structure and the bound, and lets a reader compute the error for their own transient. A single percentage would be an answer to a question about the grid this thesis happens to use, and Section 5.3 shows how badly that goes.

Both results are stated over $T_e \geq 2 \text{ eV}$. Section 5.10 shows that this is not a hedge: it is the conservative boundary adopted above the band in which the model's own assumptions fail together, and two unrelated pieces of physics place the boundary in the same place.

What this chapter takes as given. From Chapter 3: the rate equation $\dot{\mathbf{n}} = \mathbf{L}\mathbf{n} + \mathbf{S}n_{\text{ion}}$ and its two timescales τ_{relax} and τ_{slow} ; the QSS elimination and its Schur complement; the two reference states R^{QSS} and R^{CRE} and the two error measures ε^{QSS} and ε^{CRE} built on

them; the two-supply decomposition $\mathbf{n}_F = \mathbf{a} n_{1s} + \mathbf{c} n_{\text{ion}}$; the ground-fed fraction f_m , its logistic form, and the identity $d \ln R / d \ln b_1 = f_3 - f_4$. From Chapter 1: the operating range $T_e = 1\text{--}10\text{ eV}$, $n_e = 10^{12}\text{--}10^{15}/\text{cm}^3$, and the fact that an edge-localised mode (ELM) disturbs a divertor for of order $100\text{ }\mu\text{s}$. That duration is the aggressive end of the range; the value extrapolated to the ITER pedestal is $506\text{ }\mu\text{s}$ [17, Sec. 5], and Section 5.8 carries both. The benchmark point $T_e = 2.947\text{ eV}$, $n_e = 1.389 \times 10^{14}/\text{cm}^3$, grid index [23, 5], runs through this chapter as it runs through the last.

5.1 What happens after a temperature step

Take a plasma sitting quietly at collisional–radiative equilibrium (CRE), the steady state of the full 43-level system, and change the electron temperature abruptly at $t = 0$, holding n_e fixed. Three things then happen, on three separated timescales, and the middle one is where this thesis lives.

For the first few nanoseconds the excited manifold rearranges. Every one of the 42 excited levels is a reactive intermediate in the chemical-engineering sense: it holds almost no population, it is created and destroyed thousands of times faster than the species feeding it changes, and it therefore reaches a steady state with respect to that species almost immediately. At the benchmark point this takes $\tau_{\text{relax}} = 2.28\text{ ns}$.

For the next several microseconds nothing further appears to happen, and this is the interesting part. The excited populations have settled, but they have settled onto the *old* ground-state population, because the ground state holds essentially all the atoms and cannot turn over in nanoseconds. The system occupies a state that is steady in its fast variables and stale in its slow one. Call it *partial equilibrium* (PE). At the benchmark point it lasts until the ground state has moved, on $\tau_{\text{slow}} = 22.7\text{ }\mu\text{s}$, so the plateau is some four orders of magnitude wider than the relaxation that built it.

Only then does the plasma reach the new collisional–radiative equilibrium that a lookup table would have assigned to it the instant the temperature changed. Figure 5.1 shows all three phases on a single logarithmic time axis, at the benchmark and at the coldest grid point.

The ELM that motivates this thesis lasts of order $100\text{ }\mu\text{s}$, long against τ_{relax} by a factor of about 4×10^4 and comparable to τ_{slow} . Comparable cuts both ways: over the 448 defended pairs with a plateau window τ_{slow} runs from $1.2\text{ }\mu\text{s}$ to 89 ms , and at 237 of them it is shorter than $100\text{ }\mu\text{s}$, so a spectrometer integrating over an ELM records the reservoir’s relaxation there and not a frozen plateau; at every one of the 45 pairs where the averaged error exceeds 10 % it is longer than $195\text{ }\mu\text{s}$ (`validation/window_membership_exposure/`). The statement that an exposure spends its integration time looking at the plateau holds only where τ_{slow} exceeds the exposure. The plateau is the state whose error this chapter measures first; what an exposure of finite length records is the subject of Section 5.8.

5.1.1 Two error measures, and why the step size is part of each

Section 3.5 fixed the two reference states against which an observable can be judged. Two further quantities are needed here, and both are properties of a temperature *step* rather than of a single operator, which is why neither appears in Chapter 3.

Let the plasma sit at the collisional–radiative equilibrium of $\mathbf{L}(T_e, n_e)$, and let the electron temperature step to $T_e + \Delta T_e$ at $t = 0$ with n_e held fixed. Write $\mathbf{L}^-$ and $\mathbf{L}^+$ for




Figure 5.1: What the plateau looks like: the excited states keep up, the reservoir does not. The shell ratio $R(t) = n_3/n_4$ after a step in T_e at fixed n_e , on a logarithmic time axis long enough to hold both timescales. Three phases are visible. R rises on τ_{relax} as the excited manifold re-equilibrates against a ground state that has not yet moved; it sits at the partial-equilibrium value R^{PE} ; then it decays on τ_{slow} as the ground state, and with it the ionisation balance, finally responds. The excited states are fast enough to track the new temperature within nanoseconds. The reservoir that feeds them is not. In (a) the curve crosses the R^{PE} line at $0.7\tau_{\text{relax}}$, peaks near $1.2\tau_{\text{relax}}$ and settles back onto it, because $n = 4$ relaxes faster than $n = 3$ and the ratio briefly exceeds its partial-equilibrium value. (a) The benchmark, $T_e = 2.947 \rightarrow 3.09$ eV at $n_e = 1.4 \times 10^{14} \text{ cm}^{-3}$: R moves from 0.7778 through 0.8201 to 0.7711, a plateau error of 6.36%, with $\tau_{\text{relax}} = 2.24 \times 10^{-9}$ s and $\tau_{\text{slow}} = 1.85 \times 10^{-5}$ s ($M = 8243$). (b) The coldest grid point, $T_e = 1.00 \rightarrow 1.05$ eV at $n_e = 5.2 \times 10^{13} \text{ cm}^{-3}$: 1.2391 through 1.6712 to 1.2050, 38.7%, with $\tau_{\text{relax}} = 6.24 \times 10^{-9}$ s and $\tau_{\text{slow}} = 0.233$ s ($M = 3.74 \times 10^7$), five orders of magnitude of separation between the two clocks. Every timescale quoted belongs to the *post-step* operator, the one that governs the relaxation; the same grid point read from its unstepped operator gives a different M , and the two must not be interchanged. R^{PE} is computed independently by a single linear solve against that operator and agrees with the two-channel construction used elsewhere in this chapter to 2×10^{-16} . The plateau is flat to 0.17% peak to peak in (a). Curves are eigen-propagations of the full 43-state operator, cross-checked against `scipy.linalg.expm` inside the plateau window, worst disagreement 1×10^{-7} .

the operators before and after, in the notation of Section 3.3. The *step error*

$$\varepsilon_{\text{step}} = \left| \frac{R^{\text{CRE}}(T_e, n_e)}{R^{\text{CRE}}(T_e + \Delta T_e, n_e)} - 1 \right| \quad (5.1)$$

is the dimensionless distance between the two tabulated answers: what the reading would change by if the plasma re-equilibrated instantly. The *plateau error*

$$\varepsilon_{\text{plateau}} = \left| \frac{R^{\text{PE}}}{R^{\text{CRE}}(T_e + \Delta T_e, n_e)} - 1 \right|, \quad R^{\text{PE}} \equiv R^{\text{QSS}}(T_e + \Delta T_e, n_e; u^-), \quad (5.2)$$

is the distance between the tabulated answer and the state the plasma occupies once the excited manifold has settled but before the ground state has moved: the post-step excited response of Eq. (3.28), evaluated on the *pre-step* reservoir.

The third argument is $u^- = n_{1s}^-/n_{\text{ion}}$ and not b_1^- , and the difference is a factor of 24. What is frozen across the step is the ground-state *population*, which is what transport and ionisation would have to change. The departure coefficient is not frozen with it, because $b_1 = u/(Z_1(T_e)n_e)$ carries the Saha coefficient and Z_1 moves when T_e does. Holding b_1^- instead of u^- at the benchmark gives $\varepsilon_{\text{plateau}} = 0.00267$ rather than 0.0636, because b_1^{CRE} shifts by only 1.3 % across one interval, from 956.5 to 968.9, while u shifts by 23.6 %. The code freezes the population (`verify_reservoir_gain.py`, line 143), so the computed plateau errors correspond to u^- held fixed. Both are dimensionless, both are non-negative, and both are zero for a zero step.

Keeping them separate matters because they are not close. Across the operating grid $\varepsilon_{\text{plateau}}$ exceeds $\varepsilon_{\text{step}}$ at every one of the 680 point-and-direction pairs that possess a plateau, with a median ratio near 12 and a minimum of 1.45 (`verify_plateau_gridmap.py`). A diagnostician who pictures the reading sliding smoothly from one tabulated value to the other has not merely misjudged the size of the excursion; they have the wrong shape of transient.

That comparison carries a caution about the ratio itself. Because $\log(\varepsilon_{\text{plateau}}/\varepsilon_{\text{step}}) = \log \varepsilon_{\text{plateau}} - \log \varepsilon_{\text{step}}$ identically, any correlation between the amplification and $\varepsilon_{\text{step}}$ is partly tautological, and the amplification blows up wherever $\varepsilon_{\text{step}}$ passes through zero. The apparent grid maximum of $1271\times$ sits at the point of smallest $\varepsilon_{\text{step}}$ on the grid, where $\varepsilon_{\text{plateau}}$ is itself below the median (`verify_plateau_gridmap.py`). $\varepsilon_{\text{plateau}}$ is the primary quantity throughout this chapter and the amplification is reported only as a distribution.

Both quantities are defined by a step, so neither has a value until the step is named. L is tabulated only on the temperature grid, so every step used here is an integer number of grid intervals, and one interval is a factor 1.048113, that is +4.81 % heating or −4.59 % cooling (`validation/plateau_gridmap/plateau_gridmap.txt`, lines 10–11 and 18). Every $\varepsilon_{\text{plateau}}$ quoted in this chapter carries the step that produced it.

5.1.2 The closure everyone worries about is not the dominant error

Before the maps, the negative result that reframes them.

ε^{QSS} of Eq. (3.49) asks whether the excited states have kept up with the reservoir they actually see. Integrating the full 43-state system through the step and evaluating both errors along the same trajectory gives, at the benchmark point, $\varepsilon^{\text{QSS}} = 8.66 \times 10^{-6}$ against $\varepsilon^{\text{CRE}} = 6.34 \times 10^{-2}$; at $[0, 4]$, the grid point carrying the largest ε^{CRE} anywhere, $\varepsilon^{\text{QSS}} = 6.73 \times 10^{-9}$

against $\varepsilon^{\text{CRE}} = 3.869 \times 10^{-1}$. Both pairs are read from the columns `max_track_w30_shell` and `cre_start_w30_shell` of `validation/trajectory_census/trajectory_census.csv` (`verify_trajectory_census.py`), where the quantity this chapter calls ε^{QSS} is recorded under the name `max_track`.

At those two points the two errors differ by 3.9 and 7.8 decades, and they differ most at the point where the diagnostic error is worst. Over the 338 window-passing cells among the 392 candidate heating cells the gap runs from 2.9 to 9.7 decades (Figure 5.2b), so it is a distribution rather than a single number and should not be quoted as one. Neither figure is zero: the excited-state closure is an approximation and it carries a measurable residual. The point is its size relative to the other error, not its vanishing. The adiabatic elimination of Section 3.4 is doing its job to a part in 10^8 precisely where the tabulated answer is off by 39 % per +4.81 % temperature step.

The reason is the spectrum. Section 3.3 found one slow mode and then a band, with nothing in between: at $\mathbf{L}[1,4]$, the post-step operator at the worst point, the three slowest eigenvalues are -4.29 , -1.60×10^8 and -3.44×10^8 1/s, a gap of 3.7×10^7 (`verify_residual_claims.py`, `validation/residual_claims/`). An excited population cannot lag a reservoir that moves 10^7 times more slowly than it can respond.

Figure 5.2 puts both errors on the same axes over the whole grid. The closure residual never exceeds 4.4×10^{-5} anywhere on the analysed grid; the distance to the tabulated equilibrium runs from 1.66 % to 38.7 % over the same cells.

The failure is therefore not in the elimination but in the closure chosen for the eliminated variable. A two-parameter table sets $b_1 = b_1^{\text{CRE}}(T_e, n_e)$, which asserts that the plasma has finished changing. The rest of this chapter is about what that assertion costs.

5.2 The error maps

Figure 5.3 maps $\varepsilon_{\text{plateau}}$ over the operating grid for a heating step of one grid interval, +4.81 % in T_e at fixed n_e . The grid is 50 temperatures by 8 densities, and each cell was evaluated by an exact linear solve rather than by a fit: the superposition $\mathbf{n}_F = \mathbf{a} n_{1s} + \mathbf{c} n_{\text{ion}}$ on which Eq. (5.2) rests holds to a relative residual of 3.075×10^{-14} over the whole grid (`plateau_gridmap.txt`, line 19).

Three numbers anchor the map, and each is stated with the set it is an extremum over, because this project has twice been caught quoting a regime-restricted extremum globally.

- Over all 680 analysed pairs, the largest value is $\varepsilon_{\text{plateau}} = 38.69$ % per +4.81 % step, at $[0, 4]$: $T_e = 1.000$ eV, $n_e = 5.18 \times 10^{13}/\text{cm}^3$. It sits on the cold edge of the grid, so the temperature range is truncated there (`plateau_gridmap.txt`, lines 41–43). Section 5.10 shows that this value is not quantitative and that the point is not a divertor state.
- Over the range this thesis defends, $T_e \geq 2$ eV, the largest value is $\varepsilon_{\text{plateau}} = 18.07$ % per +4.81 % step, at $[15, 3]$: $T_e = 2.02$ eV, $n_e = 1.93 \times 10^{13}/\text{cm}^3$ (`verify_eps_gridmap.py`).
- At the benchmark point $[23, 5]$, $\varepsilon_{\text{plateau}} = 6.36$ % per +4.81 % step (`verify_eps_gridmap.py`, and `verify_plateau_gridmap.py` to five digits: 0.063612).

Cooling is not the mirror image of heating, and the asymmetry is mostly a step artifact rather than physics: one grid interval down is -4.59 % rather than +4.81 %,




Figure 5.2: The reversal: the approximation that is doubted holds, the assumption that is not examined fails. Two errors in the same observable, the shell ratio $R = n_3/n_4$, measured along the same trajectory after a +4.81% step in T_e at fixed n_e , and reported as the maximum over the timescale-separated plateau window $30\tau_{\text{relax}} < t < \tau_{\text{slow}}/30$. The *quasi-steady-state closure residual* is $|R(t)/R^{\text{QSS}}(n_g(t)) - 1|$, where R^{QSS} is rebuilt at each instant from the excited-state block alone, $n_E = -L_{EE}^{-1}(L_{Eg}n_g + S_E)$, evaluated at the *instantaneous* ground density rather than at any tabulated value. The *CRE distance* is $|R(t)/R^{\text{CRE}} - 1|$, the distance to the equilibrium-ionisation-balance state $n = -L^{-1}S$ that a two-parameter (T_e, n_e) lookup table returns. (a) Both, for the 338 window-passing cells among the 392 candidate heating cells, across the 8 density columns. The closure residual runs from 4.6×10^{-11} to 4.4×10^{-5} , so the closure everyone worries about is never worse than one part in 2.2×10^4 anywhere on the analysed grid: it is 8.66×10^{-6} at the benchmark ($T_e = 2.947$ eV, $n_e = 1.4 \times 10^{14}$ cm $^{-3}$), one part in 1.2×10^5 , and 6.73×10^{-9} at the coldest point [0, 4], one part in 1.5×10^8 . The CRE distance over the same cells runs from 1.66% to 38.7%: 6.34% at the benchmark and 38.7% at the coldest point. The assumption nobody examines therefore fails at the tens-of-percent level exactly where the closure everybody examines is exact to within one part in 10^5 or better. (b) The ratio of the two over the grid. The gap is 2.9 to 9.7 decades wide, 3.9 decades at the benchmark and 7.8 decades at the coldest point; it is not a single number and should not be quoted as one. Hatched cells (54 of 392) have no timescale-separated plateau window, $M \leq 900$, an imposed cut rather than a measured floor. Trajectories are eigen-propagations of the full 43-state operator, sampled at 24 points across each window and checked against `scipy.linalg.expm` at the window midpoint of every cell, worst disagreement 3×10^{-9} . Values below 2 eV are optically thin upper estimates: Lyman trapping is not switched on here.




Figure 5.3: **Where the equilibrium-reservoir closure costs most.** Fractional error $\varepsilon_{\text{plateau}}$ in the $n = 3/n = 4$ shell population ratio committed while the excited manifold has reached partial equilibrium but the ground state has not yet responded, after a +5% step in T_e at fixed n_e , the *heating* direction. Cooling is not the mirror image: the same transition between the two coldest grid temperatures, at $n_e = 1.9 \times 10^{13} \text{ cm}^{-3}$, gives 37.1% heated and 28.6% cooled, and the grid maximum moves from $[0, 4]$ (1.00 eV, $5.2 \times 10^{13} \text{ cm}^{-3}$) to $[1, 3]$ (1.05 eV, $1.9 \times 10^{13} \text{ cm}^{-3}$), where it is 28.6%. Hatched cells have no timescale-separated plateau window ($M \leq 900$, 54 of the heated cells): a plateau value is not defined there and none is plotted. The plain grey strip at the right-hand edge is the hottest temperature row, from which a +5% step cannot reach a new grid index. The three marked cells, keyed below the axes, are the grid maximum (38.7% at $T_e = 1.00 \text{ eV}$, $n_e = 5.2 \times 10^{13} \text{ cm}^{-3}$), the largest value at $T_e \geq 2 \text{ eV}$ (18.1% at $T_e = 2.02 \text{ eV}$, $n_e = 1.9 \times 10^{13} \text{ cm}^{-3}$), and the thesis benchmark point ($T_e = 2.947 \text{ eV}$, $n_e = 1.39 \times 10^{14} \text{ cm}^{-3}$, 6.36%). Below 2 eV the values shown are upper estimates: switching on Lyman trapping divides the grid maximum by 2.5–3.3 depending on the assumed slab thickness, and below 1.15 eV it also moves which density column carries the maximum (Figure 5.7). Nothing on the cold edge of this map is quantitative.




Figure 5.4: **The density dependence, all 8 columns.** $\varepsilon_{\text{plateau}}$ against n_e for 7 temperatures, heating, +5% step. Every density column computed is shown; stars mark each row’s maximum, open symbols mark points with no plateau window. Each row has a single interior maximum, but it is *broad and shallow in $\log n_e$* : across the peak column and its two neighbours, a factor 7.2 in density, ε varies by only 6–22%, so the location of the maximum is not resolved to better than about one grid interval, a factor 2.68. Over the full three decades the same row falls by a factor 2.2–6.6: the structure is a real crest in density, but a wide-topped one, and the word ridge should not be read as implying a sharp or well-localised line. The column at $n_e = 5.2 \times 10^{13} \text{ cm}^{-3}$ carries the row maximum in 3 of 49 heating rows, all of them the coldest; at every warmer temperature the maximum sits at $1.9 \times 10^{13} \text{ cm}^{-3}$. Both columns must be reported: quoting either alone misstates where the reservoir-closure error is largest, and below 1.15 eV the position moves across three columns once Lyman trapping is included (Figure 5.7), so no maximum location is claimed there at all.

and once the two are matched per grid point and normalised, the median heat/cool ratio is 1.036 (the matching and normalisation behind this ratio are not recorded: `verify_residual_claims.py` reproduces both error maps but obtains 1.111 when each error is divided by its own $|\Delta \ln T_e|$, and reports the difference rather than adopting a construction chosen to fit). What does not normalise away is the location of the maximum, which moves from $[0, 4]$ under heating to $[1, 3]$ under cooling, where $\varepsilon_{\text{plateau}} = 28.58 \%$ and the point is no longer on a grid edge (`plateau_gridmap.txt`, lines 59–61).

5.2.1 The density structure is a broad crest, not a line

Figure 5.4 cuts the map along the density axis. Every temperature row has a single interior maximum in n_e : the error is small at both ends of the density range and largest in between. Section 5.6.1 explains why. What matters first is how sharply that maximum is located, because the temptation is to read a ridge line off a colour map and quote a density.

It is not sharply located. Across the peak column and its two neighbours, a factor 7.2 in density, $\varepsilon_{\text{plateau}}$ varies by only 6 to 22 % depending on the row, while over the full three decades the same row falls by a factor 2.2 to 6.6 (`make_ch5_figures.py`, injected into the caption of Figure 5.4). The crest is real and it is wide-topped. The grid spacing is a factor 2.683 per column, so nothing finer than that is resolved in any case.

This chapter therefore quotes the density of maximum error as the range 7×10^{12} to $5 \times 10^{13} / \text{cm}^3$, which is the pair of grid columns bracketing the modal maximum

at $1.93 \times 10^{13}/\text{cm}^3$, and not as a single value. Two further facts support the range and neither supports a point value. The column at $5.18 \times 10^{13}/\text{cm}^3$ carries the row maximum in 3 of the 49 heating rows, all of them the coldest, and $1.93 \times 10^{13}/\text{cm}^3$ carries it in the rest; per-row argmax over rows 0 to 15 is [4, 4, 4, 3, 3, 3, 3, 3, 3, 3, 3, 3, 3, 3, 3] (`verify_residual_claims.py`, `validation/residual_claims/`). And a parabolic fit in $\ln n_e$ through the three columns around the maximum places its vertex at $1.65 \times 10^{13}/\text{cm}^3$, a 15% shift off the grid node. That fit is now performed on every row (`verify_crest_subgrid.py`, stamped to `validation/crest_subgrid/crest_vertices.csv`), and the result needs two qualifications the single number does not carry. Over the 49 rows with an interior maximum the vertex runs from 1.45×10^{13} to $3.71 \times 10^{13}/\text{cm}^3$ with median $1.62 \times 10^{13}/\text{cm}^3$, so $1.65 \times 10^{13}/\text{cm}^3$ is the median rather than the value, and the row-to-row spread is a factor 2.6. And a parabola through three points is an interpolation with no residual to report: refitting through five points where five exist moves the vertex by between -3.4 and $+9.6\%$, median 0.27% . **A sub-grid fit locates the crest to about 10% at best**, which is why this chapter reports the crest as a range

Reporting a single density would also overstate what this result is about. The crest belongs to the $n = 3$ and $n = 4$ pair specifically. An H_α/H_γ diagnostic, which uses $n = 3$ against $n = 5$, has its worst density at $7.2 \times 10^{12}/\text{cm}^3$, a factor 2.68 lower (`verify_residual_claims.py`, `validation/residual_claims/`). The correct form of words is that the H_α/H_β ratio is least reliable near a certain density, never that the Balmer diagnostic is.

There is no interior maximum in temperature anywhere in the defended range: the error falls monotonically along the crest column (Section 5.6.4), and the apparent peak near 2eV is the floor of the scope restriction of Section 5.10, not a feature of the plasma.

5.3 Why a percentage is not a result

Every number in the previous section is attached to a $+4.81\%$ temperature step, and 4.81% is not a physical quantity. It is what one gets from 50 logarithmically spaced grid points between 1 and 10eV. An examiner is entitled to ask what happens for an ELM that changes the temperature by 10% , and a chapter whose headline is a percentage has no answer.

The answer, over the grid steps available, is that the error grows fast:

Table 5.1: $\varepsilon_{\text{plateau}}$ against the size of the imposed heating step, in grid intervals: the grid maximum with its location, and the same-point series at three fixed operating points. Values from `validation/reservoir_gain/reservoir_gain.csv` (`verify_reservoir_gain.py`, heating direction). The one- and two-interval maxima sit at [0, 4] and the four-interval maximum at [0, 5], all on the cold edge, and none is quantitative (Section 5.10).

step	step in T_e	grid max. of $\varepsilon_{\text{plateau}}$	at [0, 4]	at [15, 3]	at [23, 5]
1	+4.81 %	38.7 % at [0, 4]	38.7 %	18.1 %	6.36 %
2	+9.85 %	87.4 % at [0, 4]	87.4 %	38.2 %	13.9 %
4	+20.68 %	190.4 % at [0, 5]	187.3 %	81.1 %	32.9 %

A factor 4.3 in fractional step size moves the error at a fixed operating point by a

factor 4.8 at $[0, 4]$ (0.3869 to 1.8732), 4.5 at $[15, 3]$ and 5.2 at the benchmark $[23, 5]$; the grid maximum moves by a factor 4.9 only because its location shifts from $[0, 4]$ to $[0, 5]$ at four intervals. Quoted bare, 38.7% is a statement about a grid, and the sentence “*the reservoir-closure error reaches 39%*” is not falsifiable because it does not say what it is 39% of; calling it the quasi-steady-state error, as is common, misnames it as well (Section 5.1.2).

The way out is not a better choice of step. It is to report the response function instead of one of its values, which is what Section 5.4 does. That the error more than doubles between the first and second rows of Table 5.1, while the quantity introduced there moves by 2%, is the reason for preferring the second object to the first.

5.4 The exact response to a finite step

5.4.1 The split has exactly two channels

Everything in this section and the two that follow rests on a partition made in Chapter 3: one slow state, the ground state $1s$, and 42 fast ones. That partition deserves an argument rather than an assertion, because the $2s$ level is a candidate for a second slow state.

The $2s$ level has no single-photon dipole route to the ground state (Section 2.2) and decays instead by two-photon emission at 8.229 1/s , a lifetime of 0.12 s against roughly a nanosecond for the $2p$ level beside it. If $2s$ held population on the slow timescale it would be a second reservoir: the supply sum $\mathbf{n}_F = \mathbf{a} n_{1s} + \mathbf{c} n_{\text{ion}}$ of Eq. (3.33) would gain a third term, the logistic of Eq. (3.40) would not merely be displaced but would change shape, and the bound of Section 5.5 would not follow.

In a plasma $2s$ is not depopulated by radiation. Proton impact mixes it into $2p$, which then decays promptly, and that route carries between 99.92% and 99.99% of the $2s$ loss budget across the grid, the floor falling at 10 eV and the ceiling at the cold edge (`make_story_figures.py`). At the thinnest and coldest point on the grid, $T_e = 1\text{ eV}$, $n_e = 10^{12}/\text{cm}^3$, the total $2s$ loss rate is $|L_{2s,2s}| = 1.2564 \times 10^9\text{ 1/s}$ against a slow eigenvalue $|\lambda_0| = 0.0149\text{ 1/s}$, a separation of 8×10^{10} .

That argument leans on a proton-impact rate coefficient. A second argument needs none. Delete the ℓ -mixing entirely and leave $2s$ with nothing but two-photon emission: 8.229 1/s is still 553 times faster than $|\lambda_0|$ at the same corner. **On this grid $2s$ cannot be slow under any assumption about ℓ -mixing whatever.** The density at which proton ℓ -mixing out of $2s$ falls to the two-photon rate, so that the collisional route stops dominating, solves $n_e \sum_i K_{i,2s}^{\ell\text{-mix}} = 8.229/\text{s}$ and lies between 6.6×10^3 and $1.06 \times 10^4/\text{cm}^3$ over the temperature grid, median $7.9 \times 10^3/\text{cm}^3$. That is 8.10 decades below the bottom of the operating range.

The direct spectral test agrees and supplies the number to quote. Over all 400 grid points the smallest ratio of $|\text{Re } \lambda|$ over the excited block $\mathbf{L}_{FF}$ to $|\lambda_0|$ is 86.5, at $T_e = 10\text{ eV}$, $n_e = 10^{15}/\text{cm}^3$, and the slow eigenvector is the ground state’s at every point.

That last statement needs its norm, because the two norms of Section 3.3 disagree about it more sharply than anywhere else in this thesis. In the *raw* eigenvector the largest component is the ground state at all 400 grid points, with the $2s$ component at most 2.05×10^{-3} of it and typically 10^{-5} : the population the slow mode moves sits in $1s$ and nowhere else. In the *population-scaled* eigenvector $2s$ reaches 0.996 of the ground state’s entry at the coldest, thinnest corner. **That second figure is not evidence of a second reservoir; it is the opposite.** It says the fractional response of $2s$ tracks the

ground state's to within 0.4 %, which is what being slaved to a single reservoir means. An independent reservoir would move on its own timescale and show up as a *difference*, not as agreement. There is no near-degenerate second slow mode. That point and that value are the same stress point and essentially the same quantity as the timescale-separation minimum $M = 86.8$ of Eq. (3.18), computed from the excited block rather than from the full matrix, so the two are a consistency check on each other and not two results. **The worst-case separation is 86.5, and that is the figure that describes the grid;** the benchmark value $M = 9982$ describes one point (Section 3.3).

All four measurements in this subsection are produced by `src/validation/verify_2s_not_slow.py` and stamped in `validation/slow_subspace/`, which carries one row per grid point and raises if the largest raw component of the slow eigenvector is ever anything but the ground state. The two-photon rate is the one number there not read from the pipeline, whose radiative array holds $A(2s \rightarrow 1s) = 0$ (Section 2.2); the matrix this thesis solves is therefore more pessimistic about $2s$ than the atom is.

5.4.2 A better variable than the departure coefficient

Chapter 3 worked in $b_1 = n_{1s}/(Z_1 n_e n_{\text{ion}})$, because that is the variable in which the observable's response has the closed form of Eq. (3.35). Here it is better to use the reservoir itself,

$$u \equiv \frac{n_{1s}}{n_{\text{ion}}} = b_1 Z_1(T_e) n_e, \quad (5.3)$$

a dimensionless ground-state population per unit ion density. At fixed (T_e, n_e) the two are proportional, so $d \ln u = d \ln b_1$ and the sensitivity of Section 3.5.4 carries over unchanged. The reason for the change of variable is that a temperature step moves Z_1 : comparing a pre-step b_1 with a post-step one would mix a change in the reservoir with a change in the Saha–Boltzmann normalisation, and u has no normalisation to move. Comparing pre- and post-step b_1 directly gives a contamination in $\Delta \ln b_1$ comparable to the quantity being measured.

Assumption. Equation (5.3) treats n_{ion} as a fixed reservoir, unchanged by the step, as throughout this thesis (Section 3.1). This removes the freedom for the ionisation balance to respond, and it is what makes u numerically equal to the ground-state population in the solved system. Chapter 6 returns to what a mobile ion population would change, and bounds it: the total bound population per unit n_{ion} at the benchmark is 8.925×10^{-4} , so freezing n_{ion} costs at most 0.09 % (`verify_residual_claims.py`, `validation/residual_claims/`).

5.4.3 The decomposition

Equation (3.35) makes the post-step observable a ratio of two affine functions of u , so the logarithmic sensitivity of Section 3.5.4 is $d \ln R / d \ln u = f_3 - f_4$, the difference of the two ground-fed fractions. Write u^- for the reservoir the plasma arrives with, the *pre-step* equilibrium value $u^{\text{CRE}}(T_e^-)$ that the excited manifold is frozen onto during the plateau, and u^+ for the *post-step* value $u^{\text{CRE}}(T_e^+)$ that the lookup table assumes. Then R^{PE} is the post-step operator evaluated at u^- , and R^{CRE} is the same operator at u^+ . Integrating between the two gives, with no approximation,

$$\ln \frac{R^{\text{PE}}}{R^{\text{CRE}}} = \int_{\ln u^+}^{\ln u^-} (f_3 - f_4) d \ln u = \bar{S} \Delta \ln u, \quad (5.4)$$

where

$$\Delta \ln u \equiv \ln \frac{u^+}{u^-} \quad (5.5)$$

is the reservoir displacement, written post-step over pre-step so that Δ carries its usual sense of new minus old, and

$$\bar{S} \equiv \frac{1}{\Delta \ln u} \int_{\ln u^+}^{\ln u^-} (f_3 - f_4) d \ln u \quad (5.6)$$

is the mean sensitivity over the interval the reservoir traverses. Both factors are dimensionless.

The sign convention, stated once and used throughout. The ground-fed fraction of $n = 3$ exceeds that of $n = 4$ at every point of this grid, so the integrand of Eq. (5.6) is positive everywhere and only the direction of travel changes sign. Heating depletes the ground state, so $u^+ < u^-$, the interval runs forward, and $\Delta \ln u < 0$ while the integral is positive. Cooling refills it, so $\Delta \ln u > 0$ while the interval runs backwards and the integral is negative. **In both directions $\bar{S} < 0$, and in both directions $G < 0$,** the latter because $\Delta \ln u$ and $\Delta \ln T_e$ always carry opposite signs. Measured over all 2288 rows of `reservoir_gain.csv` this holds without exception: $\bar{S}$ runs -0.4963 to -0.0355 and G runs -14.52 to -2.64 .

The product $\bar{S} \Delta \ln u$ is therefore positive for heating and negative for cooling, and $\varepsilon_{\text{plateau}}$ is positive in both because Eq. (5.7) takes a modulus: heating gives $e^+ - 1 > 0$ and cooling gives $1 - e^- > 0$. Nothing observable depends on the orientation, since reversing it flips $\Delta \ln u$ and $\bar{S}$ together and leaves $\varepsilon_{\text{plateau}}$ unchanged, which is why the tables report $|\bar{S}|$ and $|G|$. The convention is spelled out because it must match the script that computes $\Delta \ln u$: the opposite orientation flips the sign of G while leaving every error unchanged.

Since $f_3, f_4 \in [0, 1]$ by Section 3.5.4, and $\bar{S}$ is an average of their difference, $|\bar{S}| < 1$ everywhere. That is the universal bound and it follows from positivity alone. The sharper $|\bar{S}| \leq \tanh(|\Delta|/4)$ of Eq. (3.42) is a separate statement resting on the unit-width logistic, and Section 5.5 keeps the two apart deliberately: $f_m \in [0, 1]$ yields only the weaker. Exponentiating gives the plateau error:

$$\varepsilon_{\text{plateau}} = \left| e^{\bar{S} \Delta \ln u} - 1 \right| \quad (5.7)$$

A single scalar error has become a product of two quantities that can be measured independently: how strongly the observable responds to its reservoir, and how far the reservoir moved. Nothing in Eq. (5.4) says which of them varies across the operating map, and Section 5.6 determines that.

5.4.4 The reservoir gain, and why it is the object to report

The second factor still contains the step. Splitting it,

$$\Delta \ln u = G \Delta \ln T_e, \quad G \equiv \frac{\Delta \ln u}{\Delta \ln T_e} \xrightarrow{\Delta \ln T_e \rightarrow 0} \left. \frac{\partial \ln u^{\text{CRE}}}{\partial \ln T_e} \right|_{n_e}, \quad (5.8)$$

defines the *reservoir gain* G : the factor by which a fractional change in electron temperature is amplified into a fractional change in the equilibrium ground-state population. It is a

secant, taken along the CRE curve between the two endpoints of the step, and it is defined that way deliberately, so that Eq. (5.9) below is exact rather than first-order; it tends to the partial derivative on the right as the step closes. It is dimensionless, it is negative throughout the grid because heating ionises the ground state away, and unlike $\varepsilon_{\text{plateau}}$ it is very nearly a property of the operating point rather than of the experiment done to it. Substituting into Eq. (5.7),

$$\varepsilon_{\text{plateau}} = \left| \exp(\bar{S} G \Delta \ln T_e) - 1 \right|, \quad \lim_{\Delta \ln T_e \rightarrow 0} \frac{\varepsilon_{\text{plateau}}}{\left| \Delta \ln T_e \right|} = \left| \bar{S} G \right|. \quad (5.9)$$

The small-step limit on the right is the useful form: it is the diagnostic error per unit fractional temperature excursion, with the excursion divided out. It is written as a ratio rather than as $d\varepsilon_{\text{plateau}}/d \ln T_e$ for a reason that is easy to miss. $\varepsilon_{\text{plateau}} = |e^z - 1|$ has a corner at $z = 0$: its one-sided derivatives there are $+1$ and -1 , so the two-sided derivative does not exist and a limit written that way is not well posed at the point it is being taken. The ratio of moduli is, and it is the quantity actually wanted. The *signed* logarithmic response is smooth and gives the same content without the modulus,

$$\lim_{\Delta \ln T_e \rightarrow 0} \frac{\ln(R^{\text{PE}}/R^{\text{CRE}})}{\Delta \ln T_e} = \bar{S} G, \quad (5.10)$$

which is the form to differentiate if anything downstream needs a derivative. In fact Eq. (5.10) holds as an exact identity at every finite step, not merely in the limit, because $\bar{S}$ and G are both defined as secants over the step actually taken; writing it as a limit is weaker than the truth and the limit is kept only to make the comparison with a derivative explicit.

At one grid interval $|\bar{S}G|$ is 1.31 at the benchmark, 3.54 at $[15, 3]$ and 6.96 at $[0, 4]$, computed from the values in Table 5.2. **Those are one-interval secants, not the zero-step limit.** Extrapolating linearly in the number of intervals to $k \rightarrow 0$ gives 1.25, 3.66 and 7.41, so the secant sits 5.1, 3.5 and 6.4 % from the tangent at the three points, and the two should not be quoted interchangeably. Read the middle one as: in the worst conditions this thesis defends, a 1 % temperature transient costs the Balmer inversion about 3.5 % in the line ratio.

Table 5.2: The reservoir gain G at three grid points, for steps of one, two and four grid intervals. G moves by 5 to 6 % across a fourfold change in step size, the secant closing on its tangent, quantified below, while $\varepsilon_{\text{plateau}}$ itself changes by a factor 4.5 to 5.2 (Table 5.1). All nine values are the column G of `validation/reservoir_gain/reservoir_gain.csv` (`verify_reservoir_gain.py`), heating direction, and are reproduced from it to the digits printed by `verify_headline_provenance.py`.

point	(T_e, n_e)	$k = 1$	$k = 2$	$k = 4$	spread
benchmark [23, 5]	2.947 eV, $1.389 \times 10^{14}/\text{cm}^3$	-5.736	-5.634	-5.439	5 %
crest [15, 3]	2.02 eV, $1.93 \times 10^{13}/\text{cm}^3$	-7.766	-7.617	-7.333	6 %
cold edge [0, 4]	1.00 eV, $5.18 \times 10^{13}/\text{cm}^3$	-14.434	-14.131	-13.551	6 %

Provenance, and what is not established. The three $k = 1$ entries of Table 5.2 can be checked against independent artifacts, and all three pass. One grid interval is

$\Delta \ln T_e = \ln(10)/49 = 0.0469915$, so $|G| \Delta \ln T_e$ predicts $|\Delta \ln u| = 0.2695, 0.3649$ and 0.6783 at the three points, against measured values of 0.2695 (recovered below from $\varepsilon_{\text{plateau}}$ and $\bar{S}$ at the benchmark), 0.3649 (`verify_ridge_mechanism.py`, as recorded in a project working note) and 0.6783 (`plateau_gridmap.txt`, line 41). Table 5.2's nine values are reproduced to four significant figures by `verify_reservoir_gain.py`, whose output is stamped in `validation/reservoir_gain/reservoir_gain.csv`.

The residual step-dependence is the secant closing on its tangent, and it is not scatter. If G were a derivative it could not take three values at one point, and the three values are not noise around one: $|G|$ falls *linearly* in the number of grid intervals k . Fitting $|G| = A - Bk$ through the $k = 1$ and $k = 4$ entries and predicting the withheld $k = 2$ entry gives $5.637, 7.622$ and 14.140 at the three points of Table 5.2, against measured $5.634, 7.617$ and 14.131 . Carried across every heating point that carries all three step sizes, the withheld prediction is right to a median of 0.057% and a worst case of 0.066% at 314 of 314 points (`verify_secant_linearity.py`, reducing the stamped `validation/reservoir_gain/reservoir_gain.csv`). The intercepts $A = 5.835, 7.910$ and 14.729 are then estimates of the derivative itself. **So the 5 to 6% is a systematic and understood interval effect, not an instability in the coefficient**, and it is one to two orders of magnitude smaller than the factor 8.44 by which $\varepsilon_{\text{plateau}}$ moves over the same range.

Figure 5.5 maps both coefficients and their product over the grid. The two carry different information and different scopes: $|\bar{S}|$ belongs to the line pair and would change if a different α/β ratio were chosen, while $|G|$ belongs to the plasma and would not. Panel (c) is the product, the coefficient this chapter reports in place of a percentage.

Step dependence is not the same for all three, and the ordering is the point. Across step sizes of one, two and four grid intervals at the 736 points carrying all three, the max-to-min spread is

quantity	median	spread	worst
$ G $	1.0555		1.0674
$ \bar{S} $	1.1853		1.7366
$ \bar{S}G $	1.1625		1.8492
$\varepsilon_{\text{plateau}}$	3.7319		8.4437

so $|G|$ is the near-step-independent quantity, varying by at most 6.7% , and the product is not: it moves by up to 85% because the sensitivity $|\bar{S}|$ carries most of the residual step dependence (`verify_reservoir_gain.py`). The same table refutes any description of the product as near-step-independent. What justifies reporting the product rather than a percentage is weaker but still decisive: it is stable to a factor 1.85 where $\varepsilon_{\text{plateau}}$ moves by 8.44 , so a coefficient quoted with its step is a far more portable number than an error quoted with its step.

5.4.5 A worked number at the benchmark

The chain closes arithmetically, and it is worked with signs rather than magnitudes so that the convention above can be checked rather than trusted. At $[23, 5]$ the mean sensitivity over the step is $\bar{S} = -0.2288$ (`verify_reservoir_gain.py`, `validation/reservoir_gain/reservoir_gain.csv` column `Sbar`, heating, $k = 1$; tabulated in Table 5.4 at the *post-step* operator $[24, 5]$, the




Figure 5.5: The two structural coefficients, and the invariant they multiply to. The plateau error obeys $\varepsilon = |\exp(\bar{S}G d \ln T_e) - 1|$ exactly, which splits it into a part belonging to the observable and a part belonging to the plasma. (a) $|\bar{S}| = |f_3 - f_4|$, the sensitivity of the $n = 3/n = 4$ ratio to the ground-state reservoir, where f_n is the ground-fed fraction of shell n . (b) $|G| = |d \ln u / d \ln T_e|$, the reservoir gain, measuring how far the ground population moves for a given change in temperature. (c) Their product $|\bar{S}G|$, which is the small-step limit of $\varepsilon / |d \ln T_e|$ and is the coefficient this chapter reports in place of a percentage that depends on the step chosen. It is written as a limiting ratio and not as $d\varepsilon / d \ln T_e$ because $\varepsilon = |e^z - 1|$ has a corner at $z = 0$, where the two-sided derivative does not exist. G itself is a secant over the step taken, not a local derivative, so it carries a weak and systematic step dependence. All three are heating, one grid index in T_e ($d \ln T_e = 0.0470$), over all 392 cells. Ranges quoted with their scope, because the scope changes them: over the 338 cells drawn here that pass the plateau-window test, $|\bar{S}|$ runs 0.065 to 0.482 and $|G|$ runs 2.64 to 14.52, giving a product of 0.36 to 6.96; over all 2288 rows of the source table, which include cooling and larger steps, $|\bar{S}|$ runs 0.0355 to 0.4963 while $|G|$ is unchanged at 2.64 to 14.52. At the benchmark $|\bar{S}| = 0.229$, $|G| = 5.74$ and the product is 1.31. G is what makes this reframing worth anything: across step sizes of one, two and four grid indices at the 736 points carrying all three, $|G|$ varies by at most 6.74% (median 1.0555) while ε over the same points varies by up to a factor 8.44. Hatched cells (54 of 392) fail the plateau-window test $M > 900$; their values are drawn rather than deleted, since $\bar{S}$ and G are properties of the operator and remain defined there even where a plateau does not.

one the bound of Eq. (5.11) refers to), the gain is $G = -5.736$, and one heating interval is $\Delta \ln T_e = \ln(10)/49 = +0.0469915$. Two negatives and a positive give

$$\overline{S} G \Delta \ln T_e = (-0.2288) \times (-5.736) \times (+0.0469915) = +0.061672, \quad e^{+0.061672} - 1 = +0.063613,$$

against the 0.063612 that `verify_plateau_gridmap.py` obtains for the same point by an exact linear solve. Three quantities from three different artifacts multiply to the recorded error to four significant figures. The arithmetic is performed here in the text and is not itself a run.

5.4.6 Use the exact form, not the linearisation

The first-order form $\varepsilon_{\text{plateau}} \approx |\overline{S} \Delta \ln u|$ is what one would write down first, and it is adequate for a small step and wrong for a large one. At the four-interval step it returns 59.4% where Eq. (5.7) gives 81.1% at [15, 3] (`verify_ridge_mechanism.py`). The exponential costs nothing and this chapter uses it throughout.

The linearisation does not err in one direction only, so it cannot be quoted as a bound; which way it errs depends on which linearisation is meant. For the *endpoint* form computed at `verify_plateau_gridmap.py`, line 212, the ratio of measured to predicted error over the heating map runs 0.8417 to 1.3891 with median 1.0310: the linearised expression understates in 219 of 392 heating steps and overstates in the rest, and at the benchmark it overstates, with ratio 0.9459 (`verify_residual_claims.py`, `validation/residual_claims/`, which reproduces the range, the count of 219 and the benchmark value; over the 392 heating steps it obtains a median of 1.018, so the set the 1.031 is taken over is not recorded). **Any statement about the linearisation must name which linearisation it is about.**

A final caution about what the decomposition can and cannot be used for. Equation (5.4) is an identity, so comparing its two sides tests nothing: a check that $\varepsilon_{\text{plateau}}$ equals $|e^{\overline{S} \Delta \ln u} - 1|$ can fail only through arithmetic error. What the identity permits is an *attribution*. The two factors are separately measurable, so one may ask which of them carries the structure the maps show. That is a weaker kind of statement than a derivation, and it is the kind made in Section 5.6.

5.5 A bound that no atomic dataset can exceed

Section 3.5.4 showed that in the logarithmic reservoir variable both ground-fed fractions are the same unit-width logistic curve, displaced by $\Delta = \ln[(a_3/a_4)/(c_3/c_4)]$, and that their difference therefore has the closed-form maximum $\max |f_3 - f_4| = \tanh(|\Delta|/4)$ of Eq. (3.42). Since $\overline{S}$ of Eq. (5.6) is an average of $f_3 - f_4$ over an interval, it inherits the bound:

$$|\overline{S}| \leq \tanh\left(\frac{|\Delta|}{4}\right) < 1, \quad \Delta \text{ evaluated on the } \textit{post-step} \text{ operator.} \quad (5.11)$$

The qualifier is load-bearing. $\overline{S}$ averages $f_3 - f_4$ along the post-step curve, so the displacement that caps it is the post-step one. Tested over all 2288 rows of `reservoir_gain.csv`, the bound holds with **zero** exceptions when Δ is read on the post-step operator, and is violated at 101 of them, 4.4%, worst excess 0.0087, when it is read on the pre-step operator that labels each row.

Three properties make this worth stating separately from the algebra that produced it.

There are two bounds here, and only one of them is independent of the atomic data. The distinction is easy to lose and an examiner will not lose it.

The *universal* bound is $|f_3 - f_4| < 1$, which follows from $f_3, f_4 \in [0, 1]$ and therefore from nothing but the positivity of $\mathbf{a}$ and $\mathbf{c}$ established in Section 3.5.4. It holds for any non-negative rate data whatever, for any pair of shells, and it is what licenses the amplification statement below. Nothing about hydrogen enters it.

The *sharp* bound $\tanh(|\Delta|/4)$ is stronger and is not universal in the same sense. Its *functional form* is structural: the $\tanh$ is a consequence of the logistic having unit width, and the width is unity because the excited populations are exactly affine in n_{1s} , which is a property of the linear system and not of any rate coefficient (project working note, no producing script; Table 4.4). But its *numerical value* is set by $\Delta = \ln[(a_3/a_4)/(c_3/c_4)]$, and those four coefficients are the atomic data. **Change the rate set and the cap moves.** Every property of the 42-state network enters exactly there, and only there. What *is* established about Δ is invariance under two specific choices, which is a narrower claim than robustness to the rate data and should not be read as it. Δ is invariant under rescaling the n_{1s}/n_{ion} normalisation, verified exactly under factors of 10^{10} and 10^{-7} , and is insensitive to the ℓ -weighting used to bundle a shell: population sum, p -states, d , s and statistically weighted give 1.94561, 1.94474, 1.94447, 1.94471 and 1.94574 at the benchmark, a spread of 0.07 % (`verify_residual_claims.py`, `validation/residual_claims/`). The individual switching points x_3 and x_4 and the fractions f_p are *not* invariant under that rescaling; only their separation is.

Its physical content is a statement about amplification. Because $|\mathrm{d} \ln R / \mathrm{d} \ln b_1| < 1$ everywhere, integrating over any excursion gives $|\Delta \ln R| \leq |\Delta \ln b_1|$. A line ratio can never respond to the ground-state reservoir faster, in relative terms, than the reservoir itself moves. A stale reservoir cannot be amplified in logarithmic space, only attenuated.

The bound is loose in general and tight at the crest. At the benchmark $|\bar{S}| = 0.2288$ against a cap of 0.4503, so the sensitivity is at 51 % of what the atomic structure permits. At [23, 3], one column into the crest, $|\bar{S}| = 0.4261$ against a cap of 0.4394, that is 97 % (both from `verify_residual_claims.py`, `validation/residual_claims/`, with the cap taken at the post-step index as Table 5.4 does; the percentages are arithmetic on those two columns). Where the error is largest, the diagnostic is running at essentially the maximum sensitivity to the reservoir that *this rate set* permits.

It does not follow that no refinement of the rate coefficients could make that worse. Sitting at 97 % of the cap says there is little headroom *beneath* the cap; it says nothing about where the cap itself would move, and the cap moves with Δ , which is atomic data. A refinement can raise $\tanh(|\Delta|/4)$ and can move u^{peak} relative to the operating point, in either direction.

What is measured is insensitivity to two specific perturbations of the rate set. Re-building $\mathbf{L}$ with a density-consistent Debye cutoff at every grid point moves $f_3 - f_4$ by at most 0.42 %, and scaling every ℓ -mixing rate over a range of $\times 0.1$ to $\times 10$ moves it by under 0.3 % (Section 2.7). Both concern ℓ -mixing.

One refinement of the excitation dataset is now tested, and it is the one the benchmark of Section 2.8 makes available. Substituting the RMPS rate coefficients wherever they exist and re-solving the steady state moves Δ by 9.1 % on average over the 448 defended pairs and by at most 17.8 %, and moves the cap with it, by 8.4 % on average (`src/validation/verify_reservoir_gain_anderson.py`; the means over every row of the file, all three step sizes and the cold edge included, are 7.1 % and 6.5 %).

What does not move is the *headroom*. $|\bar{S}|$ and its cap fall together: the median of $|\bar{S}|$ as a fraction of its own cap goes from 80.4% to 78.2%, and its maximum from 99.9% to 99.7%. The bound is violated at none of the 338 points under either rate set. Running at 97% of the cap is therefore not an artefact of this particular rate set.

The recombination channel can be substituted too, and it behaves differently enough to be worth separating. $\alpha_{3\text{BR}}$ is not independent data: Eq. (2.6) builds it from the ionisation coefficient by detailed balance, so substituting the ionisation dataset substitutes the three-body recombination with it. The model takes ionisation from CCC TICS below $n = 9$ and from Lotz above it, so Lotz is an independent dataset already in hand, and running the whole table from Lotz instead replaces both at once. Over the states where CCC supplies the model, Lotz is larger by a factor between 1.16 and 5.32, median 2.13, a far bigger perturbation than the excitation substitution, and one that changes the ionisation loss on the diagonal of $\mathbf{L}$ as well as the source vector.

It moves Δ by 1.9% and the cap with it, and $\varepsilon_{\text{plateau}}$ by 16.4%, over the same 448 pairs. The disparity is the point, and it is sharpest in what each swap does to the headroom. Over the 338 heating pairs on which the bound was checked, the excitation swap moves $|\bar{S}|$ by 8.4% and the cap by 6.7% on average, in absolute terms, and moves them together: the ratio $|\bar{S}|/\tanh(|\Delta|/4)$ changes by 4.3% on average and by -1.4% at the median pair. The ionisation swap moves $|\bar{S}|$ by 17.8% on average inside a cap that moves by 1.5%, and it does so in both directions: the ratio changes by 17.7% on average, downward at 62 percent of the pairs and by -8.3% at the median pair, while the median of the headroom distribution itself is unchanged, 80.4% before and 80.5% after. The ionisation data do not push the operating point one way; they scramble where each pair sits beneath its cap. The two datasets act on different objects. Δ is a ratio of ratios, $(a_3/a_4)/(c_3/c_4)$, and a change in the ionisation coefficients that is similar at $n = 3$ and $n = 4$ largely cancels in it; what such a change moves is where the reservoir sits, which is the variable Section 6.4 identifies as the one that controls $|\bar{S}|$. **Refining the excitation data moves the cap; refining the ionisation data moves the operating point beneath it, pair by pair and in both directions.** Under all four rate sets the bound is violated at none of the 338 points.

Radiative recombination is the one input with no second dataset here. It is Johnson's Eq. (7), quoted as accurate to 5% against exact Karzas–Latter Gaunt factors, and scaling it by that much, a sensitivity test rather than a substitution and not to be quoted as a dataset comparison, moves $\varepsilon_{\text{plateau}}$ by 1.1% and at most 2.4%. Radiative recombination is not where this uncertainty lives.

5.5.1 What the bound does not constrain

A diagnostician does not suffer an error in R . They measure R and invert it. At known n_e , Eq. (3.44) makes the inferred temperature error the ratio of two derivatives, both taken at fixed n_e : the sensitivity of the observable to the reservoir over its sensitivity to temperature. Equation (5.11) constrains the numerator. It says nothing about the denominator, and Chapter 3 deferred the denominator to this chapter. If n_e is not independently known the relevant object is the 2×2 Jacobian of Eq. (3.45) instead, which this thesis does not measure.

The denominator vanishes. Scanning T_e at $n_e = 5.18 \times 10^{13}/\text{cm}^3$, $\varepsilon_{\text{step}}$ of Eq. (5.1) passes through a minimum of 4.9×10^{-5} near $T_e = 6.87 \text{ eV}$, with values 7.18×10^{-4} , 3.79×10^{-4} , 4.9×10^{-5} , 2.74×10^{-4} and 5.89×10^{-4} at 6.251, 6.551, 6.866, 7.197 and

7.543 eV (`verify_plateau_gridmap.py`). Since $\varepsilon_{\text{step}}$ is the fractional distance between the two tabulated answers either side of the step, a zero means R^{CRE} is stationary in temperature there: $\partial \ln R / \partial \ln T_e = 0$.

The $n = 3/n = 4$ ratio is locally non-invertible for temperature near 6.9 eV at this density. There, a bounded error in the observable produces an unbounded error in the inferred temperature, and the plateau error is not the relevant failure mode at all: the diagnostic has no temperature resolution to lose. This is a separate failure from the one the rest of the chapter maps, it is arguably the more severe of the two for a practitioner, and it is visible only because $\varepsilon_{\text{step}}$ was computed and kept rather than discarded as uninteresting.

The scan above is one density column, and the locus is not confined to it. Evaluating R^{CRE} on the full grid and locating where $\partial \ln R / \partial \ln T_e$ changes sign (`verify_crest_subgrid.py`, `validation/crest_subgrid/eps_step_locus.csv`) puts a stationary point inside the operating range in **six of the eight density columns**, at 5.18 to 8.69 eV, moving to higher temperature as the density rises. It leaves the range only at the two highest densities, $n_e \geq 3.73 \times 10^{14} / \text{cm}^3$, where R is monotone in T_e throughout and the smallest $\varepsilon_{\text{step}}$ on the column is 1.1×10^{-3} rather than the 10^{-5} to 10^{-4} found elsewhere.

Where the locus exists the diagnostic is locally blind: a temperature excursion of any size moves the tabulated ratio by almost nothing, so the inversion cannot see it at all. **That is a different failure from the one this thesis maps, and a worse one for a practitioner, because it has no transient in it and persists in steady state.** Its consequence for the inversion is measured in Section 5.5.2: 44 of the 218 heating pairs in the defended scope reach their inferred temperature only by crossing it. It also bears on the non-uniqueness of the inversion: a stationary point in $R(T_e)$ is a fold, and on either side of it two temperatures return the same ratio.

5.5.2 The inversion, measured

Equation (3.44) is a linearisation, and Chapter 3 promised the exact finite-step form wherever a number is quoted. The exact form is a root-find. At known n_e , the temperature the table returns for the plateau observable R^{PE} is every T_e^* on that density column at which $R^{\text{CRE}}(T_e^*, n_e) = R^{\text{PE}}$, and the inferred-temperature error is $\ln T_e^* - \ln T_e^+$, with T_e^+ the true post-step temperature. That root-find was run at every one of the 784 point-and-direction pairs of Section 5.1.1, with the same step rule, the same plateau-window rule and the same three linear solves (`verify_inversion_error.py`). It reproduces the stamped `plateau_gridmap.csv` with a worst difference of zero, and on the two monotone density columns the sign of every inferred error equals the sign of the observable error times the sign of $d \ln R^{\text{CRE}} / d \ln T_e$, 180 of 180, a check an inverted interpolation would fail at every pair.

In more than a third of the defended scope the table has no solution inside its own domain. Of the 448 pairs at $T_e \geq 2$ eV, 166 carry a plateau ratio below the smallest value the table holds anywhere on that density column, so no T_e^* exists within the tabulated 1 to 10 eV, and every one of them is a cooling step. That count belongs to the table's edges as much as to the plasma, and it is quoted with the domain attached wherever it appears. A further 44, all heating, place T_e^* on the far side of a stationary point of Section 5.5.1 from T_e^+ , so the inversion reached its answer only by crossing a fold. That leaves 238 pairs with a unique, fold-free answer, and those carry the numbers below.

Over those 238 pairs a one-interval temperature step of 4.81 % is read back

as a temperature error with median 47.7 % ($|\Delta \ln T_e^*| = 0.390$), a 90th percentile of 96 % and a worst case of 144 %. As an amplification of the true excursion the median is 8.3, the 90th percentile 14.3 and the largest 19.0. Including the 44 fold-crossed pairs at their nearest root raises the median to 56 % and the amplification to 9.5; the fold-free figures are the ones quoted, because a nearest root on the far side of a fold is a choice and not a measurement. Heating and cooling differ, at medians of 55 % over 174 heating pairs and 38 % over 64 cooling pairs.

What is a theorem. At fixed n_e the table is the composition $R^{\text{CRE}}(T_e) = R(T_e, u^{\text{CRE}}(T_e))$, so its slope splits by the chain rule into the manifold's own response and the response through the reservoir,

$$\Sigma \equiv \frac{d \ln R^{\text{CRE}}}{d \ln T_e} = P + SG, \quad P \equiv \left(\frac{\partial \ln R}{\partial \ln T_e} \right)_u, \quad S \equiv \left(\frac{\partial \ln R}{\partial \ln u} \right)_{T_e} = f_3 - f_4, \quad G \equiv \frac{d \ln u^{\text{CRE}}}{d \ln T_e}, \quad (5.12)$$

exactly and locally. Here S carries the natural sign of a derivative and is positive for the (3, 4) pair; the interval average $\bar{S}$ of Eq. (5.6) is the same quantity with the opposite sign, by its limits of integration, and the two must not be confused. On the plateau the reservoir is frozen, so to first order the observable moves by $\delta \ln R^{\text{PE}} \simeq P \delta \ln T_e$, while the table, which knows only equilibrium reservoirs, can account for any movement only at its own slope Σ . Inverting the plateau ratio against the table therefore returns $\delta \ln T_e^* \simeq (P/\Sigma) \delta \ln T_e$, and the error in the inferred temperature is

$$\delta \ln T_e^* - \delta \ln T_e \simeq -\frac{SG}{\Sigma} \delta \ln T_e = -\frac{SG}{P + SG} \delta \ln T_e, \quad (5.13)$$

which is Eq. (3.44) with the reservoir displacement written as $\delta \ln u = G \delta \ln T_e$. The amplification of the true excursion is $|SG/\Sigma|$, and it diverges where $\Sigma = P + SG \rightarrow 0$. The stationary point of Section 5.5.1 is therefore not a second, unrelated failure: it is the limit in which the two channels cancel exactly.

The finite step carries the same decomposition, exactly. Over one grid interval $D = \Delta \ln T_e$ the three quantities are secants, and all three come out of the three solves that give $\varepsilon_{\text{plateau}}$:

$$\bar{P} = \frac{\ln [R^{\text{PE}}/R^{\text{CRE}}(T_e^-)]}{D}, \quad \bar{SG} = \frac{\ln [R^{\text{CRE}}(T_e^+)/R^{\text{PE}}]}{D}, \quad \bar{\Sigma} = \frac{\ln [R^{\text{CRE}}(T_e^+)/R^{\text{CRE}}(T_e^-)]}{D} = \bar{P} + \bar{SG} \quad (5.14)$$

where the identity is a telescoping sum of logarithms. The reservoir secant is not a product of two separately averaged factors: by Eq. (5.4) and Eq. (5.8) it equals $-\bar{S}G$ exactly, the sign being the one just noted, and `verify_operator_slope_decomposition.py` confirms both identities to 5×10^{-15} at all 784 pairs.

What is measured. Two signs hold at every one of the 448 pairs in scope: $\bar{P} > 0$, the (3, 4) ratio rising with the temperature of the manifold at a frozen reservoir, and $\bar{SG} < 0$, which is $S > 0$ times $G < 0$, heating emptying the reservoir. Neither is derived here; both are properties of this operator over the defended range. The net slope is negative, $\bar{\Sigma} < 0$, at 359 of the 448, and the local slope at the post-step node at 360. The two channels are nearly equal and opposite: medians $|\bar{P}| = 1.07$ and $|\bar{SG}| = 1.12$, while the per-pair cancellation ratio $|\bar{\Sigma}|/|\bar{P}|$ has median 6.8 % across the 448 pairs. The median of $|\bar{\Sigma}|$ is separately 0.091; dividing the printed medians gives 8.5 %, a coarser summary of the same cancellation, and the two should not be conflated.

Direction: a local result, and its two boundary exceptions. With $P > 0$ and $SG < 0$, Eq. (5.13) places the inferred temperature on the side the plasma came from wherever $\Sigma < 0$: colder than T_e^+ on heating, hotter on cooling. The exact root-find agrees with no exception, 172 of 172 heating pairs and 64 of 64 cooling pairs with $\Sigma(T_e^+) < 0$. Where $\Sigma(T_e^+) > 0$, above the fold, the same equation predicts a hotter branch on heating. There are 46 such heating pairs in scope. At 44 the tabulated ratio does not climb by the size of the offset before the table ends at 10 eV, and the nearest root is a colder one reached across the fold, which is the fold-crossed class above; at the other 2 a colder root lies inside the step interval itself. A local slope predicts the local direction and not the existence of a root over a finite interval, so that last statement is empirical and belongs to this table and this step.

Amplification: the ratio of the reservoir term to the residual slope. A one-interval transient displaces the ratio by $|\overline{SG}|$ while the table reads the displacement back at $|\overline{\Sigma}|$, so the finite-step prediction for the amplification is $|\overline{SG}|/|\overline{\Sigma}|$, large wherever the two channels nearly cancel. Against the exact root-find over the 238 fold-free pairs it has a rank correlation of 0.93 and overstates, the median of the pair-by-pair ratio being 1.26; the local form $|\overline{SG}|/|\Sigma(T_e^+)|$ overstates by 18 %. Per column the exact median amplification runs 8.2, 10.2, 12.7, 12.7, 10.7 and 6.8 across the six columns from 10^{12} to $1.39 \times 10^{14}/\text{cm}^3$ that carry a stationary point, then 3.4 and 1.7 on the two monotone columns above; over the 56 pairs at $n_e \geq 3.73 \times 10^{14}/\text{cm}^3$ the median inferred-temperature error is 16 % and the largest 19 %. The columns that carry the fold are the columns that carry the amplification because both are the same cancellation, one at its limit. The off-table cooling cases are the same physics again: at those 166 pairs $|\overline{P}|$ exceeds $|\overline{\Sigma}|$ by a median factor 19.7, against 4.8 where a cooling root exists, and a ratio falling that much faster than the equilibrium curve leaves the curve's range.

Why $P > 0$: a decomposition, not a proof. At fixed u the shell populations are affine, so with $A_m = \partial \ln a_m / \partial \ln T_e$ and $C_m = \partial \ln c_m / \partial \ln T_e$,

$$P = \underbrace{f_3 A_3 - f_4 A_4}_{P_a} + \underbrace{(1 - f_3) C_3 - (1 - f_4) C_4}_{P_c}, \quad (5.15)$$

exactly. Over the step, with A_m and C_m as secants, the ground-fed coefficients rise with temperature and the recombination-fed coefficients fall at every pair: $A_m > 0$ and $C_m < 0$ at 448 of 448. Given those signs, $P_a > 0$ if and only if $(f_3/f_4) / (A_4/A_3) > 1$, and it is, at 448 of 448, median 2.9 and never below 1.2: the ground-fed channel is the larger share of $n = 3$ by more than A_4 exceeds A_3 . And $P_c > 0$ if and only if $[(1 - f_4)/(1 - f_3)] / (|C_3|/|C_4|) > 1$, which holds at 426 of 448, median 2.0, and fails at 22, all at $n_e = 10^{15}/\text{cm}^3$ between 2.0 and 3.4 eV, where both shells are almost entirely recombination-fed ($f_3 = 0.08$, $f_4 = 0.01$), the share factor is close to unity, and the sign is decided by $|C_3|$ exceeding $|C_4|$. The sum $P_a + P_c$ is positive at all 448 and reproduces $\overline{P}$ to a median 3.7 %, a 90th percentile of 13 % and a worst case of 35 %; the residual is the second-order part of the exact finite-step form, the logarithm of a mixture, which the script verifies to 10^{-14} . So $P > 0$ is an empirical property of the present operator over the defended range, associated with the unequal supply fractions of the two shells, and not a general consequence of $f_3 > f_4$ alone; twenty-two pairs already refuse the stronger channel-by-channel claim. Every number in this paragraph is stamped in `validation/operator_slope_decomposition/`.

Three cautions travel with these numbers. They are single-parameter: n_e is taken as known, and the two-parameter Jacobian of Eq. (3.45) is still not measured. They

are plateau values and not the time averages of Section 5.8: they say what the table returns while the reservoir is displaced, not what an exposure of a given length records. And they inherit the frozen-reservoir conditional of Chapter 6 exactly as $\varepsilon_{\text{plateau}}$ does, because they are $\varepsilon_{\text{plateau}}$ pushed through the table. One definitional choice was tested and does not matter: interpolating the table linearly in (T_e, R) instead of $(\ln T_e, \ln R)$ moves the fold-free median by 0.01 %. Every figure in this subsection is stamped in `validation/inversion_error/` against the same `L_grid` and `S_grid` hashes as the rest of the chapter.

5.6 Where the error is largest, and why

The maps of Section 5.2 show where the diagnostic error is large, not why. The structure along the density axis and the structure along the temperature axis have different causes, and reading them as one effect gets both wrong.

5.6.1 Across density: the operating point, not the span

Evaluating both factors of Eq. (5.4) at every grid point settles which of them locates the crest (`verify_ridge_mechanism.py`, over $T_e = 2\text{--}9\text{ eV}$ and all eight density columns). The column maximising $\varepsilon_{\text{plateau}}$ is the column maximising $|\bar{S}|$ in 32 of 32 temperature rows for a one-interval step and 31 of 31 for a four-interval step; the column maximising $|\Delta \ln u|$ coincides with the error maximum in 0 of 32 and 0 of 31.

Coincidence counts are weak evidence on their own, because a quantity that barely varies still has its maximum somewhere. The amplitudes settle it. Across the three decades of density $|\bar{S}|$ varies by a factor 7.11 and $|\Delta \ln u|$ by a factor 1.01. **A quantity that changes by one per cent cannot produce a crest in which the error changes by a factor of seven**, wherever its own maximum falls.

Table 5.3: Density profiles of the plateau error and of the two factors of Eq. (5.4), each normalised to its own row maximum and then taken as the median over the 32 temperature rows with $T_e = 2\text{--}9\text{ eV}$, for a one-interval temperature step. From `verify_ridge_mechanism.py`.

n_e (1/cm ³)	10^{12}	$2.7 \cdot 10^{12}$	$7.2 \cdot 10^{12}$	$1.9 \cdot 10^{13}$	$5.2 \cdot 10^{13}$	$1.4 \cdot 10^{14}$	$3.7 \cdot 10^{14}$	10^{15}
$\varepsilon_{\text{plateau}}$	0.437	0.683	0.937	1.000	0.807	0.512	0.280	0.138
$ \bar{S} $	0.445	0.689	0.938	1.000	0.813	0.520	0.285	0.141
$ \Delta \ln u $	0.997	0.995	0.993	0.991	0.990	0.992	0.996	1.000

The normalised density profiles of $\varepsilon_{\text{plateau}}$ and of $|\bar{S}|$, the first two rows of Table 5.3, agree to 2.2 % at every column, and the $|\Delta \ln u|$ row is flat to one per cent. The density structure of the plateau error is the density structure of the observable’s sensitivity to its reservoir.

That settles *which factor*. It does not settle why that factor has an interior maximum, and here the explanation the bound suggests is wrong. Equation (5.11) caps the sensitivity at $\tanh(|\Delta|/4)$, so the natural guess is that Δ itself peaks near $2 \times 10^{13}/\text{cm}^3$: that the two supply channels are most widely separated there, and the sensitivity follows. That guess was written down in advance by an independent re-derivation of Chapter 3, together with

a structural argument for why Δ should collapse at high density, namely that in local thermodynamic equilibrium (LTE) both channels would give the same Boltzmann shape and Δ would go to zero.

It is falsified. Δ is nearly flat across the density axis: 0.9655 at $10^{12}/\text{cm}^3$, 1.9403 at $1.39 \times 10^{14}/\text{cm}^3$ and 1.9368 at $10^{15}/\text{cm}^3$, at the post-step temperature $T_e = 3.089 \text{ eV}$ of the benchmark heating step, temperature index 24 (`validation/fujimoto_r1_consequence/`; the triplet belongs to index 24, not to the benchmark temperature 2.947 eV , index 23). Above $7 \times 10^{12}/\text{cm}^3$ the cap $\tanh(|\Delta|/4)$ spans only 0.394 to 0.450, a 14 % range. The grid never reaches LTE, and because the model is open, with $\mathbf{S}_F$ an externally imposed source not tied to n_{ion} by the Saha relation, the two channels cannot merge by construction (`verify_residual_claims.py`, `validation/residual_claims/`).

What sweeps instead is the operating point. Table 5.4 puts the two side by side: while the cap moves by 14 %, the position of the CRE reservoir relative to the sensitivity peak, $\ln(u^{\text{CRE}}/u^{\text{peak}})$, runs from +1.29 to -3.29 , and $|\bar{S}|$ tracks it.

Table 5.4: Why the sensitivity has an interior maximum in density, at the post-step temperature $T_e = 3.089 \text{ eV}$ (temperature index 24) of the benchmark heating step. The cap $\tanh(|\Delta|/4)$ is nearly constant while the operating point $\ln(u^{\text{CRE}}/u^{\text{peak}})$ sweeps through zero. $\varepsilon_{\text{plateau}}$ is for a one-interval, +4.81 % step. Every value is reduced from stamped artifacts by `verify_headline_provenance.py`: Δ and $\tanh(|\Delta|/4)$ from `validation/molecular_channel/` at the post-step index 24, and $|\bar{S}|$ from `validation/reservoir_gain/`, where it is the secant of the step $23 \rightarrow 24$ and is stored against index 23. The $\ln(u^{\text{CRE}}/u^{\text{peak}})$ column takes its numerator at index 23 and its denominator at index 24; evaluating it consistently at either index shifts every entry by about 0.25 and moves neither the sign change, which stays between $j = 2$ and $j = 3$, nor the maximum of $|\bar{S}|$ at $j = 3$.

j	$n_e \text{ (1/cm}^3\text{)}$	Δ	$\tanh(\Delta /4)$	$\ln(u^{\text{CRE}}/u^{\text{peak}})$	$ \bar{S} $
0	1.00×10^{12}	0.9655	0.2368	+1.2877	0.1754
2	7.20×10^{12}	1.6662	0.3940	+0.3842	0.3884
3	1.93×10^{13}	1.8857	0.4394	-0.2487	0.4261
5	1.39×10^{14}	1.9403	0.4503	-1.7695	0.2288
7	1.00×10^{15}	1.9368	0.4496	-3.2910	0.0668

The maximum sits where $\ln(u^{\text{CRE}}/u^{\text{peak}})$ changes sign. The physical reading is direct. Both f_3 and f_4 rise from zero to unity as the reservoir grows, and they switch from recombination-fed to ground-fed at different reservoir strengths, so their difference vanishes in both supply limits and peaks between them. The sensitivity curve is fixed. Raising the density slides the plasma’s own CRE reservoir along it, from the ground-fed side at low density, through the mixed-supply region, to the recombination-fed side at high density. The crest is where that slide crosses the peak.

This is a statement about where the operating point sits, not about how sensitive the diagnostic could be, and it forbids the inference that the effect disappears at high density. The cap does not close; over this grid the operating point walks off the peak, and Table 5.4 shows $\ln(u^{\text{CRE}}/u^{\text{peak}})$ falling monotonically to the grid’s top density. Whether the operating point meets the peak again beyond that density is not computed here.

5.6.2 Three tests of the mechanism

The account above predicts more than the shape of one curve, and each prediction can fail. The crest positions and the boundary level in these tests are computed from the same rate matrix that produces the maps, so the third test compares the operator with itself; the first two compare it with Griem’s criterion, which is external to the operator; none of the three is experimental.

The crest must move with the shell pair, and by a predicted amount.

The reservoir strength at which a shell switches from recombination-fed to ground-fed is a version of Griem’s LTE criterion [22, 23], which places the boundary between collisionally coupled and radiatively decoupled levels at a density scaling as $n^{-17/2}$ in the principal quantum number. Moving the diagnostic from the (3, 4) pair to the (4, 5) pair should therefore move the crest by $(5/4)^{8.5} = 6.7$. Measured: 7.2. The argmax of $|f_n - f_{n+1}|$ at $T_e = 2.02$ eV runs $3.73 \times 10^{14}/\text{cm}^3$ for (2, 3), $1.93 \times 10^{13}/\text{cm}^3$ for (3, 4) and $2.68 \times 10^{12}/\text{cm}^3$ for (4, 5), with (5, 6) and above falling off the bottom of the grid (`verify_ridge_mechanism.py`).

The absolute density must match, not just the ratio. Griem’s LTE criterion evaluated at $n = 4$ and $T_e = 2$ eV gives $2.05 \times 10^{13}/\text{cm}^3$ against a measured crest at $1.93 \times 10^{13}/\text{cm}^3$, a 6 % difference.

A second diagnostic of the same operator must put the boundary level in the same place. `verify_boundary_descent.py` computes, without reference to f_3 or f_4 , the effective principal quantum number at which the collisional and radiative rates balance. At the crest density it returns $\bar{n} = 3.08$ to 3.25 at every temperature on the grid (`validation/boundary_descent.csv`), that is, between $n = 3$ and $n = 4$, which is exactly the pair whose supply fractions differ most.

What would have refuted the mechanism, and did not occur: the crest failing to move with the shell pair, or moving the wrong way. A competing explanation was refuted outright. The suggestion that both shells become ground-fed at high density because collisional coupling dominates predicts $f_3, f_4 \rightarrow 1$; measured at $T_e = 2.947$ eV, $n_e = 10^{15}/\text{cm}^3$, they are 0.071 and 0.011, so both shells become *continuum*-fed instead (`verify_residual_claims.py`, `validation/residual_claims/`).

5.6.3 Alternatives that would produce the same appearance

An interior maximum on a coarse grid can be manufactured in several ways, and the attribution above does not by itself exclude them.

Grid resolution. The grid is spaced by a factor 2.683, so the crest location is known only to that factor. What is established is that the maximum is interior and does not move between temperature rows above 1.15 eV; locating it more precisely requires rebuilding and revalidating the rate matrix on a finer density grid, which is future work.

Sub-grid drift. A parabolic fit in $\ln n_e$ through the peak and its two neighbours puts the vertex at $3.70 \times 10^{13}/\text{cm}^3$ at 1 eV. It falls to $1.45 \times 10^{13}/\text{cm}^3$ at 5.96 eV and then rises again to $1.52 \times 10^{13}/\text{cm}^3$ at 9.54 eV, so the largest and smallest vertices differ by a factor 2.56 and the drift is not monotone in temperature: it reverses once, above about 6 eV (`verify_residual_claims.py`, `validation/residual_claims/`, reducing `validation/crest_subgrid/`). Griem’s criterion predicts a drift upward as $\sqrt{T_e}$, a factor 3.2 across the same range, which the measured vertex does not follow in either direction over the whole grid. The whole drift is smaller than one grid interval and therefore unresolved, so the crest should be described as constant to within a factor 2.68

over 1 to 10 eV, never as fixed. The likely reason for the sign is that the neutral fraction n_{1s}/n_{ion} collapses across the grid, dragging the supply crossover down in density faster than the Griem boundary drags it up, which is consistent with Section 5.6.1 but has not been separated from it.

Step-size convention. A crest whose position depended on the size of the imposed step would be an artifact of that choice. It does not: the error maximum falls on the same density column for a one-interval and a four-interval step in 31 of the 31 temperature rows both can reach, and so does the maximum of $|\bar{S}|$ (`verify_ridge_mechanism.py`). This tests steps of fixed *fractional* size; a step of fixed absolute size in electron-volts spans a different number of grid intervals at each temperature and is a different experiment, not tested here.

Truncation. The model resolves $n \leq 8$ and bundles $n = 9$ –15. The mechanism turns on the difference between the $n = 3$ and $n = 4$ supply fractions, both fully resolved, but those fractions depend on the recombination flux cascading from above, which the bundling approximates. `verify_bundling_psm20.py` (Section 6.8.3) answers a different question, whether the bundled shells would be statistical in ℓ if the mixing were present, and does not test this. The sensitivity of the crest *position* to n_{max} is untested, and cannot be tested with the atomic data this model has: it is carried as the second of the five open questions in Section 6.10.

The population set. Two of the scripts used here apply different filters: `verify_plateau_gridmap.py` excludes the 54 points with no plateau window, and `verify_ridge_mechanism.py` does not. All 54 lie at $j \geq 4$, so the ridge script’s statistics for the two highest density columns include points at which the plateau state being computed does not physically exist (`verify_residual_claims.py`, `validation/residual_claims/`). The attribution of Section 5.6.1 does not depend on those columns, but a reader recomputing the profile from the raw output should apply the filter.

5.6.4 Across temperature: a displacement effect

The temperature dependence has the opposite explanation, and this is where treating the two axes alike would go wrong.

Along the crest column the plateau error falls monotonically from 18.07 % at $T_e = 2.02$ eV to 4.55 % at 8.29 eV, a factor 3.72, both per +4.81 % step. Decomposed through Eq. (5.4), $|\Delta \ln u|$ contributes a factor 2.72 of that fall and $|\bar{S}|$ only 1.29; in logarithmic terms the displacement carries 76 % of the decay and the sensitivity 19 %. At the four-interval step the split is 72 % and 13 % (`verify_ridge_mechanism.py`).

The observable therefore does not become much less sensitive as the plasma heats. **The ground state simply moves less.** In the language of Section 5.4, the reservoir gain $|G|$ falls with temperature, from 14.4 at 1.00 eV to 5.7 at 2.95 eV (Table 5.2), and a step of fixed fractional size therefore displaces the reservoir progressively less.

Why the gain falls with temperature is not established here. $|\Delta \ln u|$ is flat to one per cent across three decades of density while varying by a factor 2.7 across temperature, which implies that $\partial \ln n_{1s} / \partial \ln T_e$ at collisional–radiative equilibrium is controlled by temperature alone. This is an observation from `verify_ridge_mechanism.py`, not a derived result. A natural candidate is the Arrhenius factor in the electron-impact ionisation rate out of the ground state, whose logarithmic temperature derivative is $\chi_H/T_e = 4.6$ at the benchmark against the measured $|G| = 5.7$ (Table 5.2), but the comparison has not been carried through and the two need not be equal.

5.6.5 Three explanations this work does not support

Three readings of the crest are natural and do not survive the numbers. Each is recorded because it is among the first a reader is likely to propose.

It is not a detachment criterion. The crest sits at $1.93 \times 10^{19}/\text{m}^3$: a factor 5.2 below the density band that published edge modelling associates with a detached ITER divertor [45], and a factor 1.6 below the modelled upstream separatrix density [2]. It falls between the two, in a region neither source characterises. Two further facts make the identification untenable rather than merely unsupported. At the citable detached band, $n_e = 10^{14}$ to $10^{15}/\text{cm}^3$ with $T_e \geq 2 \text{ eV}$, the plateau error is at most 10.4 % per +4.81 % step and the 100 μs time-averaged estimate of Section 5.8 never reaches 10 %, peaking at 7.2 %. And the crest location is proportional to the assumed neutral density, which in this model is set by a local ionisation balance with no transport and no recycling, whereas in a divertor it is set by recycling; Chapter 6 measures that proportionality and finds one decade in n_{1s} moves the crest by roughly one decade in n_e . The crest is reproducible. Its identification with a named divertor operating regime is not supported by anything in this work.

It is not a peak in temperature. The error decreases monotonically with temperature throughout the defended range (Section 5.6.4); the apparent ridge line at 2 eV in Figure 5.3 is the edge of the scope restriction.

It is not simply where b_1^{CRE} meets the sensitivity maximum. That reading accounts for the density structure but fails on the temperature axis: at the four-interval step the ratio $b_1^{\text{CRE}}/b_1^{\text{peak}}$ passes through unity at $T_e = 4.7 \text{ eV}$, where the error is 33 % and falling monotonically (`verify_plateau_gridmap.py`). The endpoint of the reservoir trajectory is the wrong object: Eq. (5.6) shows that what enters is the mean sensitivity over the whole interval traversed.

5.7 Does the transient visit the plateau?

Everything so far is algebraic. R^{PE} was built by holding b_1 at its pre-step value and solving the post-step excited block, which is a construction, not an observation. If the real trajectory never lingers near that state, the map measures a fiction.

It does linger. Propagating the full 43-state system through the step with `scipy.linalg.expm` and tracking $\varepsilon^{\text{CRE}}(t)$ (`verify_plateau_bridge.py`) gives, at the benchmark point and at [15, 3]:

- The excited manifold reaches its instantaneous QSS state after $1.94 \tau_{\text{relax}}$ at the benchmark and $3.72 \tau_{\text{relax}}$ at [15, 3].
- Once settled, the state deviation from QSS is 1.3×10^{-6} and 8.1×10^{-8} respectively. These are the settled values, not the maxima over the plateau interval, which are 471 and 1721 times larger because the interval begins in the tail of the fast transient. An external review of this work quoted the maxima; the settled column is the one that supports the claim that the manifold has equilibrated.
- Over the plateau, $\varepsilon^{\text{CRE}}(t)$ departs from the analytic $\varepsilon_{\text{plateau}}$, and the size of the departure $\max |\varepsilon^{\text{CRE}}(t)/\varepsilon_{\text{plateau}} - 1|$ depends on the interval it is measured on. On the bridge script’s measured plateau, which at [15, 3] is the shorter interval, it is 3.33 % at the benchmark and 1.28 % at [15, 3]. On the fixed $k = 30$ window $[30 \tau_{\text{relax}}, \tau_{\text{slow}}/30]$

used by the grid-wide census below it is 0.0323 at the benchmark and 0.0301 at [15, 3] (`flat_w30_shell` in `validation/trajectory_census/`). That flatness is what licenses quoting a single $\varepsilon_{\text{plateau}}$ instead of a time-dependent error, and it is measured rather than assumed.

The propagator itself was checked by running `expm` against `expm_multiply`, two independent algorithms, which agree to 1.13×10^{-10} and 3.69×10^{-12} . Two other checks that were run, the ODE residual and the semigroup identity, are tautological for this construction and this `expm` implementation and are *not* evidence.

Scope. The bridge test above is at two grid points. Running the same propagation at every analysed pair (`verify_trajectory_census.py`, output in `validation/trajectory_census/`) extends the dynamic chain full transient $\rightarrow R^{\text{QSS}}(u(t)) \rightarrow R^{\text{PE}}$ to the grid. At all 680 pairs with a plateau window, the maximum departure of the excited manifold from the instantaneous QSS manifold over the $k = 30$ window is at most 4.4×10^{-5} (median 1.1×10^{-6}) against the 2×10^{-3} tolerance, and sustained tracking begins within $4.84 \tau_{\text{relax}}$ everywhere (median 2.93). The departure of $\varepsilon^{\text{CRE}}(t)$ from the analytic plateau over the window has median 0.0022 and maximum 0.0154, at cool [1, 3]. The maximum tracking error over the window scales roughly as $0.06/M$: the median of $M \max_W \varepsilon^{\text{QSS}}$ over the 680 pairs is 0.060, with a range of 0.001 to 0.427, and the log-log correlation of $\max_W \varepsilon^{\text{QSS}}$ with $1/M$ is +0.945 (`validation/trajectory_census/`, section A), which is consistent with the singular-perturbation expectation that the slow eigenvector departs from the QSS manifold by $O(1/M)$. The eigen-propagation used for the census was checked against `scipy.linalg.expm` at every pair, on the state and on the observable. The density structure of Section 5.6.1 is a separate, grid-wide algebraic calculation.

One preregistered criterion failed. The first plateau criterion required $\varepsilon^{\text{QSS}} < 2 \times 10^{-3}$ throughout $20 \tau_{\text{relax}}$ to $0.02 \tau_{\text{slow}}$ and reached 2.778×10^{-3} at the very end of the window at [15, 3]; it was retained unchanged. It failed because the ground state had begun its expected motion while the tracking error stayed at 6.8×10^{-7} : the system had moved *along* the QSS manifold, not off it. The replacement interval definition used above was not preregistered.

5.8 How long the error lasts, and what an ELM averages

A spectrometer integrates. If the error decayed on nanoseconds it would be invisible in a $100 \mu\text{s}$ exposure, so the decay law matters as much as the amplitude.

The error decays as a single exponential with time constant τ_{slow} . Fitting $\varepsilon^{\text{CRE}}(t)$ over the decay gives $\tau_{\text{eff}}/\tau_{\text{slow}} = 1.0102$ at the benchmark and 1.0656 at [15, 3], with half-decay at 0.7002 and $0.7386 \tau_{\text{slow}}$ against $\ln 2 = 0.6931$ for a pure exponential (`verify_plateau_bridge.py`). In physical units the half-life is $12.95 \mu\text{s}$ at the benchmark and 1.10 ms at [15, 3]. Unlike the plateau duration, which is set by whatever tolerance defines the end of the plateau, the half-life carries no threshold dependence.

That gives a time-average over an event of duration τ_{drive} :

$$\left\langle \varepsilon^{\text{CRE}} \right\rangle_{\tau_{\text{drive}}} \simeq \hat{\varepsilon} \equiv \varepsilon_{\text{plateau}} \frac{\tau_{\text{slow}}}{\tau_{\text{drive}}} \left(1 - e^{-\tau_{\text{drive}}/\tau_{\text{slow}}} \right), \quad (5.16)$$

which tends to zero as $\tau_{\text{drive}} \rightarrow \infty$. Its other limit is wrong: as $\tau_{\text{drive}} \rightarrow 0$ the estimate tends to $\varepsilon_{\text{plateau}}$, whereas the true time average tends to the step error $\varepsilon_{\text{step}}$ of

the rise, 0.0283 against 0.3869 at $[0, 4]$ and 0.0087 against 0.0636 at the benchmark (`validation/trajectory_census/`), so the estimate is valid only for $\tau_{\text{drive}} \gg \tau_{\text{relax}}$. One transient precedes the plateau at the benchmark and does not enter this average: because $n = 4$ relaxes faster than $n = 3$, their relative deviations from partial equilibrium cross at $0.72 \tau_{\text{relax}}$ and the ratio n_3/n_4 overshoots R^{PE} by 0.40 % at $1.16 \tau_{\text{relax}}$ (2.6 ns), carrying ε^{CRE} to 0.0678 against $\varepsilon_{\text{plateau}} = 0.0636$ before settling by $5 \tau_{\text{relax}}$; over $t \leq 10 \tau_{\text{relax}}$ the L^2 norm of the excited-block deviation from the frozen-reservoir state never exceeds its initial value at any of the 784 pairs (maximum growth factor 0.9998), so the overshoot belongs to the ratio and not to the state. Over longer times that norm does rise again, by up to a factor 1.536, as the reservoir drift carries the state toward the new equilibrium; that is the slow decay, not a transient. The overshoot is absent at $[15, 3]$, and the transient before $50 \tau_{\text{relax}}$ moves the 100 μs average at the benchmark by -4.0×10^{-5} relative (`validation/trajectory_census/`, section A).

It is written with $\simeq$ rather than $\geq$. The reasoning that the observable cannot decay faster than the slowest mode would make Eq. (5.16) a conservative lower bound; 43-state eigen-propagation at every pair (`verify_trajectory_census.py`, `validation/trajectory_census/`) shows it is close to the true average and not a bound on it. Over the 448 pairs above 2 eV the ratio of the true 100 μs average to $\hat{\varepsilon}$ runs from 0.9791 to 1.0222 (median 0.9998, below unity at 226 of the 448), and at 506 μs from 0.9671 to 1.0377. The ratio below unity at the worst point $[0, 4]$, 0.999978, is not a quadrature artefact: raising the sampling from 400 to 1600 per decade changes the five averages of Table 4.3 by at most 4.7×10^{-6} relative and the $[0, 4]$ average by 2.7×10^{-10} (`validation/trajectory_census/`, section B). Its deficit of 2.2×10^{-5} is the net of two contributions of opposite sign: the first $50 \tau_{\text{relax}}$, during which the error is $\varepsilon_{\text{step}}$ rather than $\varepsilon_{\text{plateau}}$, lower the average by 5.2×10^{-5} relative, and the slow-mode mechanism below raises it by the remainder. $\hat{\varepsilon}$ is what its name says: a single-slow-mode estimate, within 2.2 % of the true average at 100 μs and 3.8 % at 506 μs over the grid.

The spectral gap of Section 5.1.2 is why a single exponential comes this close: the slow mode projects strongly onto the observable at every pair and no second mode contributes. The residual departure is not a transient. $\varepsilon^{\text{CRE}} = |R(t)/R^{\text{CRE}} - 1|$ is a ratio whose denominator, the $n = 4$ population, relaxes on the same τ_{slow} as the numerator. With δ the slow-mode amplitude in the $\text{H}\beta$ channel relative to its CR-equilibrium value, the ratio of true average to estimate tends to $(1 + \delta) \ln(1 + \delta)/\delta = 1 + \delta/2 - \delta^2/6 + O(\delta^3)$ for $\tau_{\text{drive}} \gg \tau_{\text{slow}}$. Heating gives $\delta > 0$ and the estimate reads low; cooling gives $\delta < 0$ and it reads high. Over the warm pairs δ runs from -0.22 to $+0.30$, which is why the extreme ratios are a cooling minimum and a heating maximum in the same density column at neighbouring temperatures, cool [49, 1] and heat [48, 1] at 100 μs .

What the demotion costs the census is two pairs. Counting the same 448 pairs on the true 100 μs average instead of on $\hat{\varepsilon}$ moves the census from 45 to 43, worst 0.1753 at heat [15, 3]. No pair enters. The two that leave, cool [22, 3] at 0.09949 and cool [23, 2] at 0.09999, sit at and just below the threshold; the second misses it by 1.5×10^{-4} relative, so the propagated count is 43, with 44 within the numerics. At 506 μs the propagated count is 24 on the shell ratio and 23 on the A -weighted line ratio. **Every census in this chapter should therefore be read as 45 on the estimate and 43 propagated.**

5.8.1 The census, inside the defensible scope

Counting the pairs whose time-averaged error $\hat{\varepsilon}$, the average of the instantaneous CR-equilibrium-reference error over $\tau_{\text{drive}} = 100 \mu\text{s}$, exceeds 10%, with a one-interval step (`validation/divertor_map/divertor_map.csv`):

Table 5.5: Breakdown census on the 100 μs time average of the model-defined instantaneous CR-equilibrium-reference error, counted on the single-slow-mode estimate $\hat{\varepsilon}$ of Eq. (5.16) with a one-interval step. The average is taken over a chosen duration τ_{drive} ; it is not a measurement model of an ELM. Each row states the set it counts over.

scope	pairs	$\hat{\varepsilon} > 10\%$	worst
all analysed pairs	680	202	0.3868 at [0, 4]
$T_e \geq 2 \text{ eV}$	448	45	0.1748 at [15, 3]
$T_e \geq 2 \text{ eV}$ and $n_e \geq 10^{14}/\text{cm}^3$	108	0	0.0717 at [15, 5]

The middle row is the result. Of the 202 pairs in the first row, 157 lie below 2 eV, inside the region Section 5.10 excludes, and 85 of them have a Lyman- α line-centre optical depth above unity over the half of a 5 cm slab, the convention Section 6.1 uses throughout (`validation/lyman_optical_depth/`, census (a)). The defensible statement is 45 of 448 pairs above 2 eV, worst 17.5% at $T_e = 2.02 \text{ eV}$, $n_e = 1.93 \times 10^{13}/\text{cm}^3$, at the 100 μs event duration these maps use; at the duration extrapolated to the ITER pedestal it is 24 of 448, worst 15.3%, at the same point (Table 5.6). At citable divertor densities there are zero breakdowns at either duration, with the worst case 7.2%. The benchmark point does not break down: its 100 μs time-averaged error is 1.2%.

An earlier note gave the first row as 201, a transcription slip; the direct recount from `divertor_map.csv` is $202 = 105 \text{ heating} + 97 \text{ cooling}$ (`verify_transport_selection.py`, which raises if any count in this table disagrees with the file), and nothing in the defensible scope depended on it.

The event duration was assumed and should not have been. $\tau_{\text{drive}} = 100 \mu\text{s}$ was fixed at the start of this work with no source; the refereed ITER extrapolation is $\tau_{\text{IR}}^{\text{ELM}} = 506 \mu\text{s}$ [17, Sec. 5] (Section 1.6), and since Eq. (5.16) averages a decaying error over the window, a longer window gives a smaller average, so 100 μs is the aggressive end rather than a neutral choice.

Table 5.6: Duration sensitivity of the census, over the 448 pairs above 2 eV at a one-interval step. The maps throughout this chapter are computed at 100 μs ; 506 μs is the extrapolation to the ITER pedestal.

τ_{drive}	$\hat{\varepsilon} > 10\%$	worst
100 μs	45	0.1748
250 μs	33	0.1664
506 μs , ITER	24	0.1533
750 μs	20	0.1435

The count nearly halves and the worst case barely moves. Over a factor 7.5 in event duration the count falls from 45 to 20 while the worst case falls from 17.5% to

14.4 %, a relative change of 18 %. At the ITER-extrapolated duration the defensible count is 24 of 448, **worst** 15.3 %, and that is the figure an ITER-relevant statement should use. The maps and every derived quantity in this chapter are computed at 100 μs and are labelled with it; the 506-microsecond census is a strict subset of the 100-microsecond one, because the estimate falls monotonically with the window at every point, so no conclusion drawn about the 45 fails to apply to the 24.

Two further sensitivities belong with the census. The count depends strongly on the threshold: at 5, 10 and 20 % it is 348, 202 and 61 over the same 680 pairs, so the count should always be quoted with the threshold that produced it. The worst case does not reference the threshold at all and is unaffected by it. And the 104 points excluded for having no plateau window contain no breakdowns whatever, all sitting at $T_e \geq 3.24\text{ eV}$ with a maximum estimate of 0.0011, so the exclusion does not hide anything; including them changes 202/680 to 202/784.

A detector records a ratio of integrated brightnesses, not the mean of the instantaneous error. A spectrometer with exposure τ_{drive} integrates each line separately and reports $\int_0^{\tau_{\text{drive}}} j_\alpha dt / \int_0^{\tau_{\text{drive}}} j_\beta dt$, a brightness-weighted mean of the instantaneous ratio, which is not the time average of ε^{CRE} that the census counts. That observable was computed at every pair with the augmented-matrix exponential, which gives $\int_0^{\tau_{\text{drive}}} \mathbf{n} dt$ without quadrature and to the round-off of the matrix exponential; at the benchmark and 100 μs it agrees with trapezoidal quadrature of $\mathbf{w} \cdot \mathbf{n}(t)$ on the sampled trajectory to 2.7×10^{-7} (`verify_trajectory_census.py`). Across the 448 pairs the error of the exposure ratio against the CR-equilibrium ratio agrees with the time average of ε^{CRE} to within 1.2 % at 100 μs (ratio of the two errors 0.9901 to 1.0114) and to within 2.2 % at 506 μs (0.9806 to 1.0217); that 2.2 % is the agreement between two definitions of the averaged error, and is distinct from the 2.2 % by which $\hat{\varepsilon}$ departs from the true average at 100 μs . The exposure ratio gives the same census: 43 of 448 at 100 μs on both the shell ratio and the A -weighted line ratio, and 24 and 23 at 506 μs . The headline 45 of 448 in Table 5.5 is the count on the estimate; the propagated true average and the exposure ratio give 43 on the shell ratio because two cooling pairs, from [22, 3] and [23, 2], whose estimates are 0.10008 and 0.10043, come out at 0.09949 and 0.09999 when propagated, both within 0.001 of the threshold (`validation/trajectory_census/`, section C). The agreement is empirical rather than structural. At 98 of the 784 pairs the signed error changes sign during the rise, so cancellation inside the exposure integral is possible in principle; all 98 have $\varepsilon_{\text{plateau}} \leq 0.064$ and none is a census member.

The plateau window itself is a free parameter, and it moves only one of the two things. The window is defined by a separation factor k , taken as 30 throughout, so that a point is analysed only where $k \tau_{\text{relax}} < \tau_{\text{slow}}/k$, that is where $M > k^2$. That choice is imposed rather than measured, and its effect was recounted at $k = 10, 20, 30, 50$ and 100 from the stamped one-interval maps (`verify_window_sweep.py`, `validation/window_sweep/`; `verify_window_membership_exposure.py`, `validation/window_membership_exposure/`). The result separates cleanly.

Magnitudes do not move at all. $\varepsilon_{\text{plateau}}$ is a frozen-reservoir linear solve evaluated at one instant and computed before the window gate is applied, so it cannot depend on the width of a window it is not averaged over; the stored maps at $k = 10, 20, 30$ and 50 carry bit-identical $\varepsilon_{\text{plateau}}$ on all 784 pairs, 0.0636 at the benchmark heating pair and 0.0830 at the cold corner. A run with a fixed +0.6 eV step (`validation/plateau_slowmode/`),

which is the four-interval +20.7 % step at the benchmark and a +60 % step at the cold corner rather than this chapter's one-interval step, gives 0.32902966 and 0.24075077 at the same two pairs, with $M = 4856$ at the benchmark; it shows the same invariance in k at magnitudes that are not this chapter's. Sampling that run's trajectory across the window and extrapolating back to $t = 0$ returned its analytic value to 0.005 % (`validation/plateau_window_check.txt`).

Membership does move. k decides which points have a resolvable plateau at all: 779, 728, 680, 596 and 452 of the 784 pairs pass at $k = 10, 20, 30, 50$ and 100, and 547, 496, 448, 364 and 221 of the 552 defended pairs. At $k = 100$ the benchmark heating pair itself is excluded, needing $M > 10^4$ and having 8243 for its post-step operator. The census numerator does not move with the denominator: 45 pairs exceed the 10 % threshold at every one of the five values of k , with the same worst case, 0.1748 at [15, 3], because every census member has M between 3.5×10^4 and 1.4×10^6 . **So k sets the denominator of every census and, over the range swept, never the numerator**, and a census quoted without its k is as incomplete as one quoted without its threshold. Every count in this chapter uses $k = 30$.

5.8.2 Two structural objections, tested

The estimate of Eq. (5.16) uses τ_{slow} , and τ_{slow} in this model is the relaxation time of an *open* system: the ground state loses atoms to an ion reservoir that is held fixed and never gives them back, so an artificially long τ_{slow} would inflate every entry in Table 5.5. Rebuilding the operator with an exactly particle-conserving 44th state and recomputing the whole census changes the count from 202 to 200 and the worst case from 38.68 % to 38.56 %; above 2 eV it changes nothing at all, 45 of 448 (`verify_residual_claims.py`, `validation/residual_claims/`, counting the 100 μs errors above the 10 % criterion in `validation/ion_closure/`). The closure factor on τ_{slow} , its table, the named falsifier and the caveat that this closure adds an ion reservoir without closing the electron balance or adding transport are in Section 4.8.

The second objection is that a step is an idealisation and a real ELM has a finite rise time. The operator was first ramped between the two ends of one grid interval, interpolated linearly and log-linearly, because the rate tables exist only at the grid nodes; the operator along a true temperature ramp, $\mathbf{L}[T_e(t)]$, is computed below. Replacing the step by a linear operator ramp of duration t_{ramp} , integrating the full 43-state system through $\mathbf{L}(t) = (1 - f) \mathbf{L}^- + f \mathbf{L}^+$ with $f = \min(t/t_{\text{ramp}}, 1)$, and reading the error at the start of the plateau window gives, at [15, 3], the fractions 0.9995, 0.9954, 0.9555 and 0.6469 of the step's plateau error for $t_{\text{ramp}}/\tau_{\text{slow}} = 10^{-3}, 10^{-2}, 10^{-1}$ and 1 (`verify_ramp_plateau.py`, `validation/ramp_plateau/`). A single-pole system driven by a ramp reaches the fraction $\text{De}(1 - e^{-1/\text{De}})$ of its step response, with $\text{De} = \tau_{\text{slow}}/t_{\text{ramp}}$, which is 0.9995, 0.9950, 0.9516 and 0.6321 at the same four ratios: within 0.5 % of the 43-state result for $\text{De} \geq 10$, and 2.3 % low at $\text{De} = 1$, where the 43-state value itself moves a further 1.7 %, from 0.6469 to 0.6576, between linear and log-linear interpolation of the operator along the ramp (`validation/ramp_plateau/`, log-linear sensitivity block; at $t_{\text{ramp}} = 10^{-2}\tau_{\text{slow}}$ the same change is $+1.6 \times 10^{-4}$). The two interpolated ramps differ in the shape of $T_e(t)$ as well as in the interpolation: the linear one has T_e linear in time, the log-linear one has $\ln T_e$ linear in time. So the fraction there is 0.65 ± 0.01 .

The operator along a true temperature ramp has since been computed. Every rate in $\mathbf{L}$ is either a closed-form function of T_e or a Maxwell average of a stored cross section, so

$\mathbf{L}$ can be built at any T_e from the pipeline’s own inputs, and the build reproduces every stored rate table, $\mathbf{L}$ and the source at all 400 grid nodes exactly (`build_L_at_Te.py`). At the log midpoint of every one-interval span the log-linear interpolant misses the true operator by 0.18 % in the ground-state column at [15, 3] and 0.36 % at the cold corner, the linear one by 0.63 % and 3.6 %: one quarter of the errors measured over the two-interval span [15, 3] $\rightarrow$ [17, 3] through the true mid node (ratios 0.25 to 0.27 in the ground-state column, τ_{slow} and u_{CRE}), so both interpolants are second order in the interval. Below about 3.4 eV the true entry lies between the two interpolants at every ground-state-column entry; above 7.9 eV both under-estimate that column. Integrating the one-interval ramp at [15, 3] along the true $\mathbf{L}[T_e(t)]$ with T_e linear in time gives 0.9956, 0.9567 and 0.6548 of the step’s plateau error at $t_{\text{ramp}}/\tau_{\text{slow}} = 10^{-2}$, 10^{-1} and 1, inside the interval between the linear ramp and the log-linear interpolant driven by the same $T_e(t)$ at every ratio and within 0.0005 of the latter; with $\ln T_e$ linear in time the last figure is 0.6571. The stamped gap of 0.0107 at $\text{De} = 1$ therefore splits into 0.0079 of linear-interpolation error, 0.0005 of log-linear error and 0.0024 of ramp shape, and the shape of $T_e(t)$ is the part no 0-D model supplies. The band 0.65 ± 0.01 stands, with the physical value between 0.655 and 0.657 according to the shape (`verify_interior_operator.py`, `validation/interior_operator/`). One caveat bounds the word true: the pipeline’s Maxwell average uses a 5000-point energy grid, and rebuilding it on 100000 points moves τ_{slow} by 1.0 % and the ground-state column of $\mathbf{L}$ by 1.9 % at this point, more than the interpolation error just measured. That is a property of the stored rates at the nodes as much as between them; its size across the grid is given in Section 2.3, and the grid is left as it is.

During a linear ramp the ground-state reservoir completes the fraction $1 - \text{De} (1 - e^{-1/\text{De}})$ of its total excursion, the lag of a single pole behind a ramp, which for large De expands to $(t_{\text{ramp}}/\tau_{\text{slow}}) (\frac{1}{2} - \frac{1}{6\text{De}} + \frac{1}{24\text{De}^2} - \dots)$. Per unit $t_{\text{ramp}}/\tau_{\text{slow}}$ the 43-state system at [15, 3] gives 0.4984 at $\text{De} = 100$ against the exact 0.4983, and 0.4841 at $\text{De} = 10$ against 0.4837; at $\text{De} = 1$ the reservoir completes 0.3737 of its excursion against the exact 0.3679, where the two-term truncation gives 0.3333; at $\text{De} = 1000$ the drift is 5×10^{-4} of the excursion and the 43-state value is 1.3 % above the exact form (`validation/ramp_plateau/`, lag-law block). The agreement is established at these four values of De only. The relevant Damköhler-like ratio is $\tau_{\text{slow}}/\tau_{\text{drive}}$, which is 2332 at [0, 4], and not $\tau_{\text{relax}}/\tau_{\text{drive}} \approx 10^{-5}$, which describes the suppression of the fast excited-manifold error and has nothing to say about the plateau; at $\text{De} = 2332$ the plateau is suppressed by 0.02 %. Under the two operator ramps and the true ramp tested, at [15, 3], the step idealisation is safe, and the ratio that makes it safe is $\tau_{\text{slow}}/t_{\text{ramp}}$, not $\tau_{\text{relax}}/t_{\text{ramp}}$.

5.8.3 Stepping temperature and density together as a two-parameter perturbation stress test

Every map above steps T_e at fixed n_e , and an ELM raises both. The test here moves the two parameters by whole grid intervals at once; it is a stress test of the two-channel framework, not a prescription for the path an ELM takes through the (T_e, n_e) plane. A spot check at the benchmark had suggested this runs in the model’s favour, the time-averaged estimate falling from 0.0117 to 0.0030 when the density steps up with the temperature, because a higher n_e shortens τ_{slow} . That spot check is reproduced exactly by the full map, and it is not representative.

The framework extends without amendment, and that is the first result. Nothing in the derivation of Eq. (5.7) refers to *what* displaced the reservoir. Running 4968 cases over steps of 0, ± 1 and ± 2 temperature indices crossed with 0 and ± 1 density indices, the two-channel superposition holds to a worst relative residual of 3.075×10^{-14} , the same figure the temperature-only map reports, and $\varepsilon_{\text{plateau}} = |\exp(\bar{S} \Delta \ln u) - 1|$ holds to 3.7×10^{-16} (`verify_joint_step_map.py`). The sensitivity is a property of the observable and the displacement is a property of the plasma, whichever axis moved.

The gain acquires a second component, and it is additive. What does not survive unchanged is the second factorisation, $\Delta \ln u = G \Delta \ln T_e$, which has one axis in it. Measuring the two gains at each starting point as secants over the single-axis steps of one grid interval, $+4.81\%$ in T_e at fixed n_e and a factor 2.68 in n_e at fixed T_e ,

$$G_T = \left. \frac{\Delta \ln u}{\Delta \ln T_e} \right|_{n_e, \Delta \ln T_e=0.0470}, \quad G_n = \left. \frac{\Delta \ln u}{\Delta \ln n_e} \right|_{T_e, \Delta \ln n_e=0.987}, \quad \Delta \ln u \simeq G_T \Delta \ln T_e + G_n \Delta \ln n_e, \quad (5.17)$$

the additive prediction reproduces the measured $\Delta \ln u$ to a median of 0.20% and 0.54% at the 90th percentile over 2716 joint steps. Only 26 of them exceed 5% , and those sit where the two contributions cancel: heating empties the reservoir while a density rise refills it, so $\Delta \ln u$ passes through zero and a relative error on a vanishing quantity is meaningless. Their median $|\Delta \ln u|$ is 0.0057 against 0.398 elsewhere.

The two gains are very different in size. $|G_T|$ runs 2.64 to 14.52, which is the same range Section 5.4.4 reports, while $|G_n|$ runs 0.044 to 0.164. **The reservoir is roughly fifty times less sensitive to density than to temperature in logarithmic terms.** It matters at all only because the density grid is coarse: one density index is $+168\%$ against $+4.81\%$ for one temperature index, a factor 35 in logarithmic step, so the two contributions end up comparable in this particular sweep. That coarseness is a limitation of the tabulation, not a statement about ELMs, and it is why the census below is quoted per step rather than as one number.

None of the tested one-index density steps breaks the diagnostic. Of the 320 window-passing cases at $\Delta n_e = -1$ index and the 296 at $+1$, with the temperature held fixed, *none* exceeds the 10% threshold. The median $\varepsilon_{\text{plateau}}$ is 2.7% and 2.6% respectively. The error this thesis is about is driven by temperature.

The census, and why the benchmark spot check misled. Counting as Table 5.5 does, over both temperature directions at one index and $T_e \geq 2 \text{ eV}$:

Table 5.7: The breakdown census under a joint step, at $\tau_{\text{drive}} = 100 \mu\text{s}$. The middle row is Table 5.5’s middle row, recomputed by an independent script and reproducing it exactly.

density step	pairs	$\hat{\varepsilon} > 10\%$	worst
-1 index, $\times 0.373$	426	44	0.1863 at 2.02 eV, $5.18 \times 10^{13}/\text{cm}^3$
0, fixed n_e	448	45	0.1748 at 2.02 eV, $1.93 \times 10^{13}/\text{cm}^3$
$+1$ index, $\times 2.68$	379	43	0.2206 at 2.02 eV, $7.2 \times 10^{12}/\text{cm}^3$

The count barely moves, at 10.3, 10.0 and 11.3% of the pairs available. **The worst case moves the other way from the spot check: it rises by 26% , from 0.1748**

to 0.2206, when the density steps up. The benchmark and the worst point respond differently because they sit on opposite sides of the density crest of Section 5.6. Stepping up from $n_e = 1.39 \times 10^{14}/\text{cm}^3$ moves away from the crest and collapses τ_{slow} from 18.5 to 5.23 μs , so the time-average wins; stepping up from $7.2 \times 10^{12}/\text{cm}^3$ moves *toward* it, and $\varepsilon_{\text{plateau}}$ rises faster than τ_{slow} falls.

The defensible statement is that the census is insensitive to a density step of this size while the worst case is not, and that under the tested +1-index density step the fixed-density map understates the worst case rather than overstating it.

5.9 Timescale separation does not predict the reservoir-closure error

A modeller deciding whether to trust a QSS-based inversion reaches for $M = \tau_{\text{slow}}/\tau_{\text{relax}}$, the stiffness ratio that licenses eliminating the fast variables in the first place. Chapter 3 measured it from 86.8 to 1.73×10^9 over all 784 pairs of the grid; over the 680 analysed pairs on which the correlations below are computed its minimum is 902 (`validation/partial_fe/`). The natural expectation is that the diagnostic is worst where M is smallest.

It is not. Once the operating point is accounted for, the controlled partial correlation has the *opposite* sign to the naive expectation (Tables 5.8 and 5.9), and the statement has to be read precisely. At fixed (T_e, n_e) and step direction there is exactly one M and one $\varepsilon_{\text{plateau}}$, so nothing is being compared *at* a fixed point. What the partial says is that, after the additive trends in $\log T_e$ and $\log n_e$ are removed, the residual variation of $\log M$ and the residual variation of $\log \varepsilon_{\text{plateau}}$ are anticorrelated. Those residuals are small: two-way fixed effects explain $R^2 = 0.998$ of $\log M$ and 0.992 of $\log \varepsilon_{\text{plateau}}$ in the heating direction (0.987 and 0.983 pooled; `validation/partial_fe/`, two-way fixed-effects R^2 block), so the partial correlates the remaining 0.2 and 0.8 percent of the variance. Two-way fixed effects absorb every additive function of T_e and of n_e , so what remains is the T_e – n_e interaction: the partial is a second-order statement about that interaction, not a statement that larger M means larger error. The tables below contradict the stronger statement that M carries no information at all.

The two extrema make the point without statistics. Over the whole grid the largest M , 1.73×10^9 , sits at the cooling step from $[1, 0]$ ($T_e = 1.048$ eV before the step, $n_e = 10^{12}/\text{cm}^3$) and carries $\varepsilon_{\text{plateau}} = 12.0\%$; the coldest corner $[0, 0]$ exists only in the heating direction, with $M = 9.39 \times 10^8$ and $\varepsilon_{\text{plateau}} = 8.30\%$, the 59th percentile of the 680 analysed pairs. The largest error, 38.7% at $[0, 4]$, sits 52 $\times$ away in density at an M smaller by a factor 46 (`validation/divertor_map/divertor_map.csv`). Those are whole-grid numbers and both extrema lie below 2 eV. Inside the defended range $T_e \geq 2$ eV the separation persists at smaller magnitude: the largest M is 2.63×10^6 at the cooling step from $[15, 0]$ with $\varepsilon_{\text{plateau}} = 7.4\%$, the largest $\varepsilon_{\text{plateau}}$ is 18.1% at the heating step from $[15, 3]$ with $M = 2.0 \times 10^5$, the two points sit on the same T_e row 19.3 $\times$ apart in density, and the error at the M maximum is a factor 2.5 below the worst (`validation/m_rank_test/`, Part 5). Where M is largest is not where the reservoir-closure error is largest.

The raw association looks strong and is a temperature proxy. Over the 680 analysed pairs, $\text{corr}(\log M, \log \varepsilon_{\text{plateau}}) = +0.76$. But $\log M$ is 93% explained by $(\log T_e, \log n_e)$ alone, and a bare Arrhenius factor e^{χ_H/T_e} built from the hydrogen ionisation potential $\chi_H = 13.6$ eV, which contains no dynamics whatever, correlates with $\log \varepsilon_{\text{plateau}}$ at




Figure 5.6: Timescale separation does not predict the reservoir-closure error. Each point is one (grid point, step direction) pair with a timescale-separated plateau window (680 of 784); M is $\tau_{\text{slow}}/\tau_{\text{relax}}$ of the post-step operator, colour is T_e , circles heating and triangles cooling. The raw association is strong, $\text{corr}(\log M, \log \varepsilon) = +0.76$, and it is a temperature proxy rather than a dynamical statement: $\log M$ is 93% explained by $(\log T_e, \log n_e)$ alone, a bare Arrhenius factor $e^{13.6/T_e}$ containing no dynamics correlates at +0.71, and controlling for $(\log T_e, \log n_e)$ drives the association to +0.27 linearly and -0.44, the opposite sign, quadratically. The two starred extrema make the same point without statistics: the largest M on the grid (1.7×10^9 , blue) carries $\varepsilon = 12.0\%$, while the largest error (38.7%, red) occurs $52\times$ away in density, at an M smaller by a factor 46. A large timescale separation supports the quasi-steady-state reduction of the fast excited states; it does not determine the error incurred by assuming that the slow ground-state reservoir has already reached CR equilibrium.

+0.71 (`make_ch5_figures.py`). Both quantities rise steeply toward low temperature for unrelated reasons, and that shared trend is nearly all of the +0.76.

Removing the shared trend reverses the sign. Residualising both $\log M$ and $\log \varepsilon_{\text{plateau}}$ on $(\log T_e, \log n_e)$ leaves a partial correlation of +0.27 over all pairs and +0.33 over the heating pairs alone; residualising on a full quadratic basis in the same two variables leaves -0.44 (`make_ch5_figures.py`, which reports the sign of the quadratic partial when it draws Figure 5.6 and does not condition the figure on that sign).

Report the sign, not the magnitude. The magnitude of the controlled partial depends on the basis chosen for the control, which is a modelling choice and not a measurement. The sign does not, and that claim is now tested rather than asserted. Sweeping four scopes, all analysed pairs, heating only, cooling only, and the defended range $T_e \geq 2$ eV, against four control bases of increasing richness, linear, full quadratic, quadratic plus an Arrhenius term $1/T_e$, and full cubic, gives Table 5.8.

Table 5.8: Partial correlation of $\log M$ with $\log \varepsilon_{\text{plateau}}$ after residualising both on a basis in $(\log T_e, \log n_e)$. Linear control is too weak to remove the shared temperature trend. Under quadratic control or richer the partial is negative in every one of the twelve combinations. Stamped in `validation/partial_correlation/`.

scope	linear	quadratic	quadratic + $1/T_e$	cubic
all analysed pairs	+0.274	−0.442	−0.503	−0.510
heating only	+0.326	−0.272	−0.705	−0.592
cooling only	+0.338	−0.329	−0.769	−0.607
$T_e \geq 2$ eV	+0.417	−0.531	−0.549	−0.533

The partial is negative in all twelve combinations with quadratic control or richer, ranging from -0.272 to -0.769 with median -0.532 , and it is negative in the defended range on every basis. Whether that sign is a property of the polynomial basis is tested directly in Table 5.9, which replaces the polynomial by saturated two-way fixed effects, one dummy per T_e row and one per n_e column. That control removes every additive function of the two variables and assumes no functional form. Under it the partial is -0.617 pooled, -0.939 heating and -0.770 cooling, at least as negative as under quadratic control in every direction. It keeps its sign when any one of the eight density columns is dropped (pooled range -0.549 to -0.783), when any one of the 35 temperature rows is dropped (-0.609 to -0.620), and against a permutation null that shuffles $\log \varepsilon_{\text{plateau}}$ within each T_e row: none of 2000 draws reaches the observed value in any direction (`validation/partial_fe/`). Enriching the basis is not a uniform test. Additive enrichment moves the partial toward the fixed-effects value (-0.531 quadratic, -0.614 additive quadratic, -0.609 additive quartic, -0.617 fixed effects, pooled), whereas adding $T_e \times n_e$ cross terms absorbs part of the interaction in which the residual lives and pushes the partial toward zero (full quartic: -0.524 pooled, -0.281 heating, -0.335 cooling; `validation/partial_fe/`, full-quartic block). That design has a condition number of 3.2×10^5 , worse than any other control in the file, so its partial is the least well-determined of the set. An earlier note recorded -0.22 to -0.70 over eight unnamed combinations, which could not be reproduced as stated; the twelve above are named.

The sign is a statement about residuals and must not be read as a ranking. At fixed T_e , M falls monotonically with n_e in all 98 (row, direction) series of the map, while $\varepsilon_{\text{plateau}}$ peaks at the density crest of Section 5.6, at $n_e = 1.9 \times 10^{13}/\text{cm}^3$ in 95 of the 98 and

Table 5.9: Partial correlation of $\log M$ with $\log \varepsilon_{\text{plateau}}$ over the analysed pairs at $T_e \geq 2$ eV, under the quadratic control of Table 5.8 and under saturated two-way fixed effects (one dummy per T_e row and one per n_e column). The null columns give the mean and standard deviation of the partial over 2000 permutations of $\log \varepsilon_{\text{plateau}}$ within each T_e row; no draw reached the observed value in any cell. Stamped in `validation/partial_fe/`.

scope	n	quadratic	null mean (sd)	two-way fixed effects	null mean (sd)
heating only	218	−0.590	+0.036 (0.068)	−0.939	−0.002 (0.079)
cooling only	230	−0.620	+0.044 (0.065)	−0.770	−0.003 (0.074)
pooled	448	−0.531	+0.015 (0.057)	−0.617	+0.000 (0.057)

one column higher in the other three (`validation/divertor_map/divertor_map.csv`). A monotone function cannot share its maximum with a unimodal one, so M cannot rank the error across density at all: in no row of either direction does the largest M coincide with the largest $\varepsilon_{\text{plateau}}$, 0 of 50 rows pooled and 0 of 35 in the defended range (`validation/m_rank_test/`, Part 3).

That argmax result is the evidence for the claim, and it must be kept separate from a weaker test that goes the other way. If the negative partial above were a ranking, the row-wise Spearman correlation of M with $\varepsilon_{\text{plateau}}$ across the density columns would be negative in most rows. It is not. Taken over all 50 temperature rows of the analysed grid the median is +0.23, with only 36 percent of rows negative, and the artifact records the prediction it was written to test as **refuted**; restricted to $T_e \geq 2$ eV the median falls to +0.07 with 49 percent of rows negative, which is not refutation but is not support either. The adverse figure is the whole-grid one and it is stated here because quoting only the defended subset would hide it. The two results are consistent, and the reason they are is the sentence above: the row-wise correlation is set by where the crest sits relative to the grid columns, which is a fact about the sampling rather than about predictive skill, whereas the partial is a statement about residuals after the operating point is removed and the argmax comparison is a statement about location. Only the last of the three is what a modeller reaching for M would actually need, and it fails at every row.

The negative result itself is not new as a general statement: Greenland reached it for arbitrary collisional–radiative systems [20, 21], quoted in Section 1.7. What this work adds is a quantification for one diagnostic and a map of where the criterion misleads.

The objection that M is computed from rate coefficients known to about 13 % (Section 2.8), so that the criterion might fail only because the tested quantity is noisy, does not hold. The R-matrix substitution of Section 2.8.2 shifts M by 11 % on average. Against a quantity spanning 7.30 decades that is a displacement of 1.08 % of the span at its largest, and it is very nearly a common factor: the Spearman rank correlation between the two versions of M over the 400 grid operators is 0.99993, M still falls monotonically with density in all 50 temperature rows, and the fraction of $\log M$ explained by $(\log T_e, \log n_e)$ alone is unchanged at 93.7 %. Every statement in this section is about the *ordering* of M and about its structure in (T_e, n_e) , and the substitution preserves both. The criterion fails for a structural reason, not a numerical one.

The reason the criterion fails here is visible in Eq. (5.7). M compares two *rates*; $\varepsilon_{\text{plateau}}$ measures a *distance*, the product of a sensitivity and a displacement, neither of which appears in M . A large M means the excited states equilibrate quickly with the reservoir, which Section 5.1.2 confirms to a part in 10^8 . It says nothing about whether the reservoir

is in the right place. The conclusion the data support is the narrower one: M supports the fast-state QSS reduction, and it is neither a sufficient nor a robust predictor of the reservoir-closure error.

5.10 Where the quantitative scope begins

Every quantitative statement in this chapter has been restricted to $T_e \geq 2$ eV. Two separate things go wrong below that boundary. Neither of them was chosen, both were measured, and they put the boundary in the same place; a third, cruder test agrees.

5.10.1 Holding the electron density fixed stops being consistent

Every operator in this chapter is evaluated at fixed n_e : the temperature steps, the density does not. That is not an approximation whose cost can be made small by care, because in a hydrogen plasma the electrons and the neutrals are not independent quantities. Quasineutrality ties them: every ground-state atom that ionises delivers one electron, so if the reservoir collapses the electron density must rise by the same absolute amount.

Whether that matters depends on how much of the hydrogen is neutral. At $T_e = 2.947$ eV the plasma is almost fully ionised and the neutral population is a small perturbation on the electron budget, so the coupling is negligible. At $T_e = 1$ eV the model's own steady state is 95 to 98 % neutral. The ground state then holds an order of magnitude more particles than the electron population does, and a modest fractional change in it is an enormous fractional change in n_e .

Imposing conservation of nuclei across a single +4.81 % temperature step makes this quantitative. Writing $x = n_{1s}^{\text{new}}/n_{1s}^{\text{old}}$ for the fractional collapse of the reservoir, the electron-density change that the fixed-density solutions imply is:

Table 5.10: Electron-density change implied by conservation of nuclei across one one-interval temperature step, inferred from the two fixed- n_e solutions; the self-consistent closed-parcel value is given in the text. x is the ratio of post-step to pre-step ground-state population at collisional–radiative equilibrium.

step in T_e	n_e (1/cm ³)	x	inferred $\Delta n_e/n_e$
1.000 $\rightarrow$ 1.048 eV	1.00×10^{12}	0.506	+20.0
1.000 $\rightarrow$ 1.048 eV	5.18×10^{13}	0.508	+13.9
2.947 $\rightarrow$ 3.089 eV	1.39×10^{14}	0.764	+0.0002

The first two rows are not a separate piece of physics from the error itself. $x = 0.506$ is $|\ln x| = 0.68$, and 0.68 is the reservoir displacement that `plateau_gridmap.txt` records at $[0, 4]$ and that Section 5.4 identifies as the source of the 38.7 %. **The displacement that produces the cold-edge error is exactly the displacement that makes the fixed-density assumption inconsistent.** Over the 392 one-step operators, 68 of them, 17 %, imply $|\Delta n_e/n_e| > 10$ %, and 36, 9 %, imply more than 100 %. The model evaluates the post-step operator at the old n_e at every one of them. Above 2 eV the heating-direction, fixed- n_e inference never exceeds 5.4×10^{-3} , its largest value there sitting at 2.02 eV, and it first falls below 10^{-3} above 2.56 eV (`make_story_figures.py`). Taken over both step directions and solved self-consistently, the correction reaches 7.98×10^{-3}

at a cooling step at $T_e = 2.024 \text{ eV}$, $n_e = 1 \times 10^{12} / \text{cm}^3$, and moves $\varepsilon_{\text{plateau}}$ there by 2.4 % (`verify_closed_nuclei.py`). Eight parts in a thousand is still small beside the tens of percent this chapter reports, but the stronger claim that the correction is negligible in both directions is refuted, and the correction is not one part in a thousand.

This closes a gap between two objections that this thesis has until now presented as separate and separately survivable. If the plasma parcel is closed, the fixed-density inference of Table 5.10 is not the answer but the symptom: solving the nuclei balance self-consistently, so that the post-step n_{1s} is evaluated at the post-step n_e , the electron density at the 1 eV row rises by a factor 1.98 to 2.05 (`verify_closed_nuclei.py`, `validation/closed_nuclei/`), against the 10 to 20 inferred; over the 392 heating steps 67 change n_e by more than 10 % and 7 by more than 100 %, against 68 and 36 inferred; the inference is not the required change. Either way every rate coefficient in a fixed-density post-step operator is evaluated at the wrong density there. If instead the parcel is fed by transport, so that n_e genuinely does stay fixed while the neutrals are consumed, then the ground-state history is set by the influx rather than by the local balance, and Section 6.2.5 measures that this enlarges the error at 165 of 180 tested cases rather than removing it. **The two pictures cannot both be held, and the cold edge needs one of them.**

The cold corner also fails a cruder test: it is not a divertor state. The model's own steady solution at $[0, 4]$ gives $n_{1s} = 1.4669 \times 10^{15} / \text{cm}^3$, which at a neutral temperature equal to the electron temperature is a neutral pressure of 235 Pa; at $[0, 7]$ it is 3280 Pa (`verify_cold_corner_pressure.py`, `validation/cold_corner/`). For contrast the benchmark gives 0.058 Pa.

Those want a comparison, and the comparison has to be one that was read rather than repeated. The figure of roughly 1 to 20 Pa for the ITER divertor design range could not be traced to any source and is not used. Read from the ITER divertor physics basis, the SOLPS-4.3 database for the $Q_{\text{DT}} = 10$ baseline runs between a lowest divertor neutral pressure of about 2 Pa and about 10 Pa at strong partial detachment [1, Secs. 3.1.3, 4]; later high-throughput SOLPS-ITER runs extend it to $\langle p_{\text{div}} \rangle > 25 \text{ Pa}$ [46, Sec. 6], citing [3]. These are the spans of modelled databases rather than design limits, and they are used here only as orders of magnitude. Against the top of the baseline database, $[0, 4]$ sits a factor 24 over and $[0, 7]$ a factor 328 over; against the extended one, factors of 9.4 and 131. The benchmark sits a factor 173 under the baseline top.

This is a third boundary, and it falls where the other two do. Of the 400 grid points, 32 imply a neutral pressure above 25 Pa and *none of them lies at $T_e \geq 2 \text{ eV}$* ; the implied pressure drops below 25 Pa at every density once $T_e > 1.6 \text{ eV}$. Particle conservation, radiation transport and now the implied gas pressure are three unrelated pieces of physics, and all three put the floor between 1.1 and 2.4 eV.

These are fixed points of a zero-dimensional box with no transport and no pumping. A real divertor's neutral density is set by recycling and pumping together with local ionisation balance, which is the objection Chapter 6 raises against the closed parcel, so the pressure is evidence that the cold corner is outside the modelled regime rather than a prediction of a pressure.

5.10.2 The plasma is not optically thin at the cold edge

The rate matrix of Chapter 3 treats the plasma as optically thin: every photon emitted escapes. For the Balmer lines that is a good assumption over the whole grid. For the Lyman series, which terminates on the ground state where essentially all the atoms are,




Figure 5.7: **Lyman trapping, and why the quantitative scope starts at 2 eV.** (a) $\epsilon_{\text{plateau}}$ against the assumed homogeneous-slab thickness D used for the Lyman escape factors, computed self-consistently (Θ_P depends on $n(1s)$, which depends on Θ_P) for all 14 Lyman channels; $D = 0$ is the optically thin matrix used everywhere else in this chapter. The grid maximum at $T_e = 1.00 \text{ eV}$ falls from 38.7% to 11.6–15.7% over $D = 1$ –20 cm: the cold-edge number is set by a slab thickness this zero-dimensional model does not contain, and must never be quoted without it. The worst point at $T_e \geq 2 \text{ eV}$ moves by 0.5% and the benchmark by 0.7% over the same twentyfold range. (b) Number of analysed pairs exceeding the 10% threshold, counted on the plateau value itself and not on the time-averaged ELM bound: 225 falls to 177 over the whole map, while above 2 eV it is 63 against 62, unchanged to within one pair. (At $D = 20 \text{ cm}$ the analysed set is 678 pairs rather than 680: two lose their plateau window.) The $T_e \geq 2 \text{ eV}$ restriction is therefore not a hedge; it is the working boundary adopted above the band where the one neglected process most likely to invalidate the result has been shown not to, with two of the 24 crossings at $D = 20 \text{ cm}$ lying above it.

it fails badly at the cold edge. The Doppler line-centre optical depth computed from the model’s own CRE solution (table and convention in Section 6.1) gives a Lyman- α depth of 80.3 per centimetre at $[0, 4]$ against 0.00163 at 2.95 eV, so at the cold edge the Einstein coefficients in the matrix are wrong, as effective decay rates, by about 1.8 orders of magnitude, reaching about three orders at the worst point on the grid.

Rebuilding the matrix with self-consistent Lyman escape factors for all 14 Lyman channels, over an assumed homogeneous slab of thickness D , and recomputing the whole map gives Figure 5.7.

Above 2 eV nothing moves. The ELM breakdown count above 2 eV stays at 45 of 448 for $D = 1, 5$ and 20 cm, and the worst case runs 0.1748, 0.1746 and 0.1740 against 0.1748 untrapped, a change of 0.5% across a twentyfold range in a parameter this zero-dimensional model does not contain. At the benchmark the plateau error moves by 0.7% over the same range.

Below it, the cold-edge numbers are inflated. At $[0, 4]$ Lyman trapping divides the plateau error by 2.5 to 3.3, from 38.7% to between 15.7% and 11.6% depending on the

assumed slab thickness, so the point still breaks down but its *magnitude* is not quotable; below 1.15 eV the column carrying the maximum wanders across three columns with D , so no maximum location is claimed there (Table 6.2 and Section 6.1).

Chapter 6 takes up the mechanism of the fall, which is not a loss of sensitivity, and the assumptions the slab geometry imports.

5.10.3 The boundaries coincide

Radiation trapping and quasineutrality enter through the Einstein coefficients and through the particle balance respectively, and they fail together: above 2 eV neither moves anything, and by 1 eV the first has divided the headline error by up to 3.3 while the second has invalidated the operator it was computed from.

Figure 5.8 draws the two thresholded boundaries on the operating plane. The quasineutrality contour at 10% sits at $T_e = 1.41$ to 1.50 eV and is almost independent of density. The optical-depth contour at $\tau_0 = 1$ sits at 1.13 to 1.98 eV at a slab thickness of 5 cm, and across the three thicknesses in the trapping sweep it spans 1.01 to 2.36 eV. Of the 24 optical-depth crossings, 2 lie above 2 eV, both at the largest thickness $D = 20$ cm and at the two highest density columns. Those crossings take the absorbing atoms at the electron temperature. Because the Doppler cross-section scales as $T_n^{-1/2}$, colder neutrals move every crossing upward, by 0.005 to 0.29 eV over the 24 (the largest at $D = 20$ cm, $n_e = 10^{15}/\text{cm}^3$, $T_n = 0.5$ eV), while Franck–Condon atoms at 3 eV move them down by 0.04 to 0.05 eV. At $T_n = 1$ and 0.5 eV the count above 2 eV becomes 3 and 4 of 24: the $D = 5$ cm crossing at the highest density, 1.98 eV at $T_n = T_e$, moves to 2.07 and 2.17 eV, so at that thickness the 2 eV boundary is exceeded by at most 0.17 eV and in one density column (`verify_tn_boundary.py`, `validation/tn_boundary/`).

The $T_e \geq 2$ eV restriction is therefore not a hedge. Two independent tests, each with a threshold that was chosen, place the failure of this model’s own assumptions in the band from 1 to 2 eV, the implied neutral pressure of Section 5.10.1 agrees as a third and cruder test, and the restriction is the conservative working boundary adopted above that band; it is where the neglected process most likely to overturn the result has been shown not to. Every headline number in this chapter sits inside it. The cold-edge values are reported because they exist and because suppressing them would be worse, and they are reported as what they are: the asymptotic behaviour of a model outside its own domain of validity.

5.11 What this chapter establishes

The approximation this thesis set out to test is not the one that fails: along a full 43-state trajectory the quasi-steady-state closure is accurate to 6.73×10^{-9} at the grid point carrying the largest diagnostic error. What does not keep up is the ground state, and the assumption a two-parameter lookup table makes about it is that the plasma has finished changing.

What that assumption costs has a closed form, Eqs. (5.7) and (5.9), in which the reservoir gain rather than a percentage is the object worth tabulating. In the small-step limit the diagnostic error per unit fractional temperature excursion is $|\bar{S}G|$: about 1.3 at the benchmark point and about 3.5 at the worst heating pair in the defended range.

Pushed through the table, that observable error becomes an inferred-temperature error an order of magnitude larger than the transient that caused it. Inverted at known density,




Figure 5.8: Two independent constraints put the same floor under the quantitative scope. Both boundaries are measured contours at thresholds that were chosen, 10% for the density change and $\tau_0 = 1$ for the optical depth, and they come from unrelated physics. *Quasi-neutrality.* Imposing nuclei conservation across one grid step in T_e , the electron density must change by $|\Delta n_e/n_e| = |\Delta n_g|/n_{\text{ion}}$, taken from the two CR-equilibrium solves either side of the step. This zero-dimensional model holds n_e fixed, so wherever that requirement is large the model is contradicting itself. Of the 392 one-step operators on the grid, 68 require more than 10% and 36 require more than 100%. At $T_e = 1$ eV the fixed-density inference gives 10.1 to 20.0; solved self-consistently for a closed parcel the density rises by a factor 1.98 to 2.05, and the counts become 67 and 7 (`validation/closed_nuclei/`). At $T_e \geq 2$ eV it never exceeds 0.0054. The 10% contour sits at $T_e = 1.41$ to 1.50 eV and is almost independent of density. *Optical depth.* The Lyman- α line-centre optical depth over the half slab, $\tau_0 = n(1s) \sigma_0 D/2$, with σ_0 evaluated at $T_{\text{at}} = T_e$ from the Doppler width $\sqrt{2kT/m}$, giving $\sigma_0 = 5.474 \times 10^{-14} \text{ cm}^2$ at 1 eV. The optically thin matrix used throughout this thesis is only defensible where $\tau_0 < 1$. That contour runs from $T_e = 1.13$ to 1.98 eV at the slab thickness $D = 5$ cm, and across the 3 thicknesses in the trapping sweep it spans 1.01 to 2.36 eV. (a) Both boundaries on the (T_e, n_e) plane, over a map of the quasi-neutrality requirement. (b) Each constraint divided by its own limit, so that a value above one means the constraint is violated and the two can share a single axis. Both boundaries fall between 1.01 and 2.36 eV. Two independent pieces of physics, particle conservation and radiation transport, place the floor in the same narrow band, and $T_e \geq 2$ eV clears both everywhere except 2 of the 3 times eight optical-depth crossings, which sit at the largest slab thickness $D = 20$ cm and the two highest density columns. The $T_e \geq 2$ eV restriction is therefore a measured boundary, not a hedge.

the plateau ratio has no solution within the table's 1 to 10 eV domain for 166 of the 448 defended pairs, every one a cooling step, a further 44, all heating, reach an answer only by crossing a fold, and over the 238 pairs with a unique fold-free answer a 4.81 % temperature step is read back with a median error of 48 %, on the side the plasma came from at every pair. The amplification is the ratio of the reservoir term to the residual slope of the table, and it is large because the manifold's own response and the reservoir's nearly cancel, to 6.8 % of either; the stationary point of Section 5.5.1 is that cancellation completed. At the two highest densities the median error is 16 %. The two-parameter inversion is not measured.

Every timescale in this chapter moves by 11 % when the R-matrix excitation set replaces the CCC one, measured in Section 2.8.2; that is the sensitivity to one substitution, not an uncertainty. It is too small to move any conclusion drawn here, all of which rest on the ordering of M rather than its value. Carried through the steady state to the observable, and combined with the substitution of the ionisation dataset that Section 5.5 describes, the two move $\varepsilon_{\text{plateau}}$ by 15.5 % on average over the 448 defended pairs and by at most 25.4 %, always downward at the points this chapter quotes: 6.36 % to 4.80 % at the benchmark, 18.1 % to 16.4 % at the ridge. The ionisation dataset is the larger of the two levers, at 16.4 % against 9.8 %, and the benchmark point is the least robust of the three. The split between the two structural coefficients is sharp. Under the excitation substitution the reservoir gain G is nearly invariant, moving 1.1 % on average, while the sensitivity $\bar{S}$ carries almost all of the change at 9.1 %: G is a property of how the reservoir responds to temperature, and $\bar{S}$ is built from the ratios a_3/a_4 and c_3/c_4 , which are the excitation data themselves.

The sensitivity factor obeys a ceiling, $|\bar{S}| < 1$, that follows from the structure of the linear system and not from any rate coefficient, and for a given dataset the sharper form $|\bar{S}| \leq \tanh(|\Delta|/4)$, where Δ is atomic data (Section 5.5); at the crest the measured sensitivity is at 97 % of it.

Across the operating range the error has an interior maximum in density, between 7×10^{12} and $5 \times 10^{13}/\text{cm}^3$, and no interior maximum in temperature. That range is a whole-grid figure and its upper end belongs to the three coldest rows; above 2 eV every row maximises on $1.93 \times 10^{13}/\text{cm}^3$ and the subgrid vertex spans a factor 1.35 (`validation/crest_subgrid/`). The density maximum is a position effect, the plasma's own reservoir sliding along a fixed sensitivity curve (Section 5.6.1); it survived three independent tests, the model does not identify it with detachment, and it is the (3,4) pair's maximum rather than the Balmer series'.

Inside the defensible scope, $T_e \geq 2 \text{ eV}$, 45 of 448 point-and-direction pairs commit a time-averaged error above 10 % over $\tau_{\text{drive}} = 100 \mu\text{s}$, the worst being 17.5 %, falling to 24 of 448 and 15.3 % at the 506 μs extrapolated to the ITER pedestal; at citable divertor densities none do, the worst being 7.2 %. Those counts are conditional on the closed-parcel reservoir history of the zero-dimensional model; Chapter 6 shows that history is not the divertor's at any of the 45 pairs.

That scope boundary is the chapter's least expected result. It was adopted as a precaution against radiation trapping and it turns out to be the boundary of self-consistency: below 2 eV the plasma is mostly neutral, and conservation of nuclei then demands an electron-density change of 10 to 20 times across the same temperature step whose effect is being measured at fixed n_e . The cold-edge numbers are asymptotic values of a model outside its own domain, and this chapter reports them as such and builds nothing on them.

Finally, the criterion a modeller would reach for does not see any of this coming: within the defended range the maxima of timescale separation and of diagnostic error sit $19.3\times$ apart in density, and their apparent correlation reverses sign once the operating point is controlled for. M compares rates; the diagnostic error is a distance.

What the model cannot say is the subject of Chapter 6: the plasma is treated as optically thin where it is not, the atoms are held where they are put, and the gas is treated as atomic where it is mostly molecular. Each of those is quantified there rather than hedged.

Chapter 6

What This Model Cannot Say

Chapter 5 established two things: a map of how wrong a tabulated H_α/H_β inversion can be while the ground state is still catching up, and a mechanism that explains the map's shape. The map has a ridge in it, the ridge has a closed-form explanation, and the explanation was tested three separate ways. At this point in the argument a result stands.

This chapter's job is to bound that result without withdrawing it. Every number in Chapter 5 came out of one 43-state matrix describing one kind of atom under one set of assumptions, and some of those assumptions are badly violated in exactly the corner of the map where the numbers are largest. An examiner will find that corner. It is better to arrive there first, with arithmetic.

The standard this chapter holds itself to. A limitation that has been *computed* is worth more than a limitation that has been hedged. Where an objection could be turned into a calculation, it was, and the calculation is reported here with the number that would have refuted the result. Two objections were expected to be decisive and were not; they are given the same space as the ones that bite. Where an objection cannot be computed inside a model with no spatial dimension, it is stated as a *conditional*, “*if* this holds, the error is as computed”, and not softened into a caveat sentence. Where a check has simply not been run, this chapter says so.

Three abbreviations recur below and are expanded here for this chapter: *CR* for collisional–radiative, the class of model built in Chapter 3; *QSS* for quasi-steady-state, the Bodenstein-style approximation that eliminates the fast excited levels; and *CRE* for collisional–radiative equilibrium, the fully settled steady state that the diagnostic lookup table assumes and that this thesis measures the cost of assuming. An *ELM* (edge-localised mode) is the repetitive burst of heat and particles that a tokamak edge throws at the divertor every few milliseconds, and it is the transient every “event duration” below refers to; $\tau_{\text{drive}} = 100\ \mu\text{s}$ throughout, the duration used in Chapter 5 and the aggressive end of the range; the ITER-pedestal extrapolation is $506\ \mu\text{s}$ [17, Sec. 5] (Section 5.8).

6.1 Optical thickness: when the plasma cannot let its own light out

6.1.1 The optically thin assumption and how it reaches the slow clock

Chapter 2 gave every downward transition an Einstein A coefficient: the probability per second that an excited atom drops to a lower level and emits a photon. Those coefficients entered the CR matrix $\mathbf{L}$ bare. Entering bare means one specific physical assumption: that every photon emitted anywhere in the plasma leaves it and never comes back. A plasma for which this is true is called *optically thin*.

The failure of this assumption has a direct analogue in reaction engineering. Consider a reaction whose product is removed by desorption from a catalyst surface. If the product leaves the reactor, the desorption step runs at its intrinsic rate. If the product cannot leave, if it re-adsorbs somewhere else in the bed and is put back onto a site, then the *effective* desorption rate is smaller than the intrinsic one, and the whole network downstream of that step shifts. Radiation trapping is that, with photons.

Here is the specific chain. The strongest transition in hydrogen is Lyman- α : an atom in the $2p$ level drops to the ground state $1s$ and emits an ultraviolet photon. If the gas around it contains many ground-state atoms, that photon has a good chance of being absorbed by one of them, which puts *that* atom back into $2p$. Nothing has left the plasma; a $2p$ population has simply been moved from one place to another. Averaged over the volume, the $2p$ level decays more slowly than its A coefficient says. The standard way to write this down is to multiply the coefficient by an *escape factor* Θ_P between 0 and 1: $\Theta_P = 1$ is a transparent plasma, $\Theta_P = 10^{-5}$ is one in which the transition is effectively switched off. The formulation used here is Holstein's population escape factor [38].

Why this reaches the populations, and therefore the result. If $2p$ lives longer, atoms pile up in $n = 2$. Atoms in $n = 2$ are only 10.2 eV from the ground state but only 3.4 eV from the continuum, so they are far easier to ionise than ground-state atoms. A longer $n = 2$ lifetime therefore means more *stepwise* ionisation, ionisation via an excited level rather than directly out of the ground state, which means fewer neutrals and a faster turnover of the ground-state reservoir. The ground-state reservoir is the whole mechanism of this thesis. Trapping is not a cosmetic correction to a line intensity; it acts on the slow clock.

6.1.2 Where the assumption fails, in numbers

The formalism needed is one line. A photon travelling a distance D through a gas of absorbers is attenuated by $e^{-\tau}$, where the *optical depth* $\tau = \sigma_0 n_{1s} D$ is dimensionless: n_{1s} is the ground-state density in $1/\text{cm}^3$, D the path length in cm, and σ_0 the absorption cross section at line centre in cm^2 . For a Doppler-broadened line, $\sigma_0 = 2.654 \times 10^{-2} f / (\sqrt{\pi} \Delta\nu_D)$, with f the oscillator strength (a dimensionless measure of how strongly the transition couples to light) and $\Delta\nu_D$ the Doppler width set by the temperature of the absorbing atoms. $\tau \ll 1$ is transparent; $\tau \gg 1$ is opaque.

The number that matters is n_{1s} , not n_e , and Chapter 3 already established the awkward fact that these two run in opposite directions: the plasma is most neutral where it is coldest. The worst point of the error map is $T_e = 1.0 \text{ eV}$, $n_e = 5.18 \times 10^{13} 1/\text{cm}^3$, the maximum of $\varepsilon_{\text{plateau}}$ over all 680 window-resolvable points of the map, cold edge

included. There the model’s own CR solution gives $n_{1s}/n_{\text{ion}} = 28.3$, so quasi-neutrality fixes $n_{1s} = 1.47 \times 10^{15} \text{ 1/cm}^3$. At 1 eV that gives $\sigma_0 = 5.47 \times 10^{-14} \text{ cm}^2$, and the Lyman- α optical depth per centimetre of path is

T_e (eV)	$\tau_{\text{Ly}\alpha}$ per cm at $n_e = 5.18 \times 10^{13} \text{ 1/cm}^3$
1.00	80.3
1.15	10.6
1.60	0.224
2.02	0.0255
2.95	0.00163

An optical depth of 80 per centimetre means a Lyman- α photon travels about a tenth of a millimetre before it is reabsorbed. Over 5 cm of plasma the line-centre depth over the half slab is 201 and the population escape factor is 1.2×10^{-3} ; roughly one photon in eight hundred gets out. Chapter 3 quotes 3×10^{-5} at a different grid point, a different slab thickness and a different density, so the two are not the same quantity and the agreement in order of magnitude should not be read as a reproduction of one by the other. The blunt reading is that the effective $A_{2p \rightarrow 1s}$ belonging in $\mathbf{L}$ at that point is wrong, self-consistently, by about 1.8 orders of magnitude, or 2.9 orders on the thin one-shot estimate (`verify_lyman_optical_depth.py`).

Two rows down the same table the problem has evaporated. At 2.02 eV the depth per centimetre is 0.0255, and at the benchmark point (2.947 eV, $1.389 \times 10^{14} \text{ 1/cm}^3$) it is 0.0039; a photon there crosses metres of plasma before anything happens to it. **The measured quantity is not “how opaque is this plasma” but “how fast does opacity switch on as the plasma cools”**, and the answer is a factor of 5×10^4 between 1 eV and 2.95 eV.

That is why the results of Chapter 5 were declared for $T_e \gtrsim 2 \text{ eV}$ in the first place (Section 3.5). It is also, uncomfortably, where the headline was living. Recounting the 100 μs breakdown census against optical depth rather than against temperature: of the 202 points at which the 100 μs time-averaged estimate of the error exceeds 10 %, an extremum count over the 680 window-resolvable points of the full grid, 157 lie below 2 eV, 85 have $\tau_{\text{Ly}\alpha}(5 \text{ cm}) > 1$, and 15 have it above 100, on the half-slab depth used above.

The counts 105 and 26 can also be obtained, and are a convention error rather than a recomputation: they are reproduced only by combining the full-path depth $\kappa_0 D$ with a Doppler width $\sqrt{kT_n/m}$, short of the correct $\sqrt{2kT_n/m}$ by $\sqrt{2}$. With the correct width the census is 85 and 15 on the half-slab depth quoted here and 100 and 23 on the full-path depth. A deuterium cross-section with the full-path depth also returns 105 and 26, since the isotope mass enters the width as the same $\sqrt{2}$, so the two errors are indistinguishable from the counts alone. All seven conventions are tabulated in `validation/lyman_optical_depth/`. Nothing in the defended range depends on the choice of width convention: with the absorbing atoms at T_e , every crossing lies below 2 eV except two, both at $D = 20 \text{ cm}$; colder atoms move the crossings upward, and Section 5.10 gives by how much.

6.1.3 Self-consistent Lyman escape factors applied to the matrix

The expectation, stated before the calculation was run, was that trapping would demolish the map: an effective A coefficient wrong by about three orders of magnitude at the worst point ought to invalidate everything computed there and nearby, and the project’s own

Table 6.1: Effect of radiation trapping on the 100 μs breakdown census. Each entry counts the grid points at which the time-averaged estimate of the CR-equilibrium-reference error, averaged over a chosen duration $\tau_{\text{drive}} = 100 \mu\text{s}$, exceeds 10 %, out of the points at which a quasi-steady-state window is resolvable; “worst” is the largest estimate within that scope. Rows differ only in the assumed slab thickness D used to compute the self-consistent Holstein escape factor applied to all 14 Lyman transitions. The first row is the optically thin matrix used throughout Chapter 5. Above 2 eV the count does not move at all across a twentyfold range of D . The $D = 20 \text{ cm}$ row has a slightly smaller denominator because two points lose their resolvable window under trapping. Source: validation/lyman_trapping/lyman_trapping.txt.

matrix	all window-resolvable	$T_e \geq 2 \text{ eV}$	$T_e \geq 2 \text{ eV}, n_e \geq 10^{14}$
optically thin (canonical)	202/680, worst 0.3868	45/448, worst 0.1748	0/108, worst 0.0717
trapped, $D = 1 \text{ cm}$	176/680, worst 0.2925	45/448, worst 0.1748	0/108, worst 0.0712
trapped, $D = 5 \text{ cm}$	170/680, worst 0.2550	45/448, worst 0.1746	0/108, worst 0.0694
trapped, $D = 20 \text{ cm}$	160/678, worst 0.2237	45/446, worst 0.1740	0/106, worst 0.0627

records warned that under strong trapping the fast relaxation time τ_{relax} rises to 2.05 μs and the timescale separation M falls by a factor 158. That expectation was wrong.

So the calculation was done. The Lyman series was given a self-consistent Holstein escape factor and the entire error map rebuilt. Self-consistent matters here: Θ_P depends on n_{1s} , and n_{1s} depends on Θ_P through the CR balance, so the two were iterated to a fixed point rather than evaluated once. The fixed point converged at all 400 grid points for every slab thickness tested, and non-convergence was made to raise an error rather than silently return the last iterate. All 14 Lyman channels were trapped, the seven resolved $nP \rightarrow 1S$ transitions and the seven bundled ones, not Lyman- α alone.

Three gates were passed before any result was looked at. The oscillator strengths were derived from this repository’s own A coefficients by inverting the Einstein relation rather than imported, and come out as $f = 0.4162, 0.0791, 0.0290, 0.0139$ for Lyman- α through Lyman- δ , the accepted literature values to four figures. The derived Lyman- α σ_0 agrees with an independent implementation to 0.059 %, after a first run that failed at 1156 % on a 4π factor; and rebuilding the matrix with trapping switched off reproduces the canonical $\mathbf{L}$ exactly, so every trapped number below is a difference against the real matrix (Section 4.5).

One new quantity had to be invented to do this at all, and it is not in the model. The escape factor depends on how far a photon must travel to get out, so it depends on the size of the plasma, the slab thickness D . A model with no spatial dimension has no D . It was therefore swept, from 1 cm to 20 cm, and never fixed; every trapped number below carries the D it was computed at. Reporting a single trapped number without its D would be reporting an assumption as a measurement.

6.1.4 Above two electron-volts, nothing moves

Table 6.1 contains the strongest robustness statement in this thesis, and it is the opposite of what was expected.

Above 2 eV the breakdown count does not change at all, 45 of 448 points in every run, and the worst case within that set moves by 0.5 %, from 0.1748 to 0.1740, across a twentyfold range in assumed slab thickness. The escape factor at

Table 6.2: The cold-edge worst point of the optically thin error map, at $T_e = 1.0$ eV and $n_e = 5.18 \times 10^{13}$ 1/cm³ under a heating step, is the maximum of $\varepsilon_{\text{plateau}}$ over all 680 window-resolvable grid points; it is traced here as a function of the assumed slab thickness D . $\varepsilon_{\text{plateau}}$ is the plateau error, the fractional error a lookup table makes while the ground state is stale; the ELM bound is its average over an event of $\tau_{\text{drive}} = 100$ μ s. The plateau error falls by a factor 2.5 to 3.3, but the diagnostic’s intrinsic sensitivity $|f_3 - f_4|$ barely moves: what collapses is how far the ground state is displaced, not how strongly the line ratio responds to it. Source: `validation/lyman_trapping/lyman_trapping.txt`.

D (cm)	$\varepsilon_{\text{plateau}}$	ELM bound	τ_{slow} (s)	M	f_3	f_4	$ f_3 - f_4 $
0 (thin)	0.38690	0.38682	2.3324×10^{-1}	3.74×10^7	0.7728	0.2862	0.4866
1	0.15748	0.15739	9.0651×10^{-2}	2.75×10^6	0.6643	0.2055	0.4589
5	0.13274	0.13260	4.6940×10^{-2}	6.13×10^5	0.5967	0.1793	0.4174
20	0.11570	0.11550	2.8500×10^{-2}	2.14×10^5	0.5139	0.1501	0.3639

the benchmark point (2.947 eV, 1.389×10^{14} /cm³) runs from 0.9986 at $D = 1$ cm to 0.973 at $D = 20$ cm. Even at the thickest slab tested, 97 % of Lyman- α photons still leave: the plasma is optically thin there in the sense that actually matters, and assuming so costs less than the resolution of the map.

This converts a disclaimer into a measurement. The $T_e \gtrsim 2$ eV restriction was chosen on an escape-factor argument before any of this was recomputed, and it sits almost exactly at the temperature above which trapping stops mattering: a guess that has now been checked, and survived. The defensible claim of Chapter 5, 45 of 448 points above 2 eV at 100 μ s and 24 of 448 at 506 μ s (Section 5.8), is invariant under the one physical process most likely to have invalidated it.

6.1.5 Below two electron-volts, the cold-edge number is inflated threefold

The same calculation is much less kind to the cold edge, and it should be, because that is where the optical depth was 80 per centimetre. Table 6.2 follows the worst point of the optically thin map through the sweep.

The 38.7 % becomes 11.6 to 15.7 %, a factor of 2.5 to 3.3, and which value inside that range it takes is set by a slab thickness the model does not contain. The point still exceeds the 10 % threshold at every thickness tested, so the qualitative statement that the cold corner breaks down survives intact; the *number* does not, and it must never appear without its trapped range and without the optical depth that produced it.

The route by which the number falls is itself a result, and it is not the simpler route. The simpler guess is that trapping desensitises the diagnostic, that a trapped plasma simply responds less. It does not. The sensitivity $|f_3 - f_4|$, which is the factor in the exact decomposition of Section 5.6 measuring how strongly the line ratio answers to its reservoir, moves by only 6 % at $D = 1$ cm and 25 % at $D = 20$ cm. What collapses is the other factor: the displacement.

The chain is the one set out at the start of this section, now with numbers on it. Trapping lengthens the $n = 2$ lifetime; the longer lifetime raises stepwise ionisation; the ground-state population falls; and the slow clock τ_{slow} shortens by a factor 8 at the worst point. At $T_e = 1.0$ eV, $n_e = 1.93 \times 10^{13}$ 1/cm³, n_{1s} drops from 2.99×10^{14} to

$7.26 \times 10^{13} \text{ 1/cm}^3$ between $D = 1$ and 20 cm. A reservoir that turns over faster is a less stale reservoir, so $|\Delta \ln x|$ roughly halves. This is exactly the two-axis attribution of Section 5.6.4 arriving from a third direction: trapping acts on the displacement, not on the sensitivity.

6.1.6 Below 1.15 electron-volts the ridge location is not determined

One more thing moves at the cold edge, and it is a location rather than a magnitude. Taking the density at which $\varepsilon_{\text{plateau}}$ is maximal, row by row: at $T_e \geq 1.6 \text{ eV}$ that maximum falls on $n_e = 1.93 \times 10^{13} \text{ 1/cm}^3$ in every run, optically thin and trapped alike, and the values themselves barely register the change: at the benchmark temperature $\varepsilon_{\text{plateau}} = 0.12166$, 0.12166, 0.12165, 0.12163 for $D = 0, 1, 5$ and 20 cm, with $|f_3 - f_4|$ changing in the fifth decimal place. Below 1.15 eV the maximum wanders across three density columns, a factor of 7, as a function of an assumed parameter. **The two coldest rows of the grid cannot support a claim about where the ridge is at all.**

That result cuts in the mechanism’s favour, which is worth saying explicitly because it is easy to miss. The ridge attribution of Section 5.6 was verified over $T_e \geq 2 \text{ eV}$, and $T_e \geq 2 \text{ eV}$ is precisely the range in which trapping has now been shown not to move the defended census within the tested escape-factor model ($D = 1$ to 20 cm, $T_n = T_e$). The agreement with Griem’s boundary criterion is therefore not an artifact of neglected opacity; it was established where opacity does not act.

6.1.7 A contradiction resolved, and what is still open

Two statements about trapping had been circulating in this project’s records: that τ_{relax} is insensitive to Lyman- α trapping to within 1 %, and that under strong trapping τ_{relax} rises to 2.05 μs with M falling 158-fold. They look contradictory. Measured directly, both are true and neither had stated its scope: at the benchmark point (2.947 eV, $1.389 \times 10^{14} \text{ /cm}^3$) τ_{relax} moves from 2.2769×10^{-9} to $2.2947 \times 10^{-9} \text{ s}$ at $D = 20 \text{ cm}$, a change of 0.78 %; at $T_e = 1.0 \text{ eV}$, $n_e = 1.93 \times 10^{13} \text{ 1/cm}^3$ it moves from 8.8842×10^{-9} to $2.4289 \times 10^{-7} \text{ s}$, a factor of 27.3. The “within 1 %” claim is a benchmark-point claim and fails by more than an order of magnitude at the cold corner. This is the same failure mode as quoting any other extremum without the set it is an extremum over, and neither sentence may travel again without its regime attached.

What this calculation does not settle should be as clear as what it does. The geometry is assumed, not derived: a homogeneous slab, uniform emission, an isotropic radiation field, a Doppler line profile, evaluated at the plasma centre. A 0-D model cannot check any of those. One assumption inside it *was* checked and is not load-bearing: the temperature of the absorbing atoms, which sets the Doppler width. Fixing it at 3 eV, the value a molecule dissociating by the Franck–Condon process would give, instead of at T_e , changes the census at $D = 5 \text{ cm}$ from 170/680 and a worst case of 0.2550 to 173/680 and 0.2653; fixing it at 0.5 eV, colder atoms returning from a surface, changes them to 166/680 and 0.2443; above 2 eV the count is 45 of 448 under all three, with the worst case moving from 0.1746 to 0.1744 (`verify_lyman_trapping.py`, `validation/lyman_trapping_tat3/` and `..._tat05/`). Continuum lowering is untouched here and remains open. Opacity in the Balmer lines themselves, which would act on the observed ratio directly rather than through the level populations, has been computed: the

largest H_α optical depth above 2 eV is 0.010, and the worst cold cell, [0, 7], reaches 0.214 (Section 4.7, `validation/balmer_optical_depth/`).

6.2 Transport: the atoms do not stay where the model puts them

6.2.1 An objection the model cannot answer from inside

This is the objection that cannot be answered inside this model, and it is a more serious threat to the cold-edge number than anything in Section 6.1. It is also the simplest to state.

6.2.2 Batch-reactor parcel against a recycling divertor

The model of Chapter 3 is a batch reactor. A fixed parcel of gas is given a temperature, and its composition evolves until it stops changing. The slow eigenvalue of that system, τ_{slow} , is how long the ground-state population takes to reach its new value, and it is long, because ionising an atom out of the ground state is slow.

A divertor is not a batch reactor. It is closer to a continuous-flow reactor with an enormous throughput, in which the residence time of a molecule is far shorter than the time it needs to react. Neutral atoms are born at the target surface when ions strike it and recombine there, which is *recycling*, and they fly back into the plasma. They do not wait around to be ionised locally; the ones that do not ionise leave, and fresh ones arrive continuously.

The comparison is not close. At the worst point of the map τ_{slow} is 233 ms: nearly a quarter of a second for a stationary atom to be ionised. A deuterium atom at 1 to 3 eV would cross a 10 cm plasma ballistically in 6 to 10 μs , but it does not travel ballistically: resonant charge exchange with the surrounding ions randomises its direction after about 1.7 cm at this density, so escape is diffusive and takes closer to 72 μs . That correction runs in this model's favour. The threshold at which the worst point would fall below 10 % is a ground-state renewal time of about 26 μs , so the result survives by a factor of 2.8 rather than failing outright. It survives only because charge exchange traps the neutrals, which is a thin margin to rest on and is stated as one. That is one grid point, and it is one this thesis does not defend: [0, 4] lies below the 2 eV scope boundary. Section 6.2.4 asks the same question at the points that do carry the result, and gets a worse answer.

In a divertor, the ground-state neutral density at a given point is not set by local ionisation balance alone. It is set by recycling and transport together with the local atomic and molecular reactions.

6.2.3 Why this bears on the result specifically

The error of Chapter 5 exists for one reason: after a temperature step, the ground-state reservoir is *stale*. Its population has not moved while the temperature has, the excited levels equilibrate against a reservoir that belongs to the old conditions, and the lookup table, which assumes the reservoir finished settling, reads the wrong answer. Everything hangs on the reservoir not being refreshed.

Neutral transport takes the reservoir’s history out of the local balance. Any process that replaces the ground-state population on a timescale shorter than τ_{slow} makes $u(t)$ something this model does not compute, and whether the error then shrinks or grows depends on what the influx supplies (Section 6.2.5); this model contains no such process, and the 26 μs renewal threshold quoted above is the map’s own measure of how fast one would have to be to matter.

6.2.4 The assumption fails hardest where the result is generated

The worst point was the wrong place to test transport: $[0, 4]$ lies below the scope boundary of Section 5.10. The pairs that carry the defended claim are the 45 above 2 eV whose 100 μs estimate exceeds 10 % (Table 5.5), and there the answer is different.

Comparing the neutral escape time with τ_{slow} pair by pair over the grid gives Table 6.3. The escape model is the one used above, stated as a formula: the atom moves at $v = \sqrt{2kT_n/m_D}$ with $T_n = T_e$, resonant charge exchange with the ambient ions randomises its direction after $\lambda = v/(n_e \langle \sigma v \rangle_{\text{CX}})$ with $\langle \sigma v \rangle_{\text{CX}} = 1.1 \times 10^{-8} \text{ cm}^3/\text{s}$, and escape from a boundary a distance L away takes L/v when $\lambda \geq L$ and $4L^2/\pi^2 D$ with $D = v\lambda/3$ otherwise. At $[0, 4]$ it returns $\lambda = 1.72 \text{ cm}$ and $\tau_{\text{esc}} = 72.3 \mu\text{s}$, the values used above (`verify_transport_selection.py`).

Table 6.3: Ratio of neutral escape time to τ_{slow} , over the 45 defended breakdown pairs and over the other 635 analysed pairs. L is the distance from the emitting parcel to the plasma boundary. “Satisfies” means $\tau_{\text{esc}} > \tau_{\text{slow}}$, that is, the atom is ionised before it leaves, which is what the closed parcel assumes.

L	set	median $\tau_{\text{esc}}/\tau_{\text{slow}}$	largest	satisfies
10 cm	the 45 breakdown pairs	0.0079	0.159	0/45
10 cm	the other 635	0.0471	384	168/635
20 cm	the 45 breakdown pairs	0.0192	0.636	0/45
20 cm	the other 635	0.129	1535	227/635

Not one of the 45 pairs that carry the defended result satisfies the assumption used to compute it, at either scale length, while a quarter to a third of the pairs that do not break down satisfy it comfortably.

A count of zero is worth little if the 45 sit just under the line, so the distance matters. The closest of them is $[18, 4]$ heating, $T_e = 2.33 \text{ eV}$, $n_e = 5.18 \times 10^{13}/\text{cm}^3$, at $\tau_{\text{esc}}/\tau_{\text{slow}} = 0.636$ for $L = 20 \text{ cm}$: its escape time would have to rise by a factor 1.57 for the parcel to close. The median of the 45 would have to rise by a factor 52. One pair is marginal and the rest are not.

Three things were varied to see whether the conclusion depends on a choice. Tripling $\langle \sigma v \rangle_{\text{CX}}$, which lengthens every diffusive escape time in proportion, moves the count from 0 to 3 of 45; dividing it by three leaves it at 0. Replacing $\sqrt{2kT_n/m}$ by the Maxwellian mean speed $\sqrt{8kT_n/\pi m}$ shortens every escape time and leaves the count at 0. Doubling L from 10 to 20 cm leaves it at 0. The finding is not an artifact of any of the three.

The selection is structural, not incidental. A pair enters the census because τ_{slow} is long, since the single-slow-mode estimate of Eq. (5.16) rises with τ_{slow} at fixed τ_{drive} .

A long τ_{slow} means slow ionisation out of the ground state. Slow ionisation is exactly the condition under which neutral transport, whose timescale does not depend on τ_{slow} at all, sets the ground-state density instead of the local atomic physics. The census therefore selects for the regime in which its own assumption is weakest, and it does so by construction rather than by accident. That is why testing the assumption at the cold corner gave the reassuring answer and testing it inside the defended scope does not.

6.2.5 What this reaches, and what it does not

The finding above is about magnitudes; it does not touch the algebra.

Untouched by transport. The two-channel split of Eq. (3.33), exact to a relative residual of 3.075×10^{-14} ; the unit-width logistic form of the ground-fed fraction; the bound $\max |f_3 - f_4| = \tanh(|\Delta|/4)$ and its corollary $|\mathrm{d} \ln R / \mathrm{d} \ln b_1| < 1$; and the accuracy of the excited-state closure, 6.73×10^{-9} at the worst point. These are properties of the operator $\mathbf{L}(T_e, n_e)$ at a fixed operating point, or theorems about its structure. The function $f_3(u) - f_4(u)$ says how the line ratio responds to whatever the reservoir does; where on that function the plasma sits is the reservoir’s business, not the operator’s. Transport changes the reservoir’s history and cannot reach any of them.

Conditional on the closed parcel. The reservoir gain G , which is a secant along the CRE curve and therefore assumes the ground state responds through collisional–radiative balance alone; $\varepsilon_{\text{plateau}}$, which is built from it; the breakdown census; the sensitivity map $\bar{S}$ and the ridge location, both evaluated at the CRE operating point u^{CRE} that transport can move (Section 6.4); and every time-averaged magnitude in Chapter 5. If recycling renews n_{1s} on a timescale shorter than τ_{slow} , G is not the gain the plasma has. These numbers are not upper estimates. A fast-exchanging reservoir with a steady source freezes the closed-parcel error rather than removing it, and a rising source enlarges it: over the defended set the true error exceeds the closed-parcel plateau at 165 of the 180 (pair, L , source-multiplier) combinations tested (`verify_open_reservoir.py`). The closed-parcel value is a reference magnitude conditional on the reservoir’s history, not a bound on it.

The thesis therefore leads with the structure and the bound and reports the magnitudes as conditional; Section 6.2.4 adds that the condition is measured to fail at the points that matter, with $\tau_{\text{esc}}/\tau_{\text{slow}}$ between 2.5×10^{-4} and 0.64 over the 45 pairs at both scale lengths (Table 6.3).

6.2.6 The claim, restated as the conditional it is

The correct form is therefore not a caveat but a conditional, and it should be written this way wherever the number appears:

If the reservoir history $u(t)$ follows the local zero-dimensional collisional–radiative dynamics used here, with no external influx, the error is as computed; during the fast-stage plateau specifically, n_{1s} is frozen to within the closure residual.

The frozen-reservoir part is exact for the plateau, because the excited manifold equilibrates in nanoseconds while the reservoir has not yet moved, and the whole conditional is the standard implicit assumption of every 0-D collisional–radiative table in use, the object this thesis set out to test. But it is an assumption about the plasma, not about the atoms, and no eigenvalue analysis inside a 0-D matrix can validate it: the two timescales

are measured to differ by a ratio $\tau_{\text{esc}}/\tau_{\text{slow}}$ between 2.5×10^{-4} and 0.64 over the defended breakdown pairs and by 3.2×10^3 at the cold worst point (Section 6.11), and are assumed not to interact. What would settle it is a two-region or transport-coupled calculation, an edge code such as SOLPS (Scrape-Off Layer Plasma Simulation) supplying n_{1s} from a recycling influx; that is outside the scope of this thesis and is named in Chapter 7 among the further work.

6.2.7 One chord, one point: the line-of-sight integration

The same absence of a spatial dimension reaches the observable itself. Every ratio in this thesis is formed from populations at a single (T_e, n_e) , whereas a spectrometer forms its ratio from two brightnesses, each integrated along a chord that crosses whatever gradients the divertor has. As Section 1.3 sets out, the measured ratio is the n_4 -weighted average of the local ratio along that chord, and it equals the local ratio only where the local ratio is constant over the emitting region. The weighting is by H_β emissivity and not by path length, so the brightest stretch of the chord sets the answer and the cold, dim stretch contributes almost nothing.

What the calculation here establishes is therefore the error a table carries at a point. Converting it into the error a real chord carries would require the emissivity profile along the sightline, which is a transport result and not an atomic one; the two-region or edge-code coupling named in Section 6.2.5 is what would supply it. The local statement is the one this thesis defends. The chord statement is conditional on n_3/n_4 being uniform over the emitting volume, and nothing here tests that.

6.3 No molecular channels, in a gas that is mostly molecules

6.3.1 An atomic-only state space

Every state in Chapter 2's state space is a level of a single hydrogen *atom*, plus a reservoir of bare protons. There is no H_2 (or D_2), no molecular ion D_2^+ , no negative ion D^- , and no reactions between them. A search of the rate assembly for any of these returns nothing at all.

6.3.2 Molecular activated recombination as a third feed stream

Chapter 3 fed every excited level from two directions, up from the ground state by electron collisions and down from the ion continuum by recombination, and the error mechanism is the difference in how the $n = 3$ and $n = 4$ shells split their supply between those two feed streams.

Molecules open a third stream. In a cold, dense, mostly neutral gas, hydrogen atoms form H_2 ; the molecules are ionised to H_2^+ , or exchange charge, and the resulting molecular ions recombine with electrons and break apart, leaving one of the fragments in an *excited* atomic level. The whole family of such routes is called MAR, molecular activated recombination. The important feature for this thesis is not that MAR exists but where it delivers: it populates $n = 3$ directly, without the atom ever having passed

through the ground-state reservoir whose staleness is the entire mechanism here. A supply channel that bypasses the reservoir does not inherit the reservoir's lag.

6.3.3 How much gas is involved

At the worst point of the map the model's own equilibrium solution is 96.6 % neutral: $n_{1s} = 1.47 \times 10^{15}$ against $n_e = 5.18 \times 10^{13} \text{ 1/cm}^3$. A hydrogen gas at 1 eV that is 96.6 % neutral is a gas in which molecules form readily. The atomic description is not being applied at the edge of its validity there; it is being applied outside it.

The size of what is being left out is documented experimentally. *Detachment* is the operating regime in which the plasma cools enough that the ion flux to the target collapses, and the regime this diagnostic is chiefly used to identify. In spectroscopic inversions of divertor Balmer emission at its onset, molecularly-induced emission is found to account for up to 60 to 70 % of the D_α light and 10 to 20 % of D_γ [25, 26]. An H_α/H_β ratio computed from atomic processes alone is not the ratio such an instrument records.

6.3.4 What survives it, and what does not

Two consequences follow, and they point in opposite directions, so both must be stated.

The *numerical* cold-edge results are atomic-model results and nothing more. Adding a third supply channel that feeds $n = 3$ preferentially would change $f_3 - f_4$, and could plausibly change it a great deal; this work offers no direct calculation of how much. A sensitivity scenario can nonetheless be constructed by imposing the emissivity fractions cited in this section. For a given upper level the molecular fraction of the emissivity equals the molecular fraction of the population, because both channels emit through the same transition probability, and molecular-assisted recombination bypasses the ground state entirely, so it dilutes the ground-fed fraction as $f_p \rightarrow f_p(1 - \phi_p)$. Taking $\phi_3 = 0.60$ from the reported 60 to 70 % of D_α , and $\phi_4 = 0.30$ as a chosen value, since the sources report no $n = 4$ fraction, the D_γ 10 to 20 % belongs to $n = 5$, and an $n = 3$ source alone at $\phi_3 = 0.60$ already puts 0.216 into $n = 4$ on this operator (`validation/molecular_channel/`), the difference of ground-fed fractions at the cold corner falls from 0.4866 to 0.1088 and at the density crest from 0.4580 to 0.0998, and the secant sensitivity $\bar{S}$ falls with it, from 0.4822 to 0.0903 and from 0.4552 to 0.0903.

The plateau error does not fall in the same proportion, because $\varepsilon_{\text{plateau}} = |e^{\bar{S}\Lambda} - 1|$ is not linear in $\bar{S}$; treating it as linear understates the reduction. Imposing the same fractions on the model's own operator, with a non-negative source depositing into $n = 3$ and $n = 4$ so that (ϕ_3, ϕ_4) is reached exactly at the plateau, the cold-corner plateau error falls from 38.7 % to 6.3 %, a **factor of 6.1**; at the upper end of the reported ranges, $(\phi_3, \phi_4) = (0.70, 0.45)$, it falls to 4.1 %, a factor of 9.4. At the crest the same two targets give factors of 5.4 and 8.1, and at the benchmark 3.2 and 4.4 (`validation/molecular_channel/`).

Two features of that calculation are worth stating, because neither is visible in the dilution formula. The two fractions cannot be chosen independently: a source large enough to make $\phi_3 = 0.60$ already drives more than 0.15 of $n = 4$ through collisional $3 \rightarrow 4$ transfer, so $(\phi_3, \phi_4) = (0.60, 0.15)$ requires a negative $n = 4$ source and is not realisable on this operator. And the reduction is not monotone in ϕ_3 . The difference $f_3(1 - \phi_3) - f_4(1 - \phi_4)$ passes through zero near $\phi_3 = 0.74$ at $\phi_4 = 0.30$, beyond which $n = 4$ is the more ground-fed of the two and the error grows again with $\bar{S}$ carrying the opposite sign. At the crest, where that zero is reachable with non-negative sources, a

two-source dilution at $\phi_4 = 0.30$ brings the plateau error down to 0.35 % at $\phi_3 = 0.75$, past the zero at $\phi_3 = 0.7355$; a single $n = 3$ source there gives 0.30 % at $\phi_3 = 0.70$ and 1.1 % again at $\phi_3 = 0.80$, with $\bar{S}$ negative. A molecular channel strong enough cancels the diagnostic's sensitivity to the reservoir rather than merely reducing it, and past the cancellation the error returns.

The scenario is crude, it inherits every uncertainty in the quoted fractions, it applies only where those fractions do, and it holds the transient fixed: a molecular reservoir would also move $u(t)$ and $n_{\text{ion}}(t)$, which an imposed source cannot.

The *structural* results, the list of Section 6.2.5, are theorems about whatever matrix $\mathbf{L}$ is supplied, and they hold for one containing molecular states fed from the same two reservoirs, which changes the numerical value of the ingredients and not the form of the result. They do not hold for molecular hydrogen as a third independent reservoir: the observable then depends on two reservoir coordinates, u and $n_{\text{H}_2}/n_{\text{ion}}$, the logistic form in u survives only at a fixed value of the second, and the error depends on two histories rather than one (Section 7.5).

What is *not* available is a calculation with molecules in it: the factors above are the model's own operator driven by an imposed source of the published strength, not a solution of a matrix containing molecular states. This is an atomic model, its cold-edge numbers are atomic-model numbers, and the scenario above is an estimate of how wrong they are rather than a correction to them.

6.4 Where the ridge sits depends on what the model assumed about neutrals

6.4.1 Mechanism and location as separate claims

Section 5.6 found an interior maximum, a ridge, in the error map at $n_e \approx 2 \times 10^{13} \text{ 1/cm}^3$, and showed that it moves with the shell pair as Griem's criterion for local thermodynamic equilibrium (LTE) predicts. The mechanism is robust; three independent tests agreed. But a mechanism and a location are different claims, and only the first of them has been established.

6.4.2 The neutral density as the controlling variable

The ridge exists because the $n = 3$ and $n = 4$ shells switch from recombination-fed to ground-fed at different reservoir strengths, so their difference vanishes when either feed dominates and peaks in between. The controlling variable is therefore not the electron density directly; it is *how much neutral gas there is to feed from*. The electron density enters only because, in this model, it determines the neutral density through the CR ionisation balance.

That is the conditionality. Change what the model assumes about neutrals and the ridge moves, whatever the electron density does.

6.4.3 Ridge location under a rescaled neutral density

Within the two-channel split the ground-fed fraction is exactly $f_p = a_1 n_g / (a_0 + a_1 n_g)$, where n_g is the ground-state density feeding the excitation channel and a_0, a_1 are the

network responses of Section 3.5.3. Scaling n_g is therefore an exact operation on the ridge, not an approximation to one. Doing it, at $T_e = 2.02\text{ eV}$ (`verify_ng_scaling.py`, `validation/ng_scaling/`):

n_{1s} scaling	density at which $ f_3 - f_4 $ peaks
$\times 0.01$	$\leq 10^{12}\text{ 1/cm}^3$ (off grid)
$\times 0.1$	$2.30 \times 10^{12}\text{ 1/cm}^3$
$\times 1$ (the model’s own CRE value)	$1.70 \times 10^{13}\text{ 1/cm}^3$
$\times 10$	$2.50 \times 10^{14}\text{ 1/cm}^3$
$\times 100$	$\geq 10^{15}\text{ 1/cm}^3$ (off grid)

One decade in the assumed neutral density moves the ridge by roughly one decade in electron density. The response is close to proportional, and the whole three-decade span of the density grid is traversed by a two-decade change in n_{1s} . The values in the table are those of the stamped scan, the operator at $[15, j]$ with its own CRE n_{1s} scaled.

The agreement with Griem’s criterion at unit scaling, 1.93×10^{13} measured against $2.05 \times 10^{13}\text{ 1/cm}^3$ predicted, is therefore an agreement evaluated at *this model’s* collisional–radiative ionisation balance, which contains no transport and no recycling. In a real divertor the neutral density is set by recycling, which is Section 6.2’s objection arriving in a different place. The sentence “the ridge sits at $n_e \approx 2 \times 10^{13}\text{ 1/cm}^3$ ” is a statement about an ionisation balance, not about a divertor, and should be quoted with that condition attached or not quoted.

The ridge is also one line pair’s and not the Balmer series’: the $(3, 5)$ pair of H_α/H_γ has its worst density at $7.2 \times 10^{12}\text{ 1/cm}^3$, a factor 2.68 lower, and the form of words this imposes is stated in Section 5.2.

6.5 The model does not identify the crest with detachment

6.5.1 The verbal match between the ridge and detachment

The identification is tempting. The ridge sits where ionising and recombining supplies compete; detachment, the regime ITER must operate in to survive its exhaust power and the transition this diagnostic is chiefly used to identify (Section 6.3), is where ionising and recombining plasma compete. The words match.

6.5.2 The ridge density against the detached-divertor band

The numbers do not. Section 5.6.5 withdrew the identification: the ridge, at $1.93 \times 10^{19}\text{ 1/m}^3$, sits a factor 5.2 below the detached-divertor band of published edge modelling [45] and a factor 1.6 below the modelled upstream separatrix density [1, 2], in a range neither source characterises.

The comparison band deserves a word, because a factor of 1.6 is not a large margin to be arguing over. Read from the primary source, the band is 3×10^{19} to $6 \times 10^{19}\text{ 1/m}^3$ across the 53 ITER SOLPS-4.3 cases [1, Sec. 3.1.1], and it is the density *at the outer midplane*, which that paper states explicitly where it plots the dependence [1, Sec. 3.2.2]. So the

quantity is settled: outer midplane, not X-point. What is not settled by a single factor is where in the band to measure from. The ridge sits 1.6 times below the *bottom* of it and 3.1 times below the top, and the precise statement is the range rather than either end. On the same flux tube the outboard-midplane and X-point densities differ by roughly 50 %, which is comparable to the width of the band itself. The conclusion does not turn on any of this, since the ridge is below the band on every reading, but the factor should not be quoted more precisely than the band it is measured against. The repository recorded [2] as “Nucl. Fusion 62 (2022)” in `src/validation/grid.py`; the DOI was right, the volume and year wrong. It is Nuclear Fusion **63**(1) 016017 (2023), companion part A at 63(1) 016016 [46]; [3] in the same block of that file genuinely is volume 62.

6.5.3 A structure without a named tokamak region

The resulting statement is uncomfortable and should be made anyway. **The structure this work found is real and reproducible, and it sits at a density this thesis cannot identify with any named region of a tokamak.** The model does not identify it with detachment or with the upstream separatrix, and calling it either would be a claim about edge transport that a 0-D atomic model cannot support. The ridge is a property of how two shells of hydrogen divide their supply. It needs no divertor geometry to produce it, and it acquires none by being described in one.

The ridge and the detached band do not overlap: restricted to $T_e \geq 2 \text{ eV}$ and $n_e \geq 10^{14} \text{ 1/cm}^3$ the census is 0 of 108 with a worst estimate of 7.2 % (Table 6.1), so within the closed-reservoir atomic model, at the densities where the diagnostic is most used, no tested case reaches the threshold.

6.6 The open ion boundary: a closed-parcel test

6.6.1 The fixed ion density in the rate equation

Section 3.1 was explicit that the ion density n_{ion} is held fixed: the recombination feed $\mathcal{S}n_{\text{ion}}$ enters as a constant source term a feed stream into a CSTR in the reactor picture, rather than as a matrix element, and the 43-state system has no 44th state for ionised atoms to accumulate in. Atoms ionise and vanish from the accounting. The slow eigenvalue λ_0 is therefore the CR ionisation rate of ground-state hydrogen against a prescribed reservoir, not the equilibration rate of a closed system.

6.6.2 Closing the manifold with a 44th ion state

If the long τ_{slow} values were an artifact of the open boundary, so would be the claim that the error outlasts an ELM. The falsifier named in Section 4.8 before the test was run is a closure factor of 2300 at $[0, 4]$, the value of $\tau_{\text{slow}}/\tau_{\text{drive}}$ there. Closing the manifold with a 44th ion state gives a closure factor with minimum 1.00, median 1.00 and maximum 41.4 over the grid, moves the ELM census from 202 to 200 of 680 and the worst case by one part in 325, and leaves the count above 2 eV at 45 before and 45 after (Table 4.2). The factor is large only where τ_{slow} already exceeds the event duration by three orders of magnitude and is unity wherever τ_{slow} approaches τ_{drive} , so where the boundary condition matters the outcome does not depend on it. The measured factor at $[0, 4]$ is 15 against the 2300 required, a cold-corner margin; inside the defended range, at the ridge, it is

1.01 against 20, and at the benchmark τ_{slow} moves by one part in 10^3 . The objection was survived everywhere it was pressed.

6.6.3 The caveat that must not be softened

The closure adds an ion reservoir. It does *not* close the electron balance, since n_e is held fixed while the ion population moves, and it adds no transport whatever. It is a test of the boundary condition on the atomic manifold, and nothing more. Section 6.2 is the objection this test does not answer, and no amount of closing the atomic system will answer it.

The two index labels for this point, $[1, 4]$ in Table 4.2 and $[0, 4]$ in the census, are the post- and pre-step operators of one heating step: the census rows carry the pre-step index, every quoted timescale belongs to the post-step operator that governs the relaxation (Chapter 5), and the two must not be interchanged. Diagonalising each gives $\tau_{\text{slow}} = 0.421\,012\text{ s}$ at $\mathbf{L}[0, 4]$ and $0.233\,239\text{ s}$ at $\mathbf{L}[1, 4]$, so the quoted $0.233\,24\text{ s}$ belongs to $\mathbf{L}[1, 4]$.

6.7 An external benchmark this model fails, and what the failure measures

6.7.1 One population coefficient agrees and one does not

Chapter 4 presented the comparison against Fujimoto’s published population coefficients as the strongest external check this model passes. It is, for one of the two coefficients. The other disagrees at low principal quantum number, and this section reports what that disagreement measures. The short answer is that it is unexplained, that it is not the ℓ closure, and that it is direct evidence about the coefficients Chapter 5 uses.

6.7.2 The recombination-fed and ground-fed coefficients r_0 and r_1

Chapter 3 split every excited population into two pieces corresponding to the two supply channels. The population coefficient $r_0(p)$ measures how strongly level p is fed from the ion continuum by recombination; $r_1(p)$ measures how strongly it is fed from the ground state by excitation. The published tabulation gives both, computed independently, decades ago, by a different method. Agreement is a real test; disagreement is a real failure.

6.7.3 The disagreement, and what it measures

The recombination-fed coefficients agree. At $n_e = 10^{12}\text{ 1/cm}^3$, r_0 matches to 1.3% at $p = 2$ and to 12 to 14% at $p = 3$ and $p = 4$. The ground-fed coefficients at low principal quantum number do not: this model’s r_1 sits a factor 8.3 low at $p = 3$ and a factor 4.4 low at $p = 4$.

Those are the two shells whose difference is the mechanism of Chapter 5, so the disagreement had to be understood rather than noted.

Two candidate explanations were eliminated before the source was consulted. If the excitation rates feeding the ground-fed channel were wrong, the deficit would follow, but

those rates were benchmarked against R-matrix-with-pseudostates calculations and the dominant $\Delta\ell = 1$ channels agree at a median of -5.0% over 120 comparisons. And a normalisation error is excluded by r_0 itself: r_0 and r_1 carry the same p -dependent factor $Z(p)$, so an error large enough to move r_1 by a factor of 8 would have moved r_0 with it.

6.7.4 The source’s closure, imposed and tested

Table 4.1(b) is indexed by $p = 2, 3, 4, 5, 7, 10, 15$. Those are principal quantum numbers, and the table carries no ℓ label anywhere. The published coefficients are shell quantities computed with the ℓ substructure bundled.

This model resolves ℓ below $n = 9$, so the closures differ. That difference is tempting as the explanation, since recomputing the model with proton ℓ -mixing removed moves $r_1(2)$ by a factor 118 and $r_1(3)$ by a factor 3.35 at $n_e = 10^{12} \text{ 1/cm}^3$. That argument does not survive its own test. Removing ℓ -mixing is the limit in which the process is absent; the tabulation assumes the limit in which it is infinitely fast, and those are opposite ends rather than a bracket. Imposing the source’s own closure directly, by bundling this model’s operator under statistical ℓ weights, moves $r_1(3)$ by 0.42% and makes the $p = 2$ agreement worse (`verify_fujimoto_bundle.py`, stamped in `validation/fujimoto_bundle/`).

The disagreement therefore does *not* measure the ℓ closure, and this thesis does not explain it. Nor does the deficit stop short of $f_3 - f_4$ on the ground that the difference never passes through a bundled quantity; bundling is not the protection. Section 4.9 makes $a_p = r_1(p) Z_p/Z_1$, so a deficit in $r_1(3)$ and $r_1(4)$ is a deficit in a_3 and a_4 themselves, and the only protection is that Δ and the cap depend on the ratio a_3/a_4 , in which a deficit common to both shells cancels. It is not common at $10^{12}/\text{cm}^3$, 8.3 against 4.4, and nearly so at $10^{15}/\text{cm}^3$, 1.8 against 1.9. Dividing a_3 and a_4 by those ratios at every grid point (`validation/fujimoto_r1_consequence/`) moves Δ by $+0.64$ and -0.08 ; over the 280 points above 2 eV the cap moves by 29% and -4% at the median, and the sensitivity at the operating point by 30% and 11% at the median, with ranges of -58 to 524% and -38 to 68% . The disagreement is direct evidence about the coefficients Chapter 5 uses; what is open is whether the deficit lies in this model or in the tabulation, and the stake is stated in Section 4.9. The open status is carried in Section 6.10.

6.7.5 Two limitations of the benchmark, stated precisely

Table 4.1(b) is computed at $T_e = 1.28 \times 10^5 \text{ K}$, which is 11.03 eV. The grid used here reaches 10 eV, so the comparison is made at the grid edge and 9.3% below the tabulated temperature.

There is no second usable case. The companion Table 4.1(a) is at $T_e = 10^3 \text{ K}$, which is 0.0862 eV, two orders of magnitude below this grid. The external check is therefore one temperature at one edge, and it should not be described as more than that.

6.7.6 What the benchmark establishes, and what it does not

What the benchmark does establish is the recombination-fed channel, which agrees, and that is worth stating as a pass rather than leaving implicit beside a disagreement. What it does not establish is the absolute normalisation of the ground-fed channel against an independent calculation, because no tabulation resolved in ℓ was available to compare against.

The sensitivity of the density maximum to the ground-fed supply remains a separate and open question, and Section 6.4 treats it on its own terms: the location responds roughly decade-for-decade to the assumed neutral density, which is a stronger reason to quote it as a range than anything the benchmark shows.

6.8 The top of the level ladder, and what the bundling costs

6.8.1 Truncation at $n_{\max} = 15$ and bundling above $n = 8$

A hydrogen atom has infinitely many bound levels, crowding together towards the ionisation limit. A computation cannot have infinitely many states, so a boundary was chosen: this model tracks levels up to $n_{\max} = 15$ and discards the rest. Above $n = 8$ it also stops resolving orbital angular momentum ℓ , treating each shell as a single bundled state populated statistically. Both are choices, and this section states what they cost.

6.8.2 Three consequences, in decreasing order of severity

The terminal shell is over-populated. Truncation removes from the terminal shell's balance the collisional excitation out of it to $n > n_{\max}$, a way out, and the de-excitation and cascade into it from those levels; the missing loss is upward excitation, not cascade, which runs downward. Against published values this model's ground-fed coefficient at $n = 15$ is high by a factor of 4.9 to 6.2, while at $n = 10$ it agrees to 4 to 20%; the budget is in Section 4.6.2. *No result in this thesis rests on the top shell:* the observable is built from $n = 3$ and $n = 4$, both fully ℓ -resolved and far from the truncation.

The bundling nonetheless enters the result indirectly. The recombination flux cascading down from the bundled shells feeds the recombination-fed channel of both $n = 3$ and $n = 4$, one of the two channels the mechanism compares, which is why Section 5.6.3 lists truncation among the alternative explanations for the ridge that this work does not exclude.

There is no ℓ -mixing inside the bundled block. The proton-impact ℓ -mixing that Section 2.7 showed to be essential for the $n = 2$ manifold is present only for $n = 2-8$; the ℓ -mixing rate array is identically zero above that. Within the bundled shells the ℓ -distribution is assumed statistical rather than computed. At high n that assumption is far more defensible than it would be at low n , because the sublevels are nearly degenerate and every collisional process mixes them, but it is assumed, not measured, and this thesis has not measured it.

6.8.3 Collisional ℓ -mixing against escape in the bundled shells

The script that quantifies this was written and, until now, never executed. Running it produced a result and also produced three defects in the script itself, which are reported here because two of them would have made its verdict worthless.

What the script could not do. It promised to test the bundled shells against the pipeline’s own ℓ -mixing array, and that is impossible: the array is identically zero for states 36 to 42, in both directions, because ℓ -mixing was computed for the resolved block alone. The file contains nothing to test. The comparison has to be made against the rate *function* rather than the rate table, and the function is the one `compute_lmix.py` used to build the resolved block, so it is the model’s own rate rather than a second one. Calling it at $n = 2$ to 8 reproduces the stored array to 1.0000, which is a wiring check and not a physical confirmation; it establishes only that the rate carried up to the bundled shells is the same function with the same arguments.

The two defects that would have inverted the answer. The script carried its own Pengelly–Seaton expression, independent of the pipeline’s, which disagrees with the rate the model actually uses by a factor of 95 at $n = 2$ rising to 5.6×10^3 at $n = 8$, growing roughly as n^2 . Used as written, it would have reported the mixing rate two to three orders of magnitude below its true value and could have returned a verdict of marginal where the answer is not close. It is deleted rather than corrected: a verification script has no business maintaining a second implementation of a rate the pipeline already computes. The script also fell back on a synthesised logarithmic temperature grid of its own when the pipeline’s grid files were absent, printing a warning and continuing, which is the failure mode described in Chapter 4: every ratio would have carried a correct-looking label attached to conditions the model was never evaluated at. It now raises.

The result. The criterion is that the ℓ -mixing rate out of the np substate, the substate the Lyman channel empties fastest, must exceed that substate’s total radiative rate. Over the seven bundled shells and all 400 grid points, the ratio is above the well-mixed threshold of 10 at every one of the 2800 combinations. The worst case is $n = 9$ at $T_e = 1$ eV, $n_e = 10^{12}/\text{cm}^3$, where the mixing rate is $1.155 \times 10^{10}/\text{s}$ against $A(9p) = 7.51 \times 10^6/\text{s}$, a margin of 1538, that is, a factor 154 clear of the threshold (`verify_bundling_psm20.py`, `outputs/bundling/bundling_report.txt`).

Two inputs to that comparison were checked rather than assumed. The Einstein coefficients $A(np \rightarrow 1s)$ for $n = 9$ to 15 are hardcoded in the script, because the model’s own radiative file resolves ℓ only to $n = 8$ and has nothing to read above it. Since $A(np \rightarrow 1s)n^3$ tends to a constant for hydrogen, the model’s resolved values and the hardcoded table must lie on one curve: they do, with a step of 0.43% across the boundary at $n = 8$ to 9 and a monotone approach to 4.13×10^9 at $n = 15$. And the ℓ -mixing coefficient carries a Debye cutoff frozen at $n_e = 10^{14}/\text{cm}^3$, because the array has no density axis. That freezing is not small: re-evaluating the cutoff at the local density changes the $n = 15$ rate by a factor 20 upward at $10^{12}/\text{cm}^3$ and a factor 30 downward at $10^{15}/\text{cm}^3$. Against the radiative criterion it does not move the verdict, because the worst point sits at the lowest density, where the local cutoff raises the rate: the worst ratio is 1538 with the freezing and 3.26×10^4 without it. At the highest densities the same correction lowers the rate, and that matters for the criterion that follows.

What this closes and what it does not. The criterion above compares mixing with radiative decay, and radiative decay is not how a high shell empties at these densities. The total loss out of a bundled state, $-L_{kk}$, is 89.5 to 99.1% collisional transfer to other shells, 0.9 to 10.5% ionisation and at most 0.19% radiative (`verify_bundling_total_loss.py`, `validation/bundling_total_loss/`). Against that total the margin is not three orders

of magnitude. With the frozen cutoff the model is built with, mixing exceeds total escape at every one of the 2800 combinations, by 26 at $n = 9$ falling to 2.85 at $n = 14$, the minimum, all at 1 eV and $10^{12}/\text{cm}^3$, and 288 combinations sit below 10. With the cutoff evaluated at the local density the ratio falls below unity at 121 of the 2800, in every bundled shell, at the two highest densities and $T_e \leq 2.95$ eV; at 1 eV and $10^{15}/\text{cm}^3$ it runs from 0.86 at $n = 9$ to 0.09 at $n = 14$. The comment in `compute_lmix.py` that the cutoff factor varies by about 30 % across the grid understates the variation already at $n = 2$, where the local-to-frozen ratio of F runs from 1.96 at $10^{12}/\text{cm}^3$ to 0.53 at $10^{15}/\text{cm}^3$ at 1 eV (`validation/bundling_total_loss/`, RESULT 4), and fails outright above $n = 9$, where $U_m > 1$ and F falls as $n_e^{-3/2}$; it is reported here and left in place. What survives is narrower than the radiative criterion suggests, and it rests on a different mechanism. Where mixing wins, the ℓ -distribution is statistical whatever the feeds do. Where collisional transfer wins, at the two highest densities below 3 eV, the distribution is set by the feeds. Adjacent-shell n -changing collisions, which carry 71 to 95 % of the collisional loss, might be supposed close to ℓ -blind at high n . They are not. In the CCC rates the model is built from, which resolve ℓ on both sides up to $n = 10$, the adjacent upward transfer out of (n, ℓ) varies across ℓ by a factor that grows with n , from 1.22 at $n = 2$ to between 2.5 and 2.8 at $n = 9$, and the downward transfer by a factor of 5 to 6 at $n = 7$ to 10, largest at the highest ℓ in every case; weighted by $2\ell + 1$ the scatter at $n = 9$ is 27 to 46 %. The transfer into a shell from a statistically populated neighbour is not statistical either: its departure coefficient runs from 0.28 at $\ell' = 0$ to 1.66 at $\ell' = n - 1$ for $n = 9$ and 10. What licenses the statistical distribution at those 121 combinations is detailed balance, not blindness. Every ℓ -resolved rate obeys $K_{\text{deexc}} = K_{\text{exc}} (\omega_i/\omega_f) e^{\Delta E/kT_e}$, so when the two neighbouring shells stand in their mutual Boltzmann ratio the statistical distribution is the fixed point of the collisional exchange whatever the ℓ -dependence of the rates: the departure of sublevel ℓ is $\theta - 1$ times a factor bounded by 0.16, with θ the Boltzmann departure of the neighbour pair. At the 34 grid points the 121 combinations occupy, the model's own populations give $|\theta - 1| \leq 6 \times 10^{-5}$ for every bundled pair, and every shell from $n = 8$ to 15 sits within 10^{-4} of Saha–Boltzmann equilibrium with the continuum: the high shells are thermalised there. The model's own $n = 8$, its highest ℓ -resolved shell, computed with the ℓ -mixing block removed, is statistical to 10^{-4} at every one of those points. The two criteria are complementary: $|\theta - 1|$ exceeds 0.01 only below $3 \times 10^{12}/\text{cm}^3$, where the locally evaluated mixing exceeds the total loss in every bundled shell by at least a factor of 341, and no grid point fails both (`verify_1_blindness.py`, `validation/1_blindness/`). The stake was not small. Had the bundled shells been non-statistical at those points, setting their radiative decay to zero would move u_{CRE} by up to 1.9 % and $f_3 - f_4$ by 1.2 %, and raising each shell's radiative rate to that of its np sublevel would move them by 21 % and 13 %, because the bundled radiative flux is 0.27 to 0.73 of the small net flow into the ground state there. None of this closes the terminal-shell over-population above, which is a different defect with a different cause, and none of it makes the ℓ -distribution inside $n = 11$ to 15 a computed quantity, since only ℓ -summed rates exist there; for those shells the statement rests on the Boltzmann and Saha departures alone.

6.9 Remaining assumptions, and the sensitivities that are known

Four assumptions remain. Each is stated with what is known about it, rather than listed as a disclaimer, and two of the four have been tested.

Maxwellian electrons: untested. Every rate coefficient in Chapter 2 is an average of a cross section over a Maxwell–Boltzmann electron energy distribution. Ionisation and excitation are driven by the high-energy tail of that distribution, which holds a small minority of the electrons. A 1 eV plasma next to a recycling surface has no guarantee of a Maxwellian tail, and a depleted or enhanced tail would change the rates that set both clocks. Not tested here.

$T_i = T_e$ in the ℓ -mixing rates: untested. The proton-impact ℓ -mixing rates of Section 2.7 are evaluated at an ion temperature equal to the electron temperature. In a detaching divertor the ions are frequently colder than the electrons. Not tested here.

Both limit the transfer of the magnitudes to a divertor, not the two-channel structure, which holds for any rate set.

The step is instantaneous: tested, and safe for a reason that is easy to get wrong. The error maps impose a sudden temperature step, whereas a real ELM has a finite rise time. The tempting ratio $\tau_{\text{relax}}/\tau_{\text{drive}} \approx 10^{-5}$ is about the wrong quantity, the fast excited-manifold step error (Section 5.8); the plateau error is governed by the slow ground-state response, for which the relevant ratio is $\tau_{\text{slow}}/\tau_{\text{drive}}$. At the worst point of the grid that ratio is 2332, but that point sits at 1 eV and the figure is not the margin this result runs on. Over the 448 defended pairs the ratio has median 0.815, and over the 45 that break down it runs from 1.95 to 288, median 11.2 (`validation/divertor_map/`). The step idealisation is therefore safe by one to two orders of magnitude at the median of that set, not by three. The governing quantity is in any case $D_e = \tau_{\text{slow}}/t_{\text{ramp}}$ rather than $\tau_{\text{slow}}/\tau_{\text{drive}}$: a linear ramp reaches $D_e(1 - e^{-1/D_e})$ of the step’s plateau (`validation/ramp_plateau/`), so any ramp with $\tau_{\text{relax}} \ll t_{\text{ramp}} \ll \tau_{\text{slow}}$ reaches the same plateau and the idealisation *is* safe. It is safe because of the slow ratio, and the fast ratio must never be quoted for or against this result. The margin is not uniform across the breakdown set, and the bottom of the range is worth stating plainly: evaluating the same lag law at $D_e = 1.95$ gives 0.78 of the step plateau, a suppression of 22 %, which is the size of the errors this chapter is discussing. At the median $D_e = 11.2$ it is 4 % and at $D_e = 288$ it is 0.2 %. Those three figures are arithmetic on the stamped ratios and the law above, not a separate run. Where the step idealisation matters most it is therefore good to a few per cent, and at the single least favourable pair it is not.

The step is in temperature only: tested, and the correction goes the wrong way. A real ELM raises n_e as well as T_e . The joint map over 4968 cases (Section 5.8.3) leaves the exact decomposition unchanged, with the gain generalising to an additive gradient accurate to 0.2 % at the median; over $T_e \geq 2$ eV the breakdown fraction moves from 10.0 to 11.3 % when the density steps up one index and the worst case from 0.1748 to 0.2206. The earlier benchmark spot check, a bound falling by a factor 4, had the right

number and the wrong sign for the grid, because the benchmark and the worst point lie on opposite sides of the density crest; the fixed-density map understates the worst case.

Two sensitivities of the reported census belong in the same place, because both are properties of a reporting choice rather than of the physics, and both would otherwise be mistaken for measurements.

The step size is one interval of the temperature grid, which displaces the reservoir by $|\ln x|$ between 0.124 and 0.682 across the grid. A four-interval step gives 0.516 to 2.563 and raises $\varepsilon_{\text{plateau}}$ at the ridge from 18.1 % to 81.1 %. On step size alone, therefore, the reported numbers *understate* the error rather than overstate it.

The 10 % breakdown threshold is doing real work: at 5 %, 10 % and 20 % the count over the same 680 points is 348, 202 and 61 (Section 5.8), so the count is not robust to the threshold, the worst case is, and “ N points break down” must carry both.

6.10 Scope: five questions this thesis leaves open

The objections computed in this chapter were all answerable from inside a 0-D collisional-radiative matrix. Five are not, and they are gathered here rather than left scattered, so that what this thesis does not determine is visible at once. Each is stated with the reason it is out of reach and with the calculation that would settle it. None is a defect of the implementation; each is a boundary of the model class.

1. The r_1 deficit is unexplained. This model’s ground-fed population coefficient sits a factor 8.3 below Fujimoto’s Table 4.1(b) at $p = 3$, the shell the mechanism runs on, and a factor 1.8 to 2.0 below it at $\log_{10}(n_e/1/\text{m}^3) = 21$; truncation (1.2 %), Lyman trapping and the source’s own statistical- ℓ closure (0.42 %) have each been eliminated (Section 6.7). What would settle it is an ℓ -resolved published tabulation or the atomic data underlying the tabulated values, neither of which was available. The deficit is reported as open, with its consequence for $f_3 - f_4$ computed in Section 4.9.

2. There is no convergence study above $n_{\text{max}} = 15$. The state space is cut at $n = 15$ because the atomic data stop there: the quantum cross sections extend to $n = 10$, shells 11 to 15 are already semiclassical Vriens–Smeets values, and no radiative or recombination coefficients exist above 15 in any form. Truncation has been tested downward, and it moves $f_3 - f_4$ by about 1 % with the sign changing across the grid (Section 4.6.2). What is untested is the sensitivity of the ridge *position* to n_{max} , which would require rebuilding the rate pipeline at $n_{\text{max}} = 20$ rather than running a validation script: new excitation, ionisation, radiative and recombination coefficients for five further shells, and a re-indexing of every array built on the 43-state layout. That is a rate-generation exercise, not an analysis one.

3. The ground-state population is solved for, not supplied. Every magnitude in Chapter 5 is conditional on n_{1s} coming from $\mathbf{Ln} + \mathbf{Sn}_{\text{ion}} = 0$; in a divertor it is set by recycling and transport, on a timescale shorter than the local ionisation balance with $\tau_{\text{esc}}/\tau_{\text{slow}}$ between 2.5×10^{-4} and 0.64 over the 45 defended breakdown pairs, median 0.0079 at $L = 10$ cm, and by 3.2×10^3 at the cold worst point, while at the benchmark transport is the slower of the two, at 3.6 (Section 6.2). The open-boundary test of Section 4.8

bounds the sensitivity without removing the conditional, and what would settle it is stated in Section 6.2.

4. There are no molecular channels in $\mathbf{L}$. At the onset of detachment a large fraction of the Balmer light is molecular in origin. The dilution estimate of Section 6.3 says how far the cold-corner number would move under one specific dilution picture, not what a molecular model would give. What would settle it is a state space containing H_2 , H_2^+ and H^- with their rate set, a different model rather than an extension of this one.

5. The electrons are assumed Maxwellian. The reasons are given in Section 6.9. The rate integrals as built do not accept an arbitrary electron energy distribution, so this is not a parameter that can be varied within the present code, and the $T_i = T_e$ assumption in the ℓ -mixing rates is untested for the same reason.

Items 3 and 4 are named in Chapter 7 among the further work. Items 1 and 2 are bounded in their consequences rather than in their causes: the thesis states what each would change if it went the wrong way, and for item 1 that stake is computed in Section 4.9.

6.11 What is left standing

6.11.1 What this model cannot say

It cannot say anything quantitative about the cold edge of its own grid. Below 2 eV the optically thin plateau error is inflated by a factor of 2.5 to 3.3 by neglected Lyman trapping; the size of the correction depends on a slab thickness the model does not contain; and below 1.15 eV even the location of the ridge is undetermined.

It cannot say what a real divertor does, because recycling renews the ground-state population there on a transport timescale the model does not contain: 72 μs against 233 ms at the worst point, a factor 3.2×10^3 , below the scope boundary; inside the defended range $\tau_{\text{esc}}/\tau_{\text{slow}}$ runs from 2.5×10^{-4} to 0.64 over the 45 defended breakdown pairs, with a median $\tau_{\text{esc}}/\tau_{\text{slow}}$ of 0.0079 over the 45 defended pairs (Table 6.3); and 3.6 at the benchmark, where transport is the *slower* of the two.

It cannot say what an instrument records at the onset of detachment, because a large fraction of that light is molecular in origin and there are no molecules in $\mathbf{L}$.

It cannot name the density at which its own ridge sits as a region of a tokamak. The ridge falls between the modelled upstream separatrix and the detached divertor band, and cannot be assigned to either by this model.

And it cannot establish the absolute normalisation of its ground-fed population coefficient against an independent calculation: r_1 sits a factor 8.3 below the standard tabulation at $p = 3$, the disagreement is unexplained, and its stake for $f_3 - f_4$ is computed in Section 4.9 (Sections 6.7 and 6.10).

6.11.2 What this model can say, with the objections computed

Above 2 eV the breakdown census is invariant under Lyman trapping across a twentyfold range of assumed slab thickness, 45 of 448 in every run with the worst case moving by 0.5 % (Section 6.1). Closing the open boundary moves the census by two points out of 680 and the worst case by one part in 325; the factor 2300 that would have refuted the

result at the cold corner measures 15, and at the ridge the factor required is 20 against a measured 1.01 (Section 4.8). The ridge mechanism was verified over $T_e \geq 2$ eV, where trapping is now shown not to move the defended census within the tested escape-factor model ($D = 1$ to 20 cm, $T_n = T_e$), so the agreement with Griem's boundary criterion is not an artifact of neglected opacity. At citable divertor densities above 2 eV the 100 μ s time-averaged error never reaches 10 % over 108 pairs, worst case 7.2 % (Section 6.5).

6.11.3 The boundary of the claim

The pattern is consistent, and it is why this chapter is not an apology. Every objection that could be computed inside a 0-D atomic model was computed. The structural result and its bound came through unchanged; the census numbers came through as conditional magnitudes, unchanged under Lyman trapping and under the closed boundary, and shown by Section 6.2.4 to rest on a reservoir history the model does not supply at the points that carry them. The objections that remain, neutral transport and molecular chemistry, are the ones that cannot be computed inside a 0-D atomic model at all, and they have been stated as conditionals rather than dissolved into caveats.

That is the boundary of the claim. What is inside it is a quantitative statement about when a spectroscopist may trust a lookup table and when she may not, and Chapter 7 puts it into the form she can use.

Chapter 7

What It Means

The physicist of Chapter 1 is still at her console. Nothing in this thesis takes her spectrometer away, alters the two wavelengths she measures, or questions the ratio she forms from them. What it changes is the moment at which she is entitled to look that ratio up. This chapter states that condition, hands over the two numbers that convert it into an error estimate for her own transient rather than for the one modelled here, and says how much of the reasoning behind it belongs to this work and how much to a literature that is sixty years old.

7.1 The excited-state closure holds; the reservoir closure does not

Chapter 1 asked how large the error in the temperature inferred from the H_α/H_β ratio at known density becomes when the plasma is disturbed faster than its ionisation balance can relax, what physical quantity sets that error, and where in the operating range it is largest. It also named a hypothesis. The quasi-steady-state (QSS) approximation eliminates several dozen dynamical variables by asserting that the excited levels equilibrate faster than conditions change, and the expectation was that some of them would fail to keep up.

They keep up. Integrating the full 43-state system through a temperature step and evaluating both error measures along the same trajectory gives $\varepsilon^{\text{QSS}} = 8.66 \times 10^{-6}$ at the benchmark point and $\varepsilon^{\text{QSS}} = 6.73 \times 10^{-9}$ at $[0, 4]$, the grid point carrying the largest diagnostic error anywhere (Section 5.1.2). The closure is at its most accurate exactly where the tabulated answer is at its worst. The spectrum explains why: given the quasi-steady manifold and $S \neq 0$, which holds at all 400 points, there is no intermediate mode for an excited population to lag on (Section 4.8, two hot dense points excepted), and the propagated trajectories confirm it directly.

What fails is the other assumption, the one Section 1.5 separated out and then left standing. A table with two axes has to dispose of a third argument, the neutral-to-ion ratio $u = n_{1s}/n_{\text{ion}}$, and the rule it uses to dispose of it is equilibrium ionisation balance. That rule asserts that the plasma has finished changing. After a disturbance it fails by an amount that decays only on τ_{slow} , and the excited manifold, having re-equilibrated in nanoseconds onto a reservoir that has not moved, reports the partial-equilibrium state faithfully. The ratio is measured correctly and the table then assigns it to the wrong plasma.

If the fast states were the problem, the repair would be a time-dependent solve of the

excited manifold, and Section 5.1.2 shows that such a solve returns an answer already correct to a part in 10^5 at the benchmark and to 4.4×10^{-5} everywhere analysed. The repair belongs instead at the closure applied to u , which is a statement about the plasma's recent history and not about its atoms. Calling the failure a breakdown of the quasi-steady-state approximation would point the effort at the one part of the calculation that is working.

7.2 A forward criterion for the equilibrium lookup

The first thing to say is what the error is not. It is not a bias that can be determined once and subtracted. Its size depends on the transient the plasma has just undergone, and a single line ratio contains no record of that transient: the same measured number is returned by an equilibrated plasma and by a displaced one. No amount of care with the spectrometer separates them, and no correction factor tabulated on (T_e, n_e) can carry the information, because the information is not a function of (T_e, n_e) .

What can be tabulated is the plasma's *response* to being displaced, and that is a property of the operating point alone. Section 5.4 factorises the plateau error exactly into a sensitivity and a displacement, and Equation (5.8) splits the displacement again into a gain and the size of the transient, leaving

$$\varepsilon_{\text{plateau}} = \left| \exp(\bar{S} G \Delta \ln T_e) - 1 \right|, \quad \lim_{\Delta \ln T_e \rightarrow 0} \frac{\varepsilon_{\text{plateau}}}{|\Delta \ln T_e|} = |\bar{S} G|. \quad (7.1)$$

This repeats Eq. (5.9), here as a procedure in which the third factor is supplied for the transient of interest. It inherits from Chapter 5 the assumption that the ion population is a fixed reservoir and that the ground-state history $u(t)$ follows the local zero-dimensional collisional–radiative dynamics with no external influx, so that during the fast-stage plateau n_{1s} is frozen to within the closure residual; this leaves the reservoir no supply from outside the parcel and is the condition Section 6.2 cannot test.

Both structural factors are dimensionless. The mean sensitivity $\bar{S}$ is the average of $f_3 - f_4$ over the interval the reservoir traverses, taken with the orientation of Eq. (5.6), which makes it negative, and it says how strongly the observed ratio answers to the ground-state reservoir. The reservoir gain $G = \Delta \ln u / \Delta \ln T_e$, the secant along the CRE curve over the step taken, says how far equilibrium ionisation balance would move that reservoir for a given fractional change in temperature. Both $\bar{S}$ and G are negative everywhere on the grid, in both step directions, so $\bar{S}G > 0$ always; it is $\bar{S}G \Delta \ln T_e$ that is positive for heating and negative for cooling, and $\varepsilon_{\text{plateau}}$ is positive in both because Eq. (7.1) takes a modulus.

The procedure. For a CR matrix of one's own, at one's own operating point:

1. Solve the excited manifold twice, once with the ion reservoir alone and once with the ground state alone, to obtain the population coefficients a_p and c_p of Eq. (1.6) for the two emitting shells. These are the same two solves any steady-state CR table already performs. Form the switching points $x_m = \ln(c_m/a_m)$ and their separation $\Delta = x_4 - x_3$; the two ground-fed fractions are then the single logistic of Eq. (3.40), displaced by Δ .

2. Take the secant of $\ln u^{\text{CRE}}$ against $\ln T_e$ at fixed density over the step of interest to get G .
3. Insert the transient one actually cares about as $\Delta \ln T_e$. The reservoir excursion is $\Delta \ln u = G \Delta \ln T_e$, and that fixes the interval over which $\bar{S}$ is the mean of $f_3 - f_4$.

This is a forward criterion: it requires an estimate of the transient from outside the spectrum, since the ratio itself carries none.

Steps 1 and 2 belong to the operating point and are done once; step 3 belongs to the event. The sensitivity straddles the two, because the curves are the operating point's and the interval is the event's, and in the small-step limit it collapses onto the operating point as well: $\bar{S}$ tends to $f_3 - f_4$ evaluated at the undisturbed reservoir, which is why $|\bar{S}G|$ can be tabulated on (T_e, n_e) alone. That division is the whole reason for reporting structure instead of a percentage. Measured across the grid, $|G|$ moves by at most 6.74 % between step sizes of one, two and four grid intervals, over the 736 point-and-direction triples for which all three step sizes were computed, while $\varepsilon_{\text{plateau}}$ at the same points changes by up to a factor 8.44. A percentage quoted without its step is not a property of the plasma; G very nearly is. Both figures are direct reductions of `validation/reservoir_gain/reservoir_gain.csv` (2288 rows, stamped 10 September 2026) by `verify_reservoir_gain.py`, written to `validation/reservoir_gain/reservoir_gain_summary.csv` one row per filter, because the same quantity over a different filter is a different number. The spread of $|G|$ is 1.039 to 1.065 over 368 heating triples and 1.039 to 1.067 over 368 cooling triples; the spread of $\varepsilon_{\text{plateau}}$ over the same points is 2.27 to 8.44 heating and 2.33 to 4.70 cooling.

The scale, before anyone computes anything. Over the 982 heating steps of that file that have a resolvable plateau window, $|\bar{S}|$ runs from 0.0649 to 0.4822; over all 2288 rows $|G|$ runs from 2.640 to 14.519. Restricted to the range this thesis defends, $T_e \geq 2 \text{ eV}$, those become 0.0649 to 0.4552 and, over the heating pairs, 2.640 to 7.822, and the product $|\bar{S}G|$ runs from 0.362 to 3.535. The upper end is the useful number to remember. In the worst conditions defended here, a 1 % temperature excursion moves the line ratio by about 3.5 % over the heating pairs; the error in the inferred temperature is that displacement divided by the table's slope, median amplification 8.3 (Section 5.5.2). At the benchmark point the ratio moves by about 1.3 %.

7.2.1 Five things that follow, and one that does not

The response has a ceiling, and it is below unity. Because both ground-fed fractions are the same unit-width logistic curve displaced by Δ , their difference cannot exceed $\tanh(|\Delta|/4)$, and $\bar{S}$ inherits that bound (Eq. (5.11)). The $\tanh$ comes from the width of the logistic, and the width is unity because the excited populations are exactly affine in n_{1s} , which is a property of the linear system rather than of any rate coefficient. Every property of the 42-state network enters only through Δ . Since $|\text{d} \ln R / \text{d} \ln b_1| < 1$ everywhere, integrating over any excursion gives $|\Delta \ln R| \leq |\Delta \ln b_1|$, so a reservoir left stale by a factor of two cannot put the line ratio out by more than a factor of two. That protection holds for every shell pair, every (T_e, n_e) and every atomic dataset; the *height* of the ceiling does not, since it is $\tanh(|\Delta|/4)$ and Δ is built from rate coefficients. The ceiling is approached most closely where the error is largest: at [23, 3], one column into the crest, the measured sensitivity already sits at 97 % of it (Section 5.5). Substituting an independent excitation dataset moves the ceiling while leaving the headroom nearly

where it was, the median sensitivity falling from 80.4 % to 78.2 % of its own cap over the 338 heating pairs of the map (Sections 2.8.2 and 5.5). Substituting the ionisation data instead does the opposite, moving the sensitivity inside a cap that barely shifts. Running at 97 % of the cap is therefore a statement about the headroom beneath the ceiling and not about where the ceiling sits: it is not an artefact of this rate set, but neither does it bound where a refinement would put the cap itself, and the two refinements this thesis can perform move the two quantities separately.

The clock is τ_{slow} , and it is far too slow to average away. The temptation is to assume that an error introduced by a fast disturbance decays with the fast timescale. It does not. The error decays on τ_{slow} as very nearly a single exponential, the estimate built on that law matching the propagated average to within 2.2 % at 100 μs and 3.8 % at 506 μs , with a half-life of 12.95 μs at the benchmark and 1.10 ms at [15, 3] (Section 5.8), against the 2.28 ns on which the excited manifold responds. A 100 μs exposure does not integrate the error away. Counted on the time-averaged estimate of Eq. (5.16), which direct 43-state propagation places within 2.2 % of the true 100 μs average over the 448 pairs above 2 eV (`validation/trajectory_census/`), 45 of 448 point-and-direction pairs above 2 eV exceed 10 %, the worst being 17.5 % at $T_e = 2.02 \text{ eV}$ and $n_e = 1.93 \times 10^{13} / \text{cm}^3$. At the 506 μs extrapolated to the ITER pedestal [17, Sec. 5] the count falls to 24 and the worst case to 15.3 %, the 24 being a strict subset of the 45. Both are fixed-density counts, and Section 5.8.3 steps the density too: the count barely moves, at 10 to 11 % of the pairs available whichever way the density goes, while the worst case rises to 22.1 % when the density steps up with the temperature. Under the tested one-index density step, a factor 2.68, the fixed-density map understates the worst case rather than overstating it. At the densities of this model's defended range where a two-parameter inversion of this pair would be attempted, $n_e \geq 10^{14} / \text{cm}^3$, none of 108 pairs exceed 10 % and the worst is 7.2 %. The conditions to distrust are the tenuous and cool ones, and the error grows toward low density and low temperature together.

Timescale separation gives no warning. A modeller deciding whether to trust a QSS-based inversion reaches for $M = \tau_{\text{slow}} / \tau_{\text{relax}}$, which is the ratio that licenses eliminating the fast variables in the first place. It does not locate the error. Inside the defended range $T_e \geq 2 \text{ eV}$ the two maxima sit on the same temperature row $19.3\times$ apart in density: at the point of largest M the plateau error is 7.4 %, the 71st percentile of the 448 defended pairs and a factor 2.5 below the worst (Section 5.9). What survives once the operating point is controlled for is not an absence of information but information with the wrong sign: residualising both $\log M$ and $\log \varepsilon_{\text{plateau}}$ on saturated two-way fixed effects in T_e and n_e leaves a partial correlation of -0.617 , so the little that M carries about the error it carries backwards, and it is a second-order statement about the T_e - n_e interaction rather than a ranking. The reason is visible in Eq. (7.1). M compares two rates, and $\varepsilon_{\text{plateau}}$ measures a distance, the product of a sensitivity with a displacement, neither of which appears in M . A large M supports the reduction of the excited states to their quasi-steady state; that they do track it is verified directly by the propagated trajectories of Section 5.7. It says nothing about whether the reservoir is in the right place. Nor is this an artefact of how well M is known: replacing the excitation rate coefficients with an independent calculation moves it by 11 % while leaving its ordering across the grid intact, rank correlation 0.99993 (Section 5.9).

A second failure, and it is the same cancellation at its limit. The bound of Eq. (5.11) constrains the error in the ratio, and a diagnostician does not suffer an error in the ratio; she inverts it. The inversion divides by $\partial \ln R / \partial \ln T_e$, and that derivative passes through zero. Scanning temperature at $n_e = 5.18 \times 10^{13} / \text{cm}^3$, the fractional distance $\varepsilon_{\text{step}}$ between the two tabulated answers either side of the step falls to 4.9×10^{-5} near 6.9 eV (Section 5.5.1), which means the tabulated ratio is stationary in temperature there. The $n = 3/n = 4$ ratio is locally non-invertible for temperature at that point, so a bounded error in the observable produces an unbounded error in the inferred temperature. It afflicts a perfectly equilibrated plasma, and it is not unrelated to the transient error. The slope of the tabulated ratio in temperature is the residual of a cancellation between the manifold's own response P and the reservoir's SG , the two equal and opposite to a median 6.8% over the defended scope (Eq. (5.12)); the fold is where that residual passes through zero, and the same residual is the denominator of every inversion. For a practitioner it is the more dangerous of the two failures, and it was found only because $\varepsilon_{\text{step}}$ was computed and kept rather than discarded. Its locus across the rest of the grid has not been mapped, and Section 5.5.1 records what would close that.

The table misreads the transient by an order of magnitude. Pushed through the table at known density, the plateau ratio has no solution within the table's 1 to 10 eV domain for 166 of the 448 defended pairs, every one a cooling step; a further 44, all heating, reach an answer only by crossing the fold; and over the 238 pairs with a unique fold-free answer a 4.81% temperature step is read back with a median error of 48% and a worst of 144%, on the side the plasma came from at every pair (Section 5.5.2). The amplification, median 8.3, is the ratio of the reservoir term to the residual slope of the table, and is large because the two responses above nearly cancel. At the two highest densities it falls to 3.2 and the median error to 16%. These are plateau values at known density; neither the time averages of Section 5.8 nor the two-parameter inversion has been carried through the table.

What does not follow. None of the above licenses a claim about a particular machine. Every number in this section is conditional on the ground-state history $u(t)$ following the local zero-dimensional collisional-radiative dynamics with no external influx, with n_{1s} frozen during the fast-stage plateau, which is the conditional Section 6.2 states and cannot test from inside a zero-dimensional model.

7.3 What is new, at the smallest size that is defensible

Section 1.7 set out what was already known. What follows is what may now be claimed that could not be claimed then, item by item.

The two-channel decomposition is not new and is not claimed. That every excited population is exactly affine in the ground-state and ion densities, Eq. (1.6), is the population-coefficient decomposition of Bates, Kingston and McWhirter, written in the Saha-normalised form $b(p) = r_0(p) + r_1(p) b(1)$ by Fujimoto and McWhirter [6, 10, 11]. Chapter 3 says so where it uses it. Everything structural in this thesis is a consequence of that 1962 identity.

The logistic form is a rewriting, and should be called one. Writing the ground-fed fraction in the logarithmic reservoir variable turns Eq. (3.38) into the single logistic of Eq. (3.40), so the two shells' fractions are the same curve displaced. This is one line of algebra applied to a published formula, and the same function appears in chemical kinetics as the Michaelis–Menten saturation fraction, where a substrate concentration plays the part of the reservoir. The defensible statement is that I did not find it stated in this form in the collisional–radiative literature surveyed for this thesis. It is not a new function and it is not a new result about hydrogen.

The bound is old, it is old in several fields at once, and only the assembly is claimed. That $\max |f_3 - f_4| = \tanh(|\Delta|/4) < 1$ follows in a few steps from two displaced logistics. A literature search conducted for this thesis found each of its parts already in print, in four separate literatures, the earliest by more than a century.

The *value* is Yule's coefficient of colligation [47]. Two unit-width logistic curves separated by Δ have constant odds ratio $\omega = e^\Delta$, and Yule's $Y = (\sqrt{\omega} - 1)/(\sqrt{\omega} + 1)$ is identically $\tanh(\frac{1}{4} \ln \omega)$. The two agree to machine precision at every Δ tested. The same number is the maximum risk difference at fixed odds ratio, attained at baseline $1/(1 + \sqrt{\omega})$.

The *location* of the maximum is Samejima's [48]: the difference of two shifted logistic ogives is the graded-response operating characteristic, and she states it is unimodal and symmetric about the midpoint. She does not give the height.

The *general theorem* is Anderson's [49]: for a symmetric unimodal density, $\sup_x [F(x) - F(x - \Delta)] = 2F(\Delta/2) - 1$. This claim is its logistic instance, and the object itself is the Lévy concentration function.

What the search did *not* find is the bound in the collisional–radiative literature, and one negative is worth recording because it was the largest remaining publication risk. Fujimoto's *Plasma Spectroscopy* Chapter 4 [11] is where the population coefficients $r_0(p)$ and $r_1(p)$ that f_p is built from are introduced, tabulated and analysed at length. **It contains no tanh, no logistic, and no bound on the difference of two population coefficients**; searching the printed chapter for each of those returns nothing. The normalisation it does establish, $r_0(p) + r_1(p) \rightarrow 1$ at high density, is the statement that the two channels exhaust the population, which is what makes f_p a fraction in the first place and is assumed rather than derived here. The bound is not implicit in that chapter.

The corollary $|\mathrm{d} \ln R / \mathrm{d} \ln b_1| < 1$ is likewise not new. That $f_p \in (0, 1)$ is the standard statement that substrate elasticities lie between zero and the Hill coefficient [50], and a general bound on the logarithmic sensitivity of a *ratio* of steady-state observables, reducing to unity when a single state carries the perturbation, is Owen and Horowitz's [51]. The companion identity $\int (f_3 - f_4) \mathrm{d} \ln b_1 = \Delta$ is Raju's unsigned area for equal-discrimination curves [52].

The qualitative content has an empirical face in this field already. Verhaegh and co-workers report that a unique solution for the recombining fraction from a measured Balmer ratio exists only for F_{rec} between about 0.15 and 0.85 [15], which is the saturation of f_p observed directly rather than derived.

What is left, and what this thesis claims, is the *assembly*: that these are one object; that it bounds the diagnostic error of a hydrogen Balmer ratio with respect to its ground-state reservoir, below unity for every positive two-channel dataset and by $\tanh(|\Delta|/4)$ for a given one; and that the bound is near-tight exactly where the error is largest. The claim is deliberately written so that finding an older statement of any single part costs it

nothing, because that has now happened four times.

The corollary requires $n_p = A_p + B_p u$ with A_p and B_p independent of u . Where that fails, effective Hill coefficients well above unity are reachable, and two Michaelian enzymes competing for a shared substrate is the standard counterexample. It does not arise here, because u is an external reservoir rather than a pool the levels compete for.

Greenland stated the negative result in general form, twenty-five years ago. That the validity criteria for a reduced CR model are conditions on eigenvectors rather than eigenvalues, and are not related to the equilibrium timescales, is Greenland’s [20, 21]; both statements are quoted in Section 1.7. This work does not improve on either. It quantifies them for one diagnostic and maps where the criterion misleads (Section 5.9; the separation of the two maxima and the sign of the residual association are in Section 7.2), which is a smaller contribution than stating the criterion.

Sawada and Fujimoto asked the adjacent question and answered it first. Section 1.7 set out their result for the same system, atomic hydrogen under an abrupt change of conditions: which level controls the approach to steady state [19]. Their answer is a clock internal to the excited manifold. The question here is what a diagnostic inversion costs while the reservoir underneath that manifold has not settled, which is a different quantity governed by a clock four orders of magnitude slower, and which can be large at the very moment the excited-level solution is exact.

The separation between the two papers is structural rather than a matter of emphasis, and the paper states it itself. Setting out their method, they write that “under the ionizing plasma condition $n(1)$ is kept constant and n_z is set equal to zero, and vice versa under the recombining plasma condition” [19, p. 5566], having already assumed that the electron density, the ion density and the electron temperature “are given” [19, p. 5565]. **The reservoir is a fixed parameter in their calculation.** It cannot be stale, because it never moves, and the question this thesis asks cannot arise inside their framework.

That is not one paper’s choice. The same calculation is the worked example of Appendix 4B of Fujimoto’s *Plasma Spectroscopy*, whose figures are captioned as quoted from that paper, and the appendix sets it up the same way: neutral atoms are immersed in a plasma at $t = 0$ and, in its words, “if $n(1)$ is kept constant, after the initial transient this system should reduce to the high-temperature ionizing plasma” treated earlier in the chapter [11, App. 4B]. **The frozen reservoir is the setting in which the textbook validity condition is derived**, which makes the gap this thesis fills a gap in the tradition rather than an oversight in one paper.

The two-channel decomposition this thesis is built on is their Eq. (3), $n(p) = R_0(p) n_z n_e + R_1(p) n(1) n_e$, and the ground-state equation whose eigenvalue is the slow clock of Chapter 3 is their Eq. (4), $dn(1)/dt = -S_{CR} n(1) n_e + \alpha_{CR} n_z n_e$. They write Eq. (4) and hold $n(1)$ fixed while the excited levels settle; Fujimoto’s Appendix 4B then treats a time-varying $n(1)$ only as a validity condition on those levels and defers its dynamics to Appendix 5B. Neither asks what a diagnostic inversion costs while $n(1)$ is still moving. That is the gap, and it is now a matter of record rather than an inference from their emphasis.

Their validity criterion for a temporal change in T_e , their Eq. (23), $|(\tau_{\text{res}}/T) E(1, p_G)/kT_e| \ll 1$ with T the time constant of the change, bounds whether the excited levels keep up. No companion condition on the reservoir appears anywhere in the paper, and none is needed there, because the reservoir is held fixed.

One apparent conflict is worth resolving before an examiner raises it. For a true step $T \rightarrow 0$, so Eq. (23) is violated, while Section 5.1.2 reports $\varepsilon^{\text{QSS}} = 8.66 \times 10^{-6}$ for exactly such a step. The two statements are about different intervals. Equation (23) governs the settling itself, and during the first few nanoseconds after a step the excited-state closure genuinely does fail. ε^{QSS} is evaluated on the plateau, after the manifold has settled, where it holds to a part in 10^5 to 10^8 . Read that way the two results are complementary rather than merely compatible, and they partition the transient between them.

Their worked case, $n_e = 10^{12}/\text{cm}^3$ and $T_e = 10 \text{ eV}$, sits on this thesis's own grid, which is what makes the comparison a comparison rather than an analogy.

Park proposed the nonequilibrium Balmer thermometer, and that idea is his.

In 1972 Park put the H_α/H_β ratio forward as an electron temperature indicator for a plasma out of Saha equilibrium, assumed quasi-steady state within the excited levels, wrote the normalised population in the same two-coefficient form with the normalised ground-state population as the third argument, and tabulated the conversion from the apparent excitation temperature to the true one [24]. Nothing here disputes that and nothing here claims it. His tables are computed in the two limits he called terminal nonequilibrium, defined as the conditions in which a further departure from equilibrium does not change the emission ratio; that definition is the statement that the sensitivity of Section 3.5.4 vanishes, which it does when both shells are fed in the same proportion. The interval between his two limits is where that sensitivity is largest, and it is where the present grid lies: the difference in ground-fed fraction does not fall below 0.046 at any of the 400 points. What is added here is the size of the error an equilibrium table incurs in that interval, its factorisation into a susceptibility and a reservoir displacement, the clock on which it decays, and the finding that timescale separation does not rank it.

What is claimable. The cost of the ionisation-balance closure hidden inside a two-parameter Balmer table, measured here for the temperature inferred at known density, is now a computable number at any operating point of a zero-dimensional model with a closed reservoir, expressed through a gain G that is very nearly a property of the operating point and a sensitivity $\bar{S}$ that depends on the step, by 22 to 30 % across a fourfold change of step at the benchmark, together with a ceiling on it whose form no atomic dataset changes and a map of where in (T_e, n_e) the coefficients are large. That is the contribution, and it is smaller than the question Chapter 1 asked.

7.4 What this work cannot claim

These are carried forward from Chapter 6 without re-argument, and they bound every statement in Section 7.2.

Everything quantitative is conditional on a frozen reservoir, and the condition is measured to fail. If neutral transport and recycling renew n_{1s} faster than the local atomic physics does, the ground state is not the stale closed-parcel one and the reservoir gain G is not the gain the plasma has. Comparing a charge-exchange-limited neutral escape time with τ_{slow} pair by pair, not one of the 45 breakdown pairs above 2 eV satisfies $\tau_{\text{esc}} > \tau_{\text{slow}}$, at either 10 cm or 20 cm to the boundary, while 27 to 36 % of the other 635 analysed pairs do (Section 6.2.4). The selection is structural: a pair enters

the census because τ_{slow} is long, and a long τ_{slow} is exactly the condition under which transport rather than ionisation sets n_{1s} . Section 6.2.5 opens the reservoir to an escape sink and a recycling source and finds that a fast-exchanging reservoir with a steady source freezes the closed-parcel error rather than removing it, and that a rising source enlarges it, the true error exceeding the closed-parcel plateau at 165 of the 180 combinations tested. The mechanism, a ratio that depends on u with u away from its equilibrium value, operates whatever sets u ; what transport removes is this model's prediction of $u(t)$. The magnitudes of Chapter 5 are therefore reference values conditional on the closed-parcel reservoir history, neither bounds on the divertor error nor estimates of it. The two-channel split, the unit-width logistic form of $f_p(u)$ and the ceiling $\tanh(|\Delta|/4)$, being properties of $\mathbf{L}(T_e, n_e)$ at fixed conditions, are untouched by it; where on that curve the plasma sits, and so the sampled $\bar{S}$, G and the ridge location, moves with whatever sets u (Sections 6.2 and 6.4). The reservoir itself has one external check already, against ADAS: u^{CRE} sits within a median factor 1.18 of ACD/SCD and $\Delta \ln u$ per step within 5.8% (Section 4.10).

It is not a divertor result below two electron-volts, and the restriction is not a hedge. Holding n_e fixed while the ground-state population changes is self-consistent only where the plasma is largely ionised. Above 2 eV the change in n_e that conservation of nuclei implies across one temperature step, inferred from the two fixed- n_e solutions in the heating direction, never exceeds 5.4×10^{-3} ; solved self-consistently over both step directions the correction reaches 7.98×10^{-3} , at a cooling step at $T_e = 2.024$ eV and $n_e = 10^{12}/\text{cm}^3$. Over the 392 heating one-step operators the fixed-density inference puts $|\delta n_e/n_e| > 10\%$ at 68 and above 100% at 36; solved self-consistently the counts are 67 and 7 (Section 5.10.1). $T_e \geq 2$ eV is the working domain the two measured boundaries support, the quasineutrality contour at 1.41 to 1.50 eV and the Lyman optical-depth contour at 1.01 to 2.36 eV across slab thicknesses (Section 5.10), and the cold-edge numbers are values of a model outside its own domain.

The crest is not a detachment criterion. The density crest sits a factor 5.2 below the detached band, and at citable detached densities the 100 μs time-averaged estimate never reaches 10% (Section 6.5). The crest is a property of the atomic system and needs no divertor geometry to produce it.

There are no molecular channels, in a gas that is largely molecular. A third supply feeding $n = 3$ preferentially would change $f_3 - f_4$. Imposing molecular emissivity fractions of 0.60 at $n = 3$ and 0.30 at $n = 4$ on the model's own operator reduces the cold-corner plateau error by a factor of 6.1, and the upper end of the reported ranges, (0.70, 0.45), by a factor of 9.4 (validation/molecular_channel/). Those figures apply to the cold corner, which lies outside the defended range in any case; at the crest the same two targets give 5.4 and 8.1. The reduction is larger than a dilution linear in $f_3 - f_4$ suggests, because $\varepsilon_{\text{plateau}}$ is exponential in the sensitivity, and it is not monotone: the diluted difference passes through zero near $\phi_3 = 0.74$, past which the error grows again with the sensitivity reversed. The estimate is built in Section 6.3 from the emissivity fractions cited there, driven through this thesis's atomic operator rather than computed from a matrix containing molecules, and it inherits every uncertainty in those fractions. It is a sensitivity scenario and not a bound: the imposed source holds $u(t)$ and $n_{\text{ion}}(t)$ fixed, which a molecular reservoir would not, and Section 7.5 states why an independent H_2 reservoir changes the form of the result rather than its coefficients.

The atomic-data figures are substitution sensitivities, not an uncertainty budget. Two of the three input channels have been swapped for independent calculations and carried through to the observable: the excitation coefficients for R-matrix values, and the ionisation coefficients, with the three-body recombination that detailed balance builds from them, for Lotz. Together they move $\varepsilon_{\text{plateau}}$ by 15.5 % over the 448 defended pairs and at most 25.4 %, and M by 11 % (Sections 2.8.2 and 5.5). Those figures are the difference between two calculations, not an uncertainty: they say how far the answer moves between the datasets that happen to exist, and a third dataset could lie outside the interval the two define. Three gaps remain in the bound. Excitation transitions reaching $n \geq 6$ have no independent counterpart. Ionisation above $n = 9$ is Lotz in both the model and the substitution, so those shells are compared with nothing; scanned instead, an eightfold reduction of that block and of the three-body recombination built from it moves $\varepsilon_{\text{plateau}}$ by 0.7 % at the median and 3.9 % at most over the defended pairs (Section 2.8.2), so that gap is one of provenance rather than of leverage. And radiative recombination has one source only, tested at its own quoted accuracy rather than against a rival, which is a weaker statement even though the 1.1 % it produces is small. The absolute normalisation of the ground-fed coefficient remains an open disagreement, a factor 8.3 below Fujimoto’s tabulation at $p = 3$ (Section 6.7).

The density maximum is a range, not a density. It lies between 7×10^{12} and $5 \times 10^{13}/\text{cm}^3$, resolved to about a factor of 3, because every density row sits above 90 % of its own maximum across roughly a decade (Section 5.11). Quoting a single density for it overstates what the grid resolves. That factor of 3 is a whole-grid figure and its upper end is set by the three coldest rows, the only ones whose maximum falls on $5.18 \times 10^{13}/\text{cm}^3$; inside the defended range every row maximises on $1.93 \times 10^{13}/\text{cm}^3$ and the subgrid vertex spans 1.45 to $1.95 \times 10^{13}/\text{cm}^3$, a factor 1.35 (`validation/crest_subgrid/`). The range is wide because of the rows this thesis excludes, not because the crest is poorly located where it is defended. It is also the (3, 4) shell pair’s maximum and not the Balmer series’ (Section 5.2).

7.5 What would settle what is open

Four calculations would change the standing of a result in this thesis. The first two are bounded work on the model as it stands; the other two are different projects.

A density grid as fine as the temperature grid. The joint map of Section 5.8.3 crosses temperature steps with density steps of one grid index, and one density index is a factor 2.683, that is +168 %, against +4.81 % for one temperature index. The reservoir is roughly fifty times less sensitive to density than to temperature in logarithmic terms, so the two contributions end up comparable in that sweep only because the density grid is that coarse. The coarseness belongs to the tabulation and not to ELMs, which is why Table 5.7 is quoted per density step rather than as one number. Regenerating $\mathbf{L}$ on a density grid spaced like the temperature grid would let the joint census be quoted at a realistic ELM density excursion, and it is the cheapest calculation that would move a number this thesis reports, though it refines the closed-parcel census rather than turning it into a statement about an ELM. It costs the same as the map that already exists.

An ionisation dataset that reaches the top of the ladder. The substitutions of Sections 2.8.2 and 5.5 bound $\varepsilon_{\text{plateau}}$ against independent data wherever independent data exist, which is $n \leq 5$ for excitation and $n \leq 9$ for ionisation. Above $n = 9$ the model and the substitution both use Lotz, so the six shells whose feed is essentially all three-body are the shells with no independent check at all, and they are the shells where Section 2.8 records Lotz exceeding CCC by a factor of four to eight, and where Chapter 6 finds the terminal shell over-populated by 4.9–6.2 times. Extending a close-coupling ionisation calculation to $n = 10$ –15, or bounding those coefficients some other way, would both close the last gap in the atomic-data bound and test whether those two observations are the same one. As it stands the link between them remains a hypothesis.

A two-region or transport-coupled calculation. The single conditional on which all of Section 7.2 rests is that the reservoir history $u(t)$ follows the local zero-dimensional collisional–radiative dynamics with no external influx, so that n_{1s} is frozen during the fast-stage plateau. Deciding it requires n_{1s} to be supplied by a recycling influx from an edge transport calculation rather than solved for from $\mathbf{L}\mathbf{n} + \mathbf{S}n_{\text{ion}} = 0$, which is a different class of calculation and not a refinement of this one. It is the item that would convert the result from a statement about a class of CR model into a statement about a machine.

Molecular channels, which may change the form and not just the coefficients. Adding molecules as states inside $\mathbf{L}$, fed from the same two reservoirs, leaves the structural results untouched, and Section 6.3 says so. Adding molecular hydrogen as a *third independent reservoir* does not. The population balance then reads $n_p = a_p n_{1s} + c_p n_{\text{ion}} + m_p n_{\text{H}_2}$, with n_{H_2} the molecular density in $1/\text{cm}^3$ and m_p the dimensionless number of atoms in level p supported per molecule. The observable then depends on two independent reservoir ratios rather than one, and the ground-fed fraction is no longer a single logistic in a single variable. The bound of Eq. (5.11) survives only at a fixed molecular ratio, with Δ replaced by a v -dependent separation, and the error then depends on two reservoir histories rather than one. Which of the two descriptions a real detaching divertor requires is a physics question this work does not answer.

The physicist at the console keeps her spectrometer, her ratio and her table. What she gains is a test she can apply before she trusts the answer: evaluate $\bar{S}$ and G at the operating point she is looking at, multiply by the fractional temperature excursion her plasma has just undergone, and compare the result with the precision she needs. If the plasma has been quiet for several times τ_{slow} , the table is telling her about the plasma. If it has not, the table is telling her about a plasma that no longer exists, and Eq. (7.1) says by how much.

Appendix A

State Ordering

Every index that appears in this thesis – $\mathbf{L}_{ij}$, a shell sum, a figure axis labelled by state, the benchmark $\mathbf{L}[23, 5]$ – refers to a position in one fixed 43-element ordering. This appendix records that ordering so any index in the text can be resolved without reading the code, states the one assumption the ordering embeds, and names the file that produces it.

A.1 What the state space is, and where it is cut

The model tracks 43 states of neutral hydrogen: the 36 orbital-angular-momentum resolved levels $n\ell$ for $n \leq 8$, and seven ℓ -bundled shells for $9 \leq n \leq 15$. The proton is not a forty-fourth state; it is held as an external reservoir of density n_{ion} , which is the open-system boundary condition discussed in Section 3.5.

The cut at $n = 8$ is placed where it is because the observable is built from $n = 3$ and $n = 4$. Both are fully ℓ -resolved and both sit well below the resolved/bundled boundary: $n = 4$ five shells below it and eleven below the truncation at $n_{\text{max}} = 15$, $n = 3$ six and twelve. Resolving ℓ costs states quadratically – n^2 sublevels per shell – and buys nothing above the shells the diagnostic reads, provided the bundled shells are ℓ -statistical, which Section 6.8.3 tests rather than assumes.

Degeneracies are hydrogenic throughout: $g_{n\ell} = 2(2\ell + 1)$ for a resolved sublevel and $g_n = 2n^2$ for a shell, the factor 2 being electron spin. Summing the resolved sublevels of each shell returns $2n^2$ exactly for every n from 1 to 8, which is as far as the resolved levels go; each bundled state carries $2n^2$ by construction. The degeneracies of all 43 states sum to $\sum_{n=1}^{15} 2n^2 = 2480$. Ionisation potentials are hydrogenic, $I_p = \chi_{\text{H}}/p^2$ with $\chi_{\text{H}} = 13.6058 \text{ eV}$; the product $I_p p^2$ runs from 13.6057995 to 13.6058006 across the table, constant to seven significant figures.

A.2 The assumption the ordering embeds

Bundling shells $9 \leq n \leq 15$ into single states asserts that the population inside each of those shells is distributed statistically over ℓ , $n_{n\ell} = (g_{n\ell}/g_n) n_n$. Nothing in the ordering enforces this; it is a modelling choice made when the rate coefficients were summed, and it is the reason no ℓ -mixing rate appears among states 36–42. Whether it is justified is a physical question, tested in Section 6.8.3 by comparing the intra-shell ℓ -mixing rate against the processes that remove population from the shell.

Table A.1: The state ordering used throughout this thesis, as written by `assemble_K_exc.py` to `data/processed/collisions/K_exc_full/state_index.csv`. Indices are zero-based, matching the arrays. States 0–35 are ℓ -resolved; states 36–42 are ℓ -bundled shells, for which ℓ is recorded as -1 and $g = 2n^2$.

idx	label	n	ℓ	g	idx	label	n	ℓ	g
0	1S	1	0	2	22	7P	7	1	6
1	2S	2	0	2	23	7D	7	2	10
2	2P	2	1	6	24	7F	7	3	14
3	3S	3	0	2	25	7G	7	4	18
4	3P	3	1	6	26	7H	7	5	22
5	3D	3	2	10	27	7I	7	6	26
6	4S	4	0	2	28	8S	8	0	2
7	4P	4	1	6	29	8P	8	1	6
8	4D	4	2	10	30	8D	8	2	10
9	4F	4	3	14	31	8F	8	3	14
10	5S	5	0	2	32	8G	8	4	18
11	5P	5	1	6	33	8H	8	5	22
12	5D	5	2	10	34	8I	8	6	26
13	5F	5	3	14	35	8J	8	7	30
14	5G	5	4	18					
15	6S	6	0	2		<i>bundled shells</i>			
16	6P	6	1	6	36	n9	9	–	162
17	6D	6	2	10	37	n10	10	–	200
18	6F	6	3	14	38	n11	11	–	242
19	6G	6	4	18	39	n12	12	–	288
20	6H	6	5	22	40	n13	13	–	338
21	7S	7	0	2	41	n14	14	–	392
					42	n15	15	–	450

This assumption has a consequence that matters for Chapter 4: the bundled sub-block of $\mathbf{L}$ is, by construction, the statistical- ℓ limit. A bundled operator cannot represent a non-statistical ℓ distribution, so any comparison against a published tabulation that is itself bundled in ℓ is comparing two models that agree on this point by construction rather than by calculation.

A.3 Two hydrogen ionisation energies

Two slightly different values of the hydrogen ionisation energy appear in the rate pipeline, and the difference is physical rather than an inconsistency:

- 13.605 693 122 990 eV, the Rydberg energy $R_\infty hc$, is the threshold used by the collision datasets (CCC, TICS) and therefore sets every collisional threshold energy.
- 13.598 434 599 702 eV, the measured ionisation energy of the hydrogen atom including reduced-mass and relativistic corrections, sets the radiative transition energies.

Both are recorded in `L_meta.csv` as `IH_collision_eV` and `IH_radiative_eV` (`assemble_cr_matrix.py` lines 76–77). They differ by 0.053 %. The hazard this creates is that a detailed-balance relation evaluated with one energy on one side and the other on the other side will fail at the few-tenths-of-a-percent level for reasons that look like a data error and are not. Both balance relations in this model use a single ladder throughout and close far below that level.

A.4 Provenance

The ordering is generated by `src/rates/assemble_K_exc.py` (line 361) and written to `data/processed/collisions/K_exc_full/state_index.csv`. It is read by `src/validation/cr_cont` which every verification script in this thesis imports rather than redeclaring the ordering locally; the number of states is taken as `L_grid.shape[-1]` and the ground index as `arg min_p n_p`, so neither is hard-coded anywhere in the analysis.

As a structural check on the assembled operator, the columns of $\mathbf{L}$ sum to $-n_e K_{\text{ion}}(p)$ – ionisation is the only process that removes population from the 43-state manifold – and this holds to a maximum relative deviation of 1.8×10^{-12} across the 43 states at $T_e = 10$ eV, $n_e = 1 \times 10^{12}/\text{cm}^3$, and 1.1×10^{-12} at $n_e = 1 \times 10^{15}/\text{cm}^3$ (`verify_fujimoto_bundle.py`, gate 2).

Appendix B

Atomic Data Sources

Chapter 2 explains what each rate in $\mathbf{L}$ represents and why it matters. This appendix records where each one came from, which script converted it, and what it is stored in, so that any coefficient in the matrix can be traced to a published source without reading the code. It also states which coefficients are *not* independent measurements, because that is the single most important thing to know when reading a validation table: a check between two quantities derived from one another cannot fail.

B.1 The stack, end to end

B.2 Which checks can fail, and which cannot

Two rows of Table B.1 are marked *not independent*, and the consequence is worth stating plainly rather than leaving for a reader to infer. De-excitation is obtained from excitation by detailed balance, and three-body recombination from ionisation by detailed balance. A numerical test that those relations hold therefore tests the arithmetic that imposed them and nothing else. Chapter 4 grades Gate A *tautological* for exactly this reason: it passes at better than 0.01 %, and it would pass on corrupted input.

The checks that can fail are those between quantities with independent origins: the excitation rates against R-matrix-with-pseudostates calculations, the raw cross sections against microscopic reversibility, the assembled operator's effective coefficients against ADAS, and the population coefficients against Fujimoto's tabulation. The first two are graded *severe* in Table 4.4 and the last two are among its four external comparisons; of those four, one fails and one is an open disagreement.

B.3 The seam at $n = 10$

The excitation data change character between $n_{\text{upper}} = 10$ and $n_{\text{upper}} = 11$: quantum close-coupling below, a semi-empirical formula above. The seam is a property of the available data and not a modelling choice. Its consequence for this thesis is confined to the bundled shells, which the observable does not read: $f_3 - f_4$ is built from $n = 3$ and $n = 4$, seven and six shells below the seam.

Table B.1: Every coefficient entering $\mathbf{L}$ or $\mathbf{S}$, with its source, the script that produced it, and the file it is stored in. Paths are relative to `data/processed/`.

quantity	source	script	stored in
Excitation, $n_{\text{upper}} \leq 10$	Convergent close-coupling cross sections [27, 28]	<code>compute_K_CCC.py</code>	<code>collisions/ccc/K_CCC_exc_table.npy</code>
Excitation, $n = 11\text{--}15$	Vriens & Smeets semi-empirical formula [29], in the form implemented by [30]	<code>compute_K_VS.py</code>	<code>collisions/vs/K_VS_exc_table.npy</code>
De-excitation	<i>not independent</i> : detailed balance applied to the excitation rates above	<code>compute_K_CCC.py</code> , <code>compute_K_VS.py</code>	<code>K_deexc_full.npy</code>
Ionisation	CCC total ionisation cross sections, Maxwell-averaged, for $n \leq 9$; Lotz’s formula [31] for the bundled $n = 10\text{--}15$, which overestimates CCC by a factor 4–8	<code>compute_K_TICS.py</code> , <code>ionization_rates.py</code>	<code>collisions/tics/ion_final.npy</code>
Radiative recombination	Johnson’s three-term Gaunt-factor formula [32], after [33]	<code>recombination_rates.py</code>	<code>recombination/alpha_RR_resolved.npy</code>
Three-body recombination	<i>not independent</i> : detailed balance applied to the ionisation rates above	<code>recombination_rates.py</code>	<code>recombination/alpha_3BR_resolved.npy</code>
Einstein A coefficients	Hoang-Binh’s exact hydrogenic radial integrals [36], distributed as ADUU [35]. The producing script names a 1993 paper, for which this project holds no reference; the discrepancy is unresolved	<code>radiative_rates.py</code>	<code>Radiative/A_resolved.npy</code>
Proton-impact ℓ -mixing	PSM20 Debye-cutoff formula [42]	<code>compute_lmix.py</code>	<code>lmix/K_lmix.npy</code>
Assembled operator	the eight rows above	<code>assemble_cr_matrix.py</code>	<code>cpymatrix/L_grid.npy</code> , <code>S_grid.npy</code> , <code>L_meta.csv</code>
ADAS SCD96, ACD96	<i>comparison only; does not enter $\mathbf{L}$</i> [8, 9]	<code>parser_adasf11.py</code> , <code>prepare_adas.py</code>	<code>adas/SCD96_interpolated.csv</code>

B.4 Two constants, and why both appear

The hydrogen ionisation energy enters the pipeline twice, with two different values, and this is deliberate. Collisional thresholds use the Rydberg energy $R_\infty hc = 13.605\,693\,122\,990$ eV, because that is the energy scale on which the CCC and TICS cross sections are tabulated. Radiative transition energies use the measured ionisation energy of the hydrogen atom, $13.598\,434\,599\,702$ eV, which includes the reduced-mass and relativistic corrections. Both are recorded in `L_meta.csv` as `IH_collision_eV` and `IH_radiative_eV`. They differ by 0.053%. Appendix A states the hazard this creates for detailed-balance tests, and Chapter 2 the physics of why both are correct in their own place.

Appendix C

Numerical Methods and Convergence

This appendix states how the linear algebra and the time integration in this thesis are actually performed, and collects the convergence and conditioning evidence that has a producing script and a stamped artifact. Results quoted in the body that lack one are identified as such here rather than repeated.

C.1 How the systems are solved

Steady states. Every quasi-steady and collisional–radiative equilibrium population in this thesis is obtained by a direct dense solve of $\mathbf{L}_{EE} \mathbf{n}_E = -\mathbf{S}_E$ or $\mathbf{L}_{EE} \mathbf{n}_E = -\mathbf{L}_{E,g} n_{1s}$, using LAPACK’s general solver through `numpy.linalg.solve`. No iterative method and no preconditioner is used, and none is needed: the excited-state block is 42×42 . The two feed channels are solved separately and their sum checked against the full solve rather than assumed (Section 4.5.1).

Timescales. The relaxation spectrum is the eigenvalue spectrum of $\mathbf{L}$, computed with `numpy.linalg.eigvals`, sorted descending on the real part. τ_{slow} and τ_{relax} are the reciprocals of the two least negative eigenvalues, and $M = \tau_{\text{slow}}/\tau_{\text{relax}}$.

The selection of “negative” eigenvalues is the one step in this procedure where a magnitude threshold does damage. A filter that discards every eigenvalue with $\lambda > -1 \text{ s}^{-1}$, on the reasoning that such values are numerical noise near zero, is wrong: at the cold corner of the grid the true slowest mode has $|\lambda| \approx 0.0149/\text{s}$, corresponding to $\tau_{\text{slow}} = 67.2 \text{ s}$, and that filter discards it and returns the second mode in its place, wrong by nine orders of magnitude. The filter used is $\lambda < 0$, which admits every decaying mode. Chapter 2 reports the error and its effect; the point here is the method, which is that no magnitude threshold is applied to an eigenvalue whose magnitude is the quantity being measured.

Time integration. Transient solutions are obtained two ways, and the two are compared rather than trusted separately. The spectral route diagonalises $\mathbf{L}$ and propagates the deviation from the post-step steady state analytically. The direct route integrates $\dot{\mathbf{n}} = \mathbf{L}\mathbf{n} + \mathbf{S}n_{\text{ion}}$ with the implicit Radau method at relative tolerance 10^{-10} and absolute tolerance $10^{-12} \max |\mathbf{n}_{\text{old}}|$, chosen because the system is stiff: the spectrum spans between eight and fifteen orders of magnitude in $|\lambda|$ depending on the grid point.

C.2 Conditioning

Non-normality matters here more than conditioning does. The operator’s numerical abscissa $\mu(\mathbf{L})$, which bounds the initial growth rate of $\|\mathbf{n}\|$, is positive everywhere on the grid while the spectral abscissa is negative everywhere: the solution can grow before it decays, even though every mode decays. Measured over all 400 grid points (`verify_operator_conditioning.py`, `validation/operator_conditioning/`):

Table C.1: Conditioning of the excited-state block and the two abscissae, over all 400 grid points.

quantity	minimum	maximum	at the benchmark
$\kappa(\mathbf{L}_{FF})$	1.48×10^3	1.74×10^5	4.20×10^4
$\mu(\mathbf{L})$ 1/s	3.40×10^8	1.51×10^{12}	1.28×10^{11}
spectral abscissa 1/s	-1.33×10^7	-1.49×10^{-2}	-4.40×10^4

A condition number of 1.7×10^5 at its worst leaves roughly eleven significant figures in double precision, against results quoted here to three or four. Conditioning is not a threat to any number in this thesis. The positive numerical abscissa is not a numerical problem either. It is a statement about transient growth of the population vector in the Euclidean norm of the coordinates used, and would change under a rescaling of those coordinates; Chapter 3 uses it in that sense, and the direct trajectory tests of Chapter 4 are the coordinate-free evidence.

C.3 Convergence tests with a stamped artifact

Table C.2: Numerical checks that have a producing script and a stamped artifact under `validation/`.

check	result	artifact
Two-channel superposition: $\mathbf{n}^{(0)} + \mathbf{n}^{(1)}$ against the full solve, 784 pairs	3.08×10^{-14}	<code>divertor_map/</code>
Column-sum closure: $\sum_i \mathbf{L}_{ij} = -n_e K_{\text{ion}}(j)$, all 43 states, at two operators	1.75×10^{-12}	<code>fujimoto_bundle/</code>
Integrator tolerance at two grid points, <code>rtol</code> = 10^{-9} against 10^{-12}	1.20×10^{-8}	<code>plateau_slowmode/</code>
Direct integration against matrix exponential, 20 log-spaced times	1.14×10^{-9} at the benchmark, worst 3.24×10^{-8} in $\ \Delta \mathbf{n}\ /\ \mathbf{n}\ $	<code>ramp_plateau/</code>
Statistical- ℓ projector identity <code>PR = I</code>	0 exactly	<code>fujimoto_bundle/</code>

Each of these is a check on the numerics, not on the physics. They establish that the

answers reported are the answers the equations give; whether the equations are the right ones is the subject of Chapter 4 and Chapter 6.

C.4 State-space convergence, and what is not established

Truncation of the level ladder was tested downward, by rebuilding $\mathbf{L}$ at successively lower $n_{\max}$ with the loss channels of the removed states returned to the diagonal of the survivors, and the result is reported in Section 4.6.2. Two limitations of that result belong here rather than in the body.

First, it is a convergence test of $f_3 - f_4$ and, under the same truncation, of u^{CRE} , Δ and τ_{slow} , not of the state space in general. Quantities that depend on the Rydberg population directly, among them the effective ionisation coefficient, are not covered by it. Neither the reservoir gain G , the secant sensitivity $\bar{S}$ nor $\varepsilon_{\text{plateau}}$ was scanned directly; each is built from quantities the scan converges, $f_3 - f_4$ at u^{CRE} and $\Delta \ln u$, which moves by at most 2.3×10^{-5} between $n_{\max} = 14$ and 15 (`validation/nmax_downward_scan/`).

Second, the numbers reported in Section 4.6.2 come from `verify_nmax_downward_scan.py` (`validation/nmax_downward_scan/`), which gates on reproducing the stamped a_p , c_p , u^{CRE} and Δ of `validation/molecular_channel/` at all 400 points before any state is removed, and on the column sums of every truncated matrix equalling those of the full one. A historical project value with no producing script, not used as evidence here and superseded by `verify_nmax_downward_scan.py`, gives these quantities with the opposite sign at the benchmark and the cold corner; the script's values are the ones reported. The trap the scan avoids is recorded in Section 4.6.2: dropping states without returning their loss channels to the diagonal moves $f_3 - f_4$ by 5 to 26% at the three named points.

Convergence *upward*, at $n_{\max} = 18$ or 20, is not attempted and cannot be, because the atomic data stop at $n = 15$. That is the second of the five open questions of Section 6.10.

Bibliography

- [1] R. A. Pitts, X. Bonnin, F. Escourbiac, H. Frerichs, J. P. Gunn, T. Hirai, A. S. Kukushkin, E. Kaveeva, M. A. Miller, D. Moulton, V. Rozhansky, I. Senichenkov, E. Sytova, O. Schmitz, P. C. Stangeby, G. De Temmerman, I. Veselova, and S. Wiesen. Physics basis for the first ITER tungsten divertor. *Nuclear Materials and Energy*, 20: 100696, 2019. doi: 10.1016/j.nme.2019.100696.
- [2] P. C. Stangeby, J. D. Lore, R. A. Pitts, J. M. Canik, and X. Bonnin. A reduced model for the ITER divertor based on SOLPS solutions for ITER $q = 10$ baseline conditions: B. A reduced model based on reversed-direction two point modeling. *Nuclear Fusion*, 63(1):016017, 2023. doi: 10.1088/1741-4326/ac9917. Published online 13 December 2022; print issue January 2023. Companion part A: Nucl. Fusion 63, 016016 (2023), doi:10.1088/1741-4326/ac9916. Cited in this repository as “Stangeby 2022, Nucl. Fusion 62”, which is incorrect.
- [3] J. D. Lore, X. Bonnin, J.-S. Park, R. A. Pitts, and P. C. Stangeby. High gas throughput SOLPS-ITER simulations extending the ITER database to strong detachment. *Nuclear Fusion*, 62(10):106017, 2022. doi: 10.1088/1741-4326/ac8a5f.
- [4] S. I. Krasheninnikov and A. S. Kukushkin. Physics of ultimate detachment of a tokamak divertor plasma. *Journal of Plasma Physics*, 83(5):155830501, 2017. doi: 10.1017/S0022377817000654.
- [5] Wei Zhao, Yali Wang, Yuzhong Jin, Li Zhao, Hongxia Zhou, Lin Nie, Guangwu Zhong, Chunja Liu, Christopher Watts, and James Paul Gunn. Preliminary design of ITER divertor Langmuir probe system. *Fusion Science and Technology*, 76(2): 79–87, 2020. doi: 10.1080/15361055.2019.1674123.
- [6] D. R. Bates, A. E. Kingston, and R. W. P. McWhirter. Recombination between electrons and atomic ions, I. Optically thin plasmas. *Proceedings of the Royal Society of London. Series A. Mathematical and Physical Sciences*, 267(1330):297–312, 1962. doi: 10.1098/rspa.1962.0101.
- [7] A. Yu. Pigarov, J. L. Terry, and B. Lipschultz. Study of the discrete-to-continuum transition in a Balmer spectrum from Alcator C-Mod divertor plasmas. *Plasma Physics and Controlled Fusion*, 40(12):2055–2072, 1998. doi: 10.1088/0741-3335/40/12/006.
- [8] H. P. Summers, W. J. Dickson, M. G. O’Mullane, N. R. Badnell, A. D. Whiteford, D. H. Brooks, J. Lang, S. D. Loch, and D. C. Griffin. Ionization state, excited populations and emission of impurities in dynamic finite density plasmas: I. The generalized collisional–radiative model for light elements. *Plasma Physics and Controlled Fusion*, 48(2):263–293, 2006. doi: 10.1088/0741-3335/48/2/007.

- [9] H. P. Summers. The ADAS user manual, version 2.6. <http://www.adas.ac.uk>, 2004. Atomic Data and Analysis Structure, University of Strathclyde.
- [10] Takashi Fujimoto and R. W. P. McWhirter. Validity criteria for local thermodynamic equilibrium in plasma spectroscopy. *Physical Review A*, 42(11):6588–6601, 1990. doi: 10.1103/PhysRevA.42.6588.
- [11] Takashi Fujimoto. *Plasma Spectroscopy*, volume 123 of *International Series of Monographs on Physics*. Clarendon Press, Oxford, 2004. ISBN 978-0-19-853028-2. doi: 10.1093/acprof:oso/9780198530282.001.0001.
- [12] B. A. Lomanowski, A. G. Meigs, R. M. Sharples, M. Stamp, and C. Guillemaut. Inferring divertor plasma properties from hydrogen Balmer and Paschen series spectroscopy in JET-ILW. *Nuclear Fusion*, 55(12):123028, 2015. doi: 10.1088/0029-5515/55/12/123028.
- [13] D. Lumma, J. L. Terry, and B. Lipschultz. Radiative and three-body recombination in the Alcator C-Mod divertor. *Physics of Plasmas*, 4(7):2555–2566, 1997. doi: 10.1063/1.872234.
- [14] G. M. McCracken, M. F. Stamp, R. D. Monk, A. G. Meigs, J. Lingertat, R. Prentice, A. Starling, R. J. Smith, and A. Tabasso. Evidence for volume recombination in JET detached divertor plasmas. *Nuclear Fusion*, 38(4):619–629, 1998. doi: 10.1088/0029-5515/38/4/311.
- [15] K. Verhaegh, B. Lipschultz, B. P. Duval, A. Fil, M. Wensing, C. Bowman, and D. S. Gahle. Novel inferences of ionisation and recombination for particle/power balance during detached discharges using deuterium Balmer line spectroscopy. *Plasma Physics and Controlled Fusion*, 61(12):125018, 2019. doi: 10.1088/1361-6587/ab4f1e.
- [16] T. Eich, B. Sieglin, A. J. Thornton, M. Faitsch, A. Kirk, A. Herrmann, and W. Suttrop. ELM divertor peak energy fluence scaling to ITER with data from JET, MAST and ASDEX Upgrade. *Nuclear Materials and Energy*, 12:84–90, 2017. doi: 10.1016/j.nme.2017.04.014.
- [17] A. Loarte, G. Saibene, R. Sartori, D. Campbell, M. Becoulet, L. Horton, T. Eich, A. Herrmann, G. Matthews, N. Asakura, A. Chankin, A. Leonard, G. Porter, G. Federici, G. Janeschitz, M. Shimada, and M. Sugihara. Characteristics of type I ELM energy and particle losses in existing devices and their extrapolation to ITER. *Plasma Physics and Controlled Fusion*, 45(9):1549–1569, 2003. doi: 10.1088/0741-3335/45/9/302.
- [18] A. Kirk, I. T. Chapman, T. E. Evans, C. Ham, J. R. Harrison, G. Huijsmans, Y. Liang, Y. Q. Liu, A. Loarte, W. Suttrop, and A. J. Thornton. Understanding the effect resonant magnetic perturbations have on ELMs. *Plasma Physics and Controlled Fusion*, 55(12):124003, 2013. doi: 10.1088/0741-3335/55/12/124003.
- [19] Keiji Sawada and Takashi Fujimoto. Temporal relaxation of excited-level populations of atoms and ions in a plasma: Validity range of the quasi-steady-state solution of coupled rate equations. *Physical Review E*, 49(6):5565–5573, 1994. doi: 10.1103/PhysRevE.49.5565.

- [20] P. T. Greenland. On the validity of collisional–radiative models. *Journal of Nuclear Materials*, 290–293:615–618, 2001. doi: 10.1016/S0022-3115(00)00601-2.
- [21] P. T. Greenland. Collisional-radiative models with molecules. *Proceedings of the Royal Society of London. Series A: Mathematical, Physical and Engineering Sciences*, 457(2012):1821–1839, 2001. doi: 10.1098/rspa.2001.0788.
- [22] Hans R. Griem. Validity of local thermal equilibrium in plasma spectroscopy. *Physical Review*, 131(3):1170–1176, 1963. doi: 10.1103/PhysRev.131.1170.
- [23] Hans R. Griem. *Principles of Plasma Spectroscopy*, volume 2 of *Cambridge Monographs on Plasma Physics*. Cambridge University Press, Cambridge, 1997. ISBN 978-0-521-45504-6. doi: 10.1017/CBO9780511524578.
- [24] Chul Park. Hydrogen line ratios as electron temperature indicators in nonequilibrium plasmas. *Journal of Quantitative Spectroscopy and Radiative Transfer*, 12(3):323–370, 1972. doi: 10.1016/0022-4073(72)90050-7.
- [25] K. Verhaegh, B. Lipschultz, C. Bowman, B. P. Duval, U. Fantz, A. Fil, J. R. Harrison, D. Moulton, O. Myatra, D. Wunderlich, F. Federici, D. S. Gahle, A. Perek, and M. Wensing. A novel hydrogenic spectroscopic technique for inferring the role of plasma–molecule interaction on power and particle balance during detached conditions. *Plasma Physics and Controlled Fusion*, 63(3):035018, 2021. doi: 10.1088/1361-6587/abd4c0.
- [26] T. A. Wijkamp, J. S. Allcock, X. Feng, B. Kool, B. Lipschultz, K. Verhaegh, B. P. Duval, J. R. Harrison, L. Kogan, N. Lonigro, A. Perek, P. Ryan, R. M. Sharples, I. G. J. Classen, and R. J. E. Jaspers. Characterisation of detachment in the MAST-U Super-X divertor using multi-wavelength imaging of 2D atomic and molecular emission processes. *Nuclear Fusion*, 63(5):056003, 2023. doi: 10.1088/1741-4326/acc191.
- [27] Igor Bray and Andris T. Stelbovics. Convergent close-coupling calculations of electron-hydrogen scattering. *Physical Review A*, 46(11):6995–7011, 1992. doi: 10.1103/PhysRevA.46.6995.
- [28] I. Bray, D. V. Fursa, A. S. Kheifets, and A. T. Stelbovics. Electrons and photons colliding with atoms: development and application of the convergent close-coupling method. *Journal of Physics B: Atomic, Molecular and Optical Physics*, 35(15):R117–R146, 2002. doi: 10.1088/0953-4075/35/15/201.
- [29] L. Vriens and A. H. M. Smeets. Cross-section and rate formulas for electron-impact ionization, excitation, deexcitation, and total depopulation of excited atoms. *Physical Review A*, 22(3):940–951, 1980. doi: 10.1103/PhysRevA.22.940.
- [30] A. Hartgers, J. van Dijk, J. Jonkers, and J. A. M. van der Mullen. CRModel: A general collisional radiative modeling code. *Computer Physics Communications*, 135(2):199–218, 2001. doi: 10.1016/S0010-4655(00)00231-9.
- [31] Wolfgang Lotz. Electron-impact ionization cross-sections and ionization rate coefficients for atoms and ions from hydrogen to calcium. *Zeitschrift für Physik*, 216(3):241–247, 1968. doi: 10.1007/BF01392963.

- [32] L. C. Johnson. Approximations for collisional and radiative transition rates in atomic hydrogen. *The Astrophysical Journal*, 174:227, 1972. doi: 10.1086/151486.
- [33] M. J. Seaton. Radiative recombination of hydrogenic ions. *Monthly Notices of the Royal Astronomical Society*, 119(2):81–89, 1959. doi: 10.1093/mnras/119.2.81.
- [34] Junjie Mao and Jelle Kaastra. Parameterization of the level-resolved radiative recombination rate coefficients for the SPEX code. *Astronomy & Astrophysics*, 587: A84, 2016. doi: 10.1051/0004-6361/201527568.
- [35] D. Hoang-Binh. A program to compute exact hydrogenic radial integrals, oscillator strengths, and Einstein coefficients, for principal quantum numbers up to $n \approx 1000$. *Computer Physics Communications*, 166(3):191–196, 2005. doi: 10.1016/j.cpc.2004.11.005. CPC Program Library ADUU_v1_0.
- [36] D. Hoang-Binh. An exact calculation of hydrogenic radial integrals and oscillator strengths, for principal quantum numbers up to N equal to about 1000. *Astronomy and Astrophysics*, 238:449–451, 1990. No DOI; ADS bibcode 1990A&A...238..449H.
- [37] W. L. Wiese and J. R. Fuhr. Accurate atomic transition probabilities for hydrogen, helium, and lithium. *Journal of Physical and Chemical Reference Data*, 38(3):565–720, 2009. doi: 10.1063/1.3077727.
- [38] T. Holstein. Imprisonment of resonance radiation in gases. *Physical Review*, 72(12): 1212–1233, 1947. doi: 10.1103/PhysRev.72.1212.
- [39] K. Behringer. Escape factors for line emission and population calculations. Technical Report IPP 10/11, Max-Planck-Institut für Plasmaphysik, Garching, Germany, 1998. No DOI. Implemented as ADAS214; see <https://www.adas.ac.uk/man/chap2-14.pdf>.
- [40] Eite Tiesinga, Peter J. Mohr, David B. Newell, and Barry N. Taylor. CODATA recommended values of the fundamental physical constants: 2018. *Reviews of Modern Physics*, 93(2):025010, 2021. doi: 10.1103/RevModPhys.93.025010.
- [41] R. M. Pengelly and M. J. Seaton. Recombination spectra: II. Collisional transitions between states of degenerate energy levels. *Monthly Notices of the Royal Astronomical Society*, 127(2):165–175, 1964. doi: 10.1093/mnras/127.2.165.
- [42] N. R. Badnell, F. Guzmán, S. Brodie, R. J. R. Williams, P. A. M. van Hoof, M. Chatzikos, and G. J. Ferland. H, He-like recombination spectra – IV. Clarification and refinement of methodology for l -changing collisions. *Monthly Notices of the Royal Astronomical Society*, 507(2):2922–2929, 2021. doi: 10.1093/mnras/stab2266.
- [43] H. Anderson, C. P. Ballance, N. R. Badnell, and H. P. Summers. An R-matrix with pseudostates approach to the electron-impact excitation of H I for diagnostic applications in fusion plasmas. *Journal of Physics B: Atomic, Molecular and Optical Physics*, 33(6):1255–1262, 2000. doi: 10.1088/0953-4075/33/6/311.
- [44] H. Anderson, C. P. Ballance, N. R. Badnell, and H. P. Summers. Corrigendum: An R-matrix with pseudo-states approach to the electron-impact excitation of H I for diagnostic applications in fusion plasmas. *Journal of Physics B: Atomic, Molecular and Optical Physics*, 35(6):1613–1615, 2002. doi: 10.1088/0953-4075/35/6/701.

- [45] C. Guillemaut, R. A. Pitts, A. S. Kukushkin, and M. O'Mullane. Radiative power loading in the ITER divertor. *Fusion Engineering and Design*, 86(12):2954–2964, 2011. doi: 10.1016/j.fusengdes.2011.07.008.
- [46] P. C. Stangeby, J. D. Lore, R. A. Pitts, J. M. Canik, and X. Bonnin. A reduced model for the ITER divertor based on SOLPS solutions for ITER $q = 10$ baseline conditions: A. Identifying options for the control parameters. *Nuclear Fusion*, 63(1):016016, 2023. doi: 10.1088/1741-4326/ac9916. Companion part A to [2]. Table 1 gives spans, not values: the divertor-averaged neutral pressure spans a factor 5.36 over this dataset. Sec. 6 records the extended database reaching $\langle p_{\text{div}} \rangle > 25$ Pa, citing [3].
- [47] G. Udny Yule. On the methods of measuring association between two attributes. *Journal of the Royal Statistical Society*, 75(6):579–652, 1912. doi: 10.2307/2340126.
- [48] Fumiko Samejima. *Estimation of Latent Ability Using a Response Pattern of Graded Scores*. Number 17 in Psychometrika Monograph Supplement. Psychometric Society, 1969.
- [49] T. W. Anderson. The integral of a symmetric unimodal function over a symmetric convex set and some probability inequalities. *Proceedings of the American Mathematical Society*, 6(2):170–176, 1955. doi: 10.1090/S0002-9939-1955-0069229-1.
- [50] Athel Cornish-Bowden. *Fundamentals of Enzyme Kinetics*. Wiley-Blackwell, Weinheim, 4 edition, 2012.
- [51] Jeremy A. Owen and Jordan M. Horowitz. Size limits the sensitivity of kinetic schemes. *Nature Communications*, 14:1280, 2023. doi: 10.1038/s41467-023-36705-8.
- [52] Nambury S. Raju. The area between two item characteristic curves. *Psychometrika*, 53(4):495–502, 1988. doi: 10.1007/BF02294403.